

CS100 Lecture 21

Inheritance and Polymorphism II

Contents

- Pure `virtual` functions and abstract classes
- Public inheritance for the "is-a" relationship (*Effective C++* Item 32)
- Inheritance of interface vs. inheritance of implementation (*Effective C++* Item 34)

Pure `virtual` functions and abstract classes

Shapes

Define classes to represent different shapes: Rectangle, Triangle, Circle, ...

Suppose we want to draw *different* shapes using a *single* function:

```
void drawShapes(ScreenHandle &screen,  
                const std::vector<std::shared_ptr<Shape>> &shapes) {  
    for (const auto &shape : shapes)  
        shape->draw(screen);  
} // Pointers in `shapes` can point to objects of different shape classes.
```

and print their information using a *single* function:

```
void printShapeInfo(const Shape &shape) {  
    std::cout << "Area: " << shape.area()  
               << "Perimeter: " << shape.perimeter() << std::endl;  
} // `shape` can be bound to objects of different shape classes.
```

What is the relationship between `Shape` and other shape classes?

Shapes

Define a base class `Shape` and let other shape classes inherit it.

```
class Shape {  
public:  
    Shape() = default;  
    virtual void draw(ScreenHandle &screen) const;  
    virtual double area() const;  
    virtual double perimeter() const;  
    virtual ~Shape() = default;  
};
```

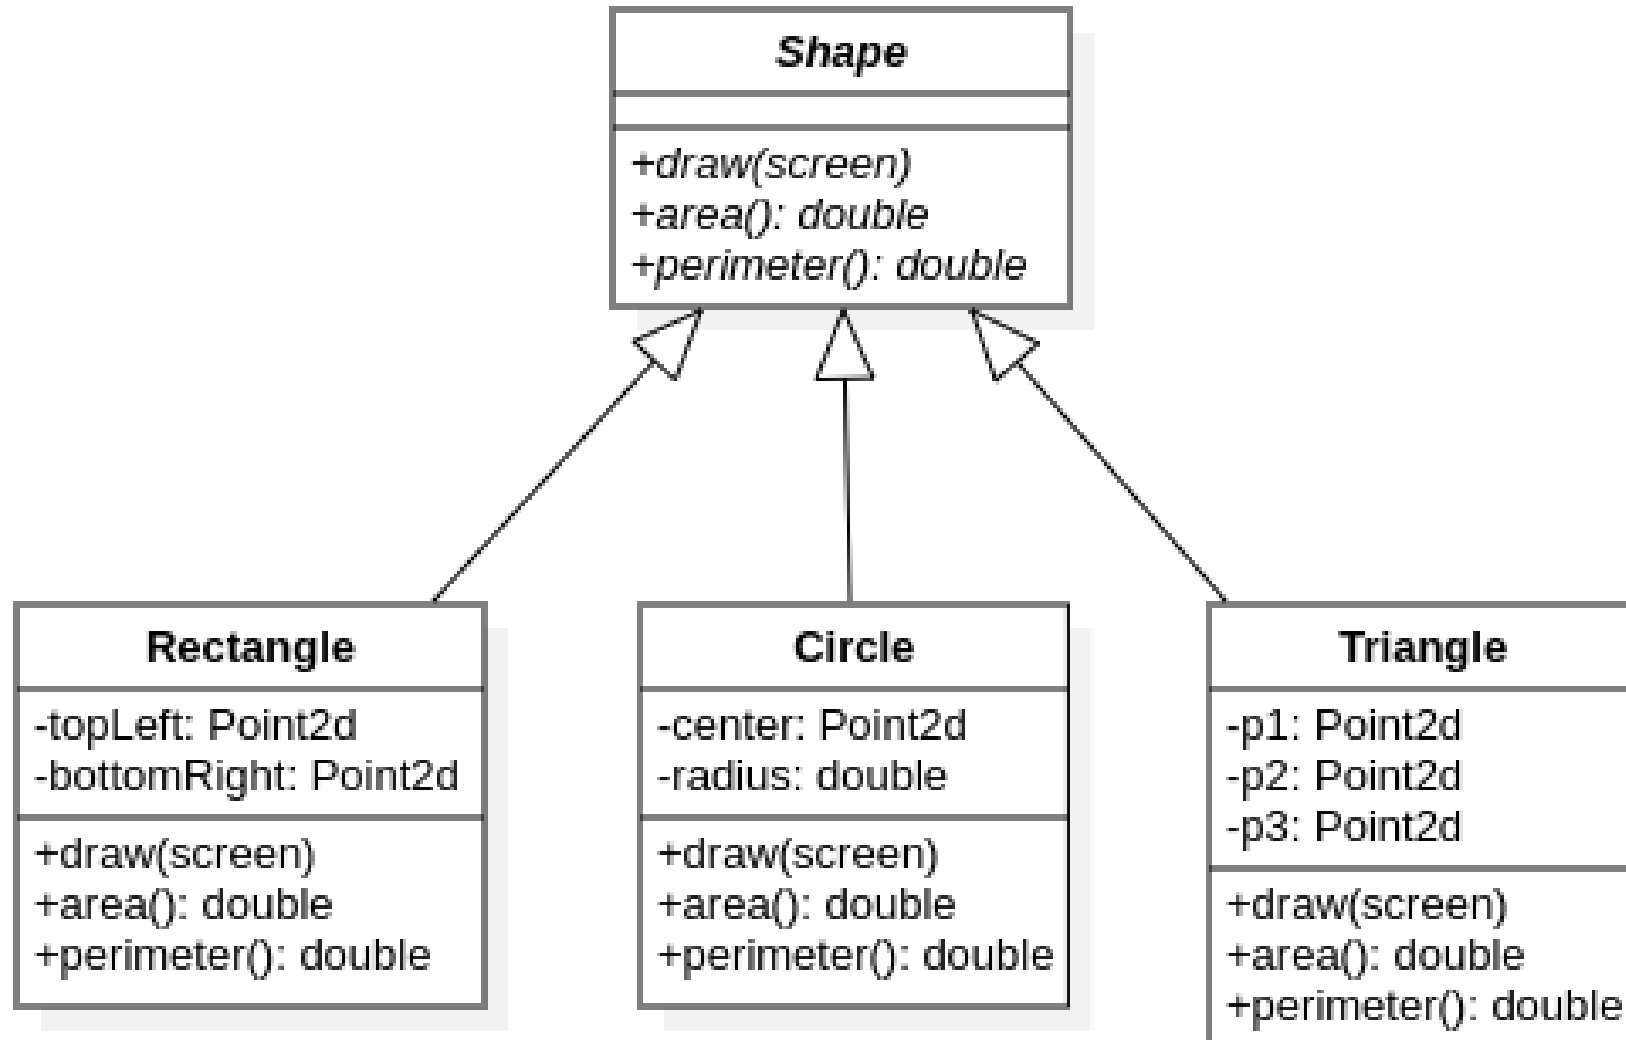
Different shape classes should define their own `draw`, `area` and `perimeter`, so these functions should be `virtual` in `Shape` (and thus can be overridden in the subclasses).

Shapes

```
class Rectangle : public Shape {
    Point2d mTopLeft, mBottomRight;

public:
    Rectangle(const Point2d &tl, const Point2d &br)
        : mTopLeft(tl), mBottomRight(br) {} // call the default ctor of `Shape`
    void draw(ScreenHandle &screen) const override { /* ... */ }
    double area() const override {
        return (mBottomRight.x - mTopLeft.x) * (mBottomRight.y - mTopLeft.y);
    }
    double perimeter() const override {
        return 2 * (mBottomRight.x - mTopLeft.x + mBottomRight.y - mTopLeft.y);
    }
};
```

Shapes



Pure **virtual** functions

How should we define `Shape::draw`, `Shape::area` and `Shape::perimeter` ?

- For the general concept "Shape", there is no way to determine the behaviors of these functions.

Pure `virtual` functions

How should we define `Shape::draw`, `Shape::area` and `Shape::perimeter` ?

- For the general concept "Shape", there is no way to determine the behaviors of these functions.
- Direct call to `Shape::draw`, `Shape::area` and `Shape::perimeter` should be forbidden.
- We shouldn't even allow an object of type `Shape` to be instantiated! The class `Shape` is only used to **define the concept "Shape" and provide required interfaces (declarations of functions that any shape class needs to support)**.

Pure `virtual` functions

If a `virtual` function does not have a reasonable definition in the base class, it should be declared as **pure `virtual`** by writing `=0` .

```
class Shape {  
public:  
    virtual void draw(ScreenHandle &) const = 0;  
    virtual double area() const = 0;  
    virtual double perimeter() const = 0;  
    virtual ~Shape() = default;  
};
```

Any class that has a **pure `virtual`** function is an **abstract class**. Pure `virtual` functions (usually) cannot be called ¹, and abstract classes cannot be instantiated.

Pure **virtual** functions and abstract classes

Any class that has a **pure virtual** function is an **abstract class**. Pure **virtual** functions (usually) cannot be called ¹, and abstract classes cannot be instantiated.

```
Shape shape; // Error.  
Shape *p = new Shape; // Error.  
auto sp = std::make_shared<Shape>(); // Error.  
std::shared_ptr<Shape> sp2 = std::make_shared<Rectangle>(p1, p2); // OK.
```

We can define a pointer or reference to an abstract class, but never an object of that type!

Pure `virtual` functions and abstract classes

An impure `virtual` function **must be defined**. Otherwise, the compiler will fail to generate necessary runtime information (the virtual table), which leads to an error.

```
class X {  
    virtual void foo(); // Declaration, without a definition  
    // Even if `foo` is not used, this will lead to an error.  
};
```

Make the interface robust, not error-prone

Besides declaring `Shape::draw`, `Shape::area`, `Shape::perimeter` as pure `virtual`, other possible designs can be:

```
class Shape {
public:
    virtual double area() const {
        return 0; // Is this good?
    }
};
```

```
class Shape {
public:
    virtual double area() const {
        throw std::logic_error{"area() called on Shape!"}; // What about this?
    }
};
```

Make the interface robust, not error-prone

```
class Shape {  
public:  
    virtual double area() const {  
        return 0;  
    }  
};
```

If `Shape::area` is called accidentally, the error will happen *silently* at run-time!

Make the interface robust, not error-prone

```
class Shape {  
public:  
    virtual double area() const {  
        throw std::logic_error{"area() called on Shape!"};  
    }  
};
```

If `Shape::area` is called accidentally, an exception will be raised at run-time.

However, a **good design should make errors fail to compile**, which can be achieved by declaring `Shape::area` as pure `virtual`.

[Best practice] If an error can be caught in compile-time, don't leave it until run-time.

Polymorphism (多态)

Polymorphism: The same function name can invoke different behaviors in different calling contexts.

- Run-time polymorphism: **dynamic binding** (contexts: types of objects being referred to by a base class pointer or reference on which the function is called).
- Compile-time polymorphism: **function overloading** (contexts: number or types of the function arguments, or `const` ness of objects on which the function is called).

Run-time polymorphism:

```
struct Shape {  
    virtual void draw() const = 0;  
};  
void drawShape(const Shape &s) {  
    s.draw();  
}
```

Compile-time polymorphism:

```
struct Rectangle {  
    void draw() const;  
    void draw(const std::string &str) const;  
};  
Rectangle r;  
r.draw();  
r.draw("rectangle");
```


Public inheritance for the "is-a" relationship

Effective C++ Item 32

Public inheritance: The "is-a" relationship

By writing that class `Sub` publicly inherits from class `Base`, you are telling the compiler (as well as human readers of your code) that

- Every object of type `Sub` *is an* object of type `Base`, but not vice versa.
- `Base` represents a **more general concept** than `Sub`, and that `Sub` represents a **more specialized concept** than `Base`.

More specifically, you are asserting that **anywhere an object of type `Base` can be used, an object of type `Sub` can be used just as well.**

- However, if you need an object of type `Sub` somewhere, an object of type `Base` can not be used.

Example: Every student *is a* person

```
class Person { /* ... */ };  
class Student : public Person { /* ... */ };
```

- Every student *is a* person, but not every person is a student.
- Anything that is true of a person is also true of a student:
 - A person has a date of birth, so does a student.
- Something is true of a student, but not true of people in general.
 - A student is enrolled in a particular school, but a person may not.

The notion of a person is **more general** than that of a student; a student is a **specialized type** of person.

Example: Every student *is a* person

The is-a relationship: Anywhere an object of type `Person` can be used, an object of type `Student` can be used just as well, **but not vice versa**.

```
void eat(const Person &p);    // Anyone can eat.
void study(const Student &s); // Only students study.
Person p;
Student s;
eat(p);    // Fine. `p` is a person.
eat(s);    // Fine. `s` is a student, and a student is a person.
study(s);  // Fine.
study(p);  // Error! `p` isn't a student.
```

Your intuition can mislead you

- A penguin *is a* bird.
- A bird can fly.

If we naively try to express this in C++, our effort yields:

```
class Bird {  
public:  
    virtual void fly();          // Birds can fly.  
    // ...  
};  
class Penguin : public Bird { // A penguin is a bird.  
    // ...  
};
```

```
Penguin p;  
p.fly();    // Oh no!! Penguins cannot fly, but this code compiles!
```

Not every bird can fly

In general, only birds with the flying ability can fly.

- There are several types of non-flying birds.

Maybe the following hierarchy models the reality much better?

```
class Bird { /* ... */ };
class FlyingBird : public Bird {
    virtual void fly();
};
class Penguin : public Bird {    // Not FlyingBird
    // ...
};
```

Not every bird can fly

Maybe the following hierarchy models the reality much better?

```
class Bird { /* ... */ };  
class FlyingBird : public Bird {  
    virtual void fly();  
};  
class Penguin : public Bird {    // Not FlyingBird  
    // ...  
};
```

- **Not necessary.** If your application has much to do with beaks and wings, and nothing to do with flying, the original two-class hierarchy might be satisfactory.
- **There is no one ideal design for every software.** The best design depends on what the system is expected to do.

What about reporting a runtime error?

To prevent the attempt to make penguins fly, another way is to redefine `Penguin::fly` to report an error at run-time:

```
void report_error(const std::string &msg); // defined elsewhere
class Penguin : public Bird {
public:
    virtual void fly() {
        report_error("Attempt to make a penguin fly!");
    }
};
```


What about reporting a runtime error?

```
void report_error(const std::string &msg); // defined elsewhere
class Penguin : public Bird {
public:
    virtual void fly() { report_error("Attempt to make a penguin fly!"); }
};
```

This does not say "Penguins can't fly." This says "Penguins can fly, but it is an error for them to actually try to do it."


To actually express the constraint "Penguins can't fly", you should **prevent the attempt from compiling**, which can be achieved by removing `Penguin::fly`.

```
Penguin p;
p.fly(); // This should not compile.
```

[Best practice] Good interfaces prevent invalid code at compile-time.

Example: A square *is a* rectangle

Should class `Square` publicly inherit from class `Rectangle` ?



Example: A square *is a* rectangle

Consider this code:

```
class Rectangle {
public:
    virtual void setHeight(int newHeight);
    virtual void setWidth(int newWidth);
    virtual int getHeight() const;
    virtual int getWidth() const;
    // ...
};

void makeBigger(Rectangle &r) {
    r.setWidth(r.getWidth() + 10);
}
```

```
class Square : public Rectangle {
    // A square is a rectangle,
    // where height == width.
    // ...
};

Square s(10); // A 10x10 square.
makeBigger(s); // Oh no!
```

Is public inheritance suitable here?

Public inheritance models "is-a", which asserts that **everything that applies to base classes must also apply to subclasses**, because every subclass object is a base class object.

However, something applicable to a rectangle (i.e., height is independent of width) is not applicable to a square whose height and width must be the same!

Thus, public inheritance is not suitable for modeling the "is-a" relationship between rectangle and square.

Public inheritance: The "is a" relationship

You need to be careful when using public inheritance to model "is a" in practice.

- To decide if public inheritance should be used, ask: **Does everything that applies to base classes are also applicable to subclasses?**

Not all "is-a" relationships can be naively modeled with public inheritance:

- "A penguin is a bird": something applicable to a bird (flying) is not applicable to a penguin.
- "A square is a rectangle": something applicable to a rectangle (independence of height and width) is not applicable to a square.

Inheritance of interface vs. inheritance of implementation

Effective C++ Item 34

Example: Airplanes for XYZ Airlines

Suppose XYZ has only two kinds of planes: Model A and Model B, and both are flown in exactly the same way.

```
class Airplane {  
public:  
    virtual void fly(const Airport &destination) {  
        // Default code for flying an airplane to the given destination.  
    }  
};  
class ModelA : public Airplane { /* ... */ };  
class ModelB : public Airplane { /* ... */ };
```

- `Airplane::fly` is declared `virtual` because in principle, different airplanes should be flown in different ways.
- `Airplane::fly` is defined, to avoid code repetition in `ModelA` and `ModelB` classes.

Example: Airplanes for XYZ Airlines

Now suppose that XYZ decides to acquire a new type of airplane, Model C, which is flown differently from Model A and Model B.

XYZ's programmers add class `ModelC` to the hierarchy, but forget to redefine the `fly` function!

```
class ModelC : public Airplane {  
    // `fly` is not overridden.  
    // ...  
};
```

This surely leads to a disaster:

```
auto pc = std::make_unique<ModelC>();  
pc->fly(PVG); // No! Attempt to fly Model C in the Model A/B way!
```


Impure virtual function: Interface + default implementation

The problem here is not that `Airplane::fly` has default behavior, but that `ModelC` was allowed to inherit that behavior without explicitly saying that it needs the behavior.

By defining an impure virtual function in a base class, we have the subclass inherit a function interface as well as a default implementation.

- Interface: Every class inheriting from `Airplane` can `fly`.
- Default implementation: If `ModelC` does not override `Airplane::fly`, it will have the inherited implementation automatically.

Separate default implementation from interface

To cut off the connection between the *interface* and its *default implementation*:

```
class Airplane {  
public:  
    virtual void fly(const Airport &destination) = 0; // pure virtual  
    // ...  
protected:  
    void defaultFly(const Airport &destination) {  
        // Default code for flying an airplane to the given destination.  
    }  
};
```

- The pure virtual function `fly` provides the **interface**: Every derived class can `fly`.
- The **default implementation** is written in an independent function `defaultFly`.

Separate default implementation from interface

If `ModelA` and `ModelB` want to adopt the default way of flying, they simply make a call to `defaultFly`.

```
class ModelA : public Airplane {
public:
    virtual void fly(const Airport &destination) {
        defaultFly(destination);
    }
    // ...
};

class ModelB : public Airplane {
public:
    virtual void fly(const Airport &destination) {
        defaultFly(destination);
    }
    // ...
};
```

Separate default implementation from interface

For `ModelC` :

- Since `Airplane::fly` is pure `virtual`, `ModelC` must define its own version of `fly`.
- If it **does** want to use the default implementation, it **must say it explicitly** by making a call to `defaultFly`.

```
class ModelC : public Airplane {  
public:  
    virtual void fly(const Airport &destination) {  
        // The "Model C way" of flying.  
        // Without the definition of this function, `ModelC` remains abstract,  
        // which does not compile if we create an object of such type.  
    }  
};
```

Still not satisfactory?

Some people object to the idea of having separate functions for providing the interface and the default implementation, such as `fly` and `defaultFly` above.

- It pollutes the class namespace with closely related function names.
 - This really matters, especially in complicated projects. Extra mental effort might be required to distinguish the meaning of overly similar names.

Read the rest part of *Effective C++* Item 34 for another solution to this problem.

Inheritance of interface vs. inheritance of implementation

Pure virtual functions specify **inheritance of interface** only.

- The base class tells the subclass: "You must support these functions, but I don't know how to implement them".

Impure virtual functions specify **inheritance of interface + default implementation**.

- The base class tells the subclass: "You must support these functions, and you can redefine (override) them or use the default implementations directly".

Non-virtual functions specify **inheritance of interface + mandatory implementation**.

- The base class tells the subclass: "You must support these functions and use my implementations".

Summary

Pure `virtual` functions and abstract classes

- A pure `virtual` function is a `virtual` function declared `= 0`.
 - Call to a pure `virtual` function is not allowed. ¹
 - Pure `virtual` functions provide interfaces and let the subclasses implement them.
- A class that has a pure `virtual` function is an abstract class.
 - We cannot create an object of an abstract class type.
 - Abstract classes are often used to represent abstract, general concepts.

Summary

Public inheritance models the "is-a" relationship.

- Everything that applies to base classes must also apply to subclasses.
- The "A bird can fly, and a penguin is a bird" example.
- The "A square is a rectangle" example.

Summary

Inheritance of interface vs. inheritance of implementation

- Pure virtual functions: inheritance of **interface** only.
- Impure virtual functions: inheritance of **interface + default implementation**.
- Non-virtual functions: inheritance of **interface + mandatory implementation**.
- In public inheritance, interfaces are always inherited.

Notes

¹ A pure `virtual` function can have a definition. In that case, it can be called via the syntax `ClassName::functionName(args)`, not via a `virtual` function call (dynamic binding).

In some cases, we want a class to be made abstract, but it does not have any pure `virtual` function. A possible workaround is to declare the destructor to be pure `virtual`, and then provide a definition for it:

```
struct Foo {  
    virtual ~Foo() = 0;  
};  
Foo::~~Foo() = default; // Provide a definition outside the class.
```

The "another solution" mentioned in page 37 is also related to this.