CS120: Computer Networks

Lecture 11. BGP

Zhice Yang

Routing Protocols

- Routing Information Protocol (RIP)
  - Algorithm: Distance Vector
- Open Shortest Path First (OSPF)
  - Algorithm: Link State
- Border Gateway Protocol (BGP)

Intradomain Routing Protocol

Interdomain Routing Protocol

The Discussion on Routing So Far …

- Routers in the network are managed by the same administrator
  - e.g., Residential building, Campus, Network of a big company, etc.
- Routers are running same routing protocol
  - e.g., OSPF or RIP.
  - These routers have certain coverage, called the routing area

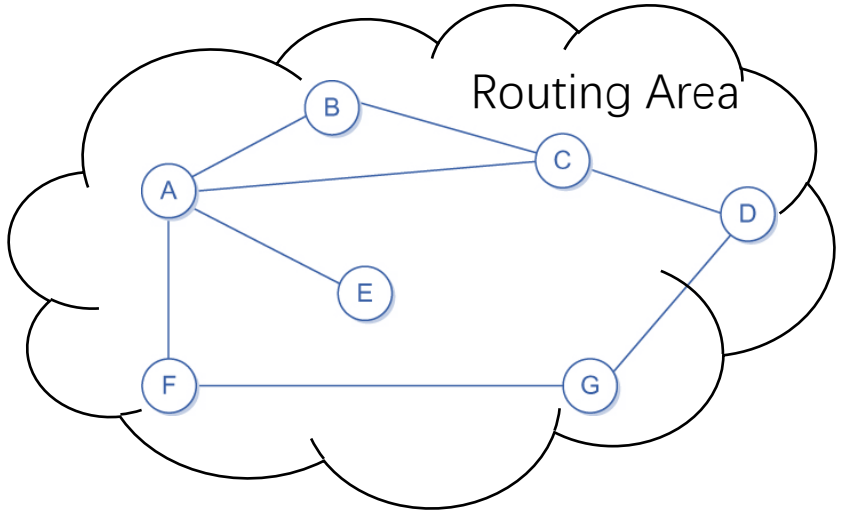

The Real Internet: Network of Network

Internet Service Providers

- End systems connect to Internet via “Consumer” ISPs (Internet Service Providers)
  - Residential, company, and university ISPs
- “Consumer” ISPs are connected “Backbone” ISPs
  - Three Major Commercial ISPs in China

  - Other ISPs
    - e.g. cernet

Network of Networks

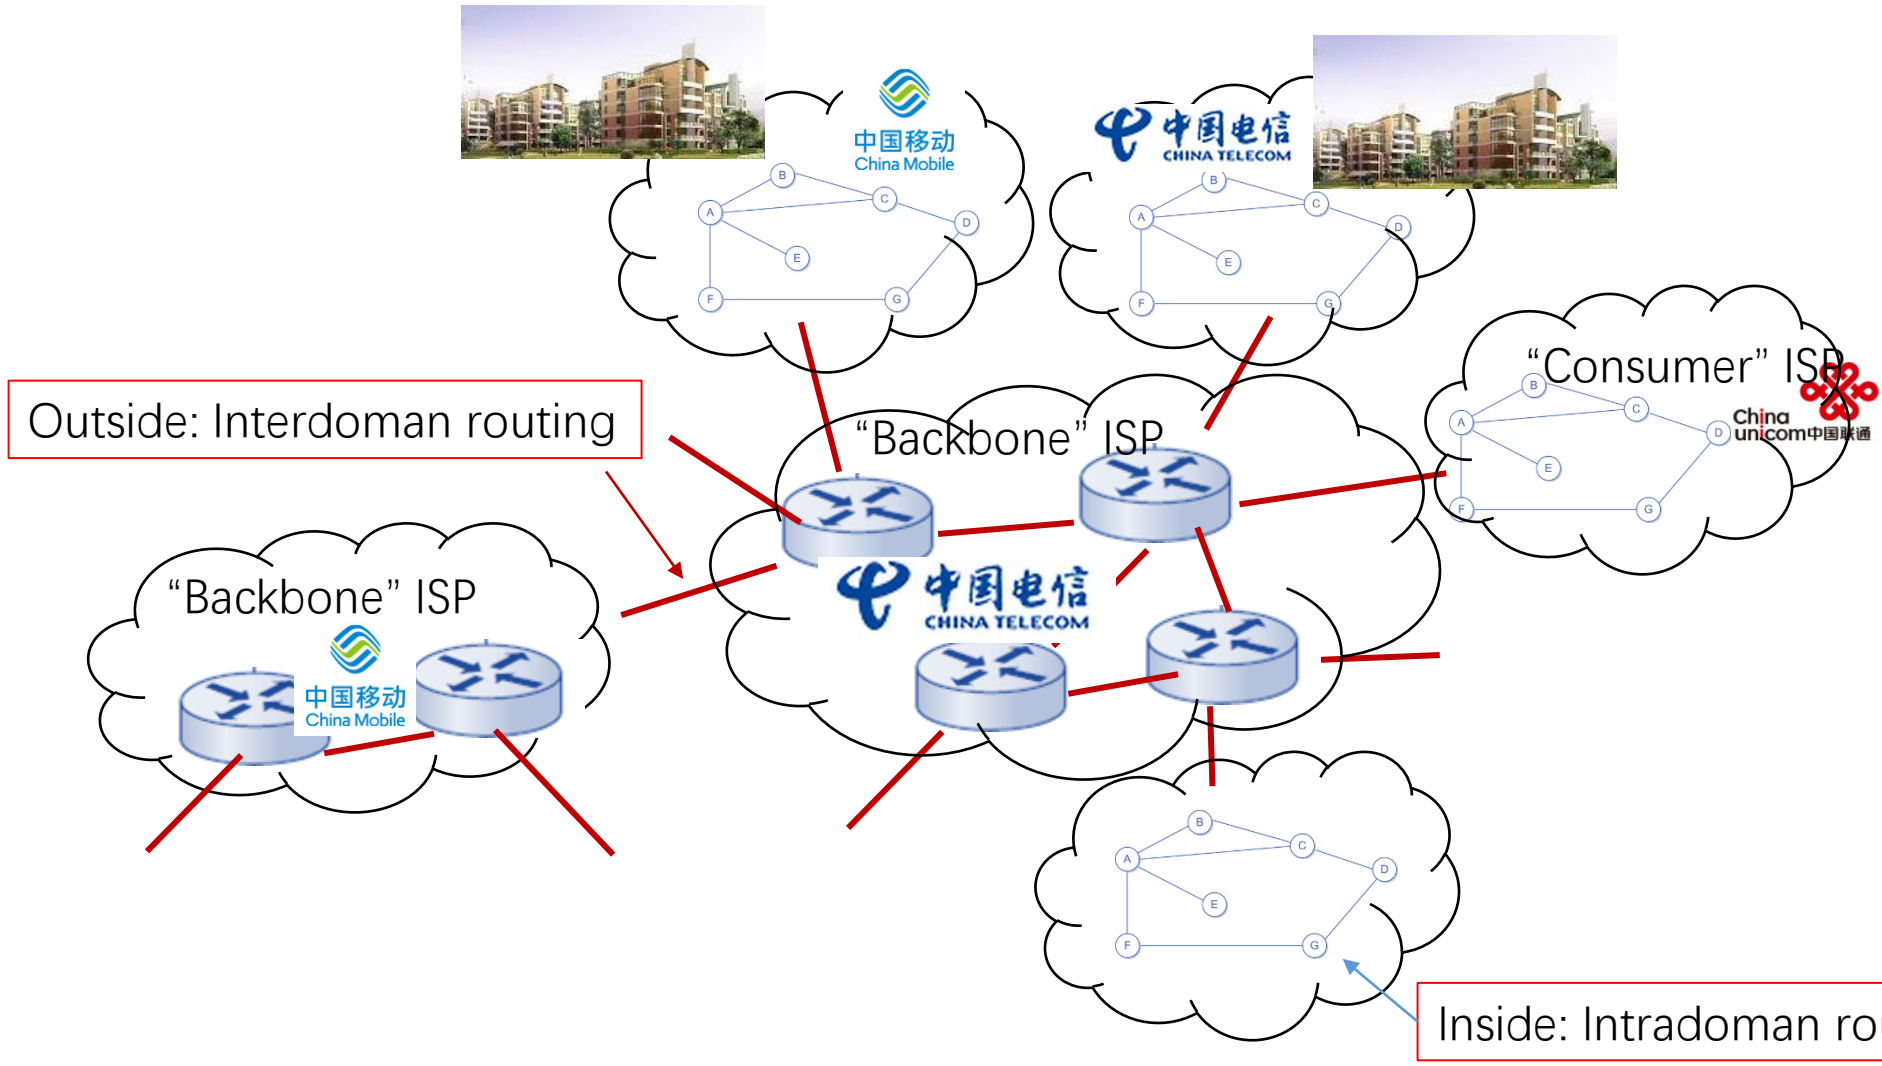

Interdomain Routing Problems

- Scalability: more than 600 million destinations
  - Storage overhead
    - Routing table
  - Calculation overhead
    - Shortest path
  - Communication overhead
    - Exchange routing information
- Routing Management: Complex Routing Polices

Interdomain Routing: New Hierarchy

- Aggregate Routers into Logical Areas: Autonomous System
- Autonomous System (AS)
  - Corresponds to an administrative domain
  - e.g. University, company, backbone network
- Routers in same AS run the same intradomain routing protocol
  - RIP, OSPF, etc.
- Routers in different AS run intrerdomain routing protocol
  - BGP, EGP
- Interdomain routing element: AS

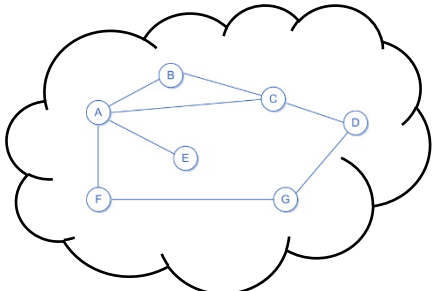

Autonomous System (AS)

Routing Hierarchy

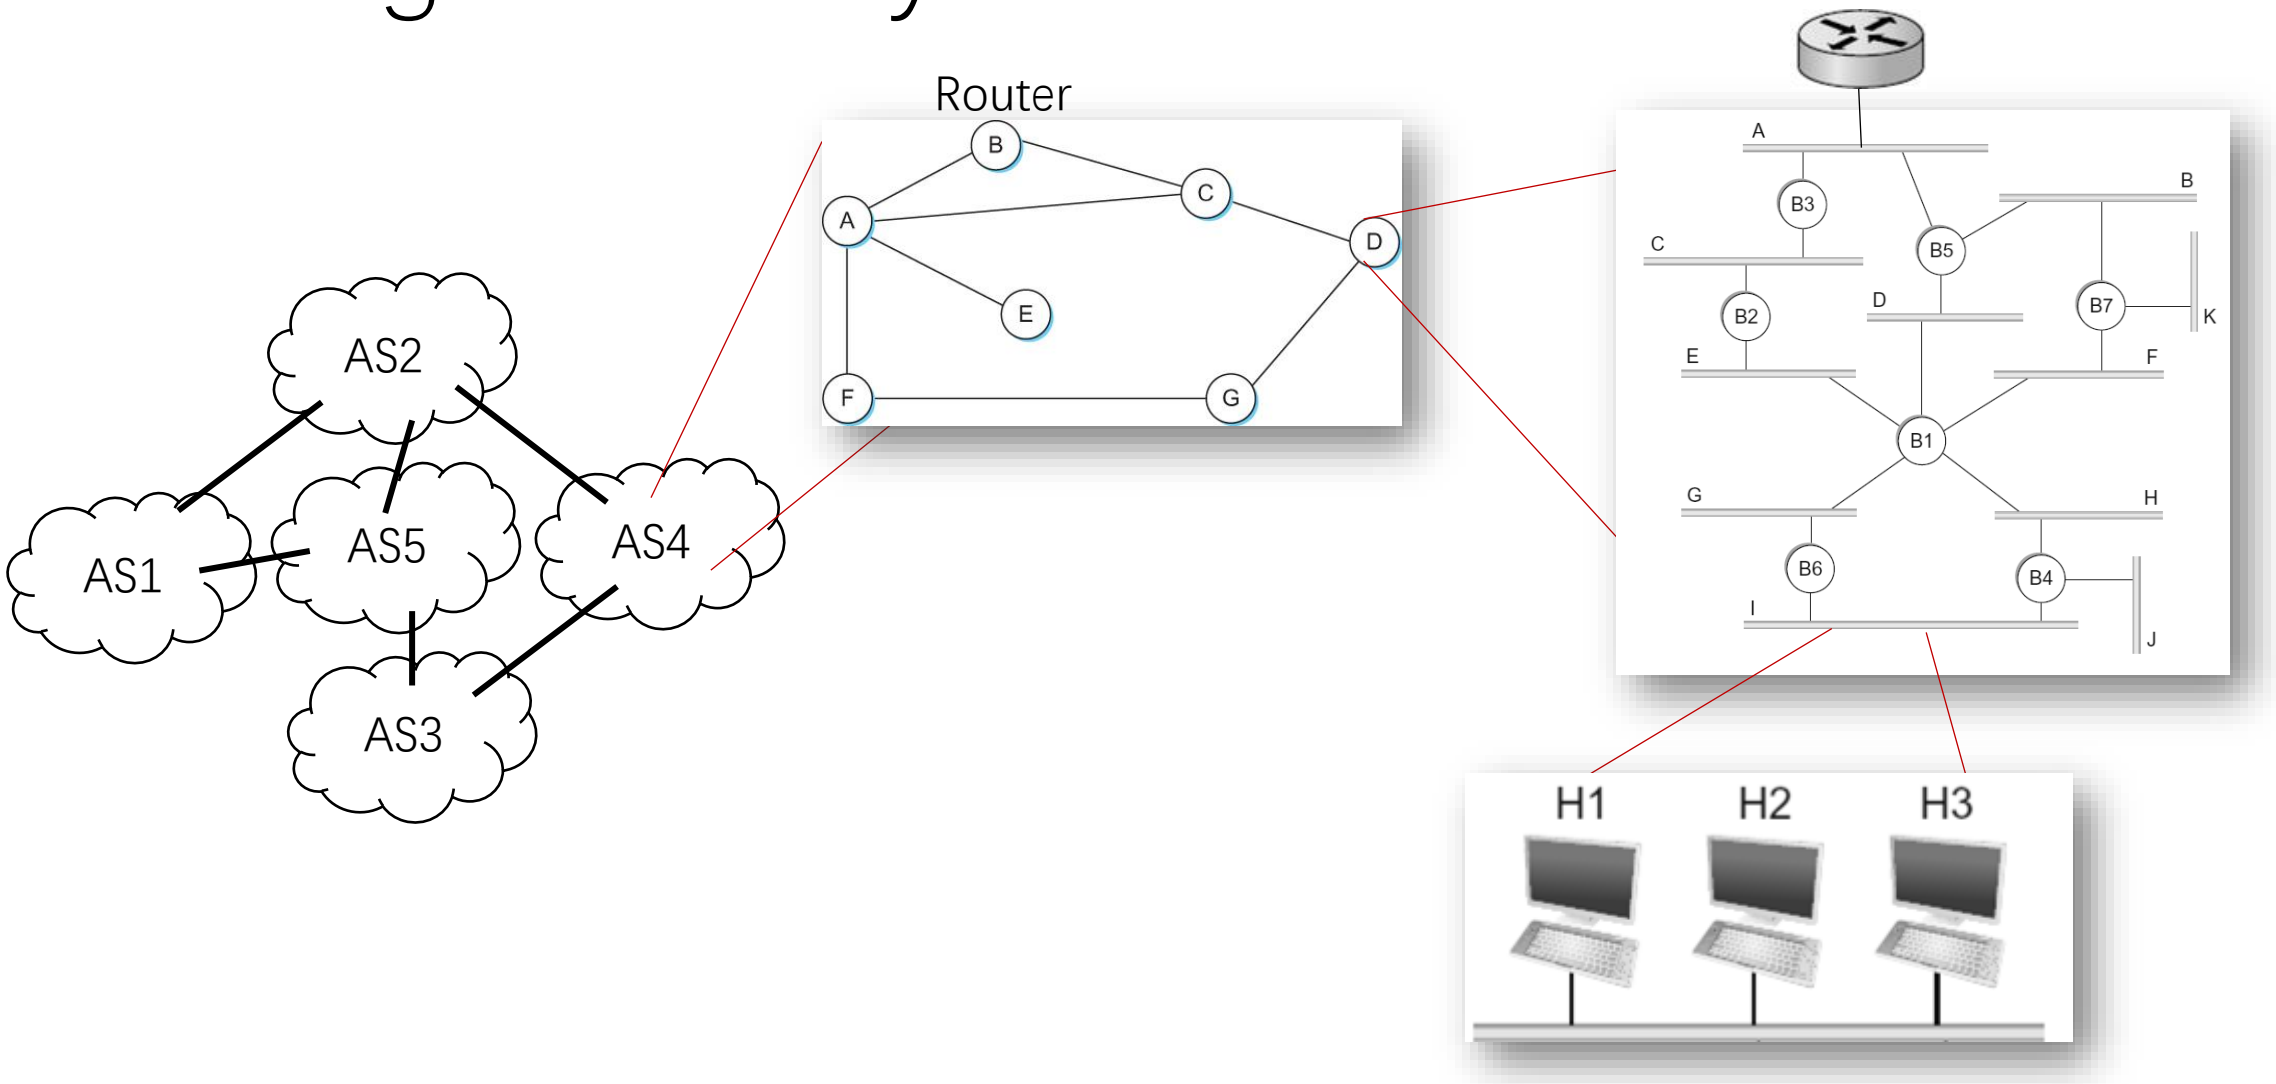

Border Gateway Protocol (BGP)

- Widely-used Interdomain Routing Protocol
- Routing Element: AS
- Routing Algorithm
  - Target on Reachability
    - Not the “shortest” route
  - Avoid Loops

Border Gateway Protocol (BGP)

- AS broadcasts route entries to neighbor ASs
  - by “border gateway” routers (on behalf of the routers in the AS)
  - Similar to RIP
  - BGP route entry
    - AS path + network prefix+ next hop
    - e.g., <AS a, AS b, AS c, ⋯> 128.96.0.0/16 12.5.6.1
  - AS number is used to detect loops

BGP: Path Vector Routing

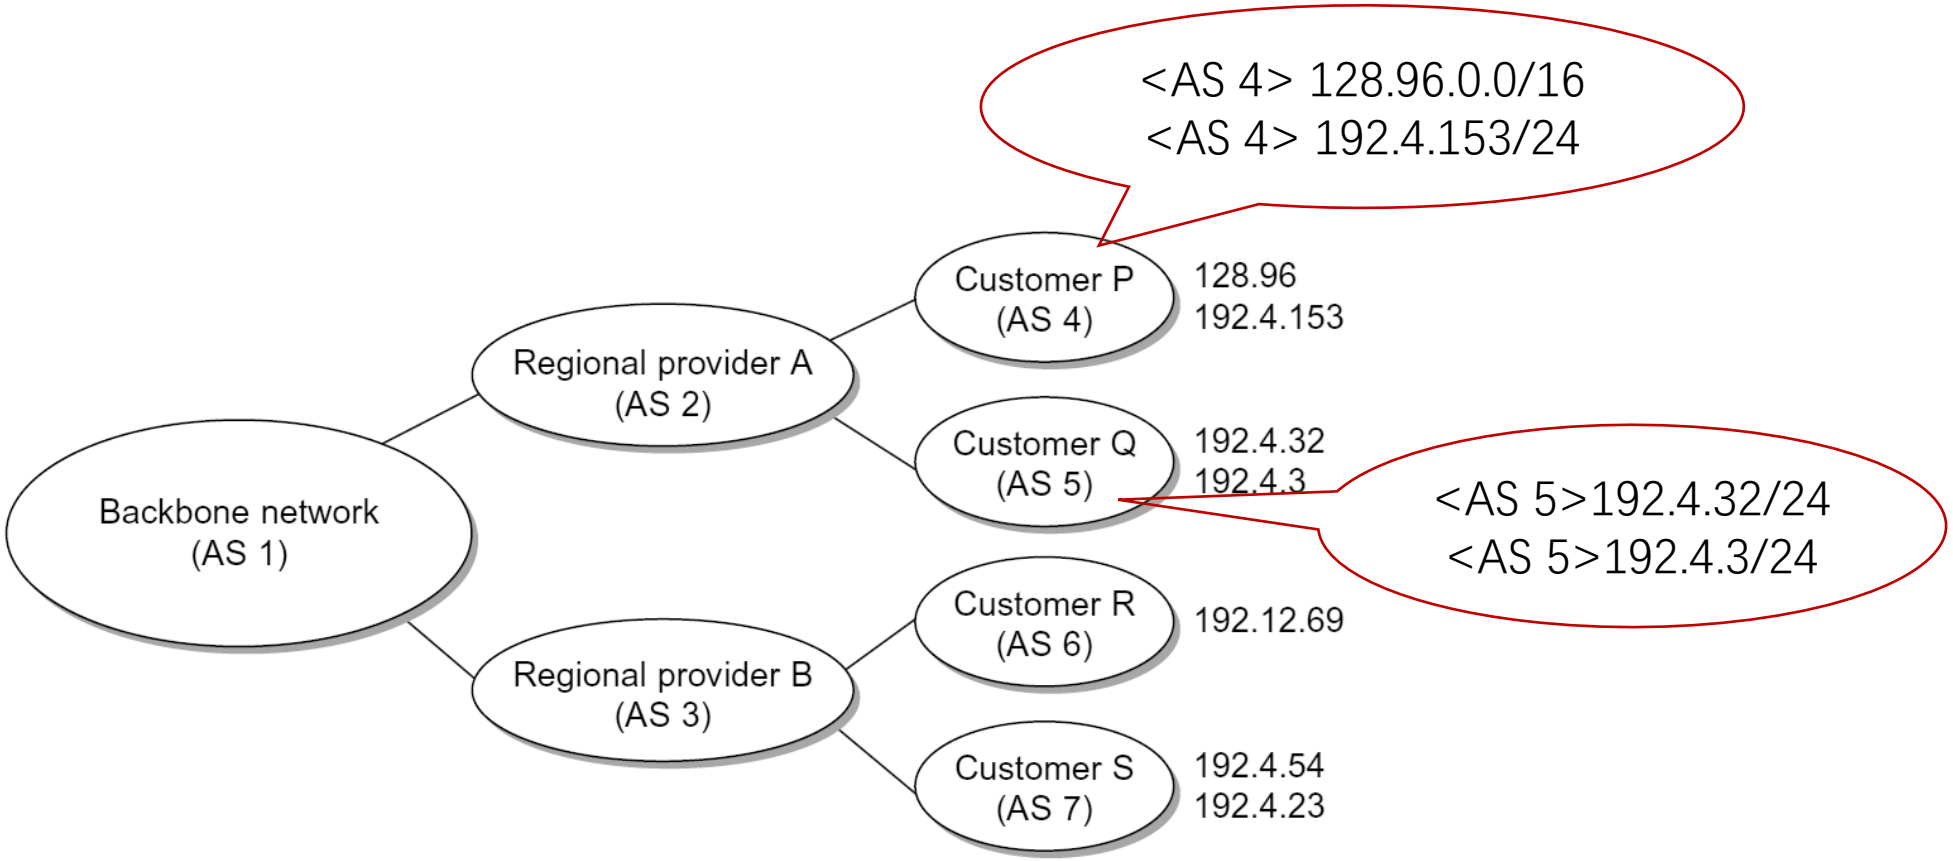

BGP: Path Vector Routing

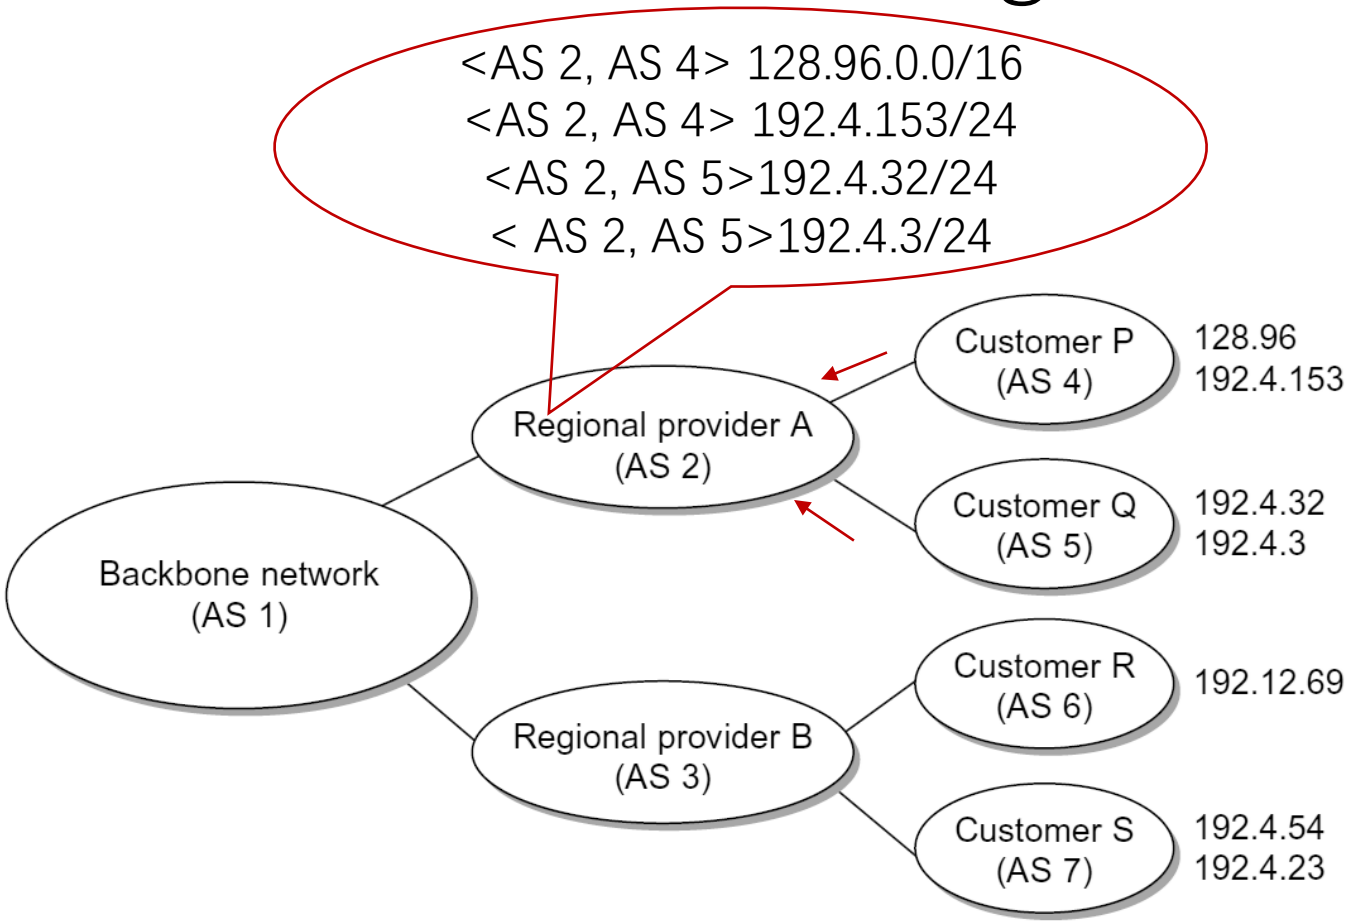

BGP: Path Vector Routing

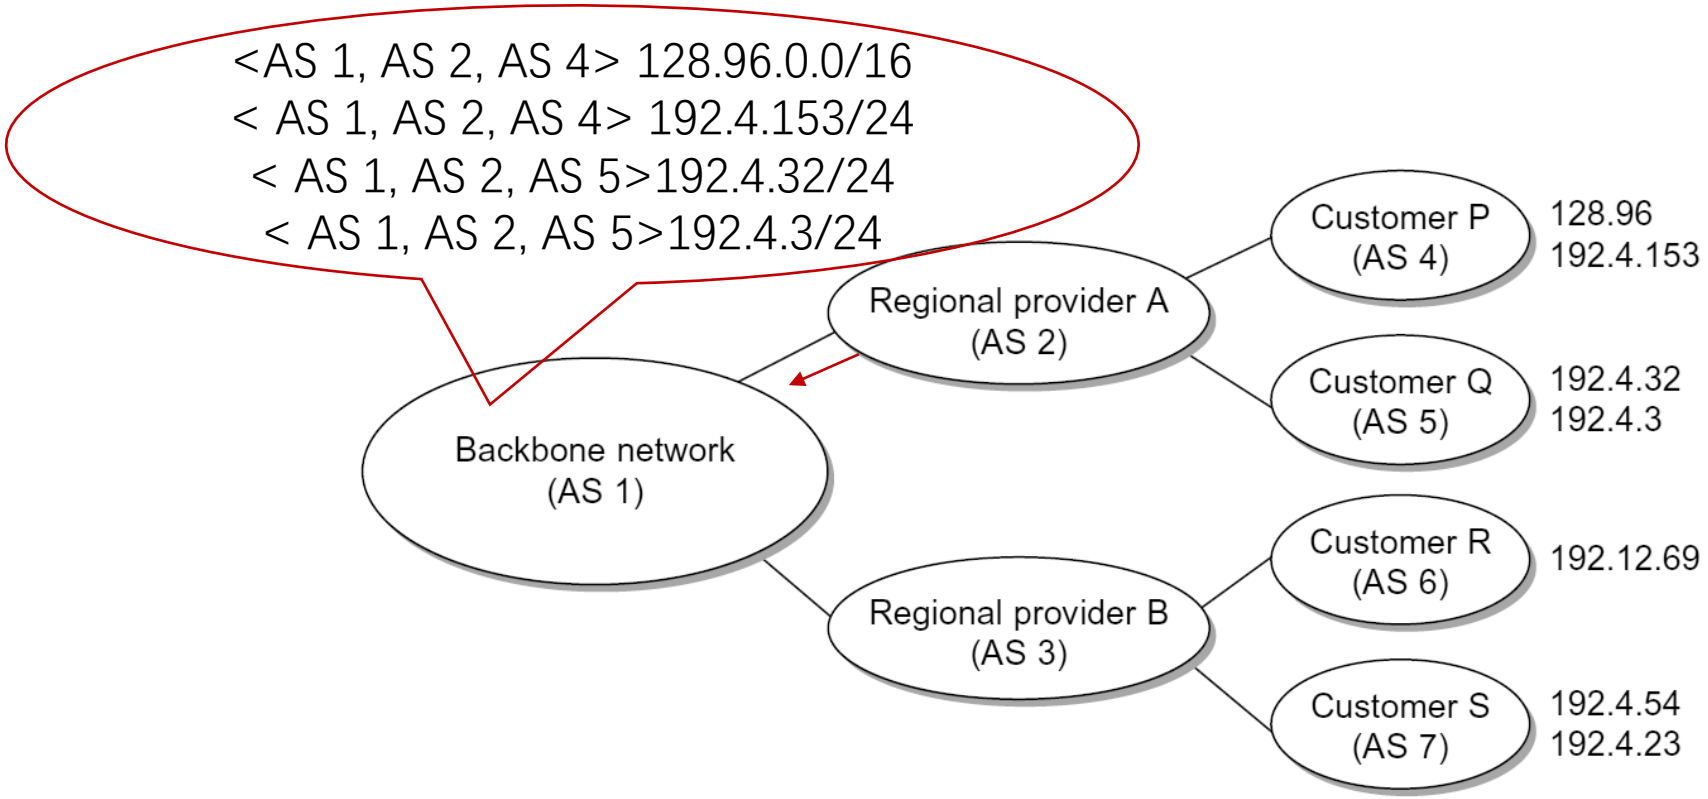

BGP: Path Vector Routing

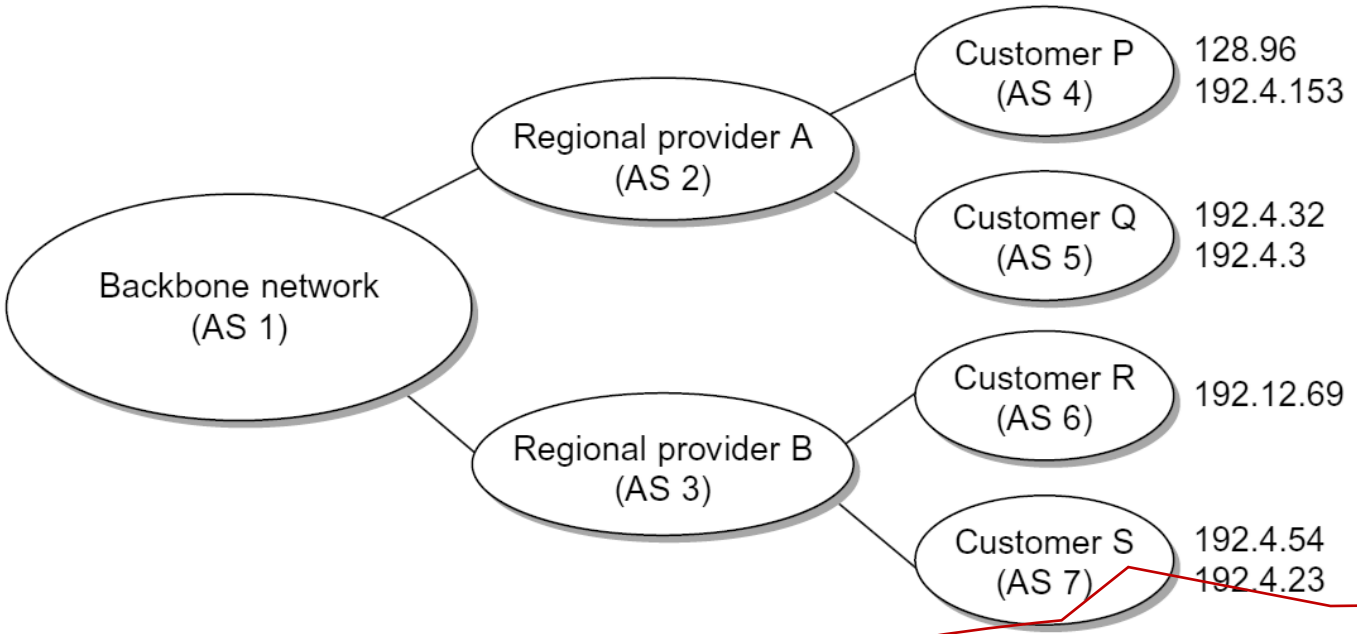

< AS 7, AS 3, AS 1, AS 2, AS 4> 128.96.0.0/16
< AS 7, AS 3, AS 1, AS 2, AS 4> 192.4.153/24
< AS 7, AS 3, AS 1, AS 2, AS 5>192.4.32/24
< AS 7, AS 3, AS 1, AS 2, AS 5>192.4.3/24

BGP: Path Vector Routing

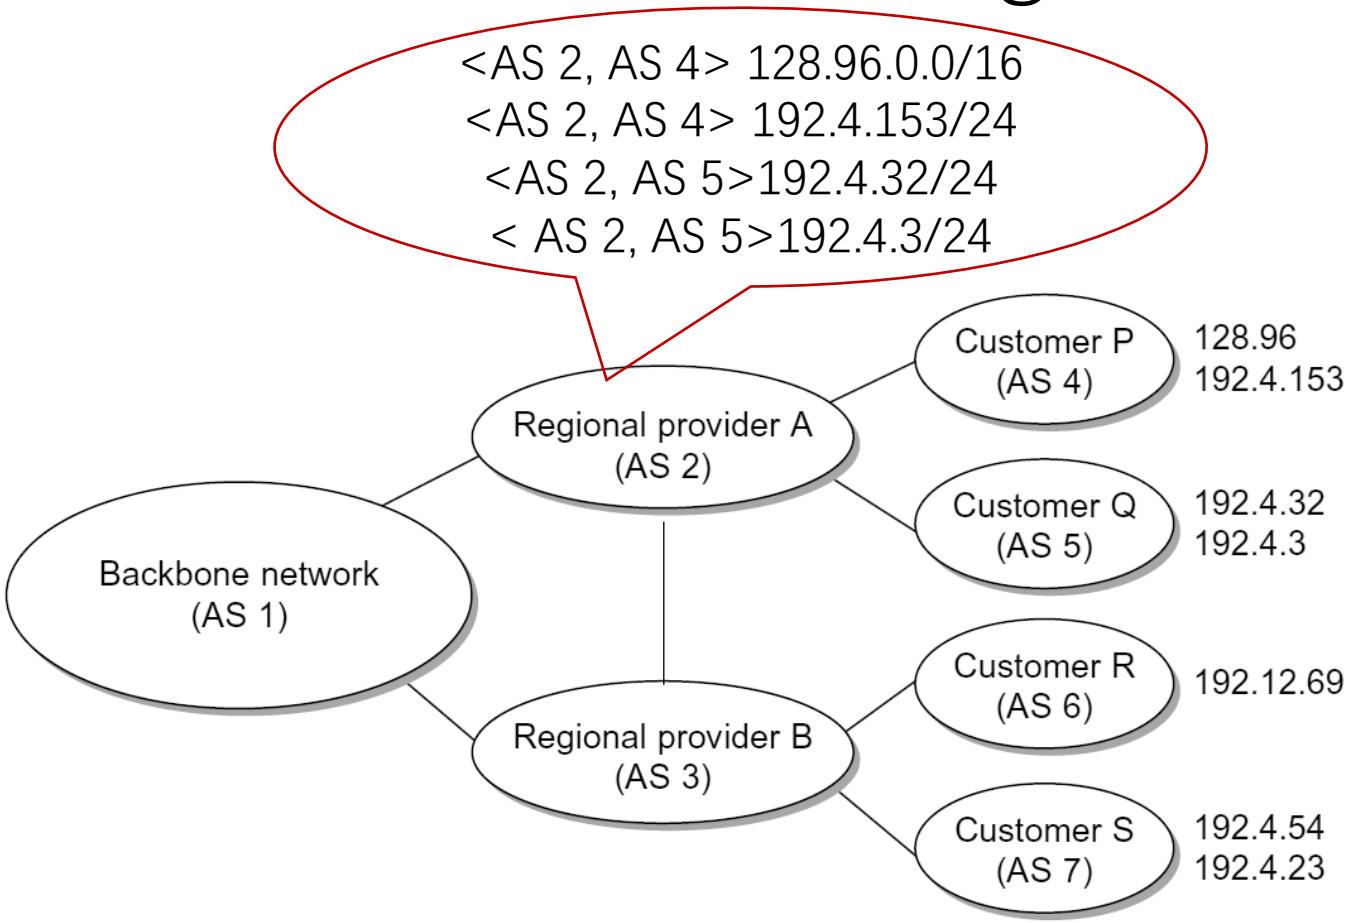

BGP: Path Vector Routing

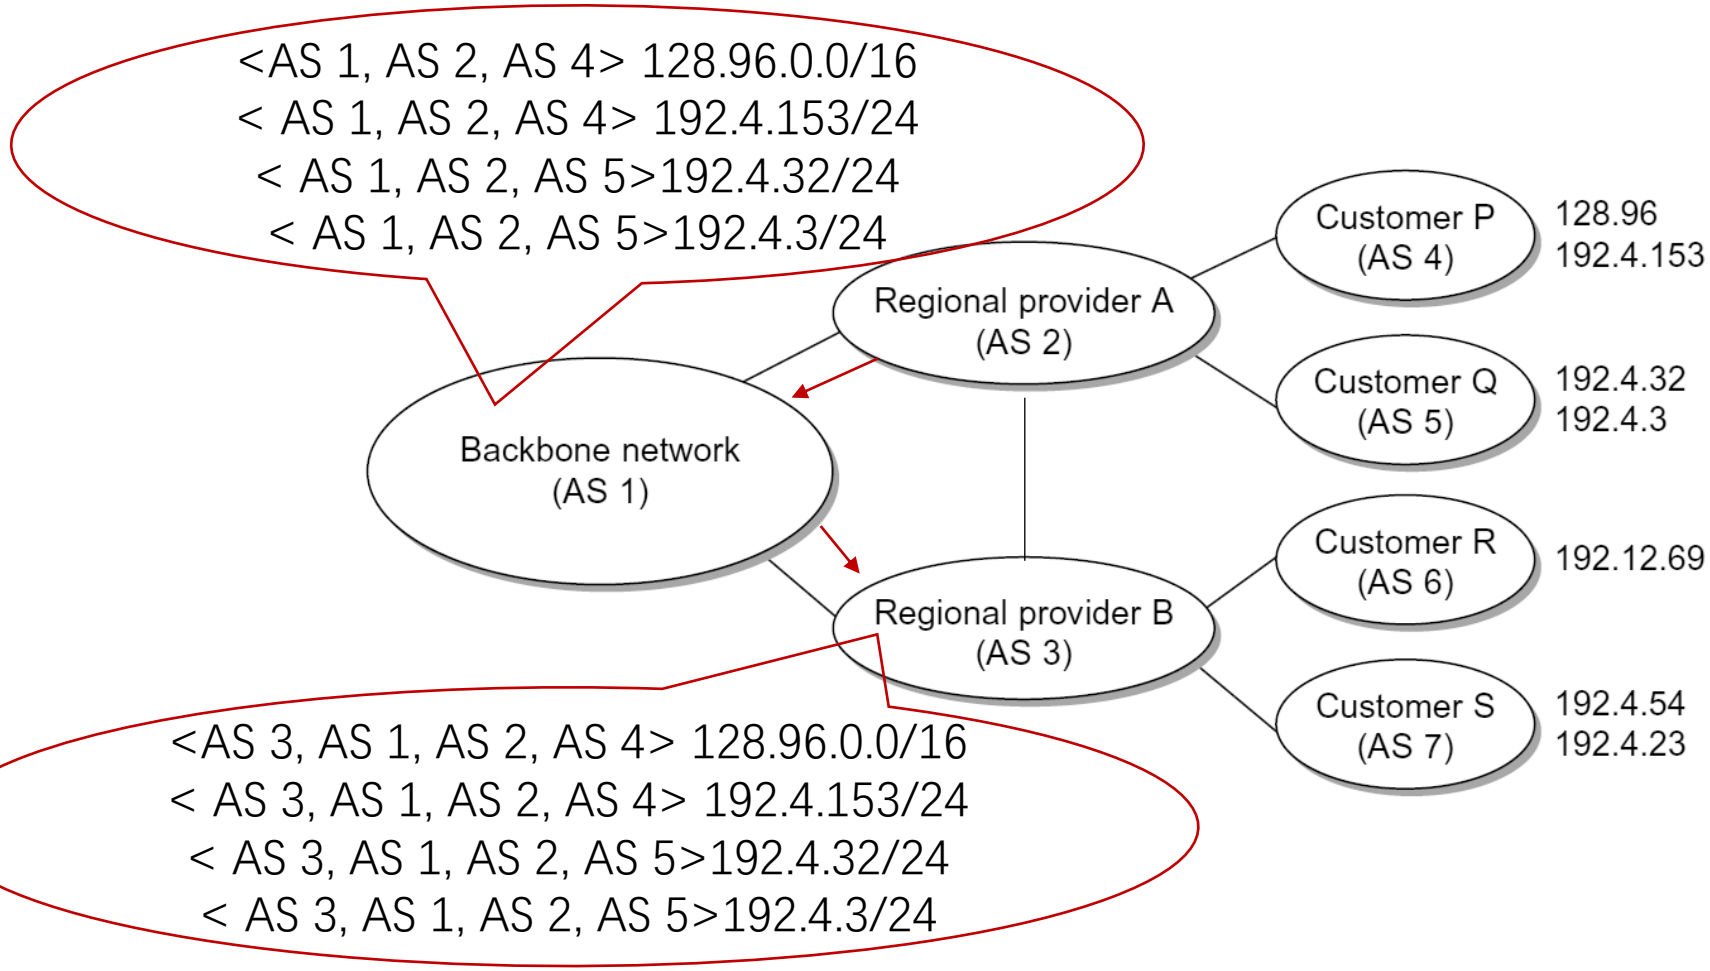

BGP: Path Vector Routing

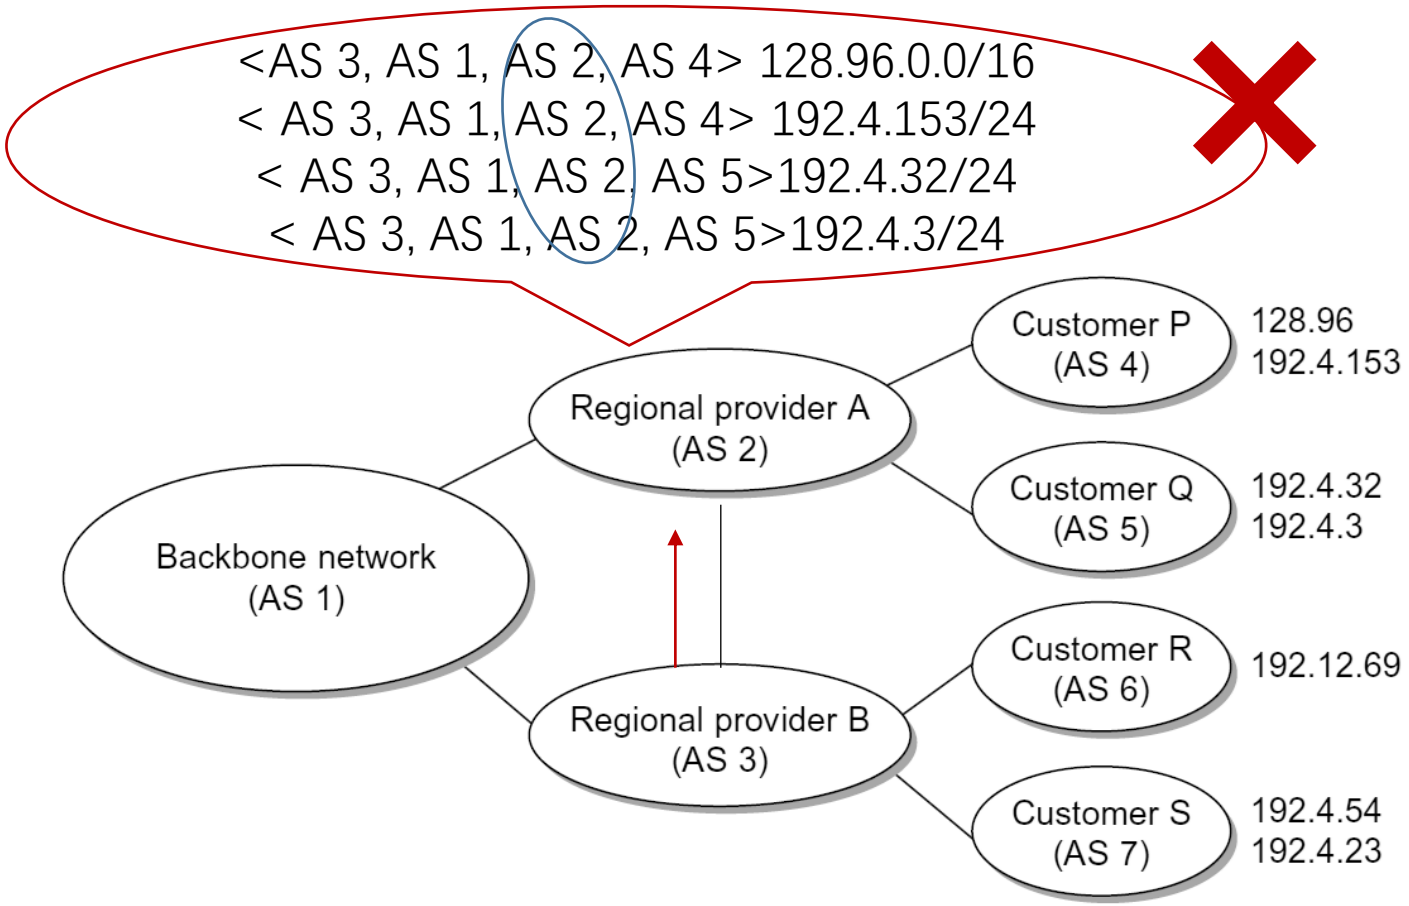

It's not a useful path

BGP: Border Router

- Border Routers: connecting more than one ASs

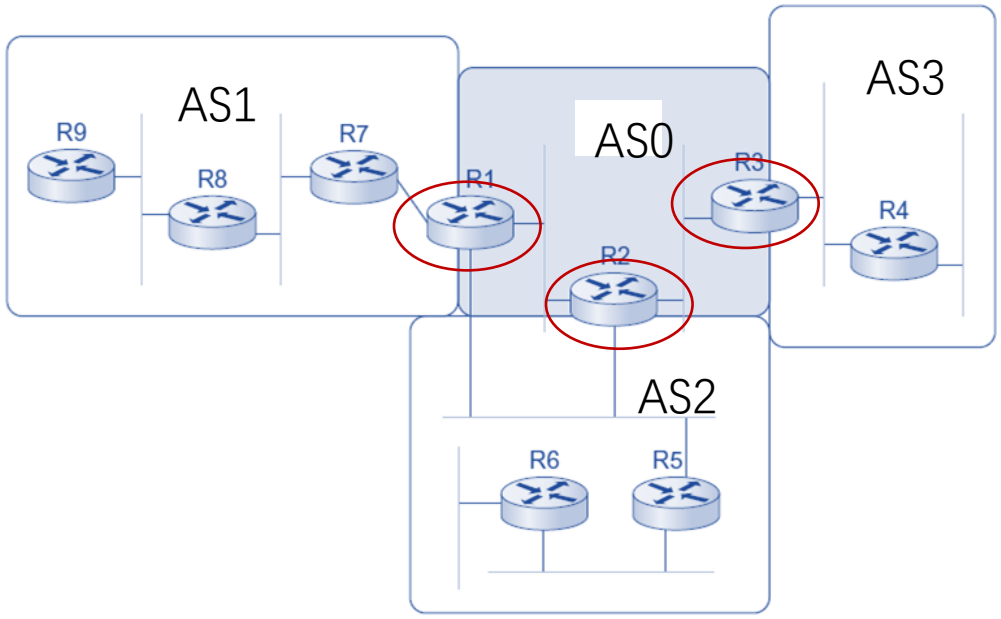

BGP: Border Router

- Border Routers: connecting more than one ASs
  - Selected and configured by AS administrators
  - Routing entries are exchanged with other Boarder Routers through exterior BGP (eBGP)
  - Routing entries are exchanged with routers within the same AS through interior BGP (iBGP)

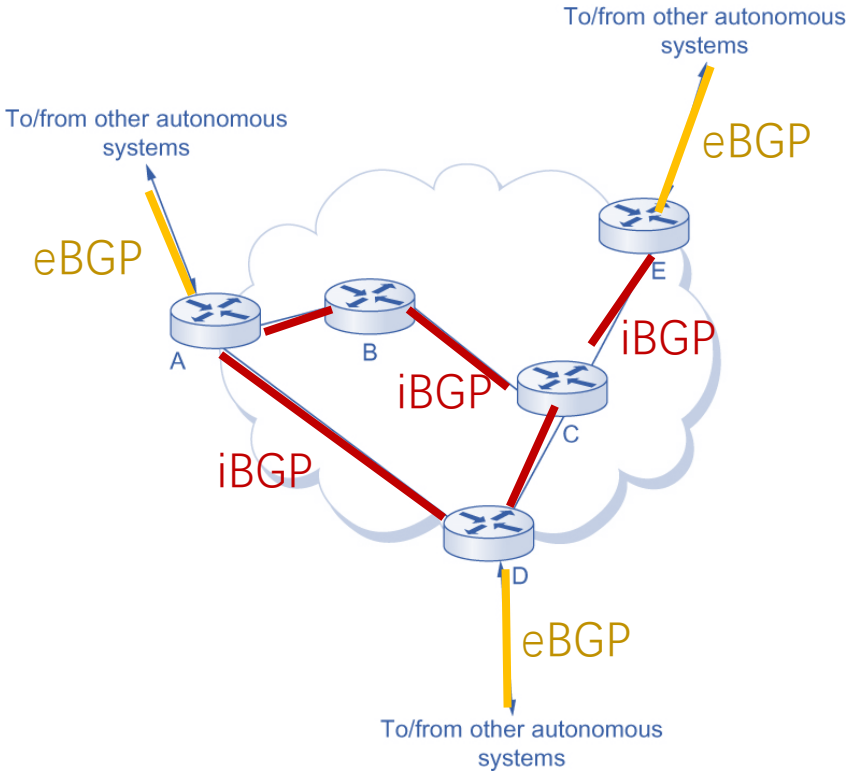

BGP: Exchanging BGP entries

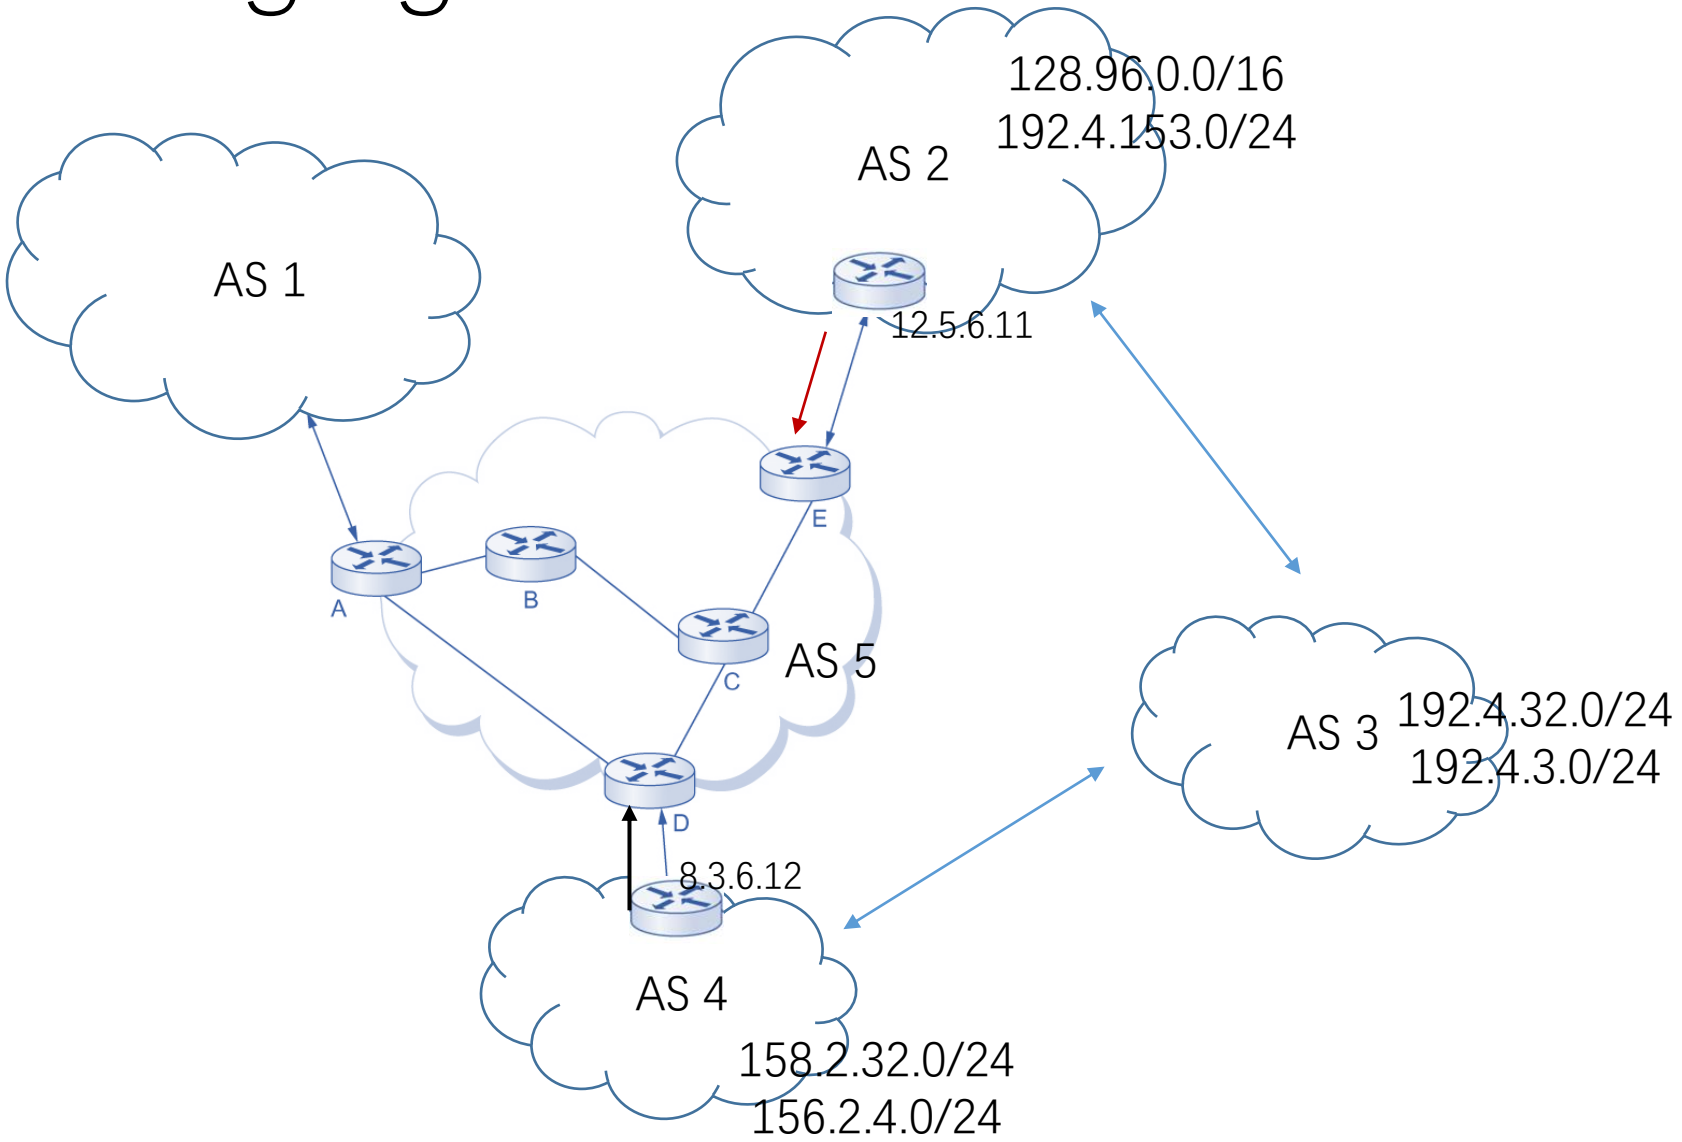

BGP: Exchanging BGP entries

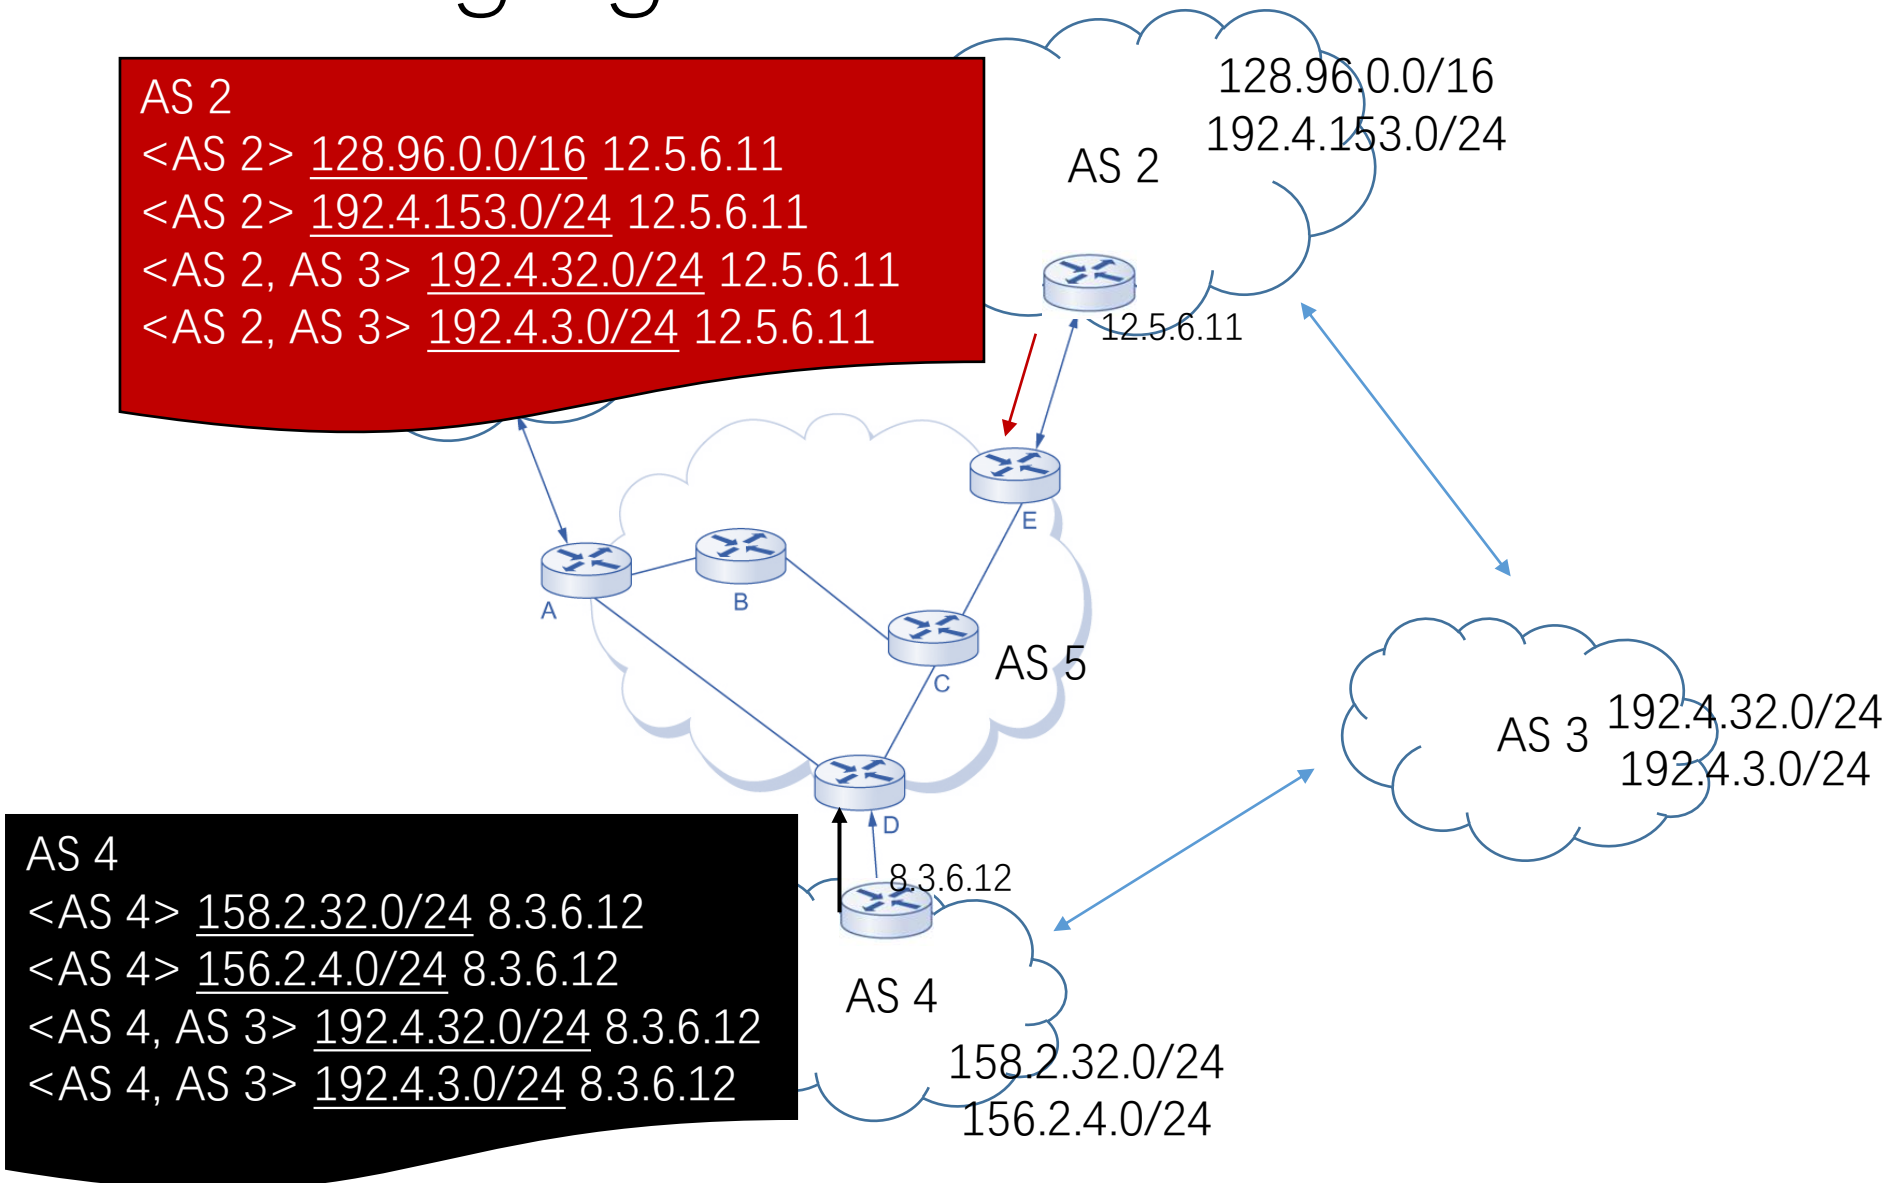

BGP: Exchanging BGP entries

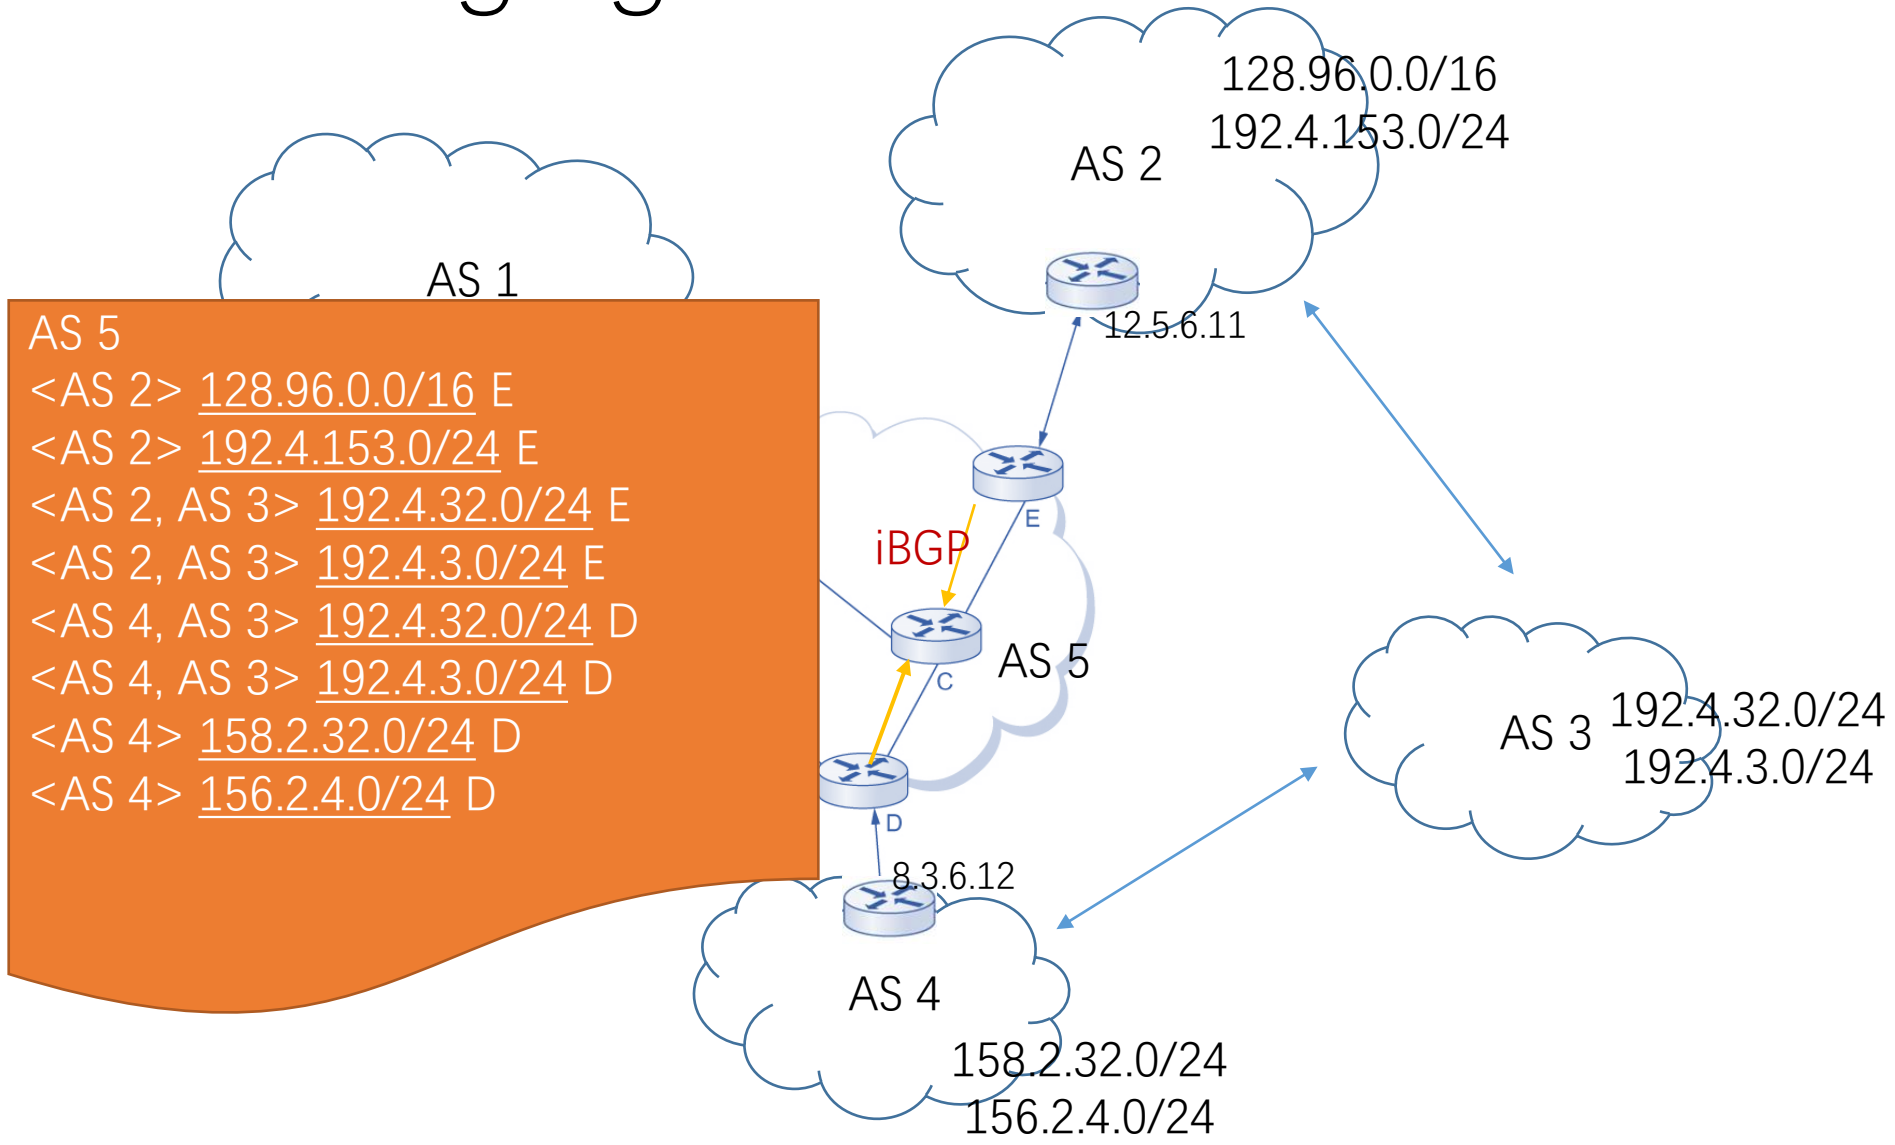

BGP: Exchanging BGP entries

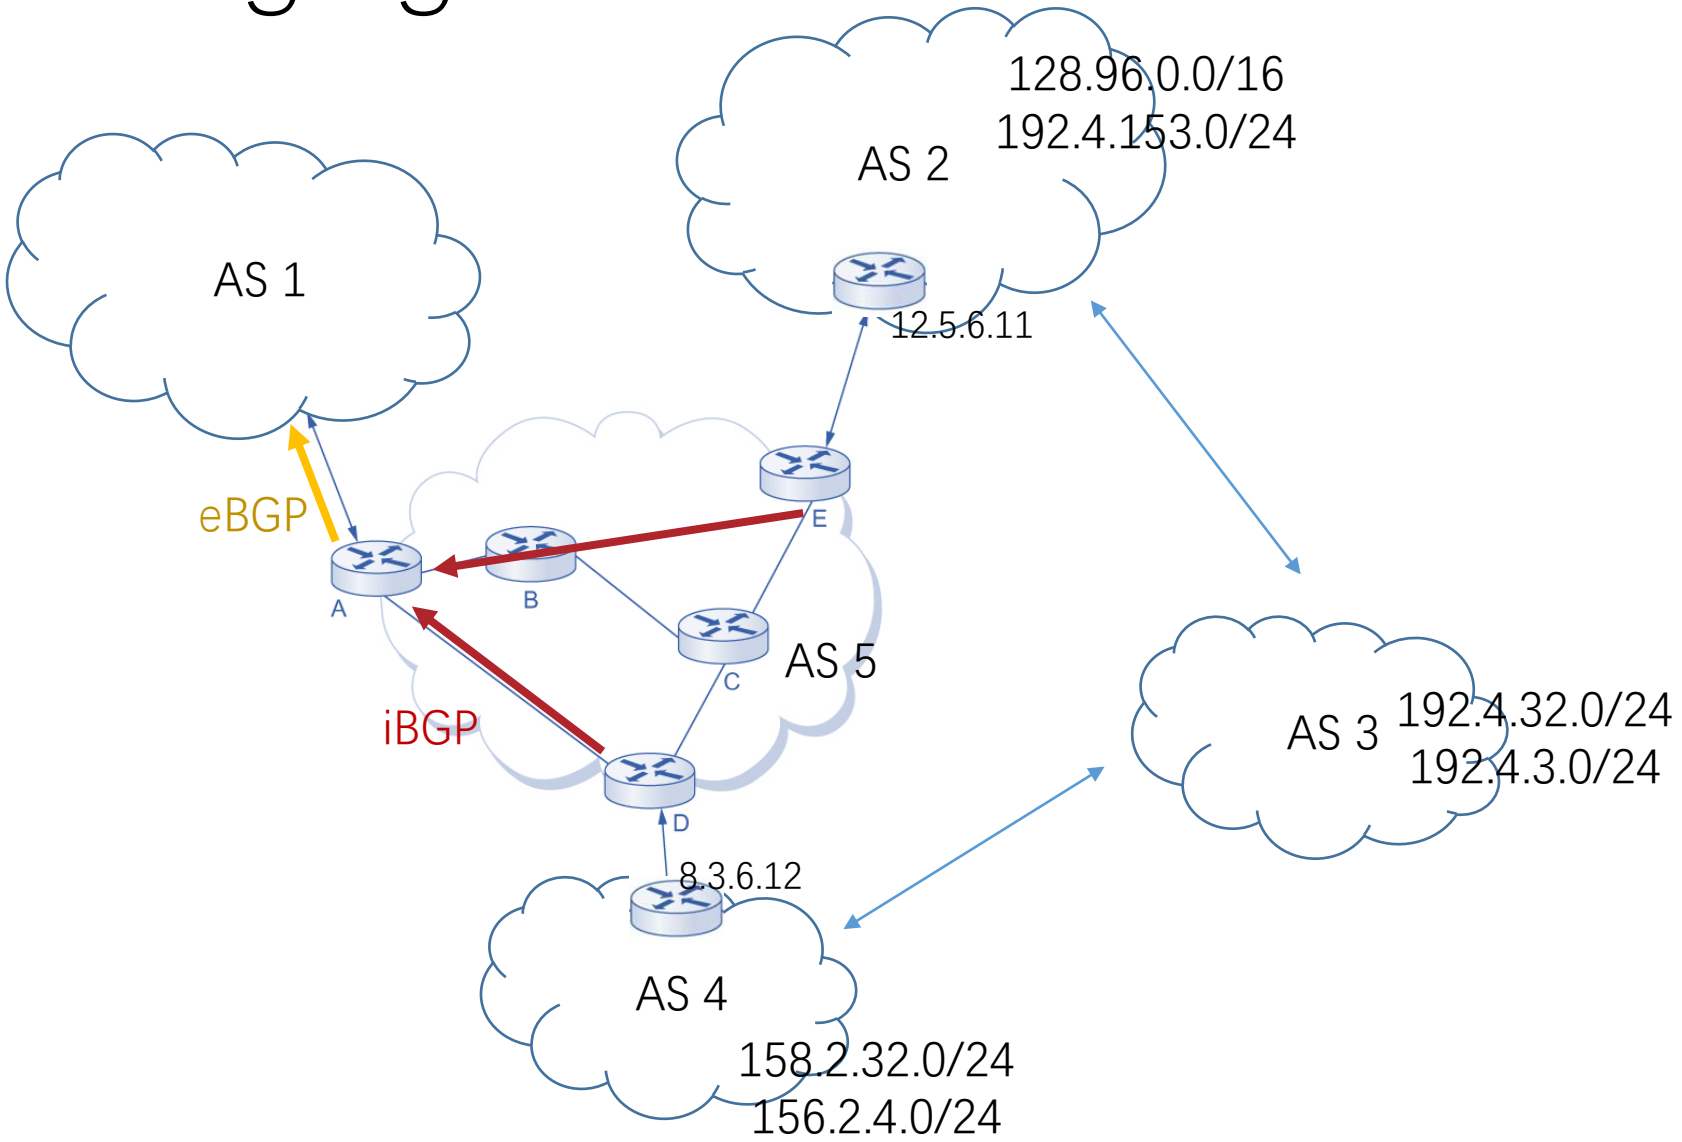

BGP: Exchanging BGP entries

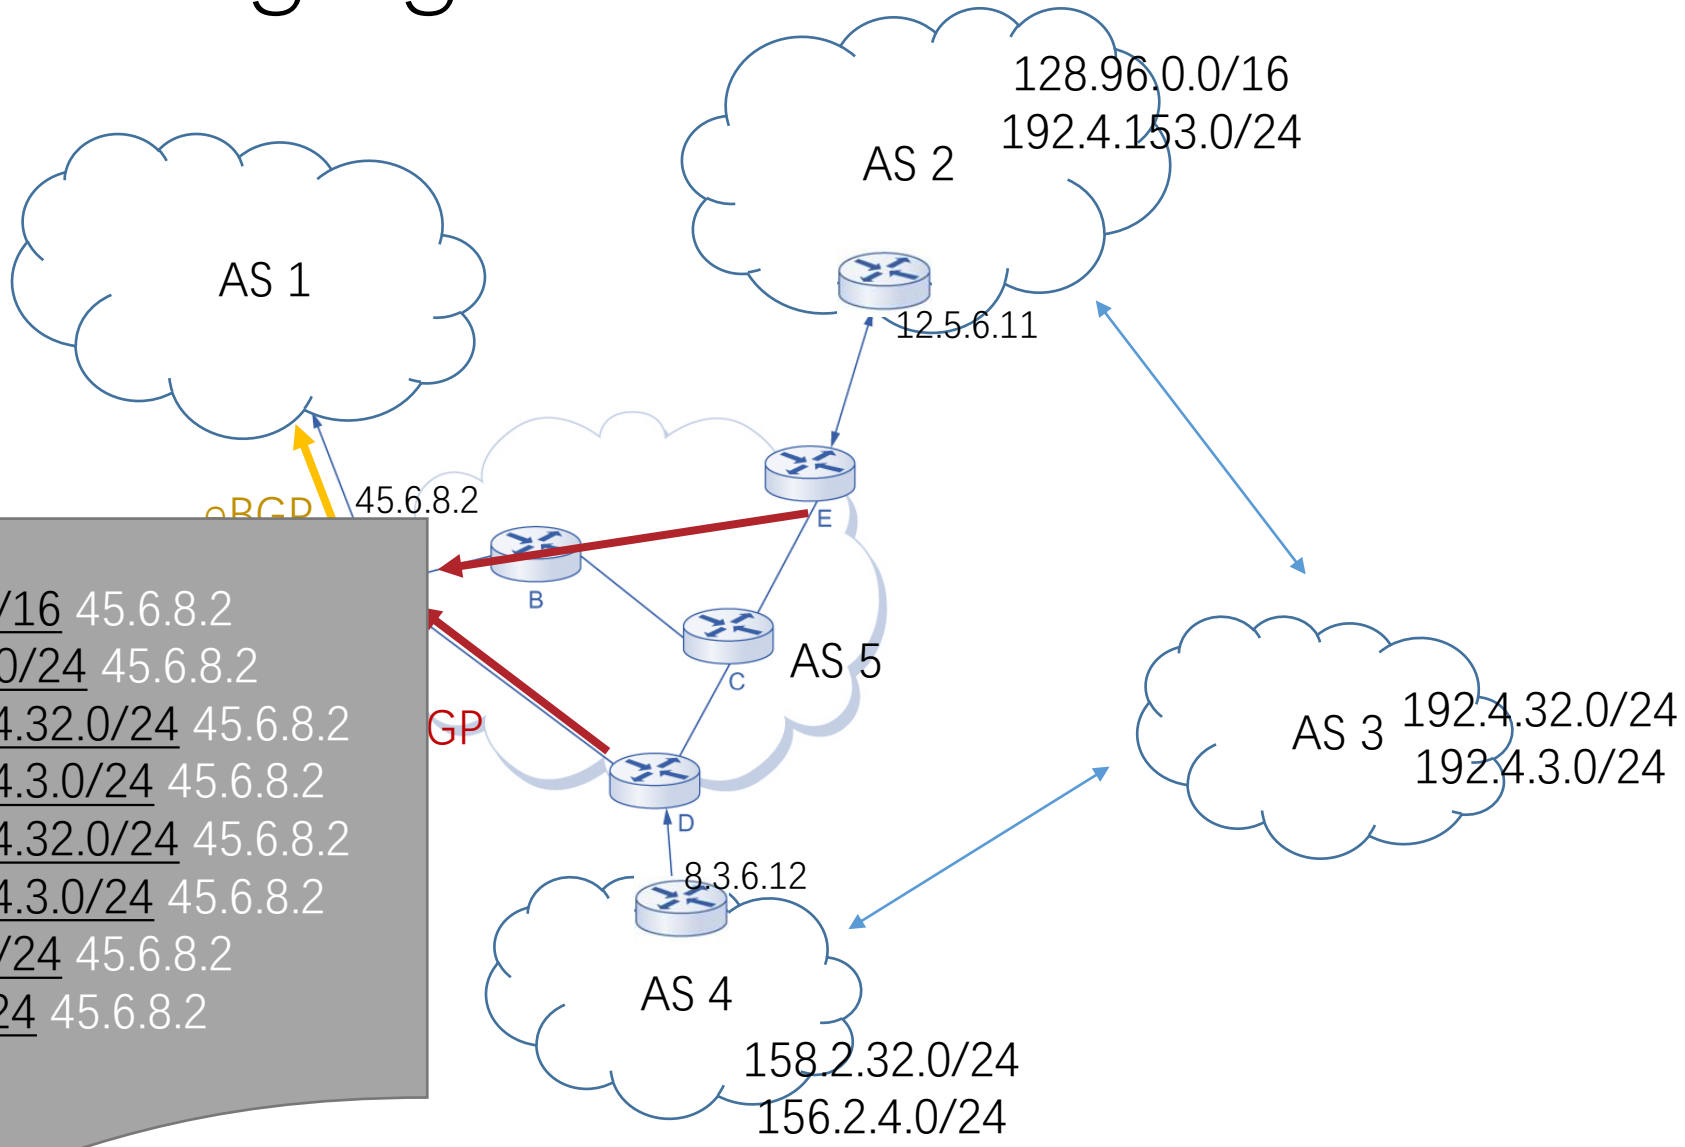

AS1
<AS 5, AS 2> 128.96.0.0/16 45.6.8.2
<AS 5, AS 2> 192.4.153.0/24 45.6.8.2
<AS 5, AS 2, AS 3> 192.4.32.0/24 45.6.8.2
<AS 5, AS 2, AS 3> 192.4.3.0/24 45.6.8.2
<AS 5, AS 4, AS 3> 192.4.32.0/24 45.6.8.2
<AS 5, AS 4, AS 3> 192.4.3.0/24 45.6.8.2
<AS 5, AS 4> 158.2.32.0/24 45.6.8.2
<AS 5, AS 4> 156.2.4.0/24 45.6.8.2

BGP: Integrate with Intradomain Routing

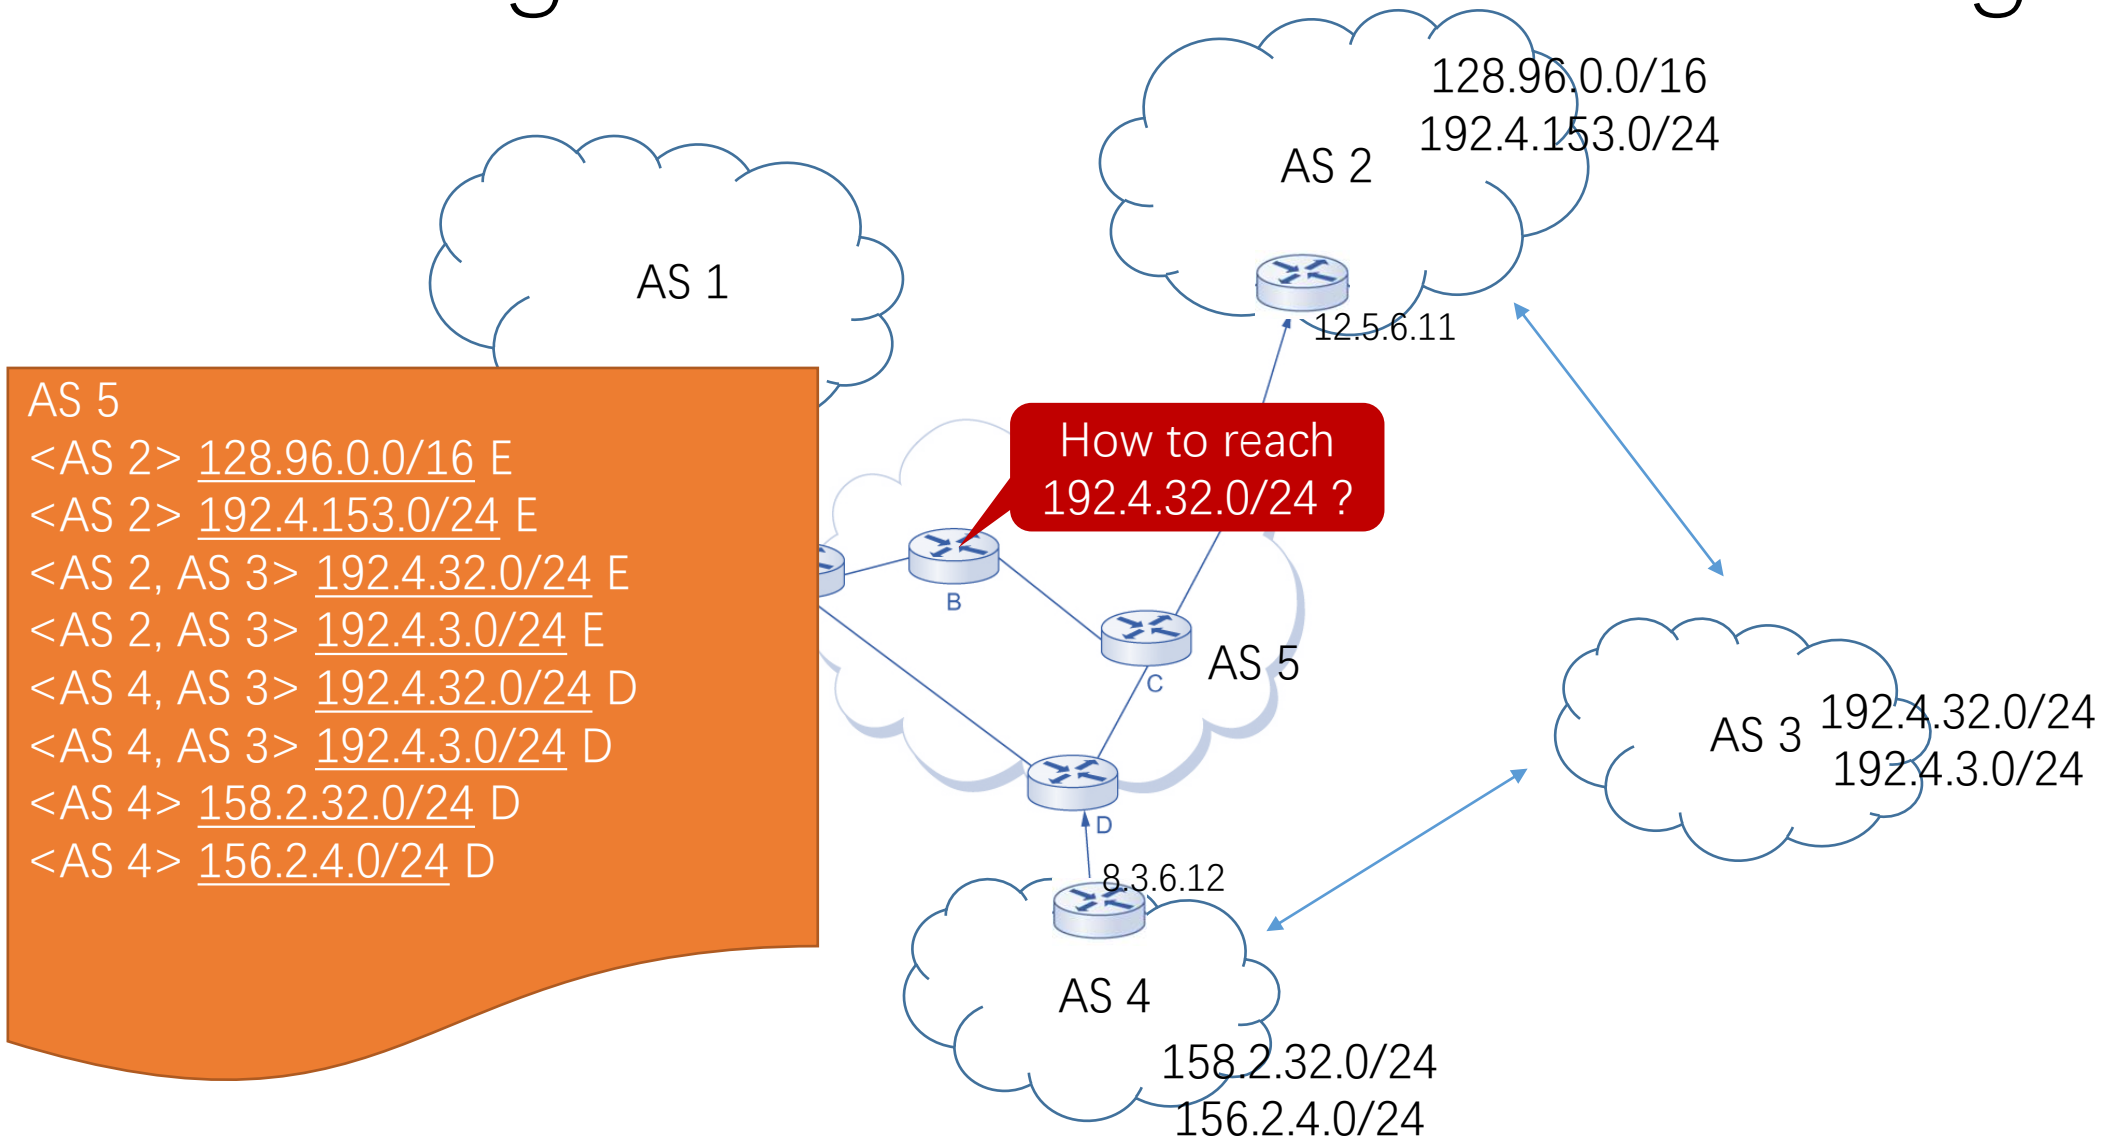

BGP: Integrate with Intradomain Routing

- Combine BGP Table and Inteadomain Routing Table

AS 5
<AS 2> 128.96.0.0/16 E
<AS 2> 192.4.153.0/24 E
<AS 2, AS 3> 192.4.32.0/24 E
<AS 2, AS 3> 192.4.3.0/24 E
<AS 4, AS 3> 192.4.32.0/24 D
<AS 4, AS 3> 192.4.3.0/24 D
<AS 4> 158.2.32.0/24 D
<AS 4> 156.2.4.0/24 D

Router B's BGP Table (obtained via iBGP)

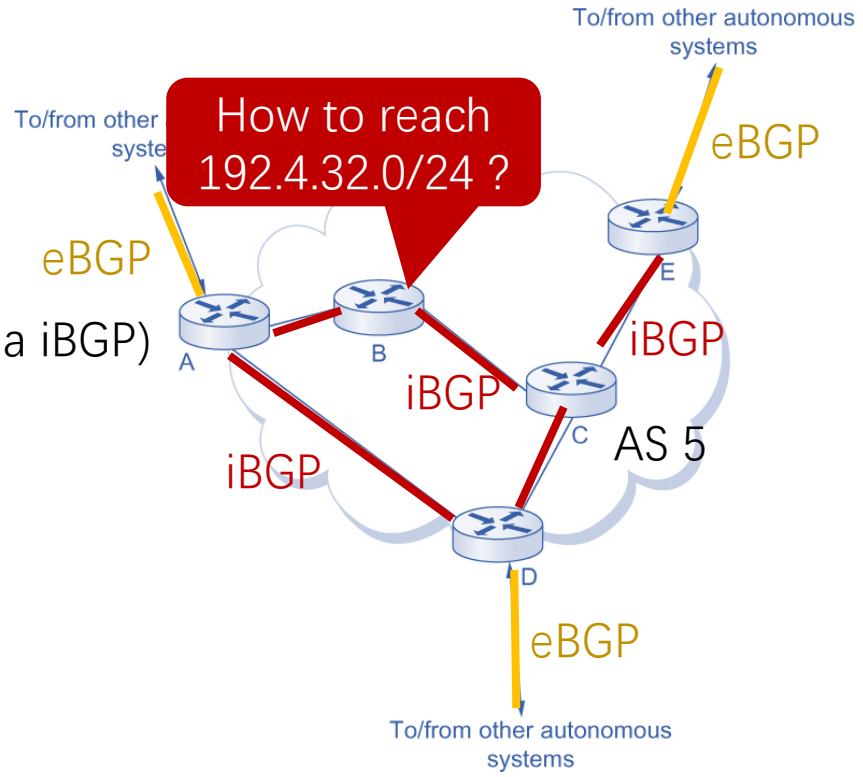

Dest	Next
A	A
C	C
D	C
E	C

Router B's Routing Table (obtained via OSPF)

BGP: Integrate with Intradomain Routing

- Combine BGP Table and Inteadomain Routing Table

AS 5
<AS 2> 128.96.0.0/16 E
<AS 2> 192.4.153.0/24 E
<AS 2, AS 3> 192.4.32.0/24 E
<AS 2, AS 3> 192.4.3.0/24 E
<AS 4, AS 3> 192.4.32.0/24 D
<AS 4, AS 3> 192.4.3.0/24 D
<AS 4> 158.2.32.0/24 D
<AS 4> 156.2.4.0/24 D

which one ?

Router B's BGP Table (obtained via iBGP)

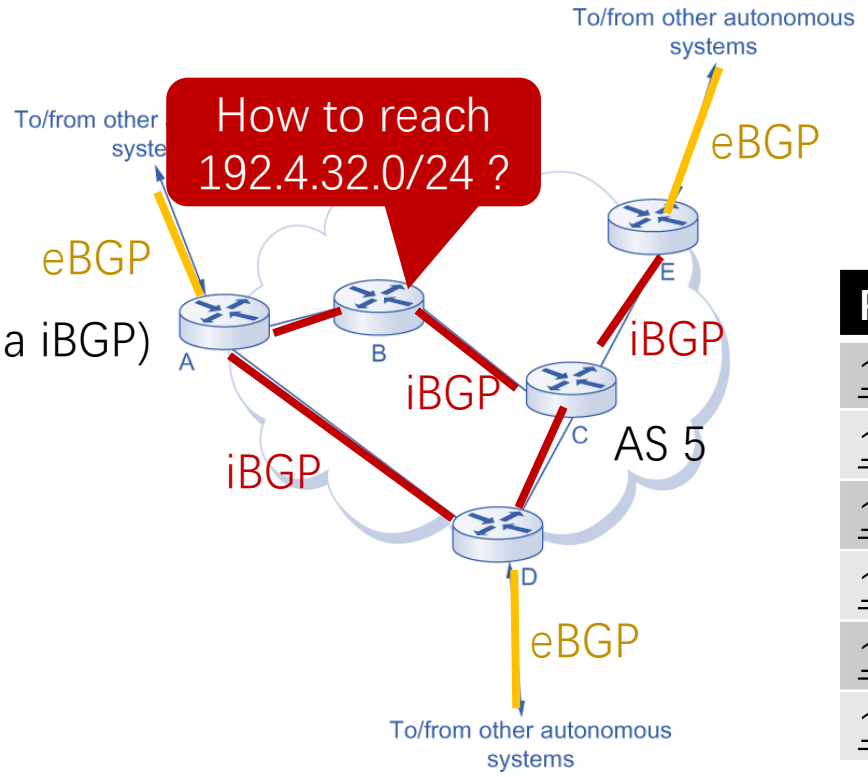

Dest	Next
A	A
C	C
D	C
E	C

Router B's Routing Table (obtained via OSPF)

Prefix	Next
128.96.0.0/16	C
192.4.153.0/24	C
192.4.32.0/24	C
192.4.3.0/24	C
158.2.32.0/24	C
156.2.4.0/24	C

Combined Routing Table for B

BGP: Integrate with Intradomain Routing

- Roughly determine the best border router for certain prefix
  - Selection Priority
    - Local Preference
    - AS hops
    - Distance to the border router
    - …

BGP: Integrate with Intradomain Routing

- Other Methods: Default Gateway Router (Static Method)
  - Inject a default router entry into all routers in the AS through intradomain routing protocol

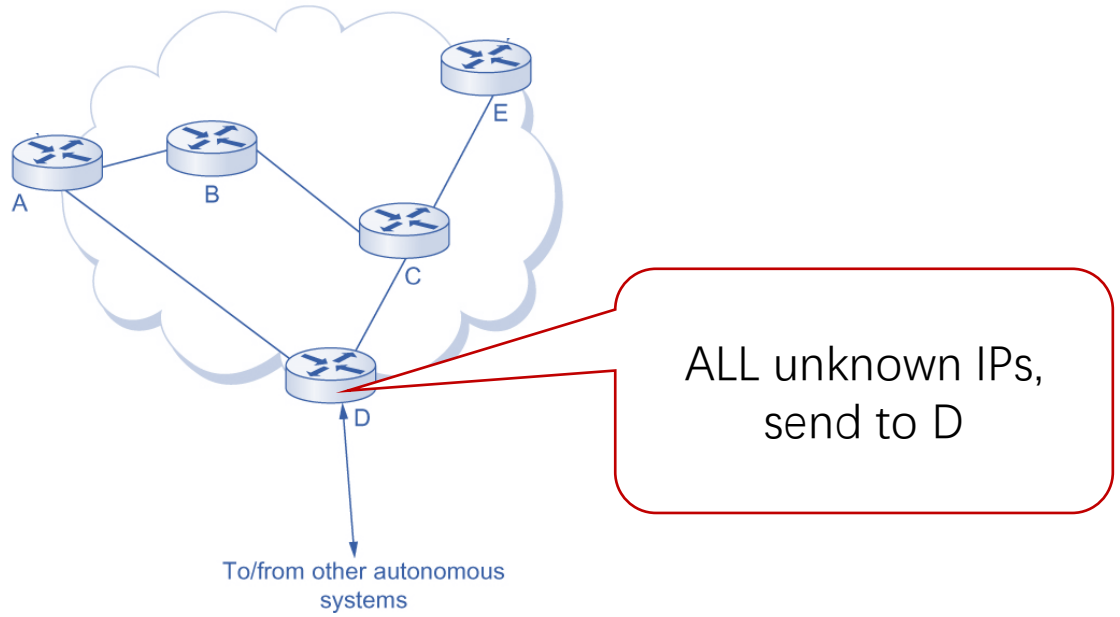

Demo

- BGP Entries
  - http://bgp.potaroo.net/as2.0/bgp-active.html
- BGP Looking Glass: check BGP entry from a certain AS to an IP address
  - e.g.,: https://www.bgp4.as/looking-glasses
- AS Number Look Up: more information about AS number
  - e.g.,: https://stat.ripe.net/app/launchpad

Example

- AS route from telia Sofia to my computer
  - https://lg.telia.net/?type=bgp&router=sfia-b2&address=59.78.171.135

Dest Network

```
59.78.171.0/24 (2 entries, 1 announced)

*BGP    Preference: 170/-201
        Source: 2.255.253.187
        Protocol next hop: 2.255.254.180
        State: <Active Int Ext>
        Local AS:  1299 Peer AS:  1299
        Age: 4d 11:16:10        Metric: 100     Metric2: 2682
        AS path: 4637 4637 4637 4637 4538 4538 24364 I  (Originator)
```

Src Router

Border Router

Src AS

AS Route

AS1299

```
        Country: EU
Registration Date: 1993-09-01
      Registrar: ripencc
          Owner: TELIANET Telia Carrier, EU
```

AS4637

```
        Country: HK
Registration Date: 1995-10-30
      Registrar: apnic
          Owner: ASN-TELSTRA-GLOBAL Telstra Global, HK
```

AS4538

```
        Country: CN
Registration Date: 2002-08-01
      Registrar: apnic
          Owner: ERX-CERNET-BKB China Education and Research Network Center, CN
```

AS24364

```
        Country: CN
Registration Date: 2005-03-03
      Registrar: apnic
          Owner: CNGI-SH-IX-AS-AP CERNET2 IX at Shanghai Jiaotong University, CN
```

Example

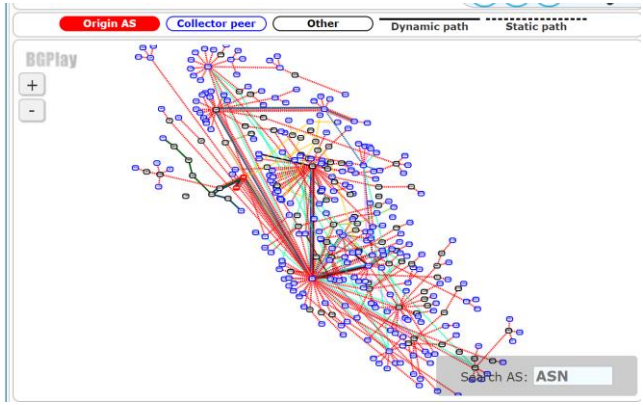

- AS route from multiple ASs
  - https://stat.ripe.net/special/bgplay#bgplay_fetch.resource=59.78.171.135

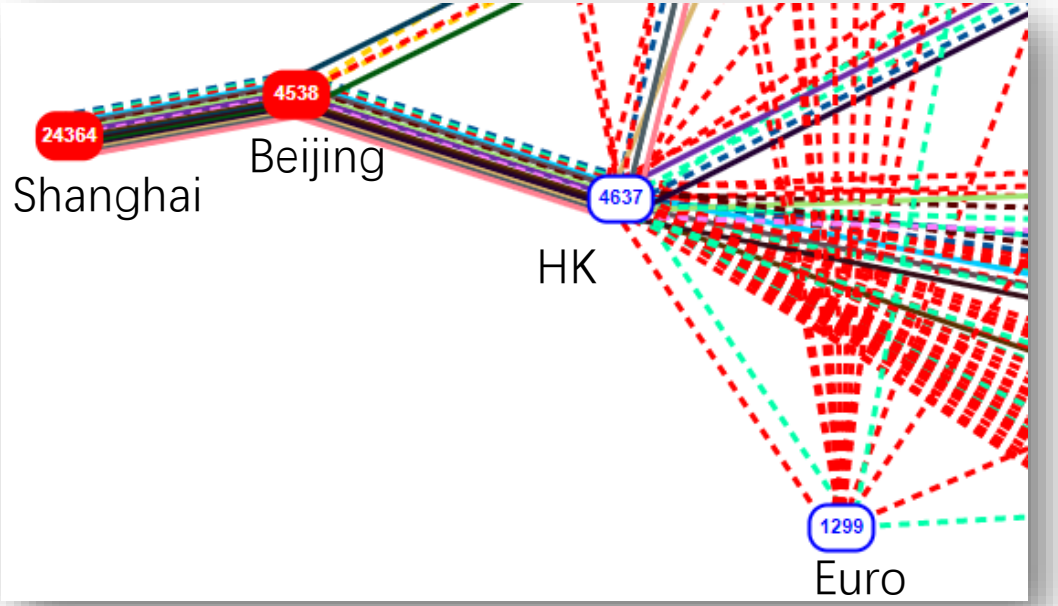


```
Network: AS1299 - Telia Carrier
Router: Sofia (sfia-b2)
Command: traceroute 59.78.171.135 as-number-lookup

traceroute to 59.78.171.135 (59.78.171.135), 30 hops max, 52 byte packets
 1  win-bb3-link.telia.net (62.115.121.36)  49.289 ms  47.015 ms  47.068 ms
 2  ffm-bb1-link.telia.net (62.115.137.202)  46.935 ms  47.882 ms  48.332 ms
 3  prs-bb3-link.telia.net (62.115.123.13)  47.737 ms  46.954 ms  47.681 ms
 4  ldn-bb3-link.telia.net (62.115.123.68)  47.194 ms ldn-bb3-link.telia.net (62.115.13
 5  ldn-b7-link.telia.net (62.115.138.151)  46.976 ms  48.680 ms  46.756 ms
 6  telstra-ic-324829-ldn-b7.c.telia.net (62.115.154.237)  47.813 ms  48.194 ms  47.57
 7  i-91.ulco-core02.telstraglobal.net (202.40.148.33) [AS  4637]  47.833 ms  51.549 ms
 8  i-0-1-1-1.gfr4-core01.telstraglobal.net (202.84.141.121) [AS  4637]  283.475 ms i-1
298.512 ms
 9  202.84.153.26 (202.84.153.26) [AS  4637]  282.290 ms  281.651 ms 202.84.157.37 (202
10  CER-0003.10026.telstraglobal.net (61.8.59.38) [AS  4637]  274.015 ms  294.241 ms  2
11  * * 101.4.114.181 (101.4.114.181) [AS  4538]  317.416 ms
12  101.4.118.121 (101.4.118.121) [AS  4538]  294.191 ms 101.4.114.238 (101.4.114.238)
13  101.4.114.58 (101.4.114.58) [AS  4538]  323.583 ms  317.161 ms  323.827 ms
14  101.4.116.85 (101.4.116.85) [AS  4538]  318.175 ms  294.344 ms  318.948 ms
15  101.4.115.106 (101.4.115.106) [AS  4538]  322.209 ms  321.473 ms  343.786 ms
16  202.112.27.2 (202.112.27.2) [AS  4538]  340.422 ms  343.047 ms  340.877 ms
17  * * *
18  * * *
19  * * *
20  * * *

(Timeout)
```

Interdomain Routing Problems

- Scalability: More than 600 million destinations
  - Storage
    - Routing Table
  - Calculation
    - Shortest Path
  - Communication
    - Exchanges Routing Information

➢Routing Management: Complex Routing Polices

Local Preference

- AS 1 can reach AS2 through AS5
  - AS5 helps AS1 to forward traffic
    - ISP negotiation
    - Priority
      - e.g., AS path prepending
  - AS5 blocks AS1 traffic to AS2
    - Does not broadcast AS2 entry to AS1

```
AS path: 4637 4637 4637 4637 4538 4538 24364
```

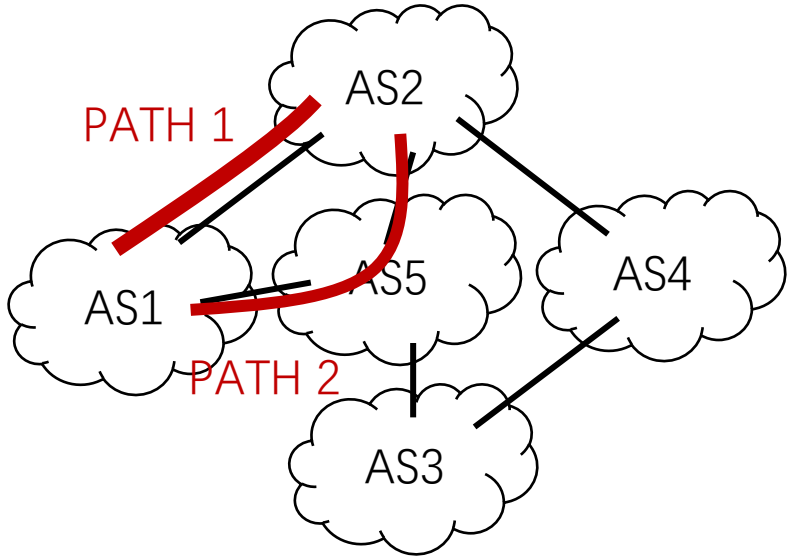

Network Neutrality

- Network Neutrality
  - ISPs supply non-discriminated IP connectivity

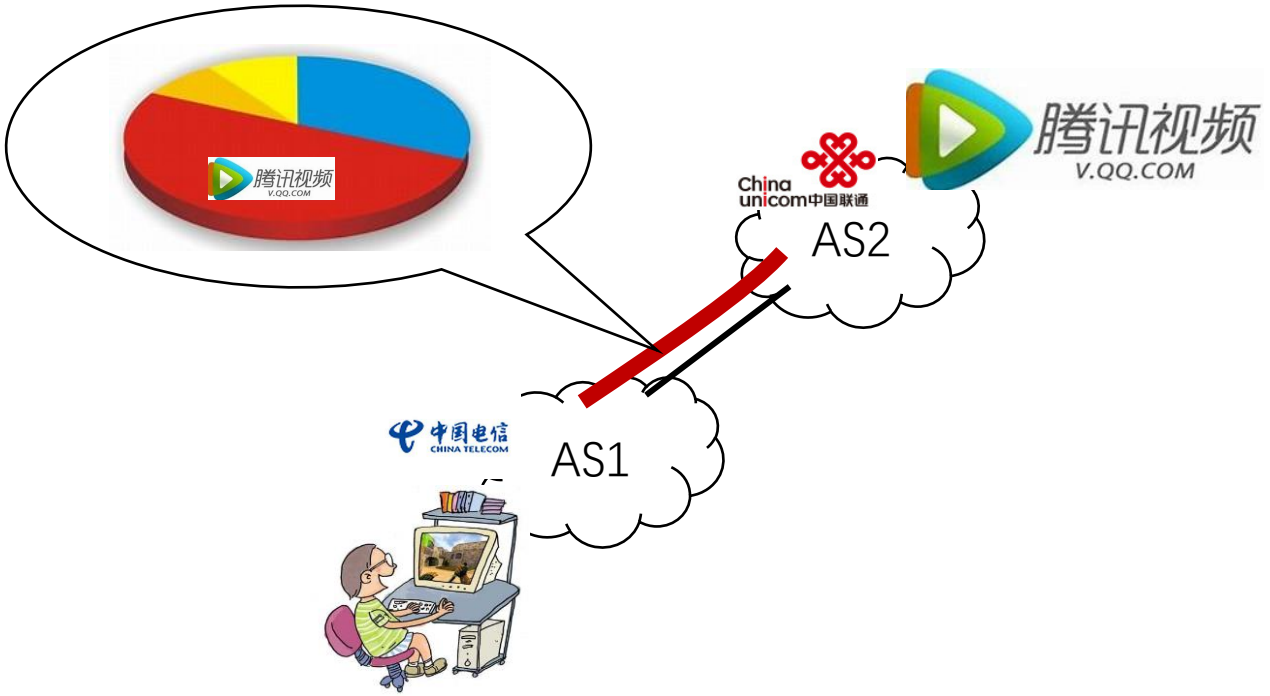

Network Neutrality

- Opposite Counterpoint
  - ISPs prioritize the access to their (often value-added) services

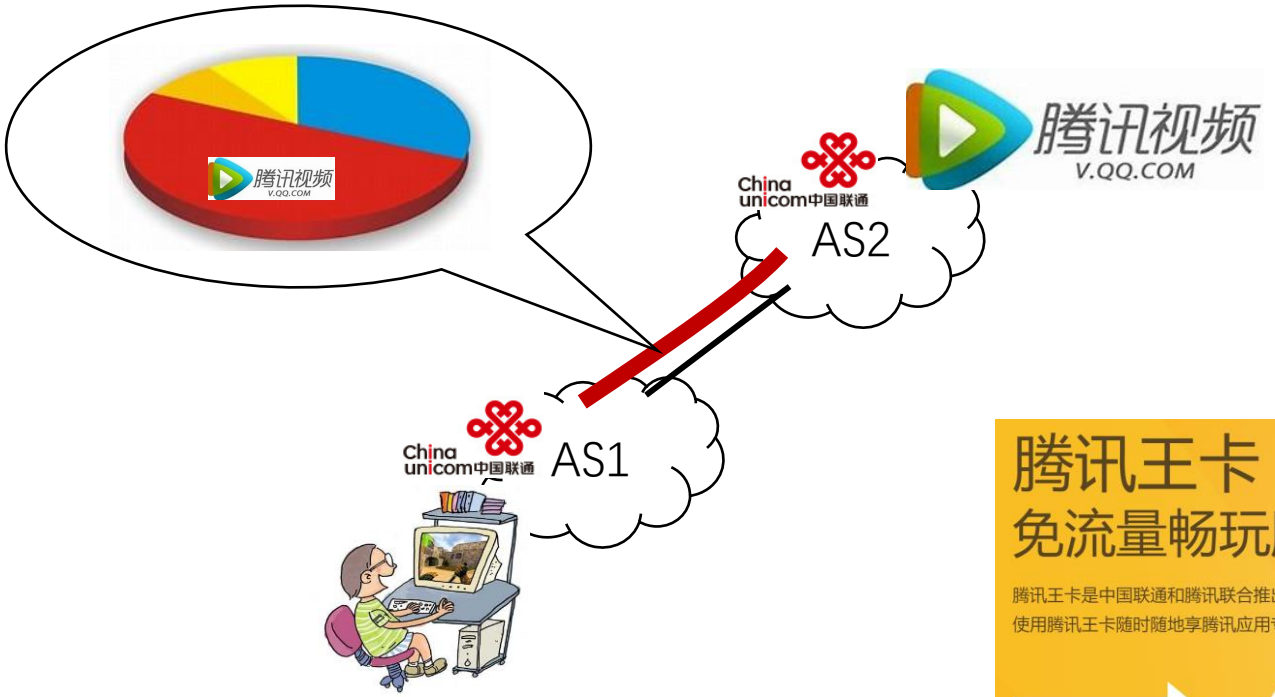

Reference

- Textbook 3.3
- Textbook 4.1
- http://www.ciscopress.com/articles/article.asp?p=24090