



CS120: Computer Networks

Lecture 18. Congestion Control 2


Zhice Yang


Congestion Control

- Host-based Congestion Control
 - Packet Loss
 - Delay
- Router-based Congestion Control
 - Queuing Discipline

Avoid Congestion

- Feedback: Network Delay







cwnd

Source 1

throughput



throughput



Queue Size

Queue
Router

Destination

100-Mbps Ethernet

Source 2

cwnd increases but rate is flat
when the network approaches congestion

TCP Vegas

- Idea: TCP source uses RTT as the sign to **avoid** network congestion
 - Compare RTT with BaseRTT
- Method:
 - Reduce rate when congestion is about to happen
 - If Actual RTT \gg Base RTT
 - Recover rate soon after bandwidth is available
 - if Actual RTT $>$ Base RTT
 - Keep certain pressure on network


TCP Vegas Algorithm

- BaseRTT: the reference for increases in RTTs
 - The minimum RTT
 - If $RTT < BaseRTT$
 - $BaseRTT = RTT$
- ExpectRate:
 - $CongestionWindow / BaseRTT$
- ActualRate:
 - ActualRTT: according to timestamps
 - $ActualRate = CongestionWindow / ActualRTT$

TCP Vegas Algorithm

- Source compares ActualRate with ExpectRate
 - if $\alpha < \text{ExpectRate} - \text{ActualRate} < \beta$
 - $\text{cwnd} = \text{cwnd}$
 - if $\text{ExpectRate} - \text{ActualRate} < \alpha$
 - $\text{cwnd}++$
 - if $\beta < \text{ExpectRate} - \text{ActualRate}$
 - $\text{cwnd}--$


TCP Vegas



Top, congestion window; bottom, expected (colored line) and actual (black line) throughput. The shaded area is the region between the α and β threshold

Two Places to Handle Network Congestion

- End hosts
 - Routers



Congestion Control



- Host-based Congestion Control
 - Packet Loss
 - Delay
- Router-based Congestion Control
 - Queuing Discipline

Queuing Discipline

- Why ?
 - Queuing discipline in routers determines how packets are transmitted (or dropped)
- Router Resource
 - Bandwidth
 - Queue Buffer

Queuing Discipline

- First-In-First-Out (FIFO)



Congestion Control

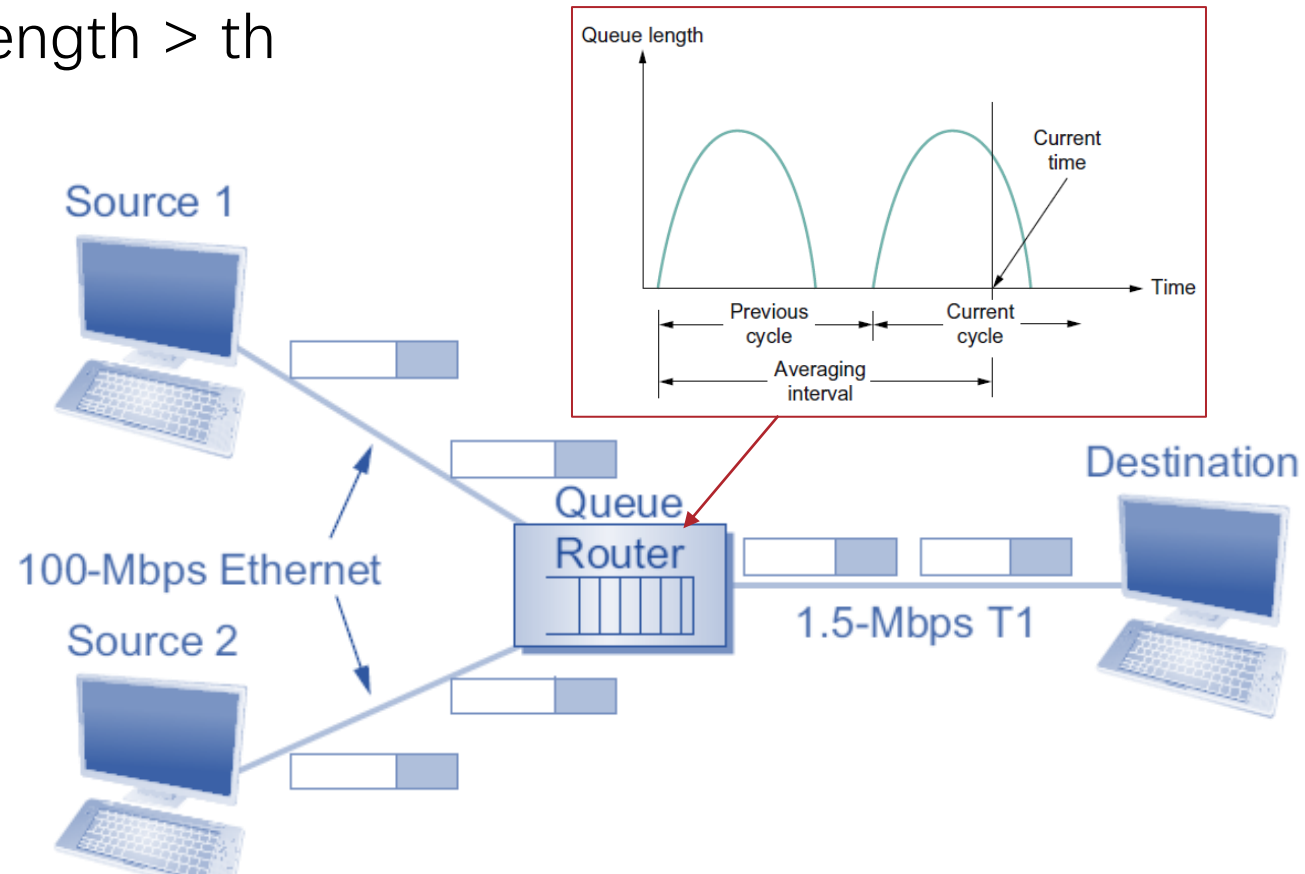
- Host-based Congestion Control
 - Packet Loss
 - Delay
- Router-based Congestion Control
 - Queuing
 - DECbit
 - Explicit Congestion Notification (ECN)
 - Random Early Detection (RED)

DECbit

- Developed for the Digital Network Architecture (DNA)
 - Before TCP/IP was “standardized”
- Idea: let routers explicitly indicate congestion
- Approach:
 - Router set congestion bit in passing packets if there is congestion
 - Destination echoes bit back to source through acks
 - Source adjusts cwnd according to congestion bit

DECbit

- Routers determine congestion
 - Monitor average queue length over last busy+idle cycle
 - If average queue length $> th$
 - set congestion bit

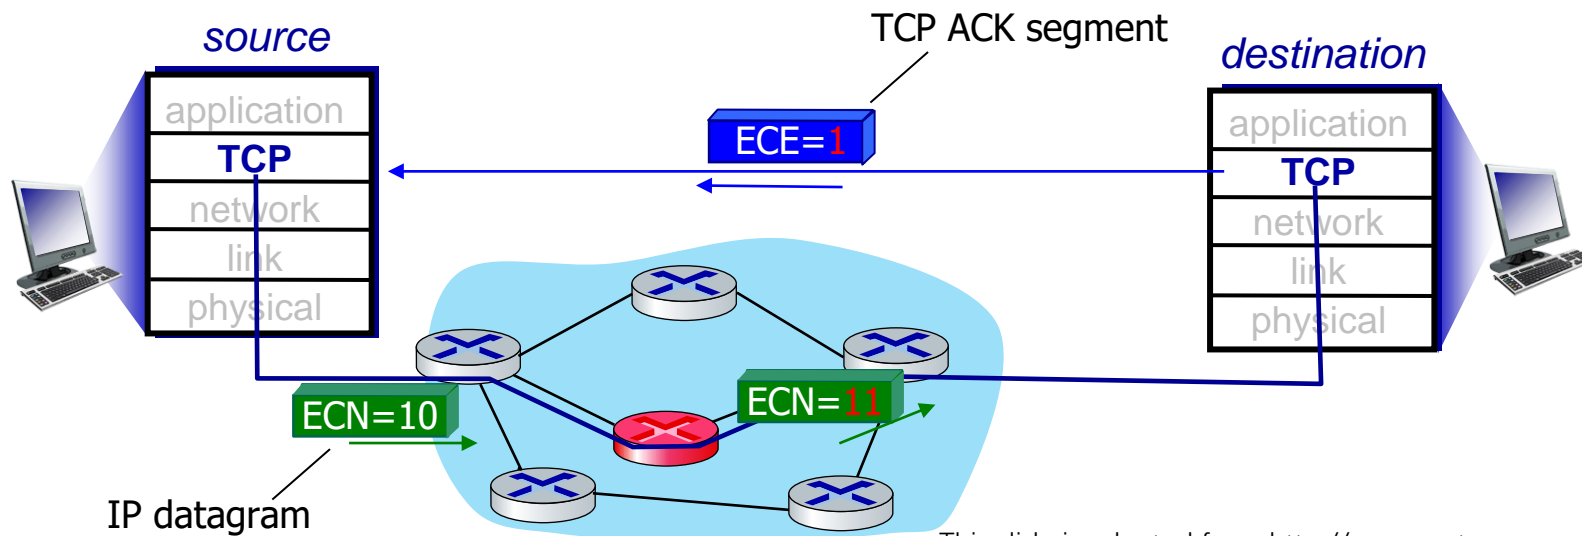


DECBit

- Destination notifies source about congestion bit
- Source reacts to congestion bit
 - If $< 50\%$ of last window's packets had bit set
 - $\text{cwnd}++$
 - If $> 50\%$ of last window's packets had bit set
 - $\text{cwnd} \times 0.875$


Explicit Congestion Notification (ECN)

- RFC 3168
 - Two bits in IP header (ECN in IP's ToS field) marked by network router to indicate congestion
 - Mechanism is similar to DECbit
 - Destination sets ECE bit (in TCP Header) on ACK segment to notify sender of congestion




Random Early Detection (RED)

- Another way is to offload a part of congestion control to routers
- Idea: let routers implicitly indicate congestion
 - Router notices that the queue is getting backlogged
 - Router randomly drops packets to signal congestion
 - Source adjusts cwnd





RED Algorithm

- Compute Average Queue Length
 - Moving average
 - $\text{AvgLen} = (1 - \text{Weight}) * \text{AvgLen} + \text{Weight} * \text{SampleLen}$



RED Algorithm

- Two queue length thresholds
 - if $\text{AvgLen} \leq \text{MinThreshold}$
 - Enqueue the packet
 - if $\text{MinThreshold} < \text{AvgLen} < \text{MaxThreshold}$
 - Drop arriving packet with probability P
 - if $\text{MaxThreshold} \leq \text{AvgLen}$ then drop arriving packet



RED Algorithm

- Computing probability P
 - $\text{TempP} = \text{MaxP} * (\text{AvgLen} - \text{MinThreshold}) / (\text{MaxThreshold} - \text{MinThreshold})$
 - $P = \text{TempP} / (1 - \text{count} * \text{TempP})$
 - Count: number of continuously queued packets without drop
- Why?
 - Drops are spaced out in time
 - When $\text{Count} == 0 \Rightarrow P = \text{TempP}$
 - When $\text{Count} == 1/\text{TempP} - 1 \Rightarrow P = 1$
 - $1/\text{TempP}$ equals to the expectation of number packets spaced between two drops

Reference

- Textbook 6.4