

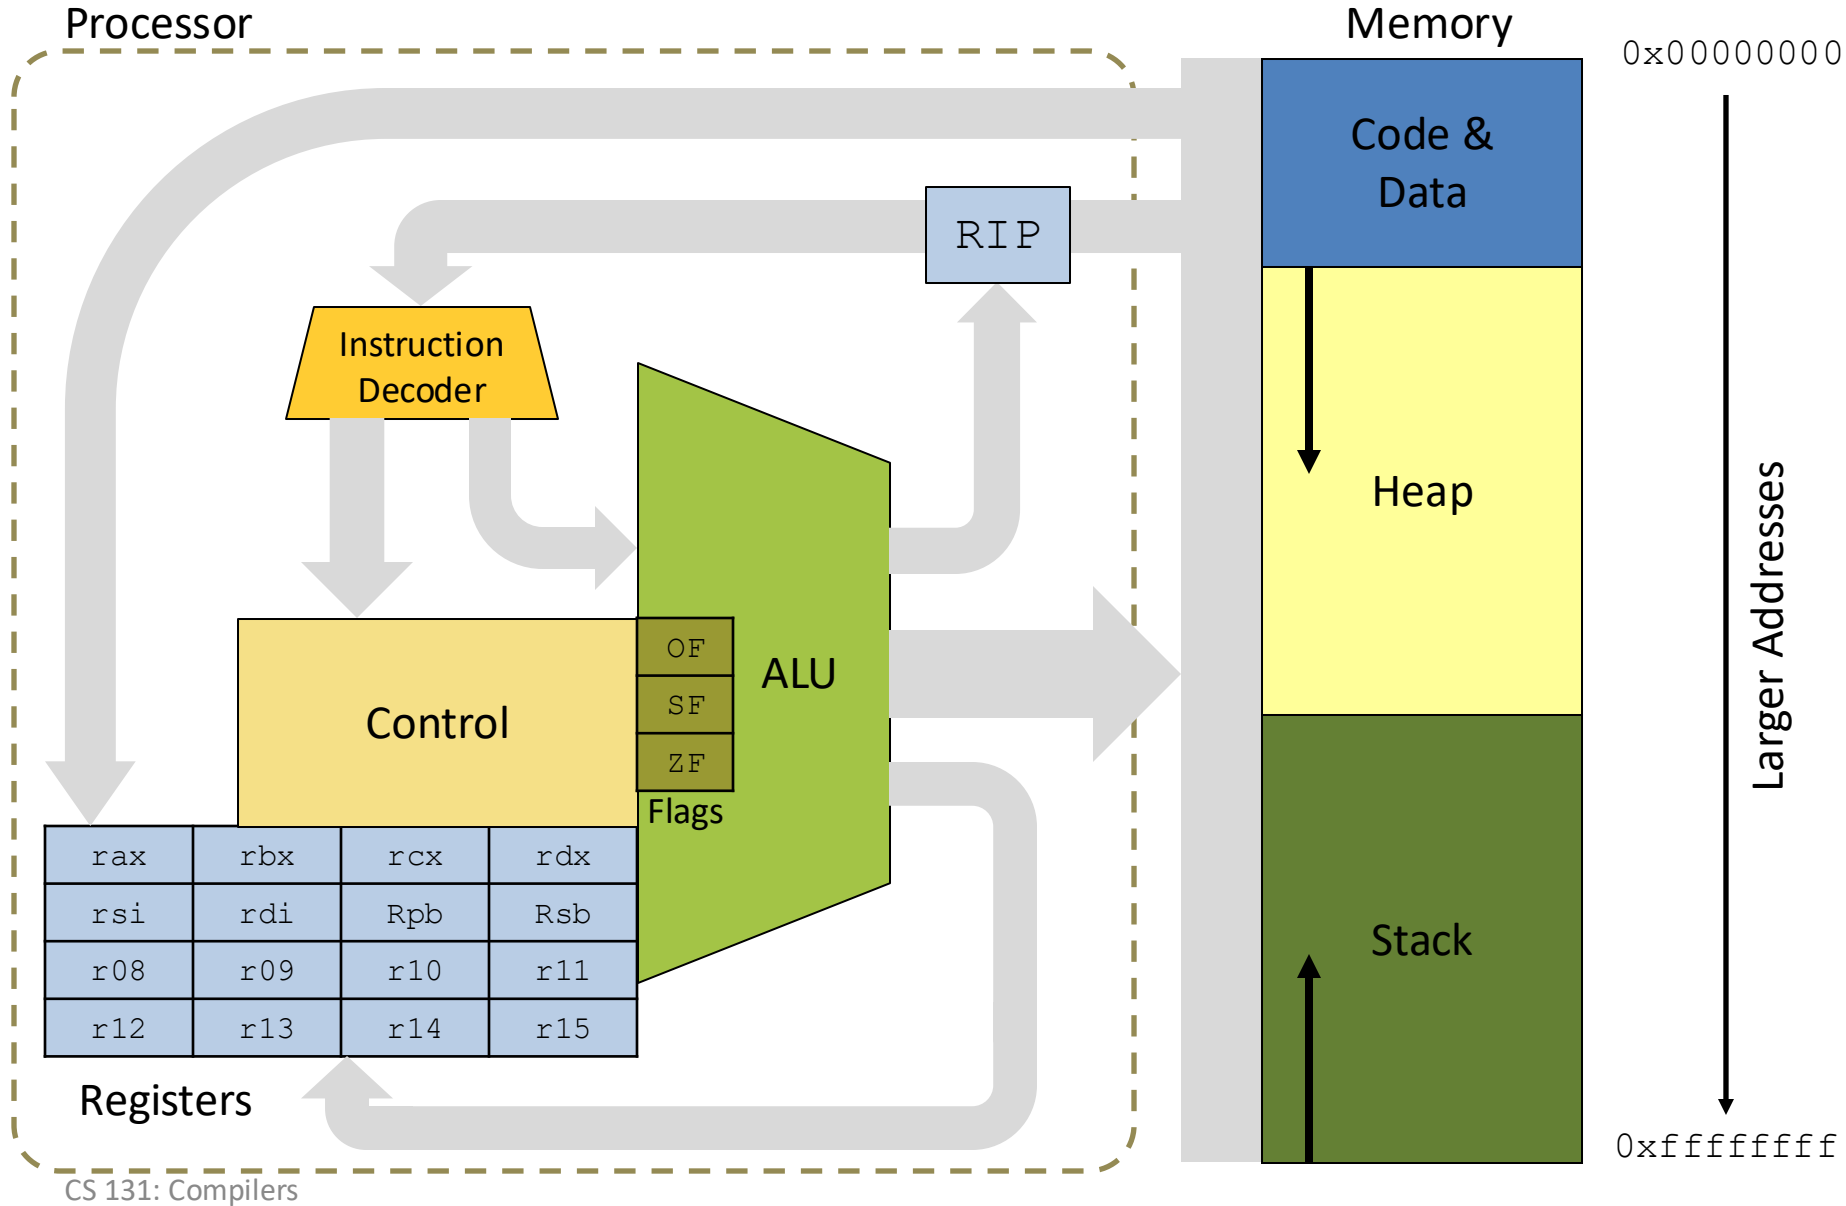
Lecture 4

# CIS 3410/7000: COMPILERS

# Announcements

- HW2: X86lite
  - Available on the course web pages soon. (look for announcement on Blackboard / Piazza).
  - Due: October 21<sup>st</sup>.
  - Pair-programming project
  - NOTE: much more difficult than hw1, so please start early!

# X86 Schematic



# X86lite State: Condition Flags & Codes

- X86 instructions set flags as a side effect
- X86lite has only 3 flags:
  - OF: “**overflow**” set when the result is too big/small to fit in 64-bit reg.
  - SF: “**sign**” set to the sign of the result (0=positive, 1 = negative)
  - ZF: “**zero**” set when the result is 0
- From these flags, we can define *Condition Codes*
  - To compare SRC1 and SRC2, compute SRC1 – SRC2 to set the flags
  - `eq` equality holds when ZF is set
  - `ne` inequality holds when (not ZF)
  - `gt` greater than holds when (not `le`) holds,
    - i.e.  $(SF = OF) \ \&\& \ \text{not}(ZF)$
  - `lt` less than holds when  $SF \neq OF$ 
    - Equivalently:  $((SF \ \&\& \ \text{not } OF) \ || \ (\text{not } SF \ \&\& \ OF))$
  - `ge` greater or equal holds when (not `lt`) holds, i.e.  $(SF = OF)$
  - `le` less than or equal holds when  $SF \neq OF$  or ZF

# Code Blocks & Labels

- X86 assembly code is organized into *labeled blocks*:

```
label1:
    <instruction>
    <instruction>
    ...
    <instruction>

label2:
    <instruction>
    <instruction>
    ...
    <instruction>
```

- Labels indicate code locations that can be jump targets (either through conditional branch instructions or function calls).
- Labels are translated away by the linker and loader – instructions live in the heap in the “code segment”
- An X86 program begins executing at a designated code label (usually “main”).

# Conditional Instructions

- `cmpq SRC1, SRC2`      Compute  $SRC2 - SRC1$ , set condition flags
- `setbCC DEST`       $DEST's\ lower\ byte \leftarrow$  if CC then 1 else 0
- `jCC SRC`       $rip \leftarrow$  if CC then SRC else fallthrough

- Example:

```
cmpq %rcx, %rax  
je __truelbl
```

Compare `rax` to `ecx`  
If `rax = rcx` then jump to `__truelbl`

# Jumps, Call and Return

- `jmp SRC`      `rip ← SRC`      Jump to location in SRC
- `callq SRC`      Push `rip`; `rip ← SRC`
  - Call a procedure: Push the program counter to the stack (decrementing `rsp`) and then jump to the machine instruction at the address given by SRC.
- `retq`      Pop into `rip`
  - Return from a procedure: Pop the current top of the stack into `rip` (incrementing `rsp`).
  - This instruction effectively jumps to the address at the top of the stack



See: runtime.c and x86.ml in lec04.zip

## DEMO: HANDCODING X86LITE




# Compiling, Linking, Running

- To use hand-coded X86:
  1. Compile main.ml (or something like it) to either native or bytecode
  2. Run it, redirecting the output to some .s file, e.g.:  
`./main.exe >> test.s`
  3. Use gcc to compile & link with runtime.c:  
`gcc -o test runtime.c test.s`
  4. You should be able to run the resulting executable:  
`./test`
- If you want to debug in gdb:
  - Call gcc with the `-g` flag too

# PROGRAMMING IN X86LITE

# 3 parts of the C memory model

- The code & data (or "text") segment
  - contains compiled code, constant strings, etc.
- The Heap
  - Stores dynamically allocated objects
  - Allocated via "malloc"
  - Deallocated via "free"
  - managed by C runtime system
- The Stack
  - Stores local variables
  - Stores the return address of a function
- In practice, most languages use this model.

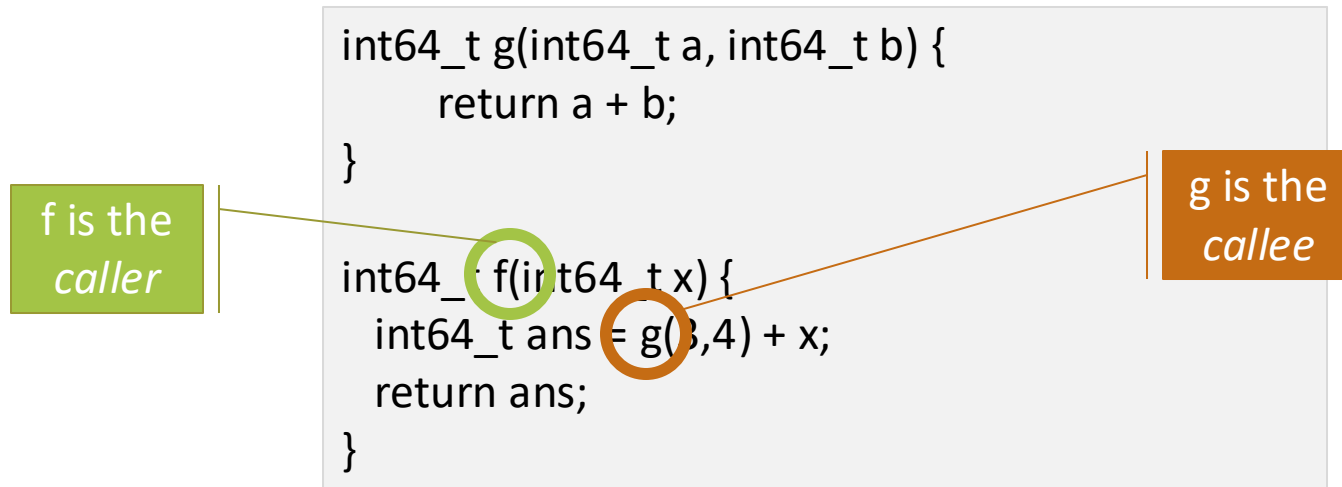


# Local/Temporary Variable Storage

- Need space to store:
  - Global variables
  - Values passed as arguments to procedures
  - Local variables (either defined in the source program or introduced by the compiler)
- Processors provide two options
  - Registers: fast, small size (64 bits), very limited number
  - Memory: slow, very large amount of space (2 GB)
    - caching important
- In practice on X86:
  - Registers are limited (and have restrictions)
  - Divide memory into regions including the *stack* and the *heap*

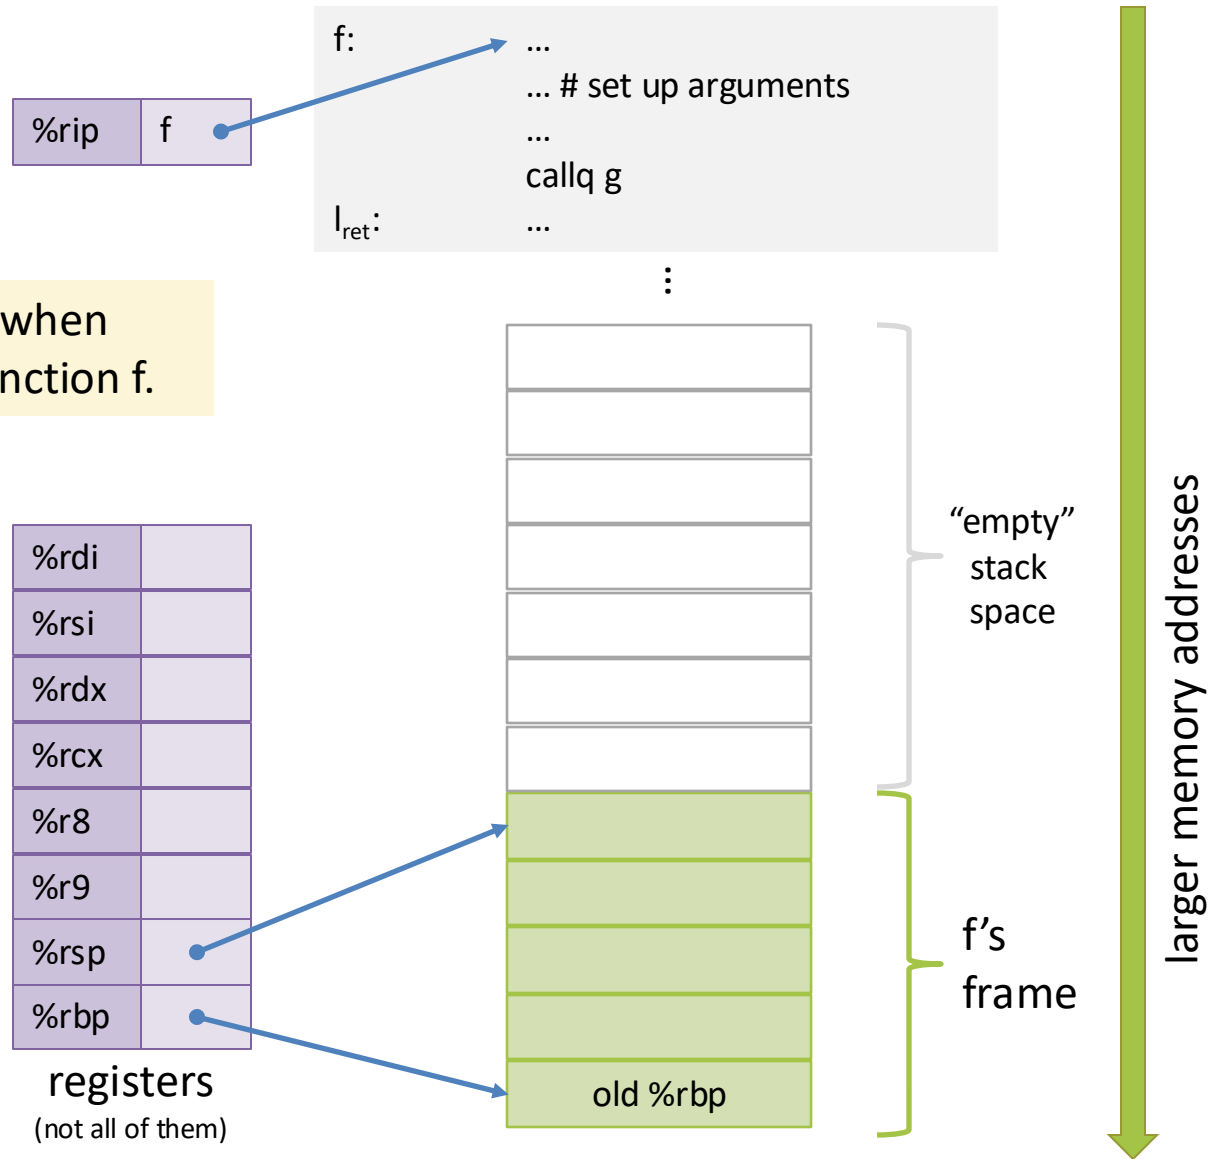
# Calling Conventions

- Specify the locations (e.g., register or stack) of arguments passed to a function and returned by the function

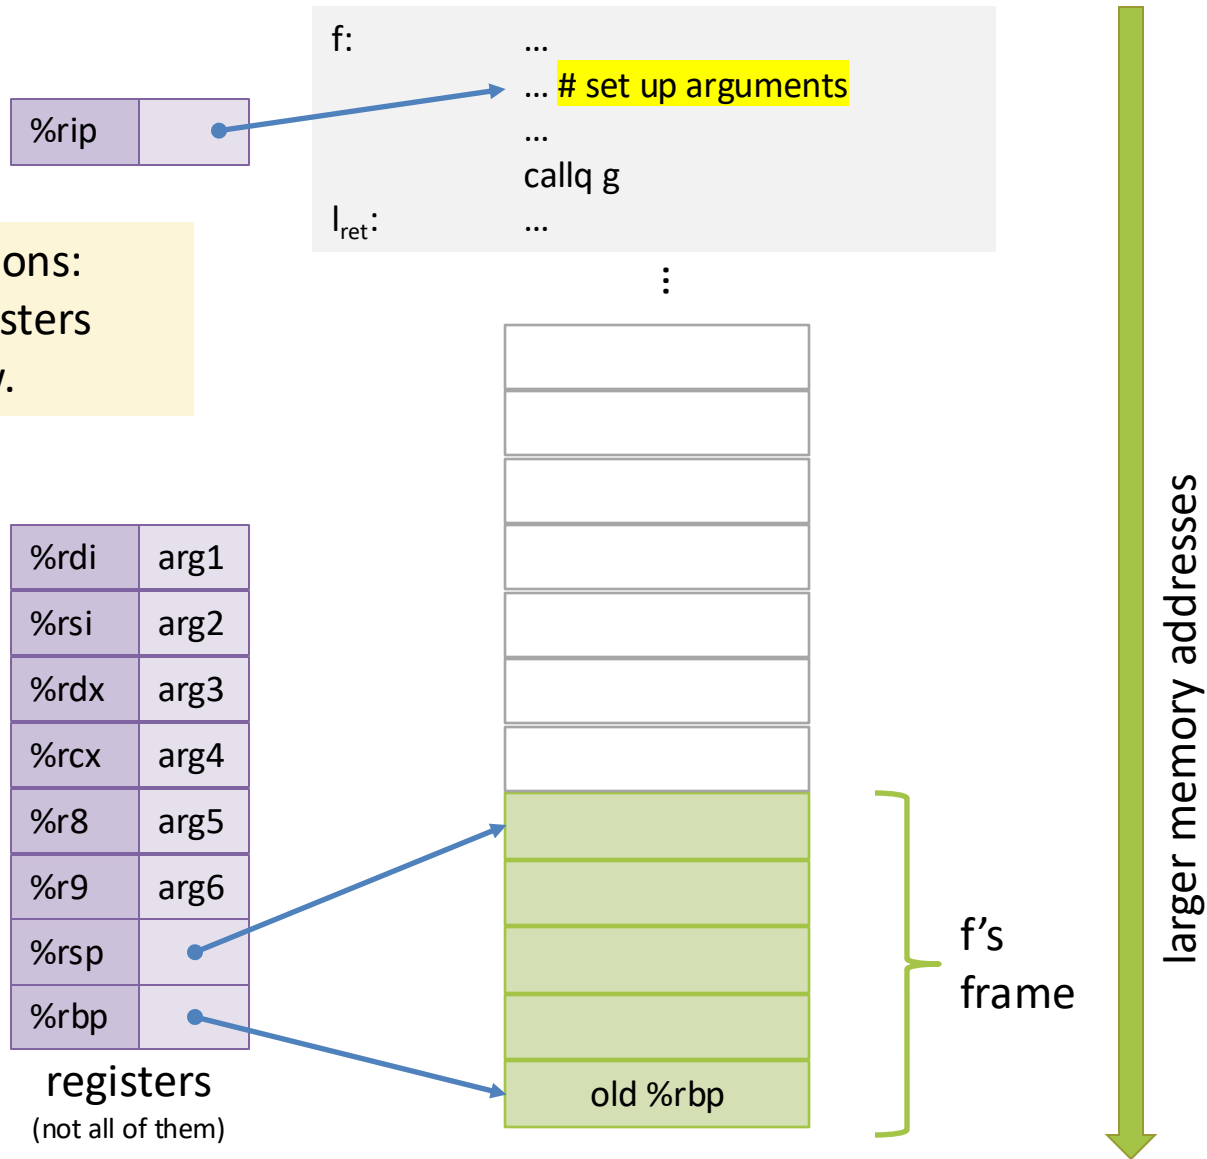


- Designate registers either:
  - **Caller Save** – e.g., freely usable by the called code
  - **Callee Save** – e.g., must be restored by the called code
- Define the protocol for deallocating stack-allocated arguments
  - Caller cleans up
  - Callee cleans up (makes supporting variable number of arguments harder)

# x64 Calling Conventions: Caller Protocol



# x64 Calling Conventions: Caller Protocol



# x64 Calling Conventions: Caller Protocol


%rip

```
f:    ...  
    ... # set up arguments  
    ...  
    callq g  
l_ret:  
    ...
```

args > 6 pushed onto  
the stack (from right to left)  
Note: %rsp changes

%rdi	arg1
%rsi	arg2
%rdx	arg3
%rcx	arg4
%r8	arg5
%r9	arg6
%rsp	
%rbp	

registers  
(not all of them)



f's  
frame

larger memory addresses



# call instruction

%rip

```
f:    ...  
    ... # set up arguments  
    ...  
    callq g  
l_ret: ...
```

To execute the call:

1. push the *return* address  
(here shown as  $l_{ret}$ )

%rdi	arg1
%rsi	arg2
%rdx	arg3
%rcx	arg4
%r8	arg5
%r9	arg6
%rsp	
%rbp	

registers  
(not all of them)



f's  
frame

larger memory addresses

# call instruction

%rip	
------	--

g:

```
pushq %rbp
movq %rsp, %rbp
subq $128, %rsp
...
```

To execute the call:  
2. set rip to address g

%rdi	arg1
%rsi	arg2
%rdx	arg3
%rcx	arg4
%r8	arg5
%r9	arg6
%rsp	
%rbp	

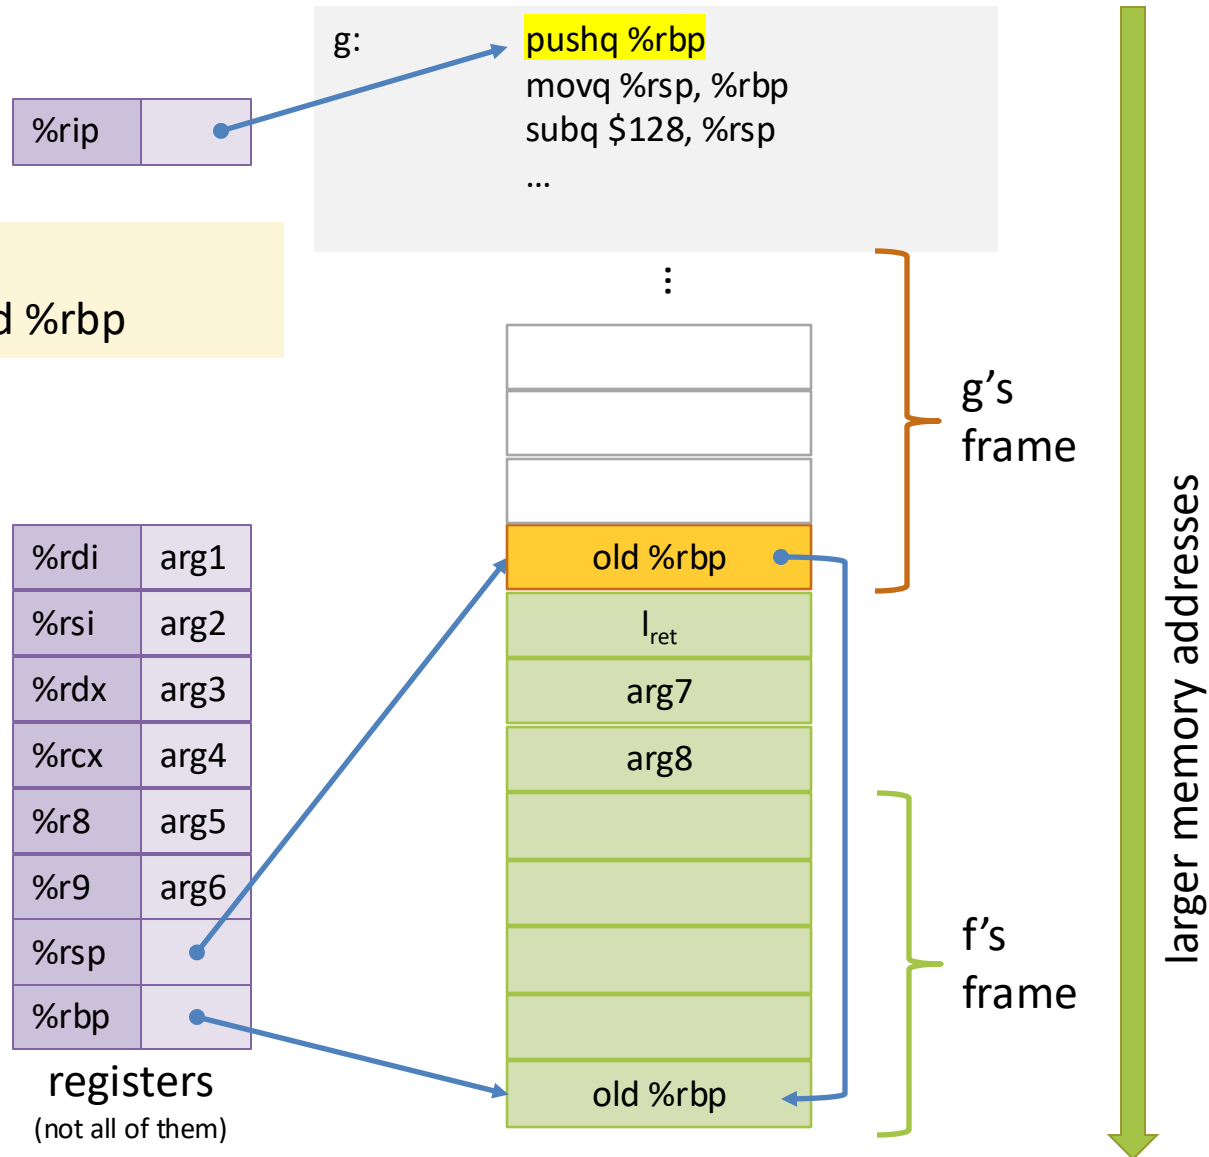
registers  
(not all of them)



f's  
frame

larger memory addresses

# callee function prologue



# callee function prologue

%rip

```
g:  pushq %rbp
    movq %rsp, %rbp
    subq $128, %rsp
    ...
```

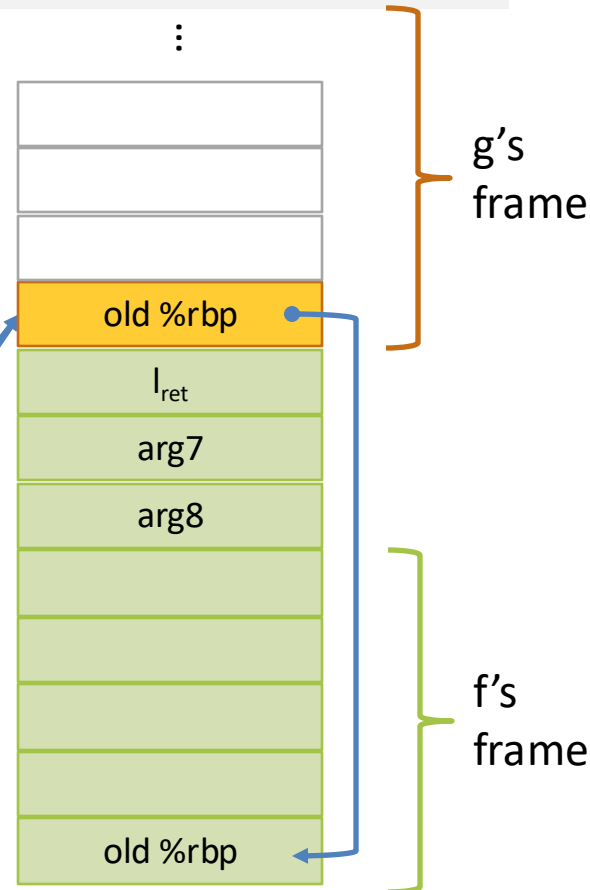
Callee protocol:

2. adjust the %rbp to point to the new “base”

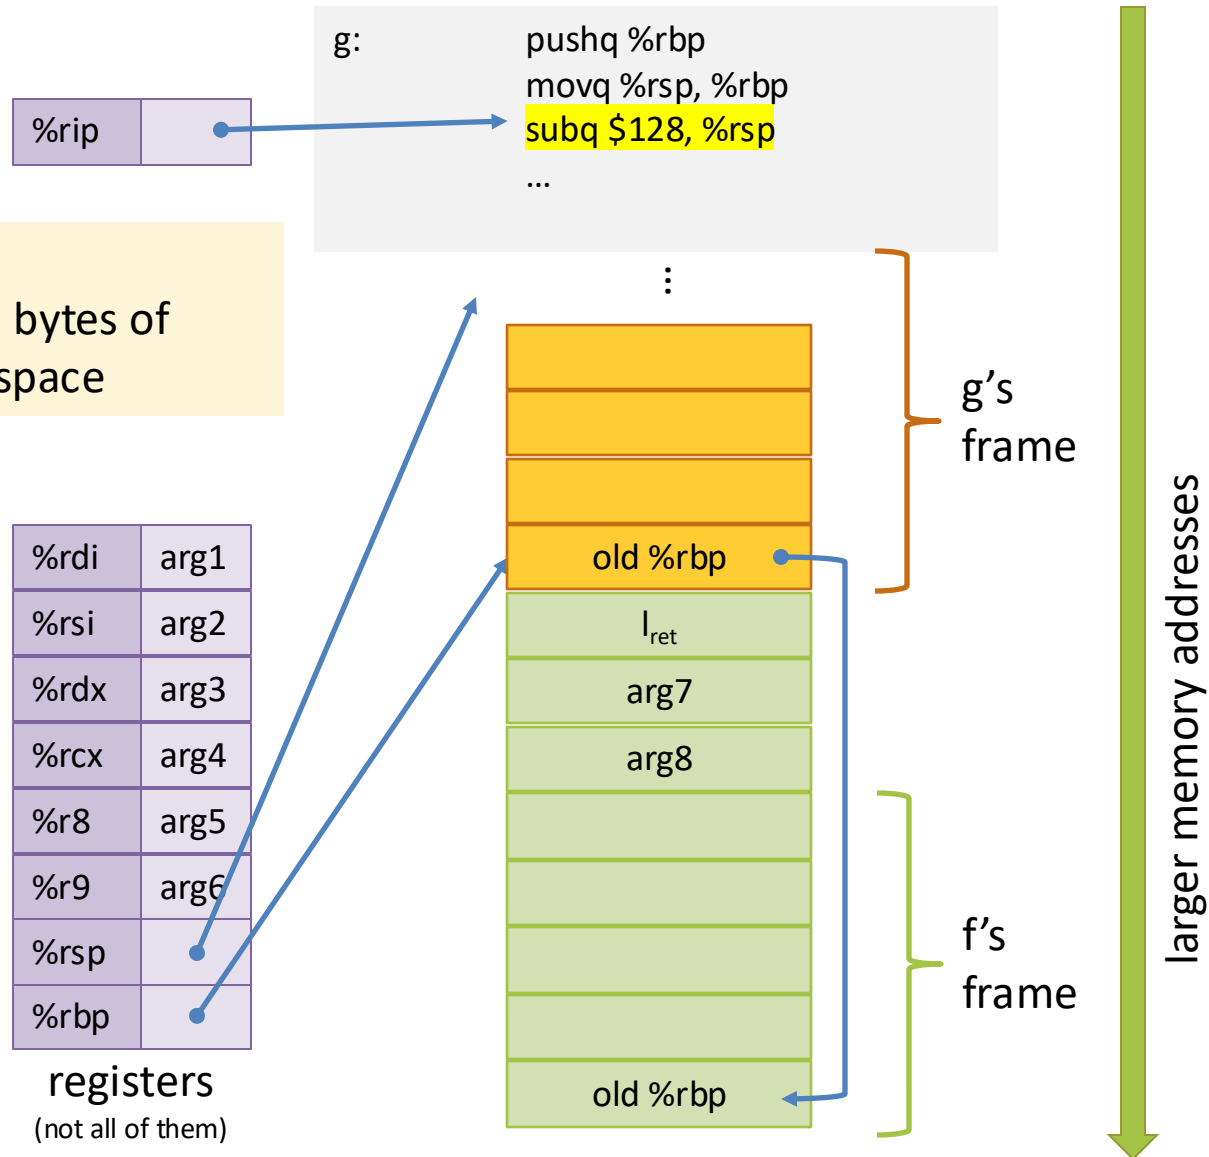
(%rbp is the “base pointer”)

%rdi	arg1
%rsi	arg2
%rdx	arg3
%rcx	arg4
%r8	arg5
%r9	arg6
%rsp	
%rbp	

registers  
(not all of them)



# callee function prologue



# callee invariants: function arguments

%rip

```
g:  pushq %rbp
    movq %rsp, %rbp
    subq $128, %rsp
    ...
```

Now g's body can run...

- its arguments are accessible either in registers or as offsets from %rbp

%rdi	arg1
%rsi	arg2
%rdx	arg3
%rcx	arg4
%r8	arg5
%r9	arg6
%rsp	
%rbp	

registers  
(not all of them)

16(%rbp)

24(%rbp)

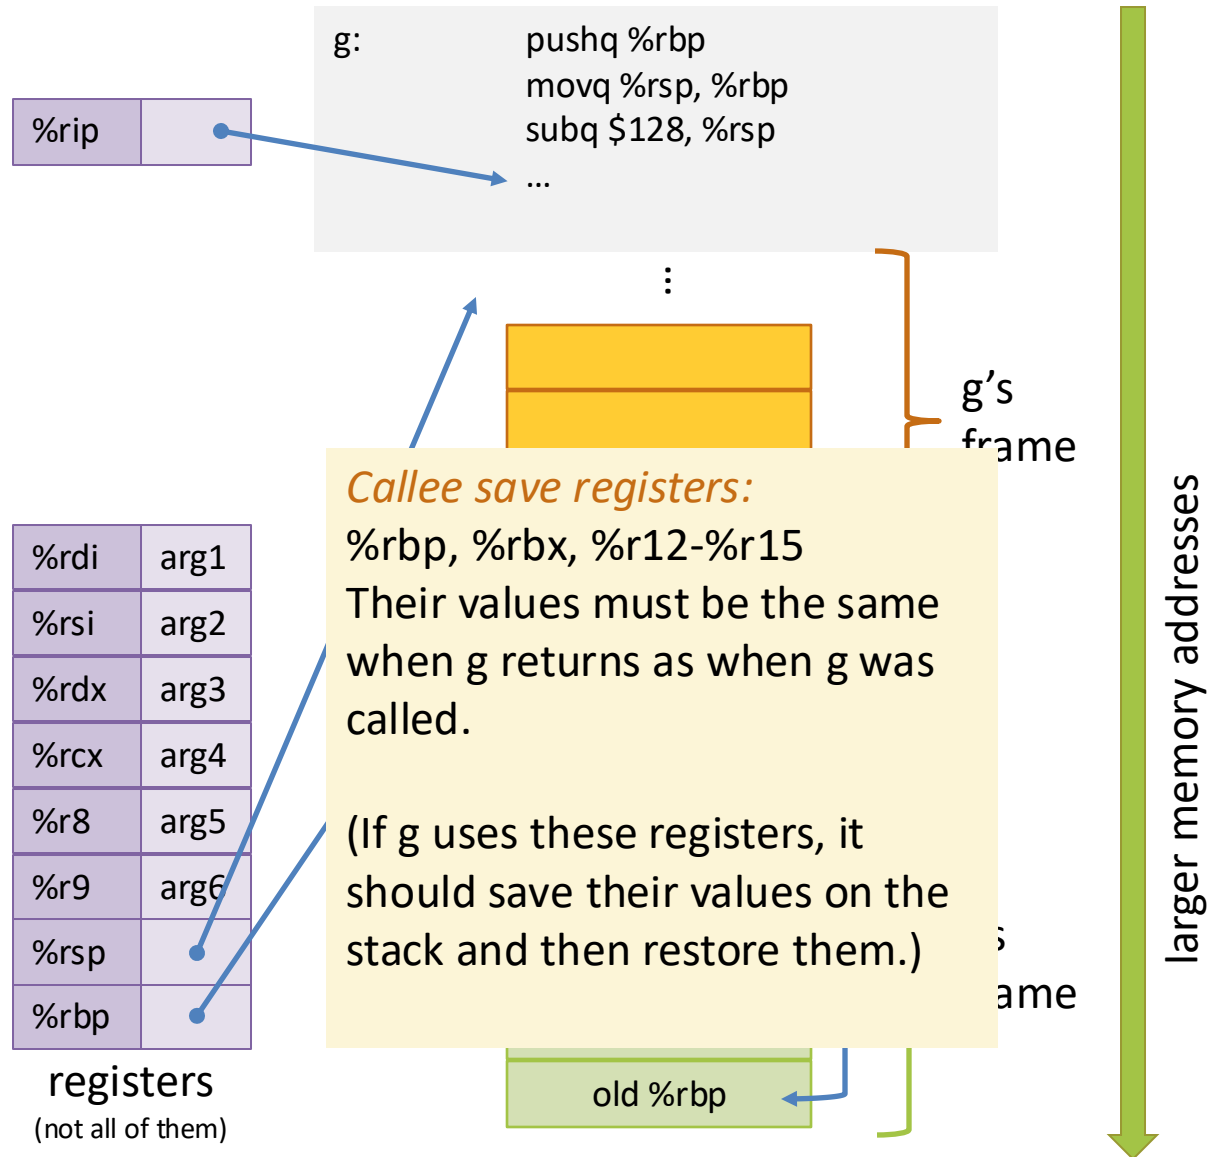


g's  
frame

f's  
frame

larger memory addresses

# callee invariants: callee save registers



# callee epilogue (return protocol)

%rip

Step 1: Move the result  
(if any) into %rax.

```
g:
...
movq ANS, %rax
addq $128, %rsp
popq %rbp
retq
```

%rax    **ANS**

%rdi    arg1

%rsi    arg2

%rdx    arg3

%rcx    arg4

%r8    arg5

%r9    arg6

%rsp

%rbp

registers  
(not all of them)



g's  
frame

f's  
frame

larger memory addresses



# callee epilogue (return protocol)

%rip

Step 2: deallocate the  
scratch space

```
g:    ...  
      movq ANS, %rax  
      addq $128, %rsp  
      popq %rbp  
      retq
```

⋮

%rax	ANS
%rdi	arg1
%rsi	arg2
%rdx	arg3
%rcx	arg4
%r8	arg5
%r9	arg6
%rsp	
%rbp	

registers  
(not all of them)



g's  
frame

f's  
frame

larger memory addresses

# callee epilogue (return protocol)

%rip

Step 3: restore the caller's %rbp

```
g:    ...  
      movq ANS, %rax  
      addq $128, %rsp  
      popq %rbp  
      retq
```

⋮

%rax    ANS

%rdi    arg1

%rsi    arg2

%rdx    arg3

%rcx    arg4

%r8     arg5

%r9     arg6

%rsp    

%rbp    

registers  
(not all of them)



g's  
frame

f's  
frame

larger memory addresses

# callee epilogue (return protocol)

%rip

Step 4: the return instruction pops the stack into %rip

```
g:    ...  
      movq ANS, %rax  
      addq $128, %rsp  
      popq %rbp  
      retq
```

⋮

%rax    ANS

%rdi    arg1

%rsi    arg2

%rdx    arg3

%rcx    arg4

%r8     arg5

%r9     arg6

%rsp    

%rbp    

registers  
(not all of them)



g's  
frame

f's  
frame

larger memory addresses

# callee epilogue (return protocol)

%rip

Step 4: the return instruction pops the stack into %rip

```
f:    ...  
    ... # set up arguments  
    ...  
    callq g  
    ...  
    .  
    .  
    .
```

l<sub>ret</sub>:

%rax    ANS

%rdi    arg1

%rsi    arg2

%rdx    arg3

%rcx    arg4

%r8    arg5

%r9    arg6

%rsp    

%rbp    

registers  
(not all of them)



f's  
frame

larger memory addresses

# back in f

%rip	
------	--

At this point, f has the result of g in %rax. It should clean up its stack as needed.

```
f:      ...  
      ... # set up arguments  
      ...  
      callq g  
l_ret:  ...
```

⋮

%rax	ANS
%rdi	arg1
%rsi	arg2
%rdx	arg3
%rcx	arg4
%r8	arg5
%r9	arg6
%rsp	
%rbp	

registers  
(not all of them)



f's  
frame

larger memory addresses

# X86-64 SYSTEM V AMD 64 ABI

- Modern variant of C calling conventions
  - used on Linux, Solaris, BSD, OS X
- Callee save: %rbp, %rbx, %r12-%r15
- Caller save: all others
- Parameters 1 .. 6 go in: %rdi, %rsi, %rdx, %rcx, %r8, %r9
- Parameters 7+ go on the stack (in right-to-left order)
  - so: for  $n > 6$ , the  $n^{\text{th}}$  argument is located at  $((n-7)+2)*8(\%rbp)$
  - e.g.: argument 7 is at  $16(\%rbp)$  and argument 8 is at  $24(\%rbp)$
- Return value: in %rax
- 128 byte "red zone" – scratch pad for the callee's data
  - typical of C compilers, not required
  - can be optimized away

# 32-bit cdecl calling conventions

- Still “Standard” on X86 for many C-based operating systems
  - Still some wrinkles about return values  
(*e.g.*, some compilers use EAX and EDX to return small values)
  - 64 bit allows for packing multiple values in one register
- All arguments are passed on the stack in right-to-left order
- Return value is passed in EAX
- Registers EAX, ECX, EDX are caller save
- Other registers are callee save
  - Ignoring these conventions will cause havoc (bus errors or seg faults)