

Lecture 19

CS131: COMPILERS

Announcements

- HW5: OAT v. 2.0
 - records, function pointers, type checking, array-bounds checks, etc.
 - Due: Wednesday, December 13th
 - Available now
 - **Start Early!**
- Teaching evaluation
 - By January 5th
 - Submit evaluation to reveal final scores.





MUTABILITY & SUBTYPING

NULL

- What is the type of null?
- Consider:
 `int[] a = null; // OK?`
 `int x = null; // not OK?`
 `string s = null; // OK?`

NULL

$G \vdash \text{null} : r$

- Null has any *reference type*
 - Null is generic
- What about type safety?
 - Requires defined behavior when dereferencing null
e.g. Java's `NullPointerException`
 - Requires a safety check for every dereference operation
(typically implemented using low-level hardware "trap" mechanisms.)

Subtyping and References

- What is the proper subtyping relationship for references and arrays?
- Suppose we have NonZero as a type and the division operation has type:
 $\text{Int} \rightarrow \text{NonZero} \rightarrow \text{Int}$
 - Recall that $\text{NonZero} <: \text{Int}$
- Should $(\text{NonZero ref}) <: (\text{Int ref})$?
- Consider this program:

```
Int bad(NonZero ref r) {  
  Int ref a = r; (* OK because (NonZero ref <: Int ref*)  
  a := 0;        (* OK because 0 : Zero <: Int *)  
  return (42 / !r) (* OK because !r has type NonZero *)  
}
```

Mutable Structures are Invariant

- Covariant reference types are unsound
 - As demonstrated in the previous example
- Contravariant reference types are also unsound
 - *i.e.*, If $T_1 <: T_2$ then $\text{ref } T_2 <: \text{ref } T_1$ is also unsound
 - Exercise: construct a program that breaks contravariant references.
- Moral: Mutable structures are *invariant*:
$$T_1 \text{ ref } <: T_2 \text{ ref} \implies T_1 = T_2$$
- Same holds for arrays, OCaml-style mutable records, object fields, etc.
 - Note: Java and C# get this wrong. They allow covariant array subtyping, but then compensate by adding a dynamic check on *every* array update!

Another Way to See It

- We can think of a reference cell as an immutable record (object) with two functions (methods) and some hidden state:

$$T \text{ ref} \simeq \{\text{get}: \text{unit} \rightarrow T; \text{set}: T \rightarrow \text{unit}\}$$

- get returns the value hidden in the state.
 - set updates the value hidden in the state.
- When is $T \text{ ref} <: S \text{ ref}$?
- Records with depth subtyping:
 - extends pointwise over each component.
$$\{\text{get}: \text{unit} \rightarrow T; \text{set}: T \rightarrow \text{unit}\} <: \{\text{get}: \text{unit} \rightarrow S; \text{set}: S \rightarrow \text{unit}\}$$
 - get components are subtypes: $\text{unit} \rightarrow T <: \text{unit} \rightarrow S$
 - set components are subtypes: $T \rightarrow \text{unit} <: S \rightarrow \text{unit}$
- From get, we must have $T <: S$ (covariant return)
- From set, we must have $S <: T$ (contravariant arg.)
- From $T <: S$ and $S <: T$ we conclude $T = S$.



See oat.pdf in HW5

OAT'S TYPE SYSTEM

OAT's Treatment of Types

- Primitive (non-reference) types:
 - int, bool
- Definitely-non-null reference types: R
 - (named) *mutable* structs with (right-oriented) *width* subtyping
 - string
 - arrays (including length information, per HW4)
- Possibly-null reference types: R?
 - Subtyping: $R <: R?$
 - *Checked downcast* syntax if?:

```
int sum(int[]? arr) {  
    var z = 0;  
    if?(int[] a = arr) {  
        for(var i = 0; i < length(a); i = i + 1;) {  
            z = z + a[i];  
        }  
    }  
    return z;  
}
```

OAT Features

- Named structure types with mutable fields
 - but using structural, width subtyping
- Typed function pointers
- Polymorphic operations: length and == / !=
 - need special case handling in the typechecker
- Type-annotated null values: `R null` always has type `R`?
- Definitely-not-null values means we need an "atomic" array initialization syntax
 - null is not allowed as a value of type `int[]`, so to construct a record containing a field of type `int[]`, we need to initialize it
 - subtlety: `int[][]` cannot be initialized by default, but `int[]` can be

OAT "Returns" Analysis

- Typesafe, statement-oriented imperative languages like OAT (or Java) must ensure that a function (always) returns a value of the appropriate type.
 - Does the returned expression's type match the one declared by the function?
 - Do all paths through the code return appropriately?
- OAT's statement checking judgment
 - takes the expected return type as input: what type should the statement return (or void if none)
 - produces a Boolean flag as output: does the statement definitely return?

Example OAT code

```
struct Base {      /* struct type with function field */
    int a;
    bool b;
    (int) -> int f
}

struct Extend {    /* structural subtype of Base via width subtyping */
    int a;
    bool b;
    (int) -> int f;
    string c;      /* added field and method */
    (int) -> int g
}

int neg(int x) { return -x; }
int inc(int x) { return x+1; }

int f(Base? x, int y){ /* function that expects a (possibly null) Base */
    if?(Base b = x){
        return b.f(y);
    } else {
        return -1;
    }
}

int program(int argc, string[] argv) {
    var s = new Extend[5]{x -> new Extend{a=3; b=true; c="hello"; f=neg; g=inc}};
    return f(s[2], -3);
}
```



STRUCTURAL VS. NOMINAL TYPES

Structural vs. Nominal Typing

- Is type equality / subsumption defined by the *structure* of the data or the *name* of the data?
- Example 1: type abbreviations (OCaml) vs. “newtypes” (a la Haskell)

```
(* OCaml: *)  
type cents = int  (* cents = int in this scope *)  
type age = int  
  
let foo (x:cents) (y:age) = x + y
```

```
-- Haskell:  
newtype Cents = Cents Integer -- Integer and Cents are  
                             -- isomorphic, not identical  
newtype Age = Age Integer  
  
foo :: Cents -> Age -> Int  
foo x y = x + y           -- Ill typed!
```

- OCaml type abbreviations are treated “**structurally**”
Haskell newtypes are treated “**by name**”

Nominal Subtyping in Java

- Example 2: In Java, Classes and Interfaces must be named and their relationships *explicitly* declared:

```
(* Java: *)
interface Foo {
    int foo();
}

class C {          /* Does not implement the Foo interface */
    int foo() {return 2;}
}

class D implements Foo {
    int foo() {return 3410;}
}
```

- Similarly for inheritance: programmers must declare the subclass relation via the “extends” keyword.
 - Typechecker still checks that the classes are structurally compatible



COMPILING CLASSES AND OBJECTS

Code Generation for Objects

- Classes:
 - Generate data structure types
 - For objects that are instances of the class and for the class tables
 - Generate the class tables for dynamic dispatch
- Methods:
 - Method body code is similar to functions/closures
 - Method calls require *dispatch*
- Fields:
 - Issues are the same as for records
 - Generating access code
- Constructors:
 - Object initialization
- Dynamic Types:
 - Checked downcasts
 - “instanceof” and similar type dispatch

Multiple Implementations

- The same interface can be implemented by multiple classes:

```
interface IntSet {  
    public IntSet insert(int i);  
    public boolean has(int i);  
    public int size();  
}
```

```
class IntSet1 implements IntSet {  
    private List<Integer> rep;  
    public IntSet1() {  
        rep = new LinkedList<Integer>();  
    }  
  
    public IntSet1 insert(int i) {  
        rep.add(new Integer(i));  
        return this;  
    }  
  
    public boolean has(int i) {  
        return rep.contains(new Integer(i));  
    }  
  
    public int size() {return rep.size();}  
}
```

```
class IntSet2 implements IntSet {  
    private Tree rep;  
    private int size;  
    public IntSet2() {  
        rep = new Leaf(); size = 0;  
    }  
  
    public IntSet2 insert(int i) {  
        Tree nrep = rep.insert(i);  
        if (nrep != rep) {  
            rep = nrep; size += 1;  
        }  
        return this;  
    }  
  
    public boolean has(int i) {  
        return rep.find(i);  
    }  
  
    public int size() {return size;}  
}
```

The Dispatch Problem



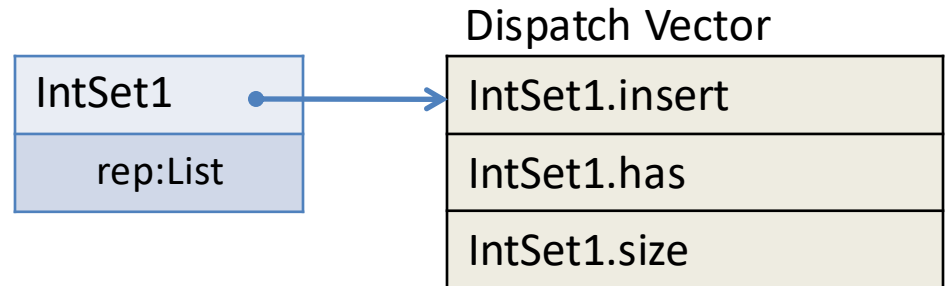
- Consider a client program that uses the IntSet interface:

```
IntSet set = ...;  
int x = set.size();
```

- Which code to call?
 - IntSet1.size ?
 - IntSet2.size ?
- Client code doesn't know the answer.
 - So objects must “know” which code to call.
 - Invocation of a method must indirect through the object.

Compiling Objects

- Objects contain a pointer to a *dispatch vector* (also called a *virtual table* or *vtable*) with pointers to method code.
- Code receiving `set:IntSet` only knows that `set` has an initial dispatch vector pointer and the layout of that vector.



Method Dispatch (Single Inheritance)

- Idea: every method has its own small integer index.
- Index is used to look up the method in the dispatch vector.

```
interface A {  
  void foo();  
}
```

Index

0

```
interface B extends A {  
  void bar(int x);  
  void baz();  
}
```

1

2

Inheritance / Subtyping:
C <: B <: A

```
class C implements B {  
  void foo() {...}  
  void bar(int x) {...}  
  void baz() {...}  
  void quux() {...}  
}
```

0

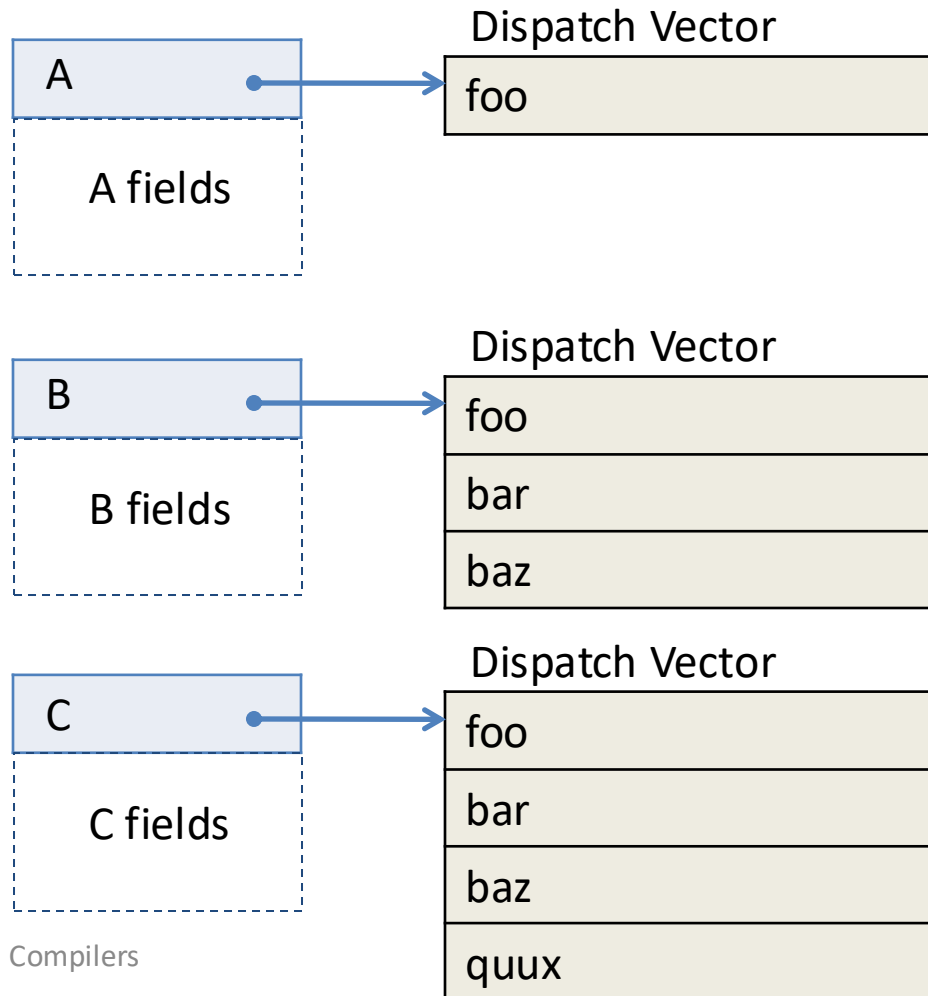
1

2

3

Dispatch Vector Layouts

- Each interface and class gives rise to a dispatch vector layout.
- Note that inherited methods have identical dispatch indices in the subclass. (*Width subtyping*)



Representing Classes in the LLVM

- During typechecking, create a *class hierarchy*
 - Maps each class to its interface:
 - Superclass
 - Constructor type
 - Fields
 - Method types (plus whether they inherit & which class they inherit from)
- Compile the class hierarchy to produce:
 - An LLVM IR struct type for each object instance
 - An LLVM IR struct type for each vtable (a.k.a. class table)
 - Global definitions that implement the class tables

Example OO Code (Java)

```
class A {  
    A (int x)           // constructor  
    { super(); int x = x; }  
  
    void print() { return; } // method1  
    int blah(A a) { return 0; } // method2  
  
}  
  
class B extends A {  
    B (int x, int y, int z){  
        super(x);  
        int y = y;  
        int z = z;  
    }  
  
    void print() { return; } // overrides A  
}  
  
class C extends B {  
    C (int x, int y, int z, int w){  
        super(x,y,z);  
        int w = w;  
    }  
    void foo(int a, int b) {return;}  
    void print() {return;} // overrides B  
}
```


Type Translation of a Class

- Each class gives rise to two implementation types:
- Object Instance Type
 - pointer to the dispatch vector
 - fields of the class
- Dispatch Vector Type
 - pointer to the superclass dispatch vector
 - pointers to methods of the class
- The inheritance hierarchy is used to statically construct the global class tables
 - which are records that have Dispatch Vector Types

Example OO Hierarchy in LLVM

Object instance types

```
%Object = type { %_class_Object* }
```

```
%_class_Object = type { }
```

```
%A = type { %_class_A*, i64 }
```

```
%_class_A = type { %_class_Object*, void (%A*)*, i64 (%A*, %A*)* }
```

```
%B = type { %_class_B*, i64, i64, i64 }
```

```
%_class_B = type { %_class_A*, void (%B*)*, i64 (%A*, %A*)* }
```

```
%C = type { %_class_C*, i64, i64, i64, i64 }
```

```
%_class_C = type { %_class_B*, void (%C*)*, i64 (%A*, %A*)*, void (%C*, i64, i64)* }
```

Class table types

```
@_vtbl_Object = global %_class_Object { }
```

```
@_vtbl_A = global %_class_A { %_class_Object* @_vtbl_Object,  
    void (%A*)* @print_A,  
    i64 (%A*, %A*)* @blah_A }
```

```
@_vtbl_B = global %_class_B { %_class_A* @_vtbl_A,  
    void (%B*)* @print_B,  
    i64 (%A*, %A*)* @blah_A }
```

```
@_vtbl_C = global %_class_C { %_class_B* @_vtbl_B,  
    void (%C*)* @print_C,  
    i64 (%A*, %A*)* @blah_A,  
    void (%C*, i64, i64)* @foo_C }
```

Class tables
(structs containing
function pointers)

Method Arguments

- Methods bodies are compiled just like top-level procedures...
- ... except that they have an implicit extra argument: `this` (or `self`)
 - Historically (Smalltalk), these were called the “receiver object”
 - Method calls were thought of as sending “messages” to “receivers”

A method in a class...

```
class IntSet1 implements IntSet {  
    ...  
    IntSet1 insert(int i) { <body> }  
}
```

... is compiled like this (top-level) procedure:

```
IntSet1 insert(IntSet1 this, int i) { <body> }
```


- Note 1: the type of “`this`” is the class containing the method.
- Note 2: references to fields inside `<body>` are compiled like `this.field`

LLVM Method Invocation Compilation

- Consider method invocation:
$$\llbracket H;G;L \vdash e.m(e_1, \dots, e_n):t \rrbracket$$
- First, compile $\llbracket H;G;L \vdash e : C \rrbracket$
to get a (pointer to) an object value of class type C
 - Call this value %obj_ptr
- Use getelementptr to extract the vtable pointer from %obj_ptr
- load the vtable pointer
- Use getelementptr to extract the address of the function pointer from the vtable
 - using the information about C in H
- load the function pointer
- Call through the function pointer, passing ‘%obj_ptr’ for this:
call (cmp_typ t) m(obj_ptr, $\llbracket e_1 \rrbracket$, ..., $\llbracket e_n \rrbracket$)
- In general, function calls may require bitcast to account for subtyping: arguments may be a subtype of the expected “formal” type

X86 Code For Dynamic Dispatch

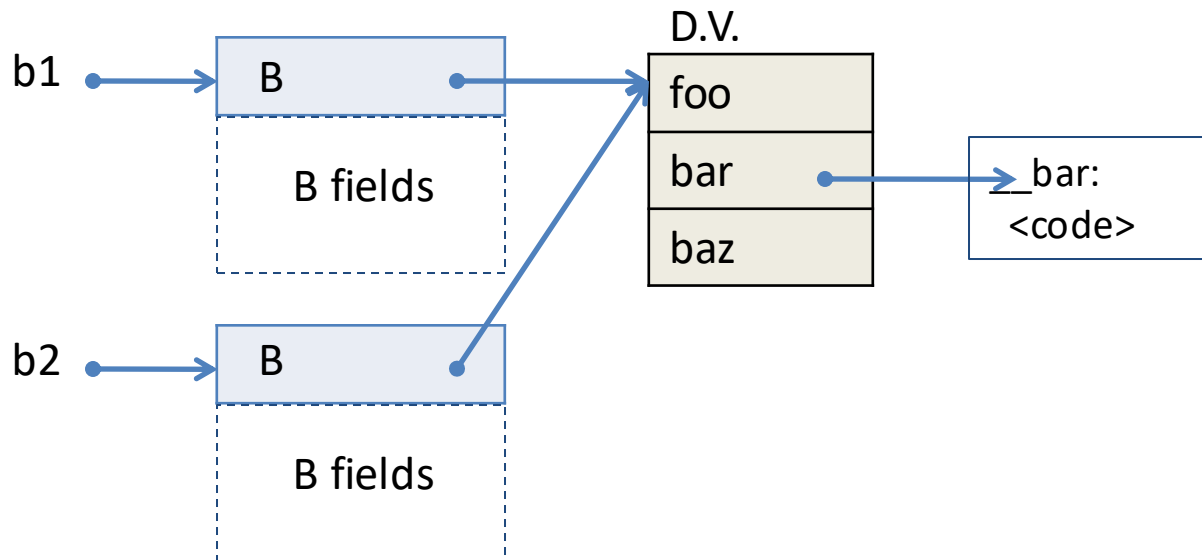
- Suppose $b : B$
- What code for $b.bar(3)$?
 - bar has index 1
 - $Offset = 8 * 1$



```
movq [[b]], %rax
movq [%rax], %rbx
movq [rbx+8], %rcx    // D.V. + offset
movq %rax, %rdi       // "this" pointer
movq 3, %rsi          // Method argument
call %rcx             // Indirect call
```

Sharing Dispatch Vectors

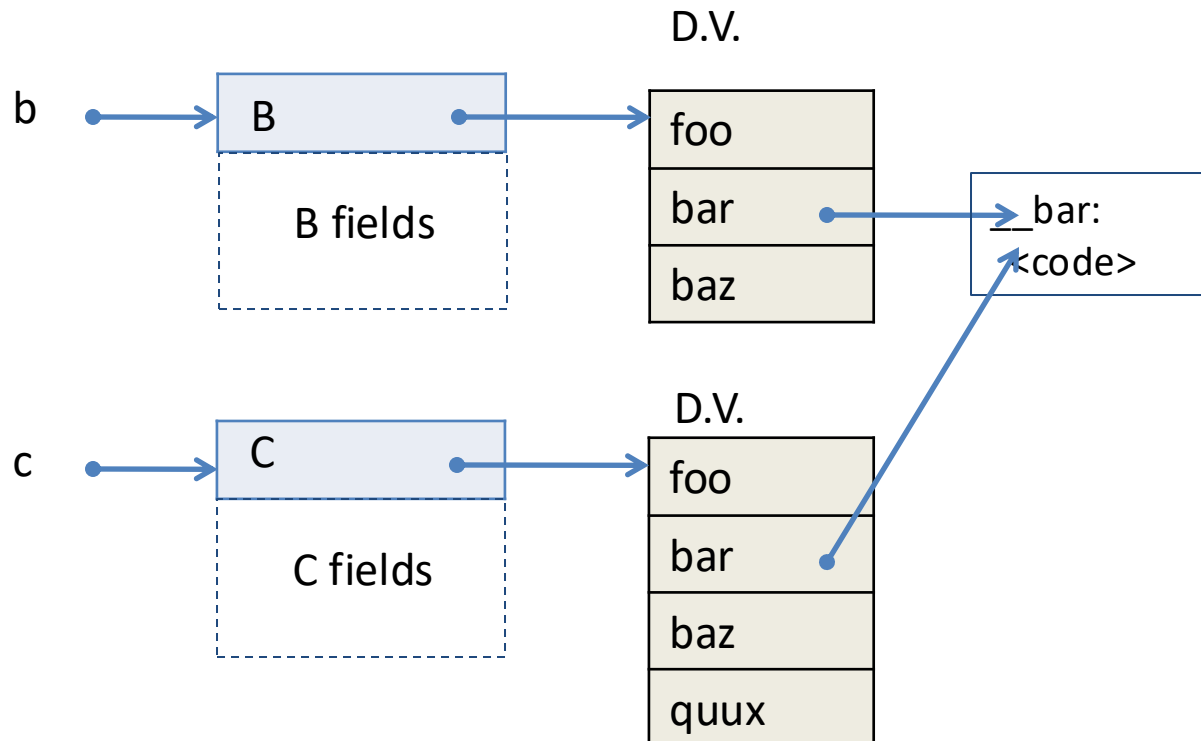
- All instances of a class may share the same dispatch vector.
 - Assuming that methods are immutable.
- Code pointers stored in the dispatch vector are available at link time – dispatch vectors can be built once at link time.



- One job of the object constructor is to fill in the object's pointer to the appropriate dispatch vector.
- Note: The address of the D.V. *is* the run-time representation of the object's type.

Inheritance: Sharing Code

- Inheritance: Method code “copied down” from the superclass
 - If not overridden in the subclass
- Works with separate compilation – superclass code not needed.



Compiling Static Methods

- Java supports *static* methods
 - Methods that belong to a class, not the instances of the class.
 - They have no “this” parameter (no receiver object)
- Compiled exactly like normal top-level procedures
 - No slots needed in the dispatch vectors
 - No implicit “this” parameter
- They’re not really methods
 - They can only access static fields of the class

Compiling Constructors

- Java and C++ classes can declare constructors that create new objects.
 - Initialization code may have parameters supplied to the constructor
 - e.g. `new Color(r,g,b);`
- Modula-3: object constructors take no parameters
 - e.g. `new Color;`
 - Initialization would typically be done in a separate method.
- Constructors are compiled just like static methods, except:
 - The “this” variable is initialized to a newly allocated block of memory big enough to hold D.V. pointer + fields according to object layout
 - Constructor code initializes the fields
 - What methods (if any) are allowed?
 - The D.V. pointer is initialized
 - When? Before/After running the initialization code?