

Lecture 20

CS 131: COMPILERS

Compiling Constructors

- Java and C++ classes can declare constructors that create new objects.
 - Initialization code may have parameters supplied to the constructor
 - e.g. `new Color(r,g,b);`
- Modula-3: object constructors take no parameters
 - e.g. `new Color;`
 - Initialization would typically be done in a separate method.
- Constructors are compiled just like static methods, except:
 - The “this” variable is initialized to a newly allocated block of memory big enough to hold D.V. pointer + fields according to object layout
 - Constructor code initializes the fields
 - What methods (if any) are allowed?
 - The D.V. pointer is initialized
 - When? Before/After running the initialization code?

Compiling Checked Casts

- How do we compile downcast in general? Consider this generalization of Oat's checked cast:

`if? (t x = exp) { ... } else { ... }`

- Reason by cases:
 - t must be either null, ref or ref? (can't be just int or bool)
- If t is null:
 - The static type of exp must be ref? for some ref.
 - If exp == null then take the true branch, otherwise take the false branch
- If t is string or t[]:
 - The static type of exp must be the corresponding string? Or t[]?
 - If exp == null take the false branch, otherwise take the true branch
- If t is C:
 - The static type of exp must be D or D? (where C <: D)
 - If exp == null take the false branch, otherwise:
 - emit code to walk up the class hierarchy starting at D, looking for C
 - If found, then take true branch else take false branch
- If t is C?:
 - The static type of exp must be D? (where C <: D)
 - If exp == null take the true branch, otherwise:
 - Emit code to walk up the class hierarchy starting at D, looking for C
 - If found, then take true branch else take false branch

“Walking up the Class Hierarchy”

- A non-null object pointer refers to an LLVM struct with a type like:

```
%B = type { %_class_B*, i64, i64, i64 }
```

- The first entry of the struct is a pointer to the vtable for Class B
 - This pointer *is* the dynamic type of the object.
 - It will have the value @vtbl_B
- The first entry of the class table for B is a pointer to its superclass:

```
@_vtbl_B = global %_class_B { %_class_A* @_vtbl_A,  
                             void (%B*)* @print_B,  
                             i64 (%A*, %A*)* @blah_A }
```

- Therefore, to find out whether an unknown type X is a subtype of C:
 - Assume C is not Object (ruled out by “silliness” checks for downcast)

LOOP:

If X == @vtbl_Object then NO, X is not a subtype of C

If X == @vtbl_C then YES, X is a subtype of C

else X == @vtbl_D, so set X to @vtbl_E where E is D's parent and goto LOOP



MULTIPLE INHERITANCE

Method Dispatch (Single Inheritance)

- Idea: every method has its own small integer index.
- Index is used to look up the method in the dispatch vector.

```
interface A {  
  void foo();  
}
```

Index

0

```
interface B extends A {  
  void bar(int x);  
  void baz();  
}
```

1

2

Inheritance / Subtyping:
C <: B <: A

```
class C implements B {  
  void foo() {...}  
  void bar(int x) {...}  
  void baz() {...}  
  void quux() {...}  
}
```

0

1

2

3

Multiple Inheritance

- C++: a class may declare more than one superclass.
- Semantic problem: *ambiguity*
class A { int m(); }
class B { int m(); }
class C extends A,B {...} // which m?
 - Same problem can happen with fields.
 - In C++, fields and methods can be duplicated when such ambiguity arises (though explicit sharing can be declared too)
- Java: a class may implement more than one interface.
 - No semantic ambiguity: if two interfaces contain the same method declaration, then the class will implement a single method
interface A { int m(); }
interface B { int m(); }
class C implements A,B {int m() {...}} // only one m

Dispatch Vector Layout Strategy Breaks

	D.V.Index
interface Shape {	
void setCorner(int w, Point p);	0
}	

interface Color {		
float get(int rgb);		0
void set(int rgb, float value);	1	
}		

class Blob implements Shape, Color {		
void setCorner(int w, Point p) {...}	0?	
float get(int rgb) {...}		0?
void set(int rgb, float value) {...}	1?	
}		

General Approaches

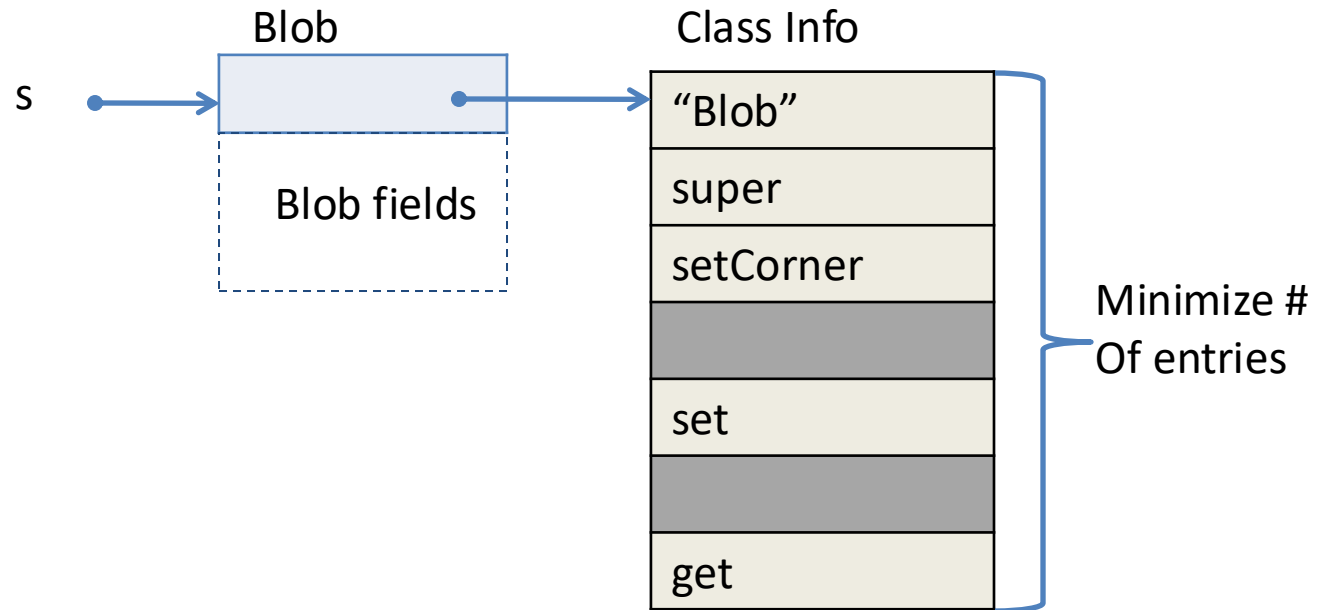
- Can't directly identify methods by position anymore.
- Option 1: Use a level of indirection:
 - Map method identifiers to code pointers (e.g. index by method name)
 - Use a hash table
 - May need to do search up the class hierarchy
- Option 2: Give up separate compilation
 - Use “sparse” dispatch vectors, or binary decision trees
 - Must know then entire class hierarchy
- Option 3: Allow multiple D.V. tables (C++)
 - Choose which D.V. to use based on static type
 - Casting from/to a class may require run-time operations
- Note: many variations on these themes
 - Different Java compilers pick different approaches to options1 and 2...

Option 2 variant 1: Sparse D.V. Tables

- Give up on separate compilation...
- Now we have access to the whole class hierarchy.
- So: ensure that no two methods in the same class are allocated the same D.V. offset.
 - Allow holes in the D.V. just like the hash table solution
 - Unlike hash table, there is never a conflict!
- Compiler needs to construct the method indices
 - Graph coloring techniques can be used to construct the D.V. layouts in a reasonably efficient way (to minimize size)
 - Finding an optimal solution is NP complete!

Example Object Layout

- Advantage: Identical dispatch and performance to single-inheritance case
- Disadvantage: Must know entire class hierarchy




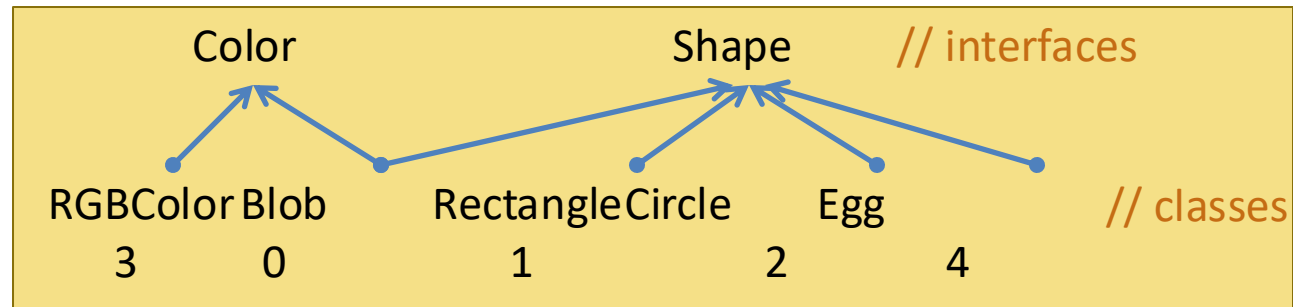
Option 2 variant 2: Binary Search Trees

- Idea: Use conditional branches not indirect jumps
- Each object has a class index (unique per class) as first word
 - Instead of D.V. pointer (no need for one!)
- Method invocation uses range tests to select among n possible classes in $\lg n$ time
 - Direct branches to code at the leaves.

Shape x;
x.SetCorner(...);



```
Mov eax, [x]
Mov ebx, [eax]
Cmp ebx, 1
Jle __L1
Cmp ebx, 2
Je __CircleSetCorner
Jmp __EggSetCorner
__L1:
Cmp ebx, 0
Je __BlobSetCorner
Jmp __RectangleSetCorner
```



Search Tree Tradeoffs


- Binary decision trees work well if the distribution of classes that may appear at a call site is skewed.
 - Branch prediction hardware eliminates the branch stall of ~10 cycles (on X86)
- Can use profiling to find the common paths for each call site individually
 - Put the common case at the top of the decision tree (so less search)
 - 90%/10% rule of thumb: 90% of the invocations at a call site go to the same class
- Drawbacks:
 - Like sparse D.V.'s you need the whole class hierarchy to know how many leaves you need in the search tree.
 - Indirect jumps can have better performance if there are >2 classes (at most one mispredict)

Option 3: Multiple Dispatch Vectors

- Duplicate the D.V. pointers in the object representation.
- Static type of the object determines which D.V. is used.


```
interface Shape {  
    void setCorner(int w, Point p);  
}
```

D.V. Index
0




```
interface Color {  
    float get(int rgb);  
    void set(int rgb, float value);  
}
```

1 0



```
class Blob implements Shape, Color {  
    void setCorner(int w, Point p) {...}  
    float get(int rgb) {...}  
    void set(int rgb, float value) {...}  
}
```




Multiple Dispatch Vectors

- A reference to an object might have multiple “entry points”
 - Each entry point corresponds to a dispatch vector
 - Which one is used depends on the statically known type of the program.

```
Blob b = new Blob();  
Color y = b; // implicit cast!
```

- Compile
Color y = b;
As
Movq $[[b]] + 8, y$





Multiple D.V. Summary

- Benefit: Efficient dispatch, same cost as for multiple inheritance
- Drawbacks:
 - Cast has a runtime cost
 - More complicated programming model... hard to understand/debug?
- What about multiple inheritance and fields?

Multiple Inheritance: Fields

- Multiple supertypes (Java): methods conflict (as we saw)
- Multiple inheritance (C++): fields can also conflict
- Location of the object's fields can no longer be a constant offset from the start of the object.

```
class Color {  
    float r, g, b; /* offsets: 4,8,12 */  
}  
class Shape {  
    Point LL, UR; /* offsets: 4, 8 */  
}  
class ColoredShape extends  
    Color, Shape {  
    int z;  
}
```



ColoredShape ??

C++ approach:

- Add pointers to the superclass fields
 - Need to have multiple dispatch vectors anyway (to deal with methods)
- Extra indirection needed to access superclass fields
- Used even if there is a single superclass
 - Uniformity

