

Lecture 24

CS131: COMPILERS

Announcements

- HW6: Analysis & Optimizations
 - Alias analysis, constant propagation, dead code elimination, register allocation
 - Due: December 30th
- Final Exam:
 - In class, Jan 2nd
 - Coverage: emphasizes material since the midterm
 - Cheat sheet: one, hand-written, double-sided, letter-sized page of notes



GENERAL DATAFLOW ANALYSIS

A Worklist Algorithm

- Use a FIFO queue of nodes that might need to be updated.

for all n , $\text{in}[n] := \emptyset$, $\text{out}[n] := \emptyset$

w = new queue with all nodes

repeat until w is empty

 let $n = w.\text{pop}()$

// pull a node off the queue

$\text{old_in} = \text{in}[n]$

// remember old in[n]

$\text{out}[n] := \bigcup_{n' \in \text{succ}[n]} \text{in}[n']$

$\text{in}[n] := \text{use}[n] \cup (\text{out}[n] - \text{def}[n])$

 if ($\text{old_in} \neq \text{in}[n]$),

// if in[n] has changed

 for all m in $\text{pred}[n]$, $w.\text{push}(m)$

// add to worklist

end

Comparing Dataflow Analyses

Liveness:

Facts: {set of uids live at a program point }

let $gen[n] = use[n]$ and $kill[n] = def[n]$

- $out[n] := \bigcup_{n' \in succ[n]} in[n']$ (backward)
- $in[n] := gen[n] \cup (out[n] - kill[n])$

Reaching Definitions:

Facts: {set of defns. that reach a program point}

let $gen[n] = \{n\}$ and $kill[n] = def[n] \setminus \{n\}$

- $in[n] := \bigcup_{n' \in pred[n]} out[n']$ (forward)
- $out[n] := gen[n] \cup (in[n] - kill[n])$

Available Expressions:

Facts: {set of rhs exps. that reach a program point}

e.g. $gen[n] = \{n\} \setminus kill[n]$ and $kill[n] = use[n]$

- $in[n] := \bigcap_{n' \in pred[n]} out[n']$ (forward)
- $out[n] := gen[n] \cup (in[n] - kill[n])$

Comparing Dataflow Analyses

Liveness:

Facts: {set of uids live at a program point }

let $gen[n] = use[n]$ and $kill[n] = def[n]$

- $out[n] := \bigcup_{n' \in succ[n]} in[n']$ (backward)
- $in[n] := gen[n] \cup (out[n] - kill[n])$

Each analysis solves constraints over some **domain** of facts.

Reaching Definitions:

Facts: {set of defs. that reach a program point}

let $gen[n] = \{n\}$ and $kill[n] = def[n] \setminus \{n\}$

- $in[n] := \bigcup_{n' \in pred[n]} out[n']$ (forward)
- $out[n] := gen[n] \cup (in[n] - kill[n])$

Available Expressions:

Facts: {set of rhs exps. that reach a program point}

e.g. $gen[n] = \{n\} \setminus kill[n]$ and $kill[n] = use[n]$

- $in[n] := \bigcap_{n' \in pred[n]} out[n']$ (forward)
- $out[n] := gen[n] \cup (in[n] - kill[n])$

Comparing Dataflow Analyses

Liveness:

Facts: {set of uids live at a program point }

let $\text{gen}[n] = \text{use}[n]$ and $\text{kill}[n] = \text{def}[n]$

- $\text{out}[n] := \bigcup_{n' \in \text{succ}[n]} \text{in}[n']$ (backward)
- $\text{in}[n] := \text{gen}[n] \cup (\text{out}[n] - \text{kill}[n])$

The "flow function"
(i.e. effect of an instruction
on the facts) can often
be defined by gen and kill.

Reaching Definitions:

Facts: {set of defs. that reach a program point}

let $\text{gen}[n] = \{n\}$ and $\text{kill}[n] = \text{def}[n] \setminus \{n\}$

- $\text{in}[n] := \bigcup_{n' \in \text{pred}[n]} \text{out}[n']$ (forward)
- $\text{out}[n] := \text{gen}[n] \cup (\text{in}[n] - \text{kill}[n])$

Available Expressions:

Facts: {set of rhs exps. that reach a program point}

e.g. $\text{gen}[n] = \{n\} \setminus \text{kill}[n]$ and $\text{kill}[n] = \text{use}[n]$

- $\text{in}[n] := \bigcap_{n' \in \text{pred}[n]} \text{out}[n']$ (forward)
- $\text{out}[n] := \text{gen}[n] \cup (\text{in}[n] - \text{kill}[n])$

Comparing Dataflow Analyses

Liveness:

Facts: {set of uids live at a program point }

let $gen[n] = use[n]$ and $kill[n] = def[n]$

- $out[n] := \bigcup_{n' \in succ[n]} in[n']$ (backward)
- $in[n] := gen[n] \cup (out[n] - kill[n])$

Backward analyses define $out[]$ in terms of $in[]$.

Reaching Definitions:

Facts: {set of defns. that reach a program point}

let $gen[n] = \{n\}$ and $kill[n] = def[n] \setminus \{n\}$

- $in[n] := \bigcup_{n' \in pred[n]} out[n']$ (forward)
- $out[n] := gen[n] \cup (in[n] - kill[n])$

Forward analyses define $in[]$ in terms of $out[]$.

Available Expressions:

Facts: {set of rhs exps. that reach a program point}

e.g. $gen[n] = \{n\} \setminus kill[n]$ and $kill[n] = use[n]$

- $in[n] := \bigcap_{n' \in pred[n]} out[n']$ (forward)
- $out[n] := gen[n] \cup (in[n] - kill[n])$

Comparing Dataflow Analyses

Liveness:

Facts: {set of uids live at a program point }

let $gen[n] = use[n]$ and $kill[n] = def[n]$

- $out[n] := \bigcup_{n' \in succ[n]} in[n']$ (backward)
- $in[n] := gen[n] \cup (out[n] - kill[n])$

Reaching Definitions:

Facts: {set of defs. that reach a program point}

let $gen[n] = \{n\}$ and $kill[n] = def[n] \setminus \{n\}$

- $in[n] := \bigcup_{n' \in pred[n]} out[n']$ (forward)
- $out[n] := gen[n] \cup (in[n] - kill[n])$

Each domain of facts comes equipped with a way of aggregating information.

Available Expressions:

Facts: {set of rhs exps. that reach a program point}

e.g. $gen[n] = \{n\} \setminus kill[n]$ and $kill[n] = use[n]$

- $in[n] := \bigcap_{n' \in pred[n]} out[n']$ (forward)
- $out[n] := gen[n] \cup (in[n] - kill[n])$



Common Features

- All of these analyses have a *domain* over which they solve constraints.
 - Liveness, the domain is sets of variables
 - Reaching defs., Available exprs. the domain is sets of nodes
- Each analysis has a notion of `gen[n]` and `kill[n]`
 - Used to explain how information propagates across a node.
- Each analysis propagates information either *forward* or *backward*
 - Forward: `in[n]` defined in terms of predecessor nodes' `out[]`
 - Backward: `out[n]` defined in terms of successor nodes' `in[]`
- Each analysis has a way of aggregating information
 - Liveness & reaching definitions take union (`U`)
 - Available expressions uses intersection (`∩`)
 - Union expresses a property that holds for *some* path (existential)
 - Intersection expresses a property that holds for *all* paths (universal)

(Forward) Dataflow Analysis Framework

A forward dataflow analysis can be characterized by:

1. A domain of dataflow values \mathcal{L}
 - e.g. \mathcal{L} = the powerset of all variables
 - Think of $\ell \in \mathcal{L}$ as a property, then “ $x \in \ell$ ” means “ x has the property”
2. For each node n , a flow function $F_n : \mathcal{L} \rightarrow \mathcal{L}$
 - So far we’ve seen $F_n(\ell) = \text{gen}[n] \cup (\ell - \text{kill}[n])$
 - So: $\text{out}[n] = F_n(\text{in}[n])$
 - “If ℓ is a property that holds before the node n , then $F_n(\ell)$ holds after n ”
3. A combining operator \sqcap
 - “If we know *either* ℓ_1 or ℓ_2 holds on entry to node n , we know at most $\ell_1 \sqcap \ell_2$ ”
 - $\text{in}[n] := \bigsqcap_{n' \in \text{pred}[n]} \text{out}[n']$



Generic Iterative (Forward) Analysis

for all n , $\text{in}[n] := \top$, $\text{out}[n] := \top$

repeat until no change

for all n

$$\text{in}[n] := \bigcap_{n' \in \text{pred}[n]} \text{out}[n']$$
$$\text{out}[n] := F_n(\text{in}[n])$$

end

end

- Here, $\top \in \mathcal{L}$ (“top”) represents having the “maximum” amount of information.
 - Having “more” information enables more optimizations
 - “Maximum” amount could be inconsistent with the constraints.
 - Iteration refines the answer, eliminating inconsistencies

Structure of \mathcal{L}

- The domain has structure that reflects the “amount” of information contained in each dataflow value.
- Some dataflow values are more informative than others:
 - Write $\ell_1 \sqsubseteq \ell_2$ whenever ℓ_2 provides at least as much information as ℓ_1 .
 - The dataflow value ℓ_2 is “better” for enabling optimizations.

Example 1: for liveness analysis, *smaller* sets of variables are more informative.

- Having smaller sets of variables live across an edge means that there are fewer conflicts for register allocation assignments.
- So: $\ell_1 \sqsubseteq \ell_2$ if and only if $\ell_1 \supseteq \ell_2$

Example 2: for available expressions analysis, larger sets of nodes are more informative.


- Having a larger set of nodes (equivalently, expressions) available means that there is more opportunity for common subexpression elimination.
- So: $\ell_1 \sqsubseteq \ell_2$ if and only if $\ell_1 \subseteq \ell_2$

\mathcal{L} as a Partial Order

- \mathcal{L} is a *partial order* defined by the ordering relation \sqsubseteq .
- A partial order is an ordered set.
- Some of the elements might be *incomparable*.
 - That is, there might be $\ell_1, \ell_2 \in \mathcal{L}$ such that neither $\ell_1 \sqsubseteq \ell_2$ nor $\ell_2 \sqsubseteq \ell_1$
- Properties of a partial order:
 - *Reflexivity*: $\ell \sqsubseteq \ell$
 - *Transitivity*: $\ell_1 \sqsubseteq \ell_2$ and $\ell_2 \sqsubseteq \ell_3$ implies $\ell_1 \sqsubseteq \ell_3$
 - *Anti-symmetry*: $\ell_1 \sqsubseteq \ell_2$ and $\ell_2 \sqsubseteq \ell_1$ implies $\ell_1 = \ell_2$
- Examples:
 - Integers ordered by \leq
 - Types ordered by $<$:
 - Sets ordered by \subseteq or \supseteq

Subsets of $\{a,b,c\}$ ordered by \subseteq

Partial order presented as a Hasse diagram.



order \sqsubseteq is \subseteq

meet \sqcap is \cap

join \sqcup is \cup

Meets and Joins

- The combining operator \sqcap is called the “meet” operation.
- It constructs the *greatest lower bound*:
 - $\ell_1 \sqcap \ell_2 \sqsubseteq \ell_1$ and $\ell_1 \sqcap \ell_2 \sqsubseteq \ell_2$
“the meet is a lower bound”
 - If $\ell \sqsubseteq \ell_1$ and $\ell \sqsubseteq \ell_2$ then $\ell \sqsubseteq \ell_1 \sqcap \ell_2$
“there is no greater lower bound”
- Dually, the \sqcup operator is called the “join” operation.
- It constructs the *least upper bound*:
 - $\ell_1 \sqsubseteq \ell_1 \sqcup \ell_2$ and $\ell_2 \sqsubseteq \ell_1 \sqcup \ell_2$
“the join is an upper bound”
 - If $\ell_1 \sqsubseteq \ell$ and $\ell_2 \sqsubseteq \ell$ then $\ell_1 \sqcup \ell_2 \sqsubseteq \ell$
“there is no smaller upper bound”
- A partial order that has all meets and joins is called a *lattice*.
 - If it has just meets, it’s called a *meet semi-lattice*.

Another Way to Describe the Algorithm

- Algorithm repeatedly computes (for each node n):
- $\text{out}[n] := F_n(\text{in}[n])$
- Equivalently: $\text{out}[n] := F_n(\prod_{n' \in \text{pred}[n]} \text{out}[n'])$
 - By definition of $\text{in}[n]$
- We can write this as a simultaneous update of the vector of $\text{out}[n]$ values:
 - let $x_n = \text{out}[n]$
 - Let $\mathbf{X} = (x_1, x_2, \dots, x_n)$ it's a vector of points in \mathcal{L}
 - $\mathbf{F}(\mathbf{X}) = (F_1(\prod_{j \in \text{pred}[1]} \text{out}[j]), F_2(\prod_{j \in \text{pred}[2]} \text{out}[j]), \dots, F_n(\prod_{j \in \text{pred}[n]} \text{out}[j]))$
- Any solution to the constraints is a *fixpoint* \mathbf{X} of \mathbf{F}
 - i.e. $\mathbf{F}(\mathbf{X}) = \mathbf{X}$

Iteration Computes Fixpoints

- Let $\mathbf{X}_0 = (\top, \top, \dots, \top)$
- Each loop through the algorithm apply F to the old vector:
 $\mathbf{X}_1 = F(\mathbf{X}_0)$
 $\mathbf{X}_2 = F(\mathbf{X}_1)$
...
- $F^{k+1}(\mathbf{X}) = F(F^k(\mathbf{X}))$
- A fixpoint is reached when $F^k(\mathbf{X}) = F^{k+1}(\mathbf{X})$
 - That's when the algorithm stops.
- Wanted: a maximal fixpoint
 - Because that one is more informative/useful for performing optimizations


Monotonicity & Termination

- Each flow function F_n maps lattice elements to lattice elements; to be sensible it should be *monotonic*:
- $F : \mathcal{L} \rightarrow \mathcal{L}$ is *monotonic* iff:
 $\ell_1 \sqsubseteq \ell_2$ implies that $F(\ell_1) \sqsubseteq F(\ell_2)$
 - Intuitively: “If you have more information entering a node, then you have more information leaving the node.”
- Monotonicity lifts point-wise to the function: $\mathbf{F} : \mathcal{L}^n \rightarrow \mathcal{L}^n$
 - vector $(x_1, x_2, \dots, x_n) \sqsubseteq (y_1, y_2, \dots, y_n)$ iff $x_i \sqsubseteq y_i$ for each i
- Note that \mathbf{F} is consistent: $\mathbf{F}(\mathbf{X}_0) \sqsubseteq \mathbf{X}_0$
 - So each iteration moves at least one step down the lattice (for some component of the vector)
 - $\dots \sqsubseteq \mathbf{F}(\mathbf{F}(\mathbf{X}_0)) \sqsubseteq \mathbf{F}(\mathbf{X}_0) \sqsubseteq \mathbf{X}_0$
- Therefore, # steps needed to reach a fixpoint is at most the height H of \mathcal{L} times the number of nodes: $O(Hn)$

Building Lattices?

- Information about individual nodes or variables can be lifted *pointwise*:
 - If \mathcal{L} is a lattice, then so is $\{f : X \rightarrow \mathcal{L}\}$ where $f \sqsubseteq g$ if and only if $f(x) \sqsubseteq g(x)$ for all $x \in X$.
- Like *types*, the dataflow lattices are *static approximations* to the dynamic behavior:
 - Could pick a lattice based on subtyping:

- Or other information:




- Points in the lattice are sometimes called dataflow “facts”

“Classic” Constant Propagation

- Constant propagation can be formulated as a dataflow analysis.


- Idea: propagate and fold integer constants in one pass:

x = 1;		x = 1;
y = 5 + x;		6;
z = y * y;		z = 36;

- Information about a single variable:
 - Variable is never defined.
 - Variable has a single, constant value.
 - Variable is assigned multiple values.

Domains for Constant Propagation

- We can make a constant propagation lattice \mathcal{L} for *one variable* like this:



- To accommodate multiple variables, we take the product lattice, with one element per variable.
 - Assuming there are three variables, x , y , and z , the elements of the product lattice are of the form (ℓ_x, ℓ_y, ℓ_z) .
 - Alternatively, think of the product domain as a context that maps variable names to their “*abstract interpretations*”
- What are “meet” and “join” in this product lattice?
- What is the height of the product lattice?

Flow Functions

- Consider the node $x = y \text{ op } z$
- $F(\ell_x, \ell_y, \ell_z) = ?$
- | | | |
|---|---|---|
| <ul style="list-style-type: none">• $F(\ell_x, \top, \ell_z) = (\top, \top, \ell_z)$• $F(\ell_x, \ell_y, \top) = (\top, \ell_y, \top)$ | } | “If either input might have multiple values the result of the operation might too.” |
|---|---|---|
- | | | |
|---|---|--|
| <ul style="list-style-type: none">• $F(\ell_x, \perp, \ell_z) = (\perp, \perp, \ell_z)$• $F(\ell_x, \ell_y, \perp) = (\perp, \ell_y, \perp)$ | } | “If either input is undefined the result of the operation is too.” |
|---|---|--|
- | | | |
|--|---|---|
| <ul style="list-style-type: none">• $F(\ell_x, i, j) = (i \text{ op } j, i, j)$ | } | “If the inputs are known constants, calculate the output statically.” |
|--|---|---|
- Flow functions for the other nodes are easy...
- Monotonic?
- Distributes over meets?




QUALITY OF DATAFLOW ANALYSIS SOLUTIONS

Best Possible Solution

- Suppose we have a control-flow graph.
- If there is a path p_1 starting from the root node (entry point of the function) traversing the nodes $n_0, n_1, n_2, \dots, n_k$
- The best possible information along the path p_1 is:
 $\ell_{p_1} = F_{n_k}(\dots F_{n_2}(F_{n_1}(F_{n_0}(T)))) \dots$
- Best solution at the output is some $\ell \sqsubseteq \ell_p$ for *all* paths p .
- Meet-over-paths (MOP) solution:

$$\bigcap_{p \in \text{paths_to}[n]} \ell_p$$



Best answer here is:

$$F_5(F_3(F_2(F_1(T)))) \sqcap F_5(F_4(F_2(F_1(T))))$$

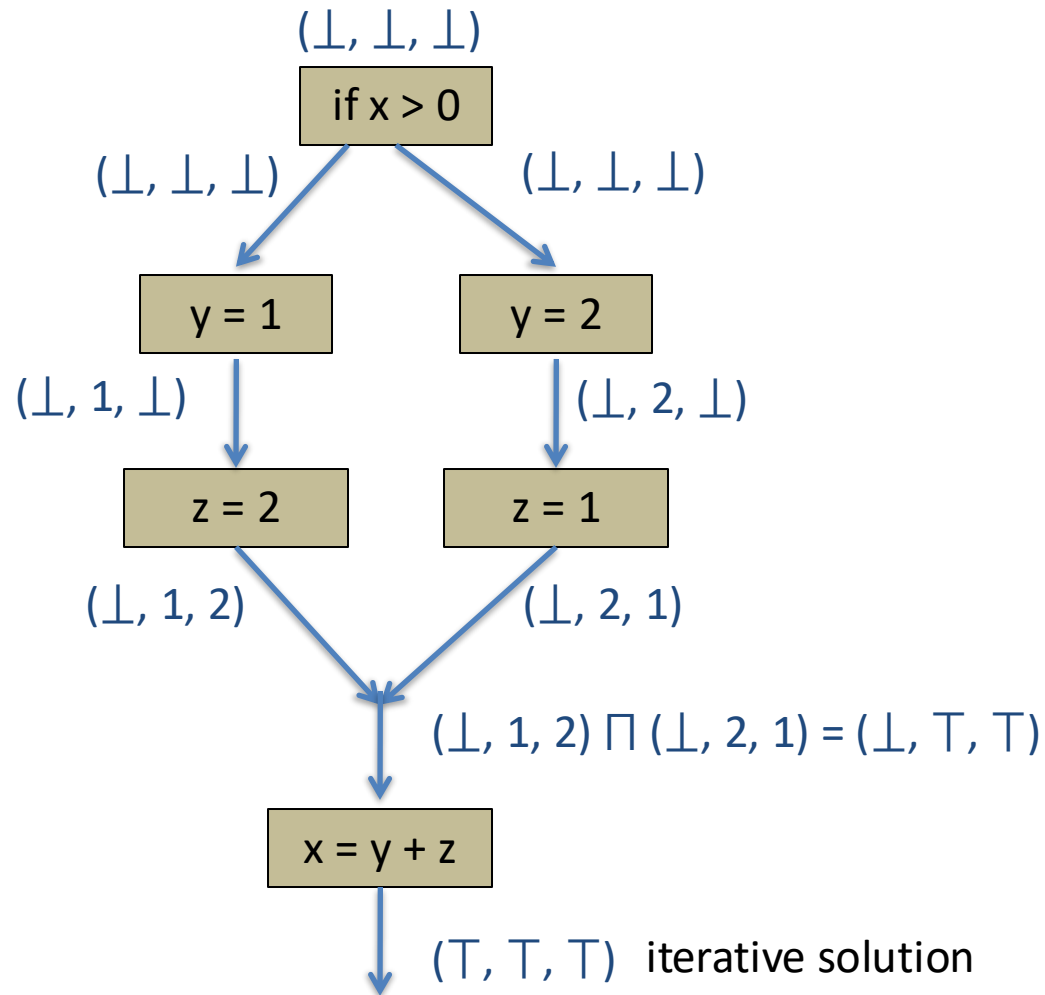
What about quality of iterative solution?

- Does the iterative solution: $\text{out}[n] = F_n(\prod_{n' \in \text{pred}[n]} \text{out}[n'])$ compute the MOP solution?
- MOP Solution: $\prod_{p \in \text{paths_to}[n]} \ell_p$
- Answer: Yes, *if* the flow functions *distribute* over \prod
 - Distributive means: $\prod_i F_n(\ell_i) = F_n(\prod_i \ell_i)$
 - Proof is a bit tricky & beyond the scope of this class. (Difficulty: loops in the control flow graph might mean there are *infinitely* many paths...)
- Not all analyses give MOP solution
 - They are more conservative.


Reaching Definitions is MOP

- $F_n[x] = \text{gen}[n] \cup (x - \text{kill}[n])$
- Does F_n distribute over meet $\sqcap = \cup$?
- $F_n[x \sqcap y]$
 - $= \text{gen}[n] \cup ((x \cup y) - \text{kill}[n])$
 - $= \text{gen}[n] \cup ((x - \text{kill}[n]) \cup (y - \text{kill}[n]))$
 - $= (\text{gen}[n] \cup (x - \text{kill}[n])) \cup (\text{gen}[n] \cup (y - \text{kill}[n]))$
 - $= F_n[x] \cup F_n[y]$
 - $= F_n[x] \sqcap F_n[y]$
- Therefore: Reaching Definitions with iterative analysis always terminates with the MOP (i.e. best) solution.

Constprop Iterative Solution




MOP Solution \neq Iterative Solution



Why not compute MOP Solution?

- If MOP is better than the iterative analysis, why not compute it instead?
 - ANS: exponentially many paths (even in graph without loops)
- $O(n)$ nodes
- $O(n)$ edges
- $O(2^n)$ paths*
 - At each branch there is a choice of 2 directions



* Incidentally, a similar idea can be used to force ML / Haskell type inference to need to construct a type that is exponentially big in the size of the program!

Dataflow Analysis: Summary

- Many dataflow analyses fit into a common framework.
- Key idea: *Iterative solution* of a system of equations over a *lattice* of constraints.
 - Iteration terminates if flow functions are monotonic.
 - Solution is equivalent to meet-over-paths answer if the flow functions distribute over meet (\sqcap).
- Dataflow analyses as presented work for an “imperative” intermediate representation.
 - The values of temporary variables are updated (“mutated”) during evaluation.
 - Such mutation complicates calculations
 - SSA = “Single Static Assignment” eliminates this problem, by introducing more temporaries – each one assigned to only once.
 - Next up: Converting to SSA, finding loops and dominators in CFGs



REGISTER ALLOCATION

Moving Towards Register Allocation

- The OAT compiler currently generates as *many* temporary variables as it needs
 - These are the %uids you should be very familiar with by now.
- Current compilation strategy:
 - Each %uid maps to a stack location.
 - This yields programs with many loads/stores to memory.
 - Very inefficient.
- Ideally, we'd like to map as many %uid's as possible into registers.
 - Eliminate the use of the alloca instruction?
 - Only 16 max registers available on 64-bit X86
 - %rsp and %rbp are reserved and some have special semantics, so really only 10 or 12 available
 - This means that a register must hold more than one slot
- When is this safe?

Register Allocation Problem

- Given: an IR program that uses an unbounded number of temporaries
 - e.g. the uids of our LLVM programs
- Find: a mapping from temporaries to machine registers such that
 - program semantics is preserved (i.e. the behavior is the same)
 - register usage is maximized
 - moves between registers are minimized
 - calling conventions / architecture requirements are obeyed
- Stack Spilling
 - If there are k registers available and $m > k$ temporaries are live at the same time, then not all of them will fit into registers.
 - So: "spill" the excess temporaries to the stack.

Linear-Scan Register Allocation

Simple, greedy register-allocation strategy:

1. Compute liveness information: $\text{live}(x)$
 - recall: $\text{live}(x)$ is the set of uids that are live on entry to x 's definition
2. Let pal be the set of usable registers
 - usually reserve a couple for spill code [our implementation uses rax, rcx]
3. Maintain "layout" uid_loc that maps uids to locations
 - locations include registers and stack slots n , starting at $n=0$
4. Scan through the program. For each instruction that defines a uid x
 - $\text{used} = \{r \mid \text{reg } r = \text{uid_loc}(y) \text{ s.t. } y \in \text{live}(x)\}$
 - $\text{available} = \text{pal} - \text{used}$
 - If available is empty: *// no registers available, spill*
 $\text{uid_loc}(x) := \text{slot } n \quad ; \quad n = n + 1$
 - Otherwise, pick r in available : *// choose an available register*
 $\text{uid_loc}(x) := \text{reg } r$

For HW6

- HW 6 implements two naive register allocation strategies:
 - none: spill all registers
 - greedy: uses linear scan
- Also offers choice of liveness
 - trivial: assume all variables are live everywhere
 - dataflow: use the dataflow algorithms
- Your job: do "better" than these.
 - To beat "greedy" on small programs – it is necessary to take into account the calling conventions
- Quality Metric - *lower* score is better:
 - total number of memory accesses
(Ind2 and Ind3 operands, Push/Pop)
 - ties broken by total number of instructions
- Linear scan is OK
 - but... how can we do better?