

Lecture 26

CS 131: COMPILERS

Announcements

- HW6: Analysis & Optimizations
 - Alias analysis, constant propagation, dead code elimination, register allocation
 - Due: December 30th
- Final Exam:
 - In class, Jan 2nd
 - Coverage: emphasizes material since the midterm
 - Cheat sheet: one, hand-written, double-sided, letter-sized page of notes



Phi nodes

Alloc “promotion”

Register allocation

REVISITING SSA

Single Static Assignment (SSA)

- LLVM IR names (via %uids) *all* intermediate values computed by the program.
- It makes the order of evaluation explicit.
- Each %uid is assigned to only once
 - Contrast with the mutable quadruple form
 - Note that dataflow analyses had these kill[n] sets because of updates to variables...
- Naïve implementation of backend: map %uids to stack slots
- Better implementation: map %uids to registers (as much as possible)
- Question: How do we convert a source program to make maximal use of %uids, rather than alloca-created storage?
 - two problems: control flow & location in memory
- Then: How do we convert SSA code to x86, mapping %uids to registers?
 - Register allocation.

Alloca vs. %UID

- Current compilation strategy:

```
int x = 3;  
int y = 0;  
x = x + 1;  
y = x + 2;
```



```
%x = alloca i64  
%y = alloca i64  
store i64* %x, 3  
store i64* %y, 0  
%x1 = load %i64* %x  
%tmp1 = add i64 %x1, 1  
store i64* %x, %tmp1  
%x2 = load %i64* %x  
%tmp2 = add i64 %x2, 2  
store i64* %y, %tmp2
```

- Directly map source variables into %uids?

```
int x = 3;  
int y = 0;  
x = x + 1;  
y = x + 2;
```



```
int x1 = 3;  
int y1 = 0;  
x2 = x1 + 1;  
y2 = x2 + 2;
```



```
%x1 = add i64 3, 0  
%y1 = add i64 0, 0  
%x2 = add i64 %x1, 1  
%y2 = add i64 %x2, 2
```

- Does this always work?

What about If-then-else?

- How do we translate this into SSA?

```
int y = ...  
int x = ...  
int z = ...  
if (p) {  
    x = y + 1;  
} else {  
    x = y * 2;  
}  
z = x + 3;
```



```
entry:  
    %y1 = ...  
    %x1 = ...  
    %z1 = ...  
    %p = icmp ...  
    br i1 %p, label %then, label %else  
then:  
    %x2 = add i64 %y1, 1  
    br label %merge  
else:  
    %x3 = mult i64 %y1, 2  
merge:  
    %z2 = %add i64 ???, 3
```

- What do we put for ???

Phi Functions

- Solution: ϕ functions
 - Fictitious operator, used only for analysis
 - implemented by Mov at x86 level
 - Chooses among different versions of a variable based on the path by which control enters the phi node.
 $\%uid = \text{phi } \langle \text{ty} \rangle \ v_1, \langle \text{label}_1 \rangle, \dots, v_n, \langle \text{label}_n \rangle$

```
int y = ...  
int x = ...  
int z = ...  
if (p) {  
    x = y + 1;  
} else {  
    x = y * 2;  
}  
z = x + 3;
```



```
entry:  
    %y1 = ...  
    %x1 = ...  
    %z1 = ...  
    %p = icmp ...  
    br i1 %p, label %then, label %else  
then:  
    %x2 = add i64 %y1, 1  
    br label %merge  
else:  
    %x3 = mult i64 %y1, 2  
merge:  
    %x4 = phi i64 %x2, %then, %x3, %else  
    %z2 = %add i64 %x4, 3
```

Phi Nodes and Loops

- Importantly, the %uids on the right-hand side of a phi node can be defined “later” in the control-flow graph.
 - Means that %uids can hold values “around a loop”
 - Scope of %uids is defined by *dominance*

```
entry:
  %y1 = ...
  %x1 = ...
  br label %body

body:
  %x2 = phi i64 %x1, %entry, %x3, %body
  %x3 = add i64 %x2, %y1
  %p = icmp slt i64, %x3, 10
  br i1 %p, label %body, label %after

after:
  ...
```


Alloca Promotion

- Not all source variables can be allocated to registers
 - If the address of the variable is taken (as permitted in C, for example)
 - If the address of the variable “escapes” (by being passed to a function)
- An alloca instruction is called promotable if neither of the two conditions above holds

entry:

```
%x = alloca i64      // %x cannot be promoted
%y = call malloc(i64 8)
%ptr = bitcast i8* %y to i64**
store i65** %ptr, %x  // store the pointer into the heap
```

entry:

```
%x = alloca i64      // %x cannot be promoted
%y = call foo(i64* %x) // foo may store the pointer into the heap
```

- Happily, most local variables declared in source programs are promotable
 - That means they can be register allocated

Converting to SSA: Overview


- Start with the ordinary control flow graph that uses allocas
 - Identify “promotable” allocas
- Compute dominator tree information
- Calculate def/use information for each such allocated variable
- Insert ϕ functions for each variable at necessary “join points”
- Replace loads/stores to alloc’ed variables with freshly-generated %uids
- Eliminate the now unneeded load/store/alloca instructions.

Where to Place ϕ functions?

- Need to calculate the “Dominance Frontier”
- Node A *strictly dominates* node B if A dominates B and $A \neq B$.
 - Note: A does not strictly dominate B if A does not dominate B or $A = B$.
- The *dominance frontier* of a node B is the set of all CFG nodes y such that B dominates a predecessor of y but does not strictly dominate y
 - Intuitively: starting at B, there is a path to y, but there is another route to y that does not go through B
- Write $DF[n]$ for the dominance frontier of node n.

Dominance Frontiers

- Example of a dominance frontier calculation results
- $DF[1] = \{1\}$, $DF[2] = \{1,2\}$, $DF[3] = \{2\}$, $DF[4] = \{1\}$, $DF[5] = \{8,0\}$, $DF[6] = \{8\}$, $DF[7] = \{7,0\}$, $DF[8] = \{0\}$, $DF[9] = \{7,0\}$, $DF[0] = \{\}$



Algorithm For Computing DF[n]

- Assume that `doms[n]` stores the dominator tree (so that `doms[n]` is the *immediate dominator* of `n` in the tree)
- Adds each `B` to the DF sets to which it belongs

for all nodes `B`

```
    if #(pred[B]) ≥ 2                                // (just an optimization)
        for each p ∈ pred[B] {
            runner := p                                // start at the predecessor of B
            while (runner ≠ doms[B])                    // walk up the tree adding B
                DF[runner] := DF[runner] ∪ {B}
                runner := doms[runner]
        }
```

Insert ϕ at Join Points

- Lift the $DF[n]$ to a set of nodes N in the obvious way:

$$DF[N] = \bigcup_{n \in N} DF[n]$$

- Suppose that at variable x is defined at a set of nodes N .

$$DF_0[N] = DF[N]$$

$$DF_{i+1}[N] = DF[DF_i[N] \cup N]$$

Let $J[N]$ be the *least fixed point* of the sequence:

$$DF_0[N] \subseteq DF_1[N] \subseteq DF_2[N] \subseteq DF_3[N] \subseteq \dots$$

That is, $J[N] = DF_k[N]$ for some k such that $DF_k[N] = DF_{k+1}[N]$

- $J[N]$ is called the “join points” for the set N
- We insert ϕ functions for the variable x at each node in $J[N]$.
 - $x = \phi(x, x, \dots, x);$ (one “ x ” argument for each predecessor of the node)
 - In practice, $J[N]$ is never directly computed, instead you use a worklist algorithm that keeps adding nodes for $DF_k[N]$ until there are no changes, just as in the dataflow solver.
- Intuition:
 - If N is the set of places where x is modified, then $DF[N]$ is the places where ϕ nodes need to be added, but those also “count” as modifications of x , so we need to insert the ϕ nodes to capture those modifications too...

Example Join-point Calculation

- Suppose the variable x is modified at nodes 3 and 6
 - Where would we need to add phi nodes?
- $DF_0[\{3,6\}] = DF[\{3,6\}] = DF[3] \cup DF[6] = \{2,8\}$
- $DF_1[\{3,6\}]$
 - $= DF[DF_0\{3,6\} \cup \{3,6\}]$
 - $= DF[\{2,3,6,8\}]$
 - $= DF[2] \cup DF[3] \cup DF[6] \cup DF[8]$
 - $= \{1,2\} \cup \{2\} \cup \{8\} \cup \{0\} = \{1,2,8,0\}$
- $DF_2[\{3,6\}]$
 - $= \dots$
 - $= \{1,2,8,0\}$
- So $J[\{3,6\}] = \{1,2,8,0\}$ and we need to add phi nodes at those four spots.

Phi Placement (Alternative)

- Less efficient than LLVM's true "dominance frontier" algorithm, but easier to understand:
- Place phi nodes "maximally" (i.e. at every node with > 2 predecessors)
- If all values flowing into phi node are the same, then eliminate it:

```
%x = phi  t %y, %pred1  t %y %pred2  ... t %y %predK
// code that uses %x
⇒
// code with %x replaced by %y
```

- Interleave with other optimizations
 - copy propagation
 - constant propagation
 - etc.

Legend of "simple" optimizations*:

LAS = load after store

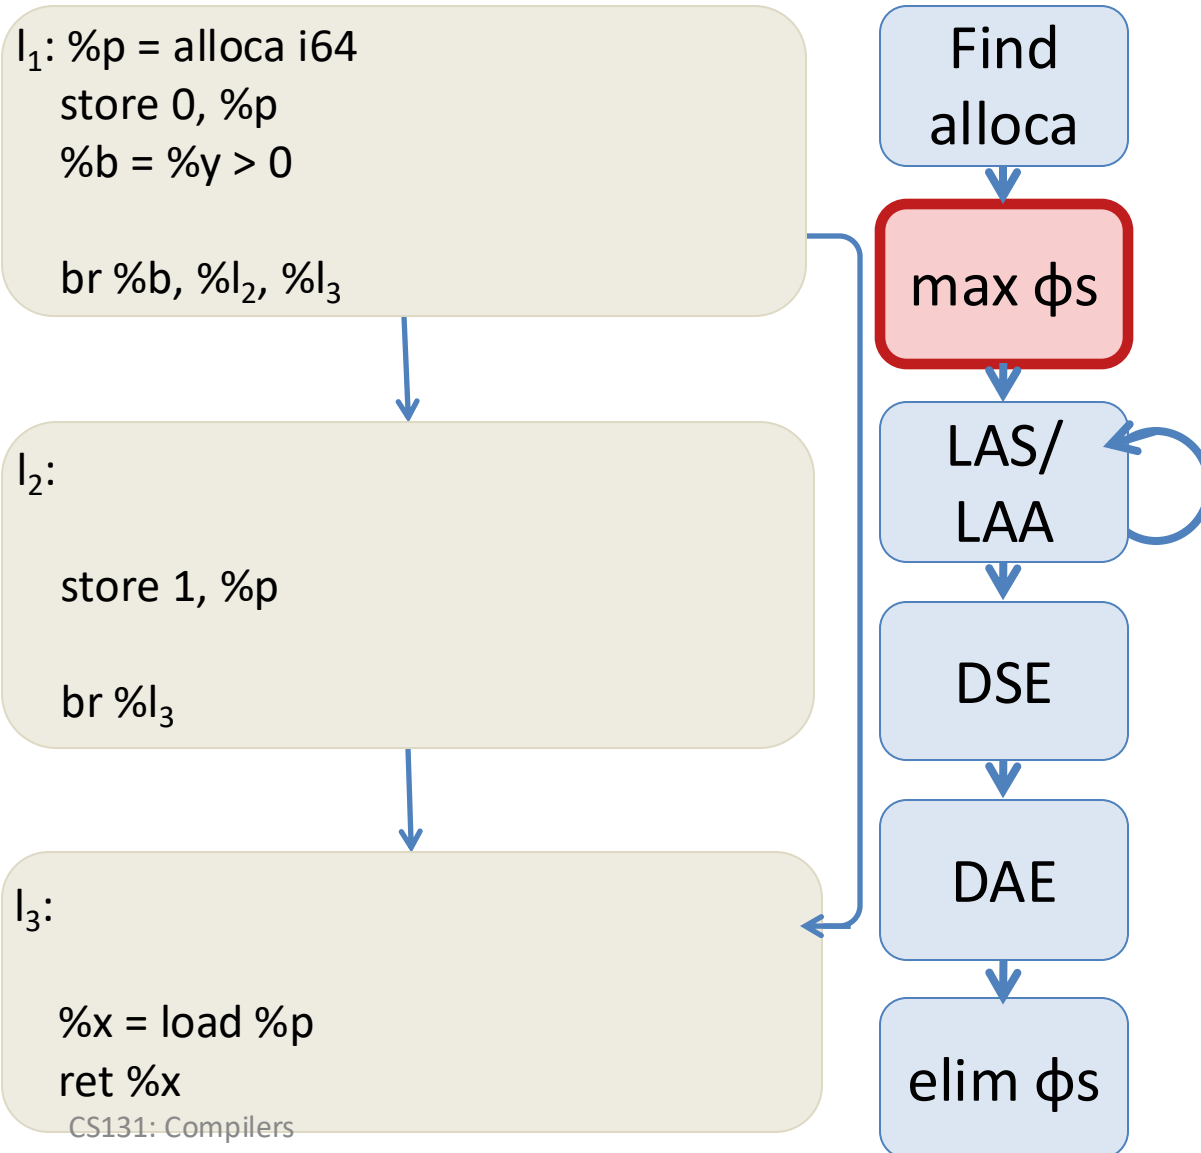
LAA = load after alloca

DSE = dead store elimination

DAE = dead alloca elimination

*nomenclature taken from LLVM IR passes

Example SSA Optimizations



- How to place phi nodes without breaking SSA?
- Note: the “real” implementation combines many of these steps into one pass.
 - Placesphis directly at the dominance frontier
- This example also illustrates other common optimizations:
 - Load after store/alloca
 - Dead store/alloca elimination

Example SSA Optimizations

l_1 : %p = alloca i64
store 0, %p
%b = %y > 0
%x₁ = load %p
br %b, %l₂, %l₃

l_2 :
store 1, %p
%x₂ = load %p
br %l₃

l_3 :
%x = load %p
ret %x

Find
alloca

max ϕ s

LAS/
LAA

DSE

DAE

elim ϕ s

- How to place phi nodes without breaking SSA?
- Insert
 - Loads at the end of each block

Example SSA Optimizations

l_1 : %p = alloca i64
store 0, %p
%b = %y > 0
%x₁ = load %p
br %b, %l₂, %l₃

l_2 : %x₃ = ϕ [%x₁, %l₁]

store 1, %p
%x₂ = load %p
br %l₃

l_3 : %x₄ = ϕ [%x₁; %l₁, %x₂: %l₂]

%x = load %p
ret %x

Find
alloca

max ϕ s

LAS/
LAA

DSE

DAE

elim ϕ s

- How to place phi nodes without breaking SSA?
- Insert
 - Loads at the end of each block
 - Insert ϕ -nodes at each block

Example SSA Optimizations

l_1 : %p = alloca i64
store 0, %p
%b = %y > 0
%x₁ = load %p
br %b, %l₂, %l₃

l_2 : %x₃ = ϕ [%x₁, %l₁]
store %x₃, %p
store 1, %p
%x₂ = load %p
br %l₃

l_3 : %x₄ = ϕ [%x₁; %l₁, %x₂: %l₂]
store %x₄, %p
%x = load %p
ret %x

Find
alloca

max ϕ s

LAS/
LAA

DSE

DAE

elim ϕ s

- How to place phi nodes without breaking SSA?
- Insert
 - Loads at the end of each block
 - Insert ϕ -nodes at each block
 - Insert stores after ϕ -nodes

Example SSA Optimizations

l_1 : %p = alloca i64
store 0, %p
%b = %y > 0
%x₁ = load %p
br %b, %l₂, %l₃

l_2 : %x₃ = ϕ [%x₁, %l₁]
store %x₃, %p
store 1, %p
%x₂ = load %p
br %l₃

l_3 : %x₄ = ϕ [%x₁; %l₁, %x₂: %l₂]
store %x₄, %p
%x = load %p
ret %x

Find
alloca

max ϕ s

LAS/
LAA

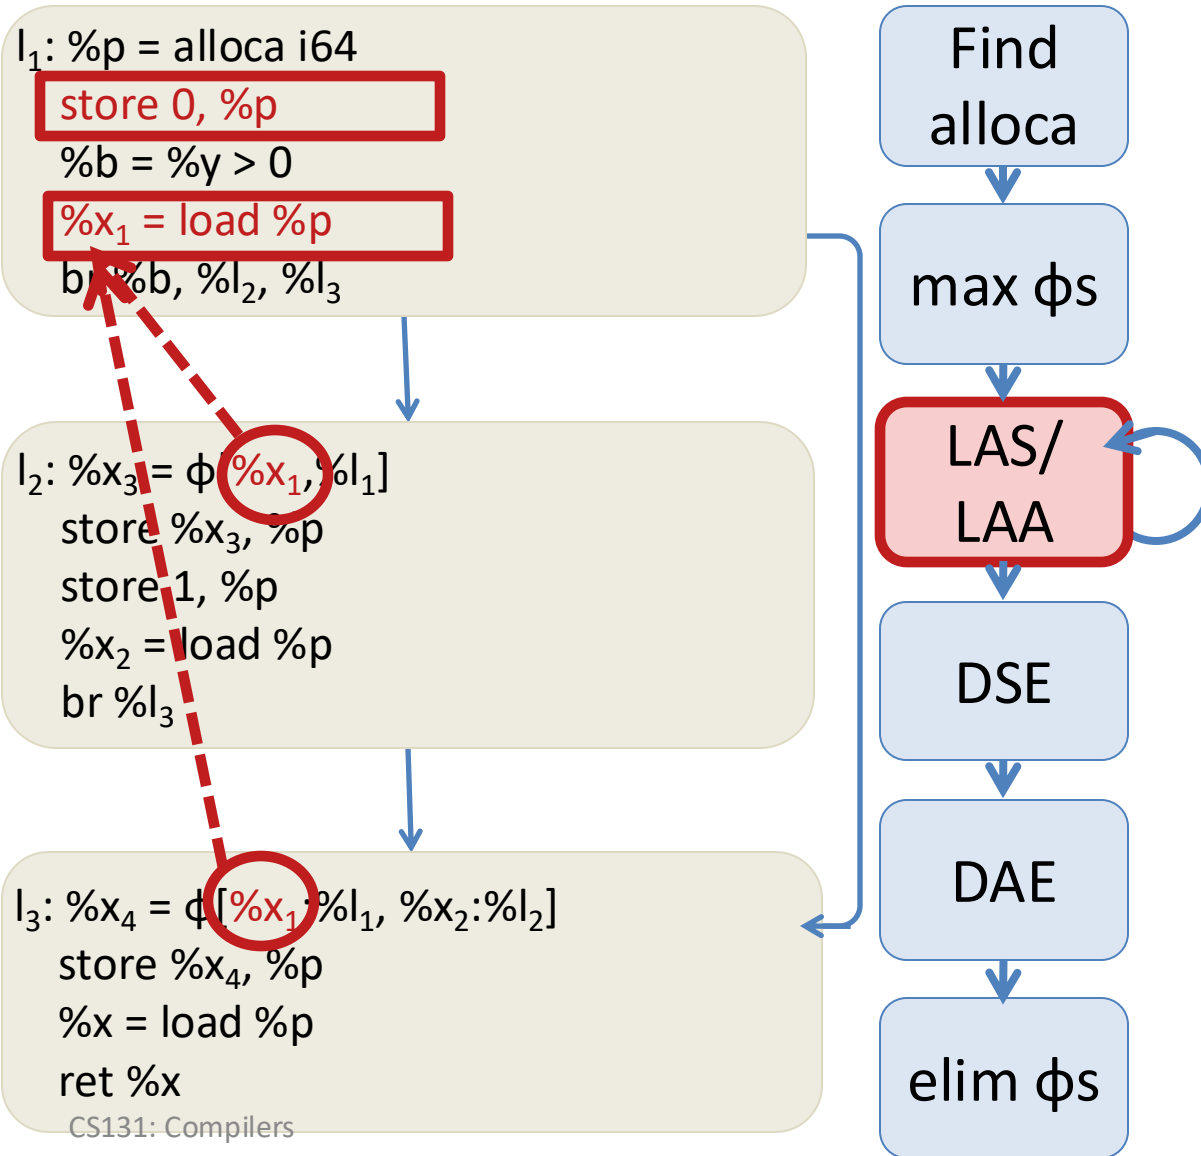
DSE

DAE

elim ϕ s

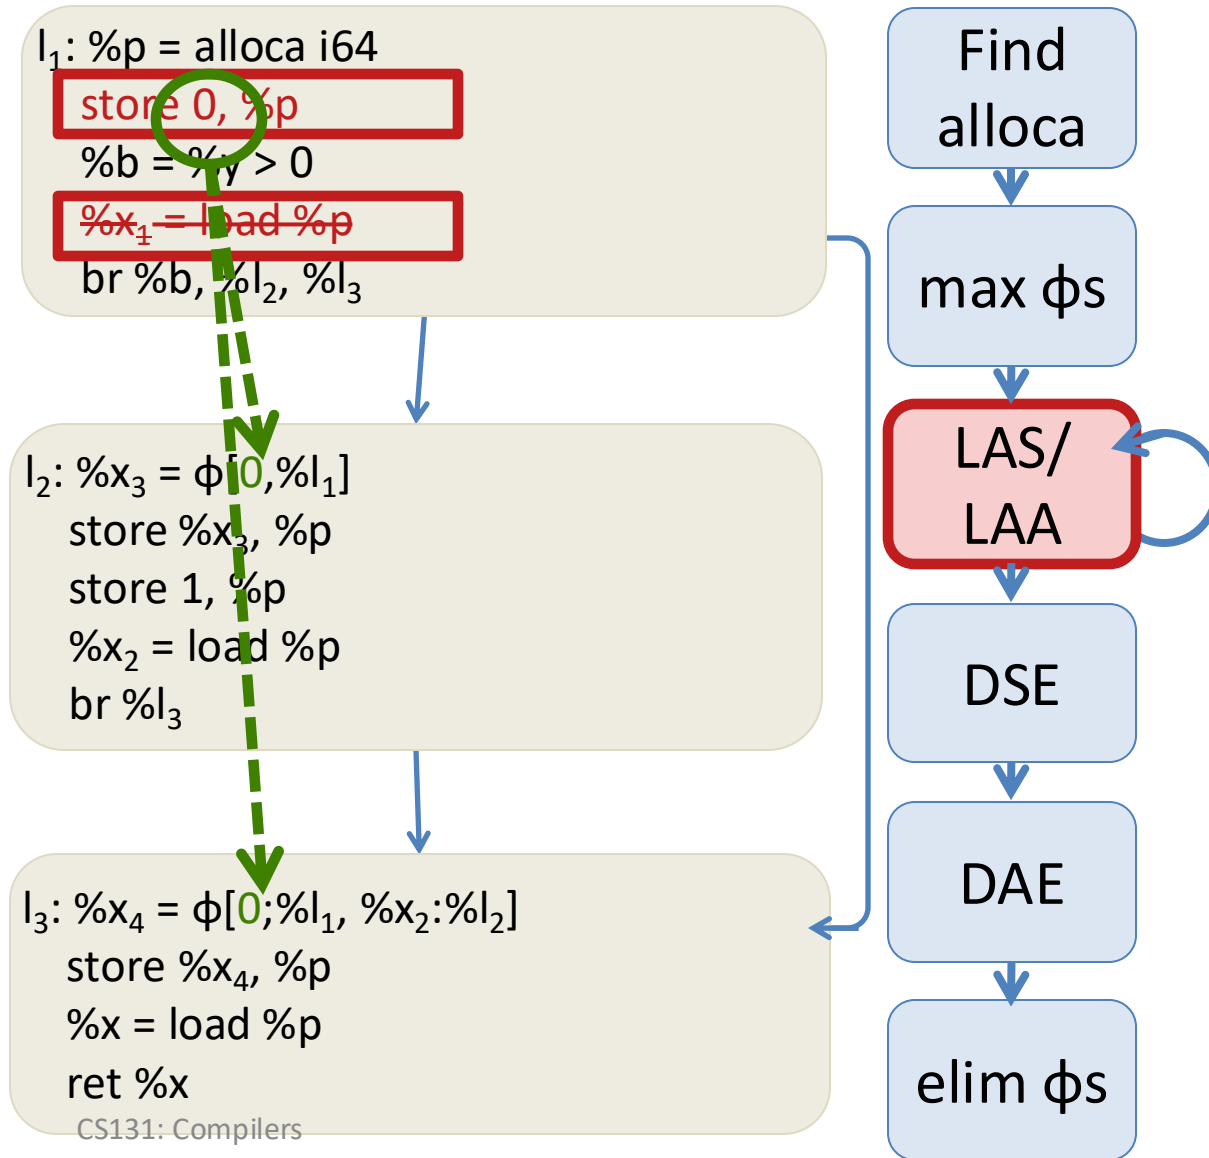
- For loads after stores (LAS):
 - Substitute all uses of the load by the value being stored
 - Remove the load

Example SSA Optimizations



- For loads after stores (LAS):
 - Substitute all uses of the load by the value being stored
 - Remove the load

Example SSA Optimizations



- For loads after stores (LAS):
 - Substitute all uses of the load by the value being stored
 - Remove the load

Example SSA Optimizations

l_1 : %p = alloca i64
store 0, %p
%b = %y > 0

br %b, %l₂, %l₃

l_2 : %x₃ = ϕ [0,%l₁]
store %x₃, %p
store 1, %p
%x₂ = load %p
br %l₃

l_3 : %x₄ = ϕ [0;%l₁, %x₂; %l₂]
store %x₄, %p
%x = load %p
ret %x

Find
alloca

max ϕ s

LAS/
LAA

DSE

DAE

elim ϕ s

- For loads after stores (LAS):
 - Substitute all uses of the load by the value being stored
 - Remove the load

Example SSA Optimizations

l_1 : %p = alloca i64
store 0, %p
%b = %y > 0

br %b, %l₂, %l₃

l_2 : %x₃ = ϕ [0;%l₁]
store %x₃, %p
store 1, %p
~~%x₂ = load %p~~
br %l₃

l_3 : %x₄ = ϕ [0;%l₁, 1;%l₂]
store %x₄, %p
%x = load %p
ret %x

Find
alloca

max ϕ s

**LAS/
LAA**

DSE

DAE

elim ϕ s

- For loads after stores (LAS):
 - Substitute all uses of the load by the value being stored
 - Remove the load

Example SSA Optimizations

l_1 : %p = alloca i64
store 0, %p
%b = %y > 0

br %b, %l₂, %l₃

l_2 : %x₃ = $\phi[0, \%l_1]$
store %x₃, %p
store 1, %p

br %l₃

l_3 : %x₄ = $\phi[0; \%l_1, 1: \%l_2]$
store %x₄, %p
%x = load %p
ret %x

Find
alloca

max ϕ s

LAS/
LAA

DSE

DAE

elim ϕ s

- For loads after stores (LAS):
 - Substitute all uses of the load by the value being stored
 - Remove the load

Example SSA Optimizations

l_1 : %p = alloca i64
store 0, %p
%b = %y > 0

br %b, %l₂, %l₃

l_2 : %x₃ = ϕ [0,%l₁]
store %x₃, %p
store 1, %p

br %l₃

l_3 : %x₄ = ϕ [0;%l₁, 1;%l₂]

store %x₄, %p

~~%x = load %p~~

ret %x₄

Find
alloca

max ϕ s

LAS/
LAA

DSE

DAE

elim ϕ s

- For loads after stores (LAS):
 - Substitute all uses of the load by the value being stored
 - Remove the load

Example SSA Optimizations

l_1 : %p = alloca i64
store 0, %p
%b = %y > 0

br %b, %l₂, %l₃

l_2 : %x₃ = $\phi[0, \%l_1]$
store %x₃, %p
store 1, %p

br %l₃

l_3 : %x₄ = $\phi[0; \%l_1, 1: \%l_2]$
store %x₄, %p

ret %x₄

Find
alloca

max ϕ s

LAS/
LAA

DSE

DAE

elim ϕ s

- Dead Store Elimination (DSE)
 - Eliminate all stores with no subsequent loads.

- Dead Alloca Elimination (DAE)
 - Eliminate all allocas with no subsequent loads/stores.

Example SSA Optimizations

l_1 : ~~%p = alloca i64~~
~~store 0, %p~~
%b = %y > 0

br %b, %l₂, %l₃

l_2 : %x₃ = $\phi[0, \%l_1]$
~~store %x₃, %p~~
~~store 1, %p~~

br %l₃

l_3 : %x₄ = $\phi[0; \%l_1, 1: \%l_2]$
~~store %x₄, %p~~

ret %x₄

Find
alloca

max ϕ s

LAS/
LAA

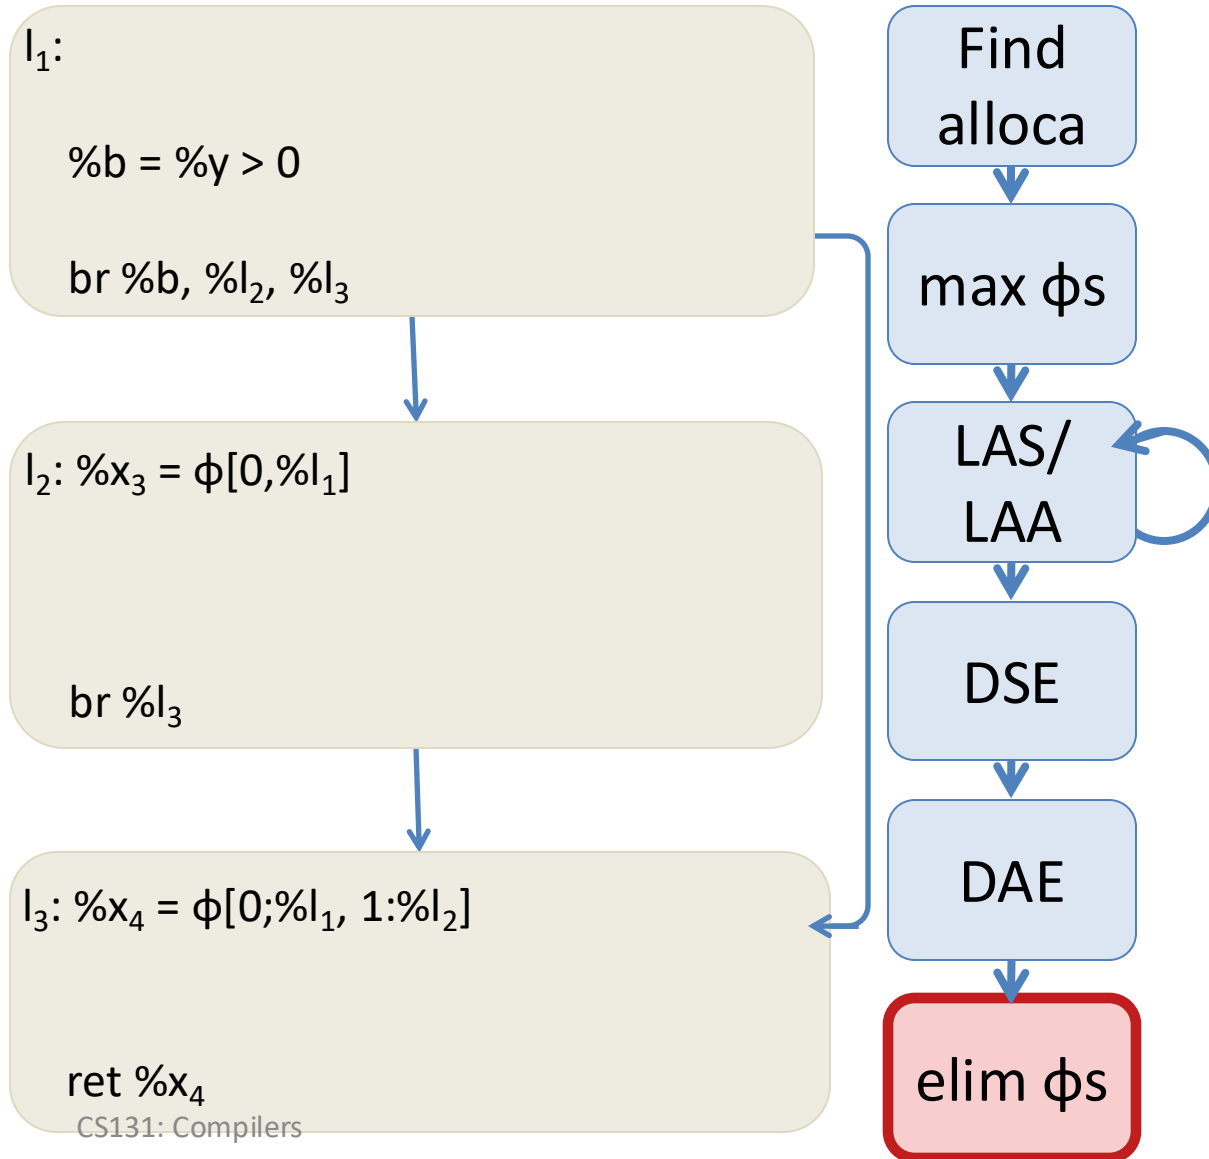
DSE

DAE

elim ϕ s

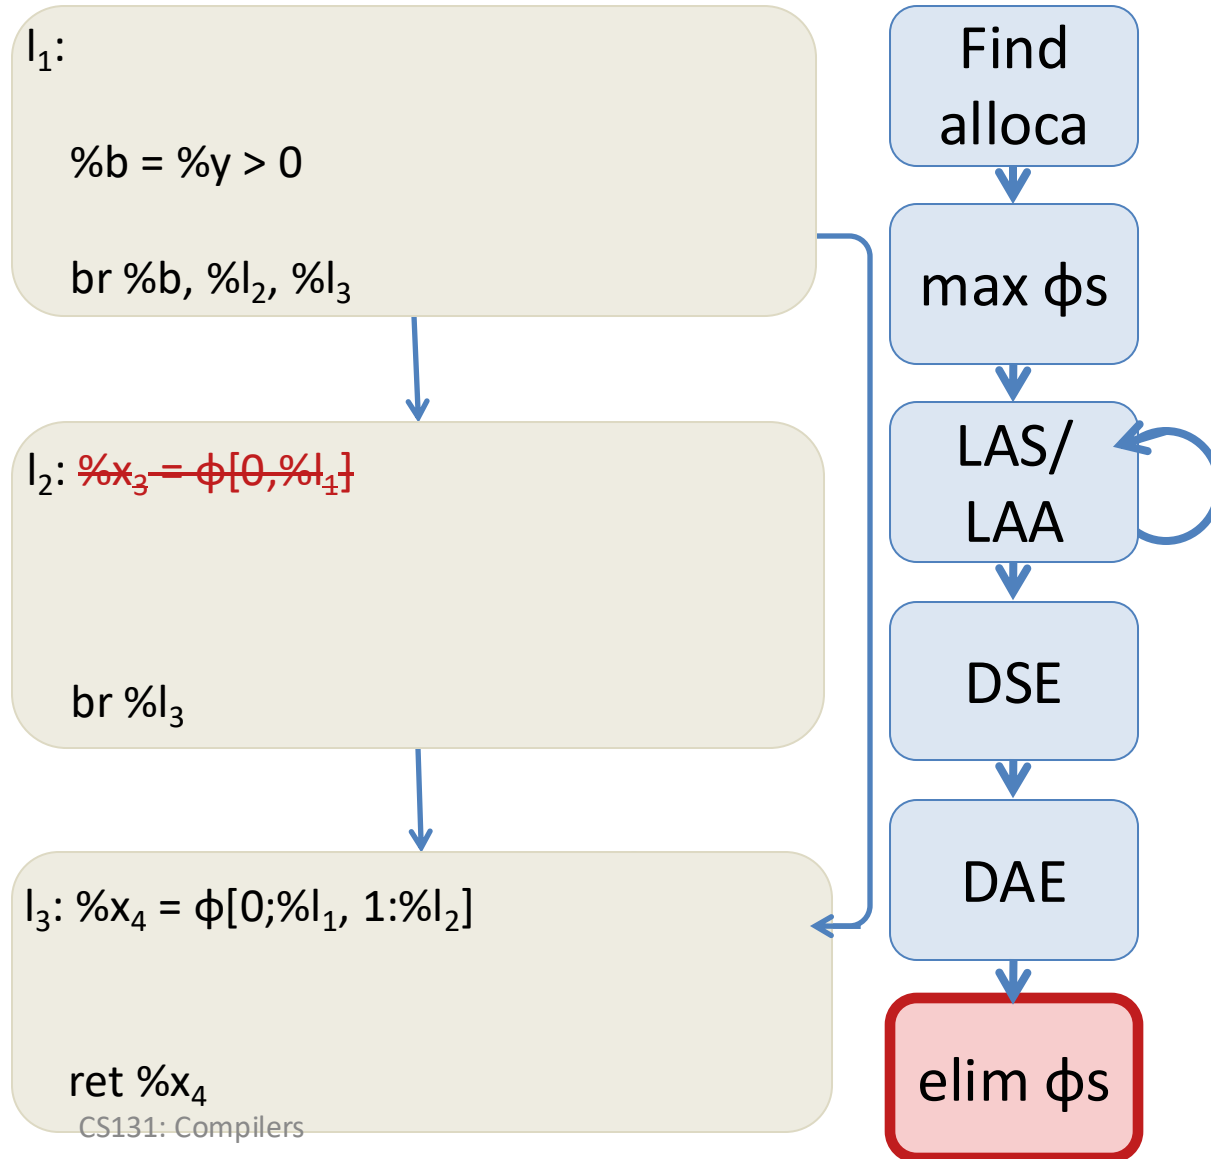
- Dead Store Elimination (DSE)
 - Eliminate all stores with no subsequent loads.
- Dead Alloca Elimination (DAE)
 - Eliminate all allocas with no subsequent loads/stores.

Example SSA Optimizations



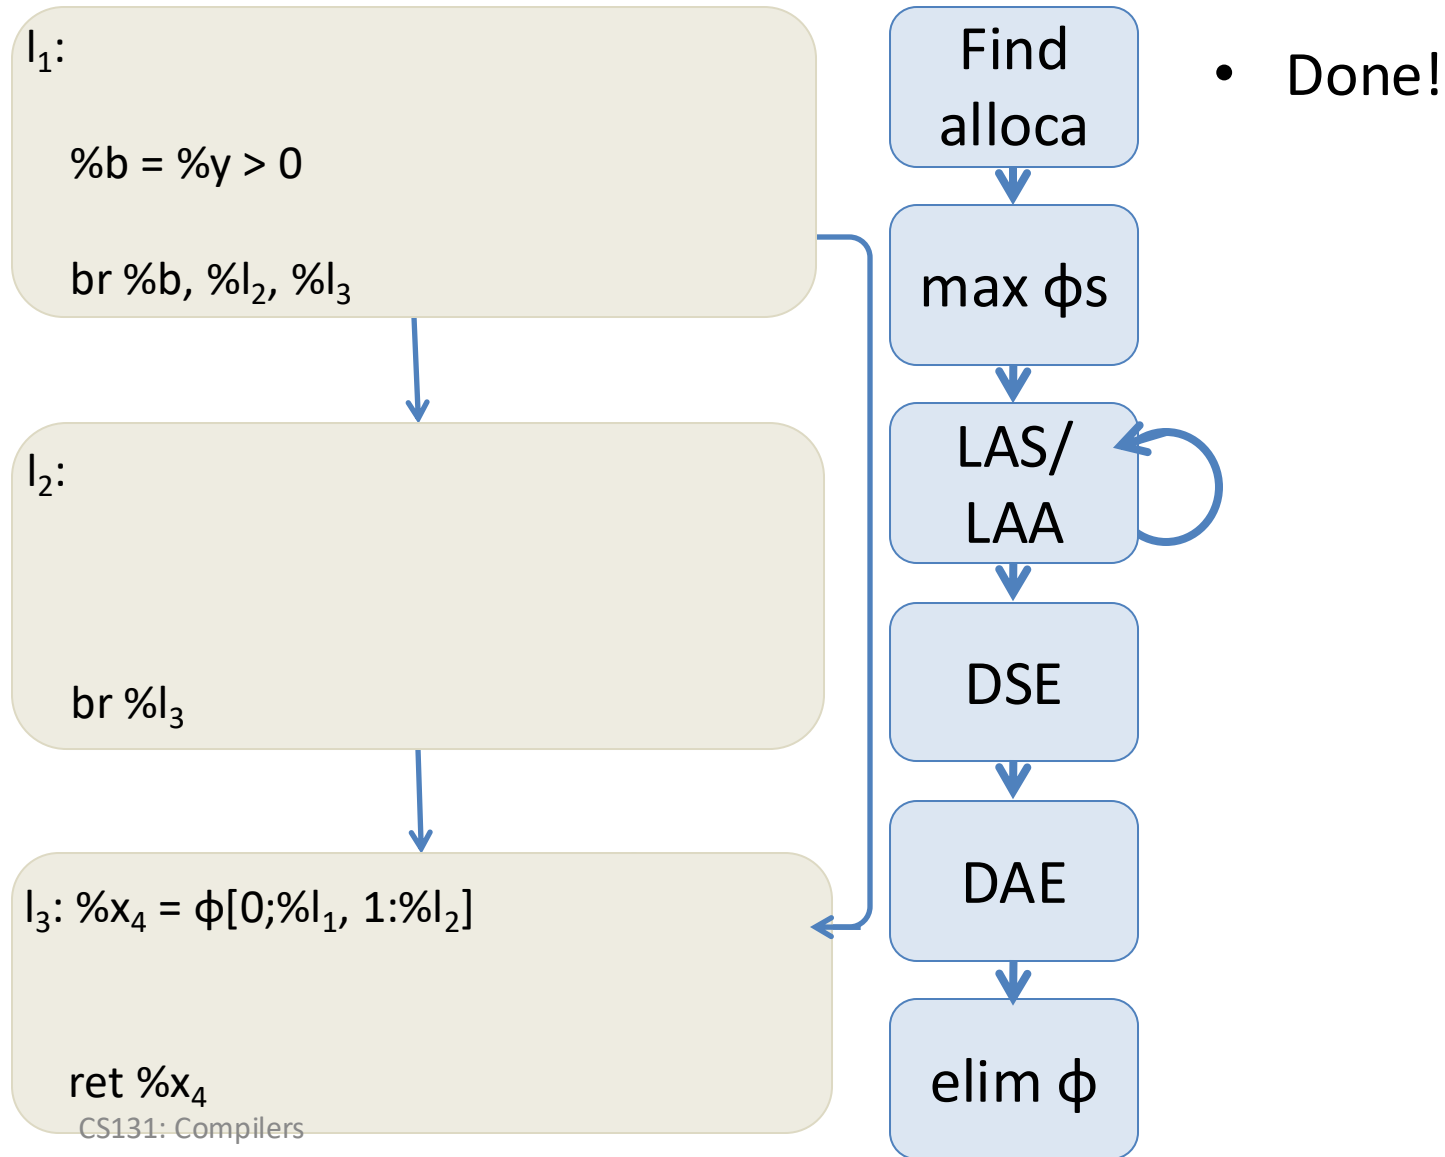
- Eliminate ϕ nodes:
 - Singletons
 - With identical values from each predecessor
 - See Aycock & Horspool, 2002

Example SSA Optimizations



- Eliminate φ nodes:
 - Singletons
 - With identical values from each predecessor

Example SSA Optimizations



LLVM Phi Placement

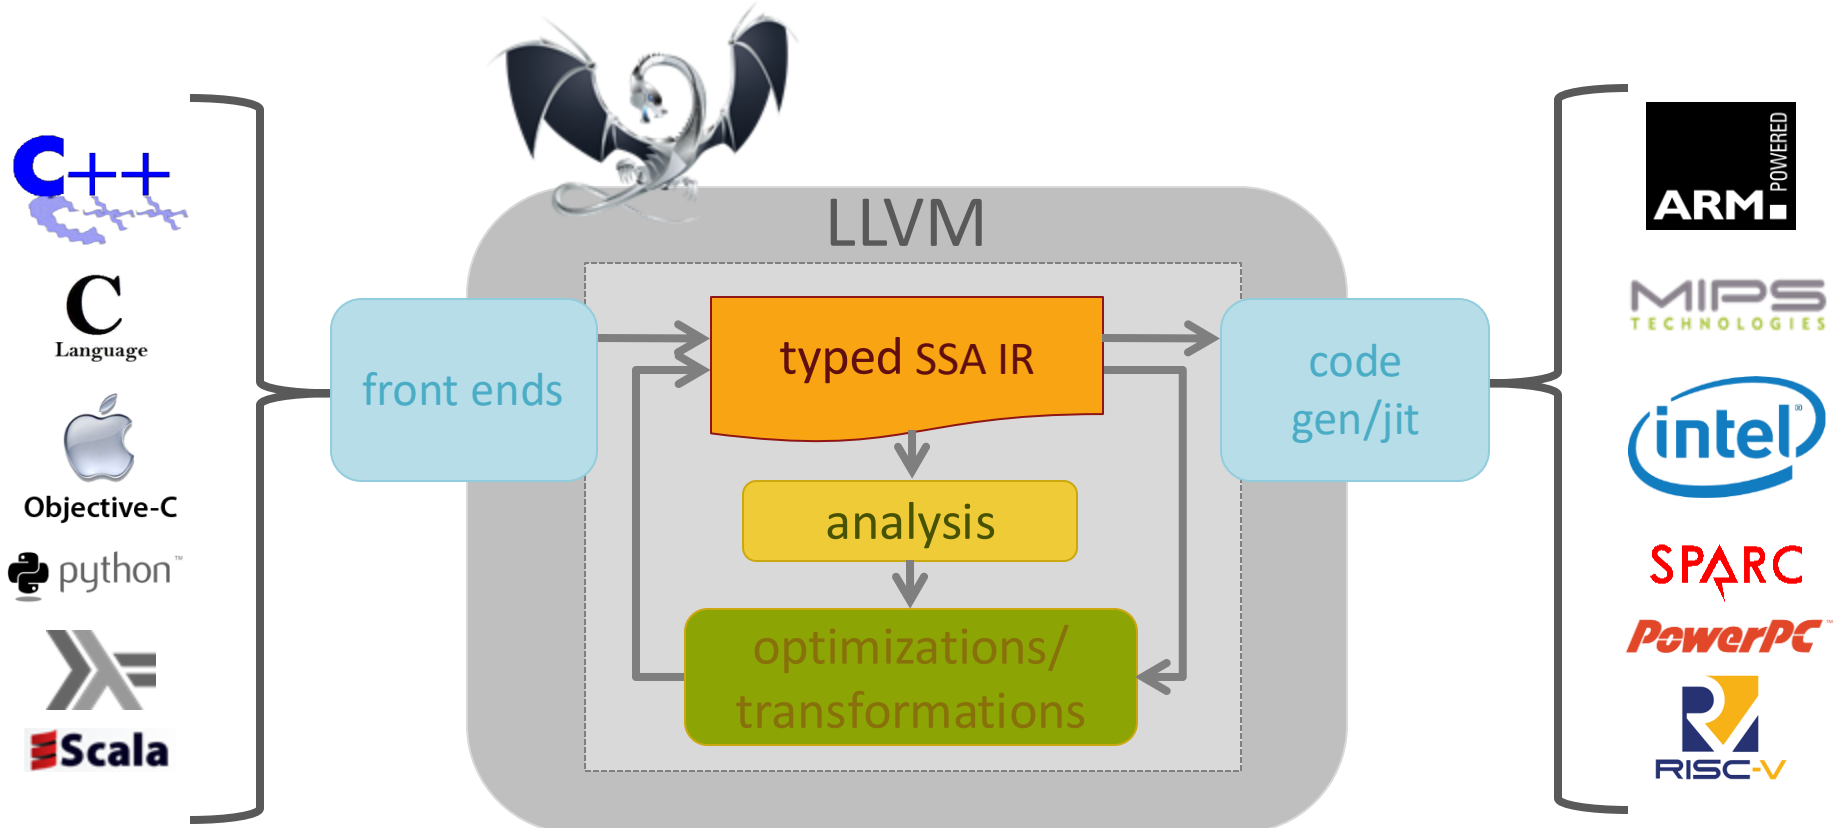
- This transformation is also sometimes called register promotion
 - older versions of LLVM called this “mem2reg” memory to register promotion
- In practice, LLVM combines this transformation with *scalar replacement of aggregates* (SROA)
 - i.e. transforming loads/stores of structured data into loads/stores on register-sized data
- These algorithms are (one reason) why LLVM IR allows annotation of predecessor information in the .ll files
 - Simplifies computing the DF



COMPILER VERIFICATION

LLVM Compiler Infrastructure

[Lattner et al.]



Other LLVM IR Features



- C-style data values
 - ints, structs, arrays, pointers, vectors
- Type system
 - used for layout/alignment/padding
- Relaxed-memory concurrency primitives
- Intrinsics
 - extend the language malloc, bitvectors, etc.
- Transformations & Optimizations

Make targeting LLVM IR
easy and attractive for
developers!

But... it's complex



LLVM Reference Manual table of contents

Abstract

This document is a reference manual for the LLVM assembly language. LLVM is a Static Single Assignment (SSA) based representation that provides type safety, low-level operations, flexibility, and the capability of representing 'all'

trinsic

Intrinsic

Intrinsics
nsics

One Example: undef

The undef "value" represents an arbitrary, but indeterminate bit pattern for any type.

Used for:

- uninitialized registers
- reads from volatile memory
- results of some underspecified operations

What is the value of %y after running the following?

```
%x = or i8 undef, 1  
%y = xor i8 %x, %x
```

One plausible answer: 0

Not LLVM's semantics!

(LLVM is more liberal to permit more aggressive optimizations)

Partially defined values are interpreted
nondeterministically as sets of possible values:

```
%x = or i8 undef, 1  
%y = xor i8 %x, %x
```

$\llbracket \text{i8 undef} \rrbracket = \{0, \dots, 255\}$

$\llbracket \text{i8 1} \rrbracket = \{1\}$

$\llbracket \%x \rrbracket = \{a \text{ or } b \mid a \in \llbracket \text{i8 undef} \rrbracket, b \in \llbracket 1 \rrbracket\}$
 $= \{1, 3, 5, \dots, 255\}$

$\llbracket \%y \rrbracket = \{a \text{ xor } b \mid a \in \llbracket \%x \rrbracket, b \in \llbracket \%x \rrbracket\}$
 $= \{0, 2, 4, \dots, 254\}$

Interactions with Optimizations

Consider:

```
%y = mul i8 %x, 2
```

versus:

$\llbracket \%x \rrbracket = \llbracket \text{i8 undef} \rrbracket$
 $= \{0, 1, 2, 3, 4, 5, \dots, 255\}$

$\llbracket \%y \rrbracket = \{a \text{ mul } 2 \mid a \in \llbracket \%x \rrbracket\}$
 $= \{0, 2, 4, \dots, 254\}$

```
%y = add i8 %x, %x
```

$\llbracket \%x \rrbracket = \llbracket \text{i8 undef} \rrbracket$
 $= \{0, 1, 2, 3, 4, 5, \dots, 255\}$

$\llbracket \%y \rrbracket = \{a + b \mid a \in \llbracket \%x \rrbracket, b \in \llbracket \%x \rrbracket\}$
 $= \{0, 1, 2, 3, 4, \dots, 255\}$



Interactions with Optimizations

Consider:

```
%y = mul i8 %x, 2
```

versus:

```
%y = add i8 %x, %x
```

Upshot: if %x is undef, we can't optimize mul to add (or vice versa)!

What's the problem?

Bug List: (12 of 435) [First](#) [Last](#) [Prev](#) [Next](#) [Show last search results](#)

Bug 33165 - Simplify* cannot distribute instructions for simplification due to undef

Status: REOPENED

Reported: 2017-05-25 02:13 PDT by Nuno Lopes

Davide Italiano 2017-05-25 08:55:40 PDT

[Comment 6](#)

With
cc Davide Italiano 2017-05-25 09:05:26 PDT

[Comment 7](#)

Test
no (unless we want to give up on some undef transformations, and special case selection
but I'm afraid others might be affected too)

B' John Regehr 2017-05-25 09:09:24 PDT

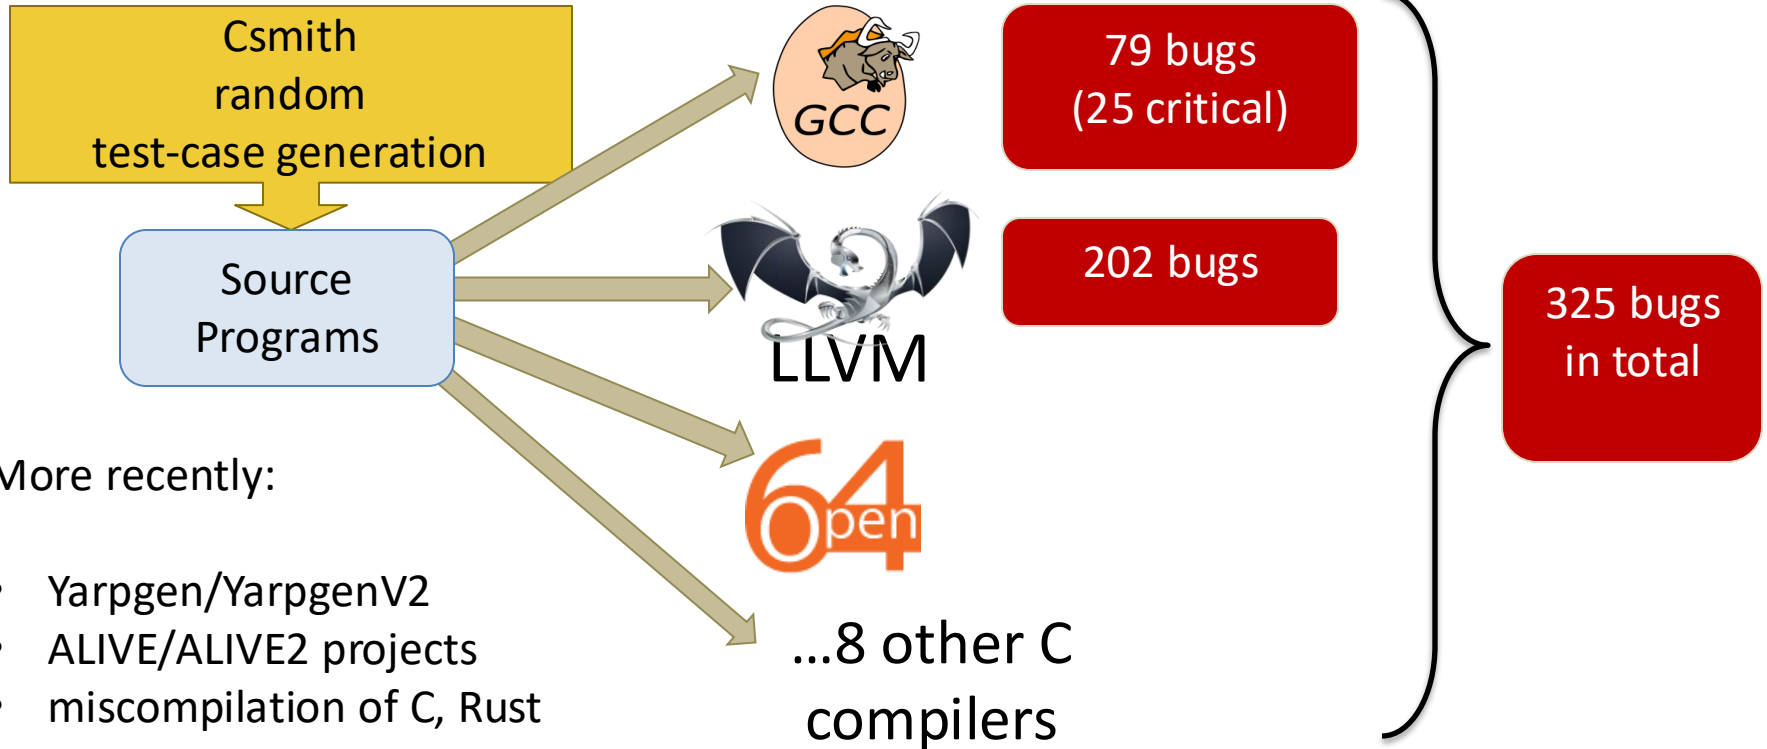
[Comment 8](#)

da
Yes, this is one of those test cases. There are so many optimization failures
Nuno has been automatically filtering out classes of mistranslation that are
to be hard to fix but I guess he decided to take a closer look at some of the

Soon I'll be able to include branches/phis in these test cases, but only for
branches due to a limitation in Alive.

Compiler Bugs

[Regehr's group: Yang et al. PLDI 2011]



More recently:

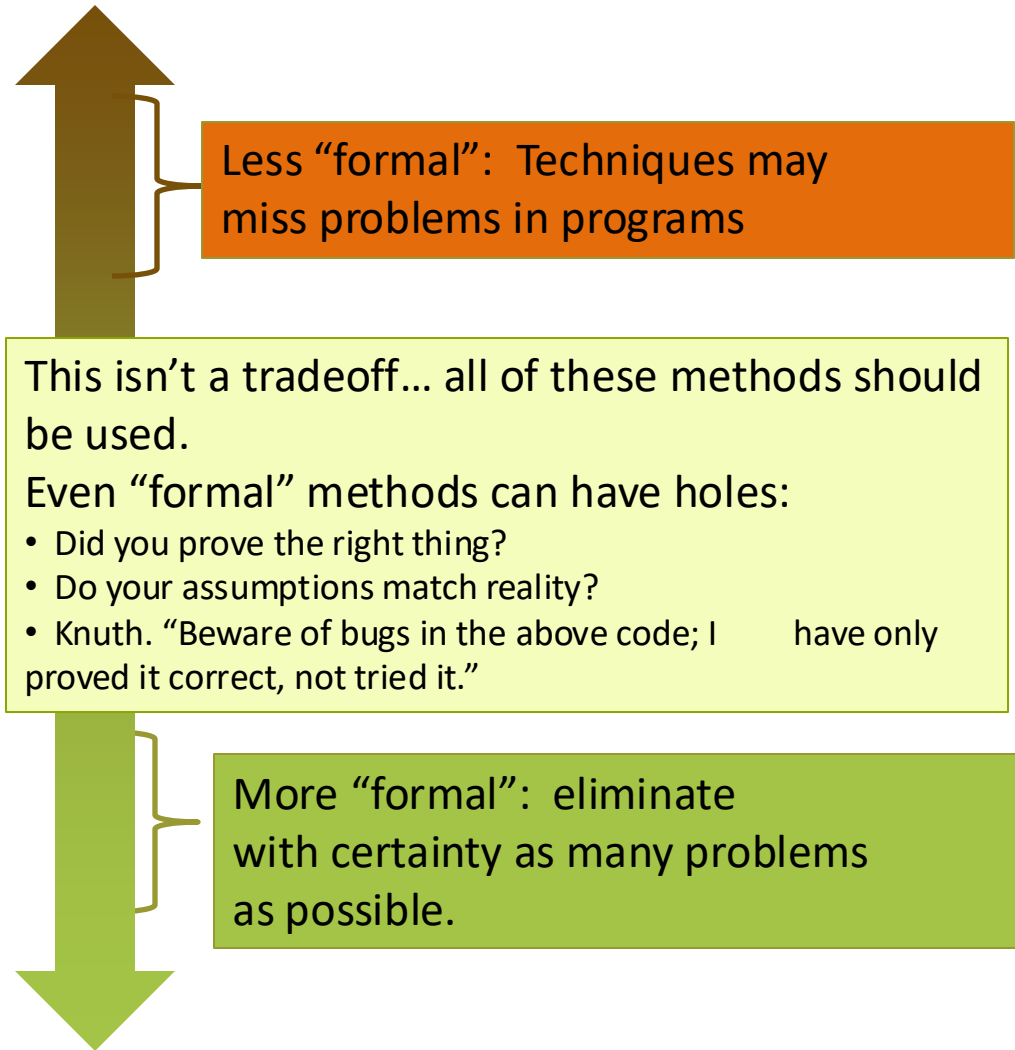
- Yarpgen/YarpgenV2
- ALIVE/ALIVE2 projects
- miscompilation of C, Rust sources [Lee et al. OOPSLA 2018]

LLVM is hard to trust
(especially for critical code)

What can we do about it?

Approaches to Software Reliability

- **Social**
 - Code reviews
 - Extreme/Pair programming
- **Methodological**
 - Design patterns
 - Test-driven development
 - Version control
 - Bug tracking
- **Technological**
 - “lint” tools, static analysis
 - Fuzzers, random testing
- **Mathematical**
 - Sound programming languages tools
 - “Formal” verification



Goal: Verified Software Correctness

- **Social**

- Code reviews
- Extreme/Pair programming

- **Methodological**

- Design patterns
- Test-driven development
- Version control
- Bug tracking

- **Technological**

- “lint” tools, static analysis
- Fuzzers, random testing

- **Mathematical**

- Sound programming languages tools
- “Formal” verification

Q: How can we move the needle towards mathematical software correctness properties?

Taking advantage of advances in computer science:

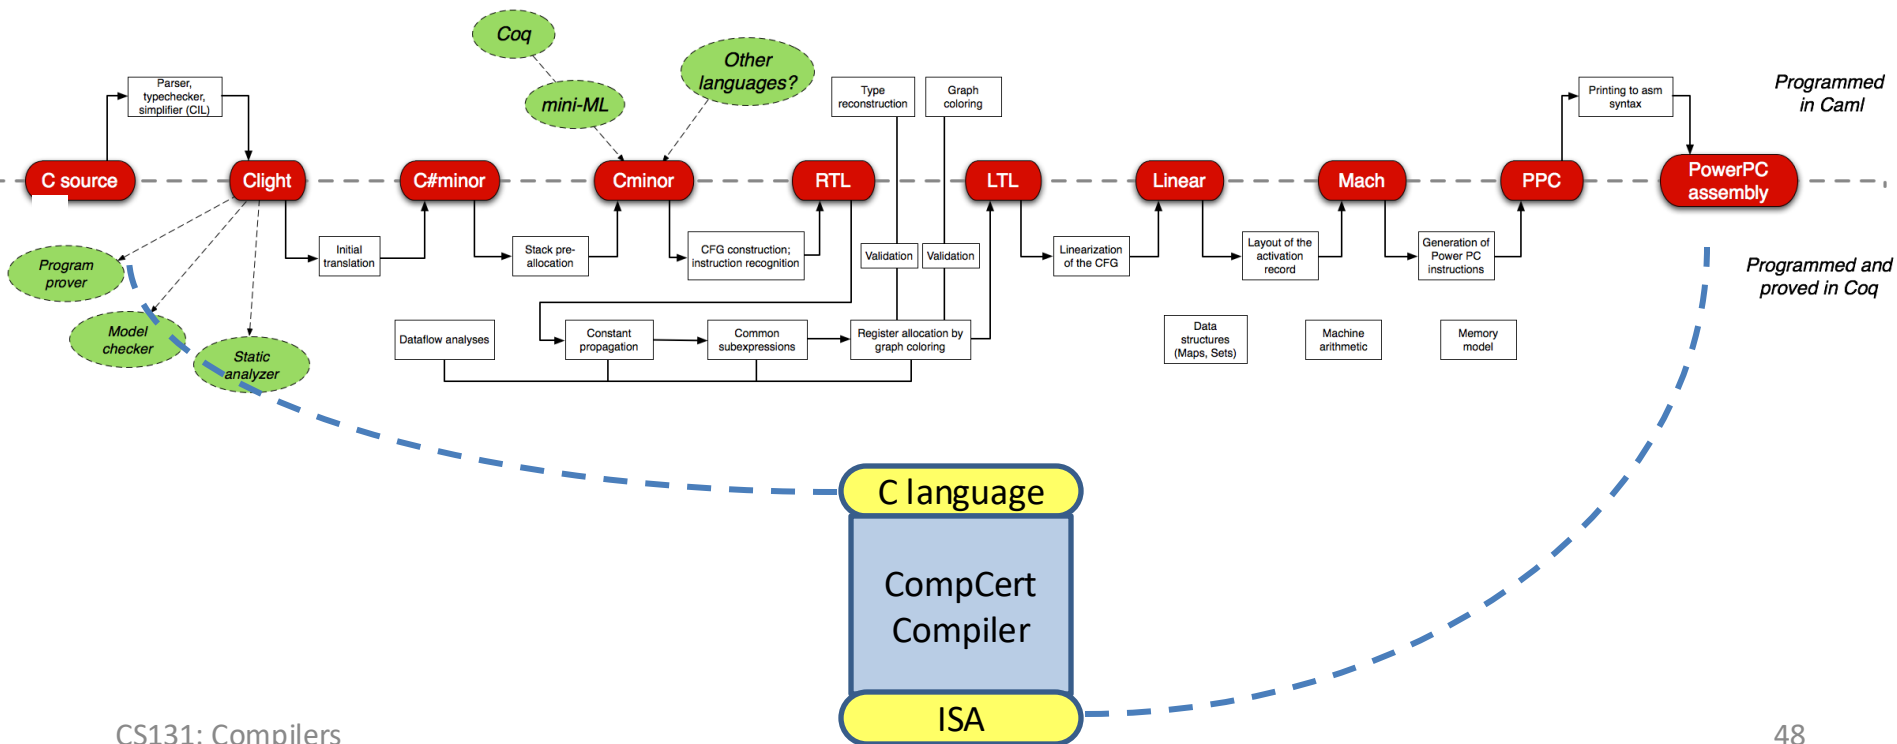
- Moore's law
- improved programming languages & theoretical understanding
- better tools:
interactive theorem provers

CompCert – A Verified C Compiler




Xavier Leroy
INRIA

Optimizing C Compiler,
proved correct end-to-end
with machine-checked proof in Coq



Csmith on CompCert?

[Yang et al. PLDI 2011]



Verification Works!

"The striking thing about our CompCert results is that the middle-end bugs we found in all other compilers are absent. As of early 2011, the under-development version of CompCert is the only compiler we have tested *for which Csmith cannot find wrong-code errors*. This is not for lack of trying: we have devoted about six CPU-years to the task. *The apparent unbreakability of CompCert supports a strong argument that developing compiler optimizations within a proof framework, where safety checks are explicit and machine-checked, has tangible benefits for compiler users.*"

– Regehr et. al 2011