

CS150A Database

Wenjie Wang

School of Information Science and Technology

ShanghaiTech University

Dec. 13, 2024

Today:

- Parallel Query Processing:

Readings:

- Database Management Systems (DBMS), Chapter 22

Review: Refinement

Refinement

- Remove Redundance: Functional Dependencies
- FD
 - Formally: An FD $X \rightarrow Y$ holds over relation schema R if, for every allowable instance r of R: $t1 \in r, t2 \in r, \pi_X(t1) = \pi_X(t2) \Rightarrow \pi_Y(t1) = \pi_Y(t2)$
 - Key/Super Key/Candidate Key
 - F^+ = closure of F:
 - the set of all FDs that are implied by F.
 - includes “trivial dependencies”
 - **Armstrong's Axioms**

Example

- **Contracts(cid,sid,jid,did,pid,qty,value), and:**
 - C is the key: $C \rightarrow CSJDPQV$
 - Proj (J) purchases each part (P) using single contract (C): $JP \rightarrow C$
 - Dept (D) purchases at most 1 part (P) from a supplier (S): $SD \rightarrow P$
- **Problem: Prove that SDJ is a key for Contracts**
- $JP \rightarrow C, C \rightarrow CSJDPQV$
 - Imply $JP \rightarrow CSJDPQV$
 - (by transitivity) (shows that JP is a key)
- $SD \rightarrow P$
 - implies $SDJ \rightarrow JP$ (by augmentation)
- $SDJ \rightarrow JP, JP \rightarrow CSJDPQV$
 - imply $SDJ \rightarrow CSJDPQV$
 - (by transitivity) (shows that SDJ is a key).
- Q: can you now infer that $SD \rightarrow CSDPQV$

Refinement

- Remove Redundance: Functional Dependencies
- FD
 - Formally: An FD $X \rightarrow Y$ holds over relation schema R if, for every allowable instance r of R : $t1 \in r, t2 \in r, \pi_X(t1) = \pi_X(t2) \Rightarrow \pi_Y(t1) = \pi_Y(t2)$
 - Key/Super Key/Candidate Key
 - F^+ = closure of F :
 - the set of all FDs that are implied by F .
 - includes “trivial dependencies”
 - **Armstrong's Axioms**
 - Compute attribute closure of X (denoted X_+) wrt F .

Attribute Closure (example)

$R = \{A, B, C, D, E\}$

$F = \{ B \rightarrow CD, D \rightarrow E, B \rightarrow A, E \rightarrow C, AD \rightarrow B \}$

- **Is $B \rightarrow E$ in F^+ ?**

$B^+ = \{B, C, D, E, \dots\}$

... Yep!

- **Is D a key for R ?**

$D^+ = \{D, E, C\}$

... Nope!

- **Is AD a key for R ?**

$AD^+ = \{A, D, E, C, B\}$

...Yep!

- **Is AD a *candidate* key for R ?**

$A^+ = \{A\}$ $D^+ = \{D, E, C\}$

...Yes!

- **Is ADE a candidate key for R ?**

No!

Refinement

- Remove Redundance: Functional Dependencies
- Functional Dependencies
- Normal Form: Boyce-Codd Normal Form (BCNF)
 - R is in BCNF if the only non-trivial FDs over R are key constraints.
- Decomposition of a Relation Scheme into BCNF
 - There are three potential problems to consider:
 - 1) May be **impossible** to reconstruct the original relation! (Lossiness)
 - Decomposition of R into X and Y is **lossless-join** w.r.t. a set of FDs F if, for every instance r that satisfies F: $\pi_X(r) \bowtie \pi_Y(r) = r$
 - Theorem: The decomposition of R into X and Y is lossless with respect to F if and only if the closure of F contains: $X \cap Y \rightarrow X$, or $X \cap Y \rightarrow Y$
 - 2) Dependency checking may require joins.
 - Decomposition of R into X and Y is dependency preserving if $(F_X \cup F_Y)^+ = F^+$
 - 3) Some queries become more expensive.

Dependency Preservation: Notes

- Critical to consider F^+ in the definition:
 - ABC , $A \rightarrow B$, $B \rightarrow C$, $C \rightarrow A$, decomposed into AB and BC .
 - Is this dependency preserving? Is $C \rightarrow A$ preserved????
- Well... F^+ contains $F \cup \{A \rightarrow C, B \rightarrow A, C \rightarrow B\}$, so...
 - $F_{AB} \supseteq \{A \rightarrow B, B \rightarrow A\}$; $F_{BC} \supseteq \{B \rightarrow C, C \rightarrow B\}$
 - So, $(F_{AB} \cup F_{BC})^+ \supseteq \{B \rightarrow A, C \rightarrow B\}$
 - Hence $(F_{AB} \cup F_{BC})^+ \supseteq \{C \rightarrow A\}$

$$(F_X \cup F_Y)^+ = F^+$$

Decomposition into BCNF

- Consider relation R with FDs F.
- If $X \rightarrow Y$ violates BCNF, decompose R into **R - Y and XY** (guaranteed to be loss-less).
 - Repeated application of this idea will give us a collection of relations that are in BCNF
 - Lossless join decomposition, and guaranteed to terminate.
- e.g., CSJDPQV, key C, $JP \rightarrow C$, $SD \rightarrow P$, $J \rightarrow S$
 - {contractid, supplierid, projectid,deptid,partid, qty, value}
 - To deal with $SD \rightarrow P$, decompose into SDP, CSJDQV.
 - To deal with $J \rightarrow S$, decompose CSJDQV into JS and CJDQV
 - So we end up with: SDP, JS, and CJDQV
- Note: several dependencies may cause violation of BCNF.
- The order in which we “deal with” them could lead to very different sets of relations!

BCNF and Dependency Preservation

- In general, **there may not be a dependency preserving decomposition into BCNF.**
- E.g., CSZ, $CS \rightarrow Z$, $Z \rightarrow C$
 - Can't decompose while preserving 1st FD; not in BCNF.
- Similarly, decomposition of CSJDPQV into SDP, JS and CJDQV is not dependency preserving (w.r.t. the FDs **$JP \rightarrow C$** , $SD \rightarrow P$ and $J \rightarrow S$).
 - However, it is a lossless join decomposition.
 - In this case, adding JPC to the collection of relations gives us a dependency preserving decomposition.
 - but JPC tuples are stored only for checking the f.d. (**Redundancy!**)

A little history

- Relational revolution
 - declarative set-oriented primitives
 - 1970' s
- Parallel relational database systems
 - on commodity hardware
 - 1980' s
- Big Data: MapReduce, Spark, etc.
 - scaling to thousands of machines and beyond
 - 2005-2015

Review: Parallel Query Processing


Why Parallelism?

- Scan 100TB
 - At 0.5 GB/sec (see lec 4):
~200,000 sec = ~2.31 days





Why Parallelism? Cont.

- Scan 100TB
 - At 0.5 GB/sec (see lec 4):
~200,000 sec = ~2.31 days
- Run it 100-way parallel:
 - 2,000 sec = 33 minutes
- 1 big problem = many small problems
 - Trick: make them independent




Two Metrics to Shoot For


- Speed-up
 - Increase HW
 - Fix workload
- Scale-up
 - Increase HW
 - Increase workload




Roughly 2 Kinds of Parallelism



: any sequential program,
e.g. a relational operator



Pipeline
scales up to pipeline depth



Partition
scales up to amount of data


We'll get more
refined soon.

Easy for us to say!


- Lots of Data:
 - Batch operations
 - Pre-existing divide-and-conquer algorithms
 - Natural pipelining
- Declarative languages
 - Can adapt the parallelism strategy to the task and the hardware
 - All without changing the program!
 - Codd's Physical Data Independence

Parallel Architectures


Shared Memory



Shared Disk



Shared Nothing
(cluster)




Shared Nothing

- We will focus on Shared Nothing here
 - It's the most common
 - DBMS, web search, big data, machine learning, ...
 - Runs on commodity hardware
 - Scales up with data
 - Just keep putting machines on the network!
 - Does not rely on HW to solve problems
 - Good for helping us understand what's going on
 - Control it in SW


Kinds of Query Parallelism

- Inter-query (parallelism across queries)
 - Each query runs on a separate processor
 - Single thread (no parallelism) per query
 - Does require parallel-aware concurrency control



Intra Query – Inter-operator


- Intra-query (within a single query)
 - **Inter-operator** (between operators)




Pipeline Parallelism

Intra Query – Inter-operator Part 2

- Intra-query
 - **Inter-operator**




Pipeline Parallelism





Intra Query - Inter-Operator Part 3

- Intra-query
 - Inter-operator




Pipeline Parallelism



Bushy (Tree) Parallelism



Intra Query – Intra-Operator

- Intra-query
 - **Intra-operator** (within a single operator)



Kinds of Query Parallelism, cont.


- Intra-query
 - Intra-operator



Partition Parallelism



Summary: Kinds of Parallelism

- Inter-Query




- Intra-Query

- Inter-Operator



Pipeline Parallelism

- Intra-Operator (partitioned)




Partition Parallelism

INTRA-OPERATOR PARALLELISM

Data Partitioning

- How to partition a table across disks/machines
 - A bit like coarse-grained indexing!




- Shared nothing particularly benefits from "good" partitioning

Parallel Scans


- Scan in parallel, merge (concat) output
- σ_p : skip entire sites that have no tuples satisfying p
 - range or hash partitioning
- Indexes can be built at each partition
- Q: Do indexes differ in the different data partitioning schemes?

Lookup by key

- Data partitioned on function of key?
 - Great! Route lookup only to relevant node
- Otherwise
 - Have to broadcast lookup (to all nodes)




Hash




Round-Robin

What about Insert?

- Data partitioned on function of key?
 - Route insert to relevant node
- Otherwise
 - Route insert to *any* node




Hash




Round-Robin

Insert to Unique Key?

- Data partitioned on function of key?
 - Route to relevant node
 - And reject if already exists




Hash




Round-Robin

Insert to Unique Key cont.

- Otherwise
 - Broadcast lookup
 - Collect responses
 - If not exists, insert anywhere
 - Else reject




Hash




Round-Robin

Remember Hashing?




Parallelize me! Hashing

- Phase 1: shuffle data across machines (h_n)
 - streaming out to network as it is scanned
 - which machine for this record?
 - use (yet another) independent hash function h_n



Parallelize me! Hashing Part 2

- Receivers proceed with phase 1 in a pipeline as data streams in
 - from local disk and network



Nearly same as single-node hashing

*Near-perfect speed-up, scale-up!
Streams through phase 1, during which
time every component works at its top
speed, no waiting.*

Have to wait to start phase 2.


Hash Join?

- Hmmmm....

If you have enough machines...


Naïve parallel hash join

- Phase 1: *shuffle* each table across machines (h_n)
 - Parallel scan streaming out to network
 - **Wait** for building relation to finish
 - Then stream probing relation through it
- Receivers proceed with naïve hashing in a pipeline *as probe data streams in*
 - from local disk and network
 - Writes are independent, hence parallel
- Note: there is a variation that has *no waiting*: both tables stream
 - Wilschut and Apers' "Symmetric" or "Pipeline" hash join
 - Requires more memory space




Parallel Grace Hash Join Pass 1

- Pass 1 is like hashing above





Parallel Grace Hash Join Pass 1 cont

- Pass 1 is like hashing above
 - But do it 2x: once for each relation being joined



Parallel Grace Hash Join Pass 2

- Pass 2 is local Grace Hash Join per node
 - Complete independence across nodes




Parallel Grace Hash Join

- Pass 1: parallel streaming
 - Stream building and probing tables through shuffle/partition
- Pass 2 is local Grace Hash Join per node
 - Complete independence across nodes in Pass 2
- Near-perfect speed-up, scale-up!
- Every component works at its top speed
 - Only waiting is for Pass 1 to end.
- Note: there is a variant that has no waiting
 - Urhan's Xjoin, a variant of symmetric hash


Parallelize me! Sorting Pass 0

- Pass 0: shuffle data across machines
 - streaming out to network as it is scanned
 - which machine for this record?
Split on value range (e.g. $[-\infty, 10]$, $[11, 100]$, $[101, \infty]$).




Parallelize me! Sorting Pass 1-n

- Receivers proceed with pass 0 as the data streams in
- Passes 1-n done independently as in single-node sorting
- **A Wrinkle: How to ensure ranges are the same #pages?!**
 - i.e. avoid data skew?




Range partitioning

- Goal: equal frequency per machine
- Note: ranges often don't divide x axis evenly
- How to choose?



Range partitioning cont.


- Would be easy if data small
- In general, can sample the input relation prior to shuffling, pick splits based on sample
- Note: Random sampling can be tricky to implement in a query pipeline; simpler if you materialize first.



How to sample a database table?
Advanced topic, we will not discuss in this class.


Parallel Sort-Merge Join

- Pass 0 .. $n-1$ are like parallel sorting above
- Note: this picture is a 2-pass sort ($n=1$); this is pass 0





Parallel Sort-Merge Join Pass 0...n-1

- Pass 0 .. n-1 are like parallel sorting above
 - But do it 2x: once for each relation, with same ranges
 - Note: this picture is a 2-pass sort (n=1); this is pass 0




Pass n (with optimization)

- Pass 0 .. n-1 are like parallel sorting above
 - But do it 2x: once for each relation, with same ranges
- Pass n: merge join partitions locally on each node




Parallel Aggregates

- Hierarchical aggregation
- For each aggregate function, need a global/local decomposition:
 - **sum**(S) = $\Sigma \Sigma (s)$
 - **count** = $\Sigma \text{count} (s)$
 - **avg**(S) = $(\Sigma \Sigma (s)) / \Sigma \text{count} (s)$
 - etc...




Parallel GroupBy

- Naïve Hash Group By
 - Local aggregation: in hash table keyed by group key k_i keep local agg_i
 - E.g. `SELECT SUM(price) group by cart;`




Parallel GroupBy, Cont.

- Naïve Hash Group By
 - Local aggregation: in hash table keyed by group key k_i keep local agg_i
 - For example, k is major, agg is $(avg(gpa), count(*))$
 - Shuffle local aggs by a hash function $h_p(k_i)$
 - Compute global aggs for each key k_i



Parallel Aggregates/GroupBy Challenge!

- Exercise:
 - Figure out parallel 2-pass GraceHash-based scheme to handle # large of groups
 - Figure out parallel Sort-based scheme




Joins: Bigger picture

- Alternatives:
 - Symmetric shuffle
 - What we did so far
 - Asymmetric shuffle
 - Broadcast join


Join: One-sided shuffle

- If R already suitably partitioned,
- just partition S, then run local join at every node and union results.



“Broadcast” Join

- If R is small, send it to all nodes that have a partition of S.
- Do a local join at each node (using any algorithm) and union results.




What are “pipeline breakers”?

- Sort
 - Hence sort-merge join can't start merging until sort is complete
- Hash build
 - Hence Grace hash join can't start probing until hashtable is built
- Is there a join scheme that pipelines?

Symmetric (Pipeline) Hash Join

- Single-phase, streaming
- Each node allocates two hash tables, one for each side
- Upon arrival of a tuple of R:
 - Build into R hashtable by join key
 - Probe into S hashtable for matches and output any that are found
- Upon arrival of a tuple of S:
 - Symmetric to R!




Symmetric (Pipeline) Hash Join cont

- Why does it work?
 - Each output tuple is generated exactly once: when the second part arrives
- Streaming!
 - Can always pull another tuple from R or S, build, and probe for outputs
 - Useful for Stream query engines!

Extensions

- Parallel Symmetric Hash Join
 - Straightforward—part of the original proposal
 - Just add a streaming partitioning phase up front
 - As in naïve hash join
- Out-of-core Symmetric Hash Join
 - Quite a bit trickier. See the [X-Join paper](#).
- Non-blocking sort-merge join
 - See the [Progressive Merge Join](#) paper



Parallel DBMS Summary

- Parallelism natural to query processing:
 - Both pipeline and partition
- Shared-Nothing vs. Shared-Mem vs. Shared Disk
 - Shared-mem easiest SW, costliest HW.
 - Doesn't scale indefinitely
 - Shared-nothing cheap, scales well, harder to implement.
 - Shared disk a middle ground
 - For updates, introduces tricky stuff related to concurrency control
- Intra-op, Inter-op, & Inter-query parallelism all possible.

Parallel DBMS Summary, Part 2

- Data layout choices important!
- Most DB operations can be done partition-parallel
 - Sort. Hash.
 - Sort-merge join, hash-join.
- Complex plans.
 - Allow for pipeline-parallelism, but sorts, hashes block the pipeline.
 - Partition parallelism achieved via bushy trees.

Parallel DBMS Summary, Part 3

- Transactions require introducing some new protocols
 - distributed deadlock detection
 - two-phase commit (2PC)
- 2PC not great for availability, latency
 - single failure stalls the whole system
 - transaction commit waits for the slowest worker
- More on this in subsequent lectures