CS150A Database

Wenjie Wang

School of Information Science and Technology

ShanghaiTech University

Dec. 23, 2024

Today:

- Analytics and ML in Data Systems:
  - Part 1
  - Data warehouse & Data Lake

Readings:

- Database Management Systems (DBMS), Chapter 25

Acknowledgement: Joseph E. Gonzalez@UCBerkeley's scourse notes

Transaction Processing vs Analytics

Online Transaction Processing (OLTP)

- Many small queries:
  - Freq. use of indexes
  - Many writes
  - Concurrency and Logging
- Managing the “Now”
  - Source of truth
- Fairly simple queries with few predicates and relations

Online Analytics Processing (OLAP) & Data Mining/ML

- Exploratory Full Table Queries
  - e.g., Agg. Sales Per Market
  - Infrequent (but bulk) writes
  - Limited transaction processing
- Recording the history
  - What was our inventory at the end of last two quarters
- Complex queries with many predicates and many relations

Analytics & ML queries:

- What was our total sales by market last quarter?
  - Summarization
- What is our predicted sales for next quarter?
  - Forecasting
- Which users will likely leave our service?
  - Churn prediction
- If a user buys X what else are they likely to buy?
  - Collaborative filtering & Recommender Systems

You embark on the journey of a data scientist …

Sales (Asia)
Sales (US)
Inventory
Advertising

Data Everywhere

- Stored Across Multiple Operational OLTP Systems
  - Different formats (e.g., currency)
    - Different schemas (acquisitions ...)
  - Mission critical
    - Serving live sales traffic
    - Managing inventory
    - … Be careful!
- Often limited historical data

We would like a consolidated, cleaned, historical snapshot of the data.

Data Warehouse

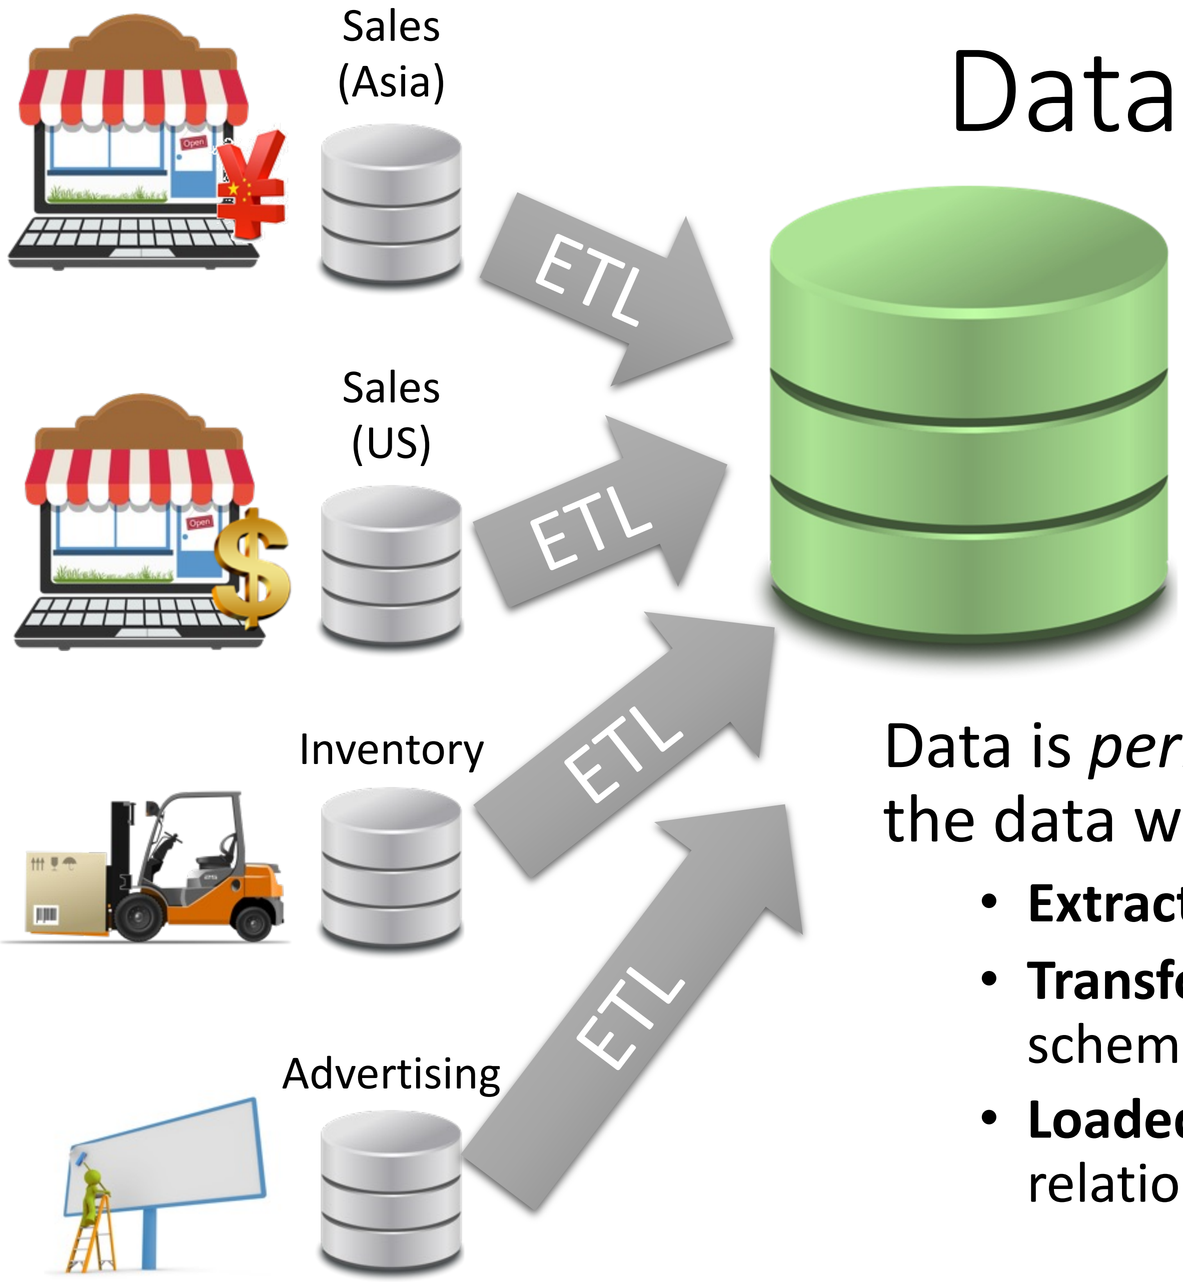

Collects and organizes historical data from multiple sources

Data is *periodically* **ETL**ed into the data warehouse:

- **Extracted** from remote sources
- **Transformed** to standard schemas
- **Loaded** into the (typically) relational system

Extracting Data from Sources

- Need to collect data from multiples sources
  - Various RDBMS vendors
  - Structured files JSON, XML

- Often done using SQL interfaces

- Validate extracted data
  - Flag corrupted records …

Transforming “Cleaning” Data

- Additional data validation and filtering

- Schema manipulation
  - Extract key fields
  - Encoding text
  - Verifying and enforcing constraints

- Data normalization (time zones, currency)

Loading Data

- Data is bulk loaded into large relations
    - Fact tables … (more on this later)
- Update:
    - Indexes
    - Metadata tables: Data about the data
        - When and how was it collected
        - Meaning of fields
    - Updating materialized views …
- Occasionally move older data to archival storage
    - Data aging

Data Warehouse

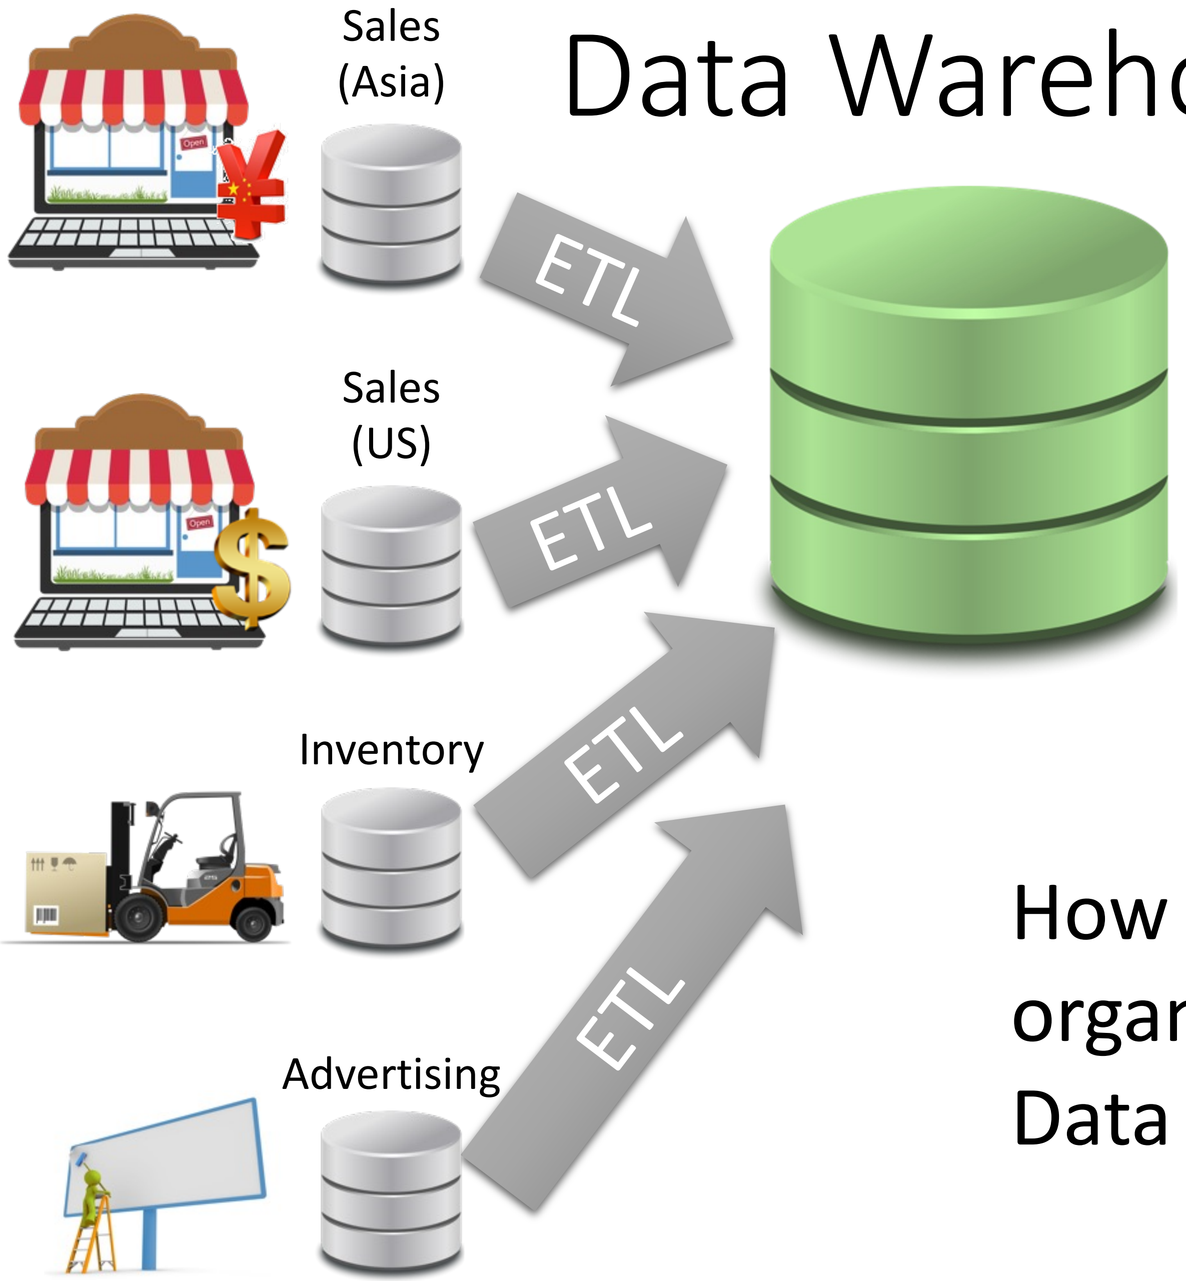

How is data organized in the Data Warehouse?

Example Sales Data:

pname	category	price	qty	date	day	city	state	country
Corn	Food	25	25	3/30/16	Wed.	Omaha	NE	USA
Corn	Food	25	8	3/31/16	Thu.	Omaha	NE	USA
Corn	Food	25	15	4/1/16	Fri.	Omaha	NE	USA
Galaxy 1	Phones	18	30	1/30/16	Wed.	Omaha	NE	USA
Galaxy 1	Phones	18	20	3/31/16	Thu.	Omaha	NE	USA
Galaxy 1	Phones	18	50	4/1/16	Fri.	Omaha	NE	USA
Galaxy 1	Phones	18	8	1/30/16	Wed.	Omaha	NE	USA
Peanuts	Food	2	45	3/31/16	Thu.	Seoul		Korea
Galaxy 1	Phones	18	100	4/1/16	Fri.	Seoul		Korea

- **Big** table: many *columns* and *rows*
  - Substantial redundancy → expensive to store and access
- Could we organize the data a little better?

Multidimensional Data Model

Sales Fact Table

pid	timeid	locid	sales
11	1	1	25
11	2	1	8
11	3	1	15
12	1	1	30
12	2	1	20
12	3	1	50
12	1	1	8
13	2	1	10
13	3	1	10
11	1	2	35
11	2	2	22
11	3	2	10
12	1	2	26

Locations

locid	city	state	country
1	Omaha	Nebraska	USA
2	Seoul		Korea
5	Richmond	Virginia	USA

Dimension Tables

Products

pid	pname	category	price
11	Corn	Food	25
12	Galaxy 1	Phones	18
13	Peanuts	Food	2

Time

timeid	Date	Day
1	3/30/16	Wed.
2	3/31/16	Thu.
3	4/1/16	Fri.

- Multidimensional “Cube” of data

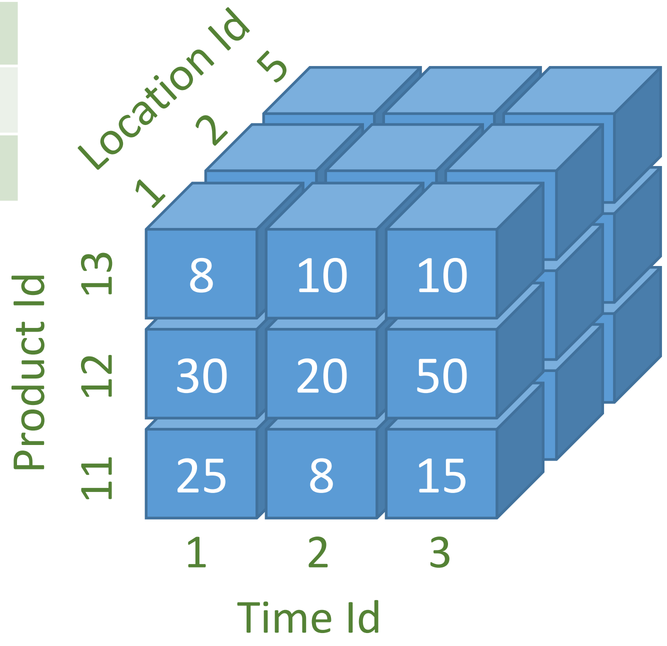

Multidimensional Data Model

Sales **Fact Table**

pid	timeid	locid	sales
11	1	1	25
11	2	1	8
11	3	1	15
12	1	1	30
12	2	1	20
12	3	1	50
12	1	1	8
13	2	1	10
13	3	1	10
11	1	2	35
11	2	2	22
11	3	2	10
12	1	2	26

Locations

locid	city	state	country
1	Omaha	Nebraska	USA
2	Seoul		Korea
5	Richmond	Virginia	USA

Dimension Tables

Products

pid	pname	category	price
11	Corn	Food	25
12	Galaxy 1	Phones	18
13	Peanuts	Food	2

Time

timeid	Date	Day
1	3/30/16	Wed.
2	3/31/16	Thu.
3	4/1/16	Fri.

- Sales Fact Table
  - Contains only foreign keys → Efficient
- Easy to manage Dimensions
  - Galaxy1 → Phablet: no need to update **Fact Table**
- Normalization
  - Minimizing redundancy

Multidimensional Data: **Star Schema**

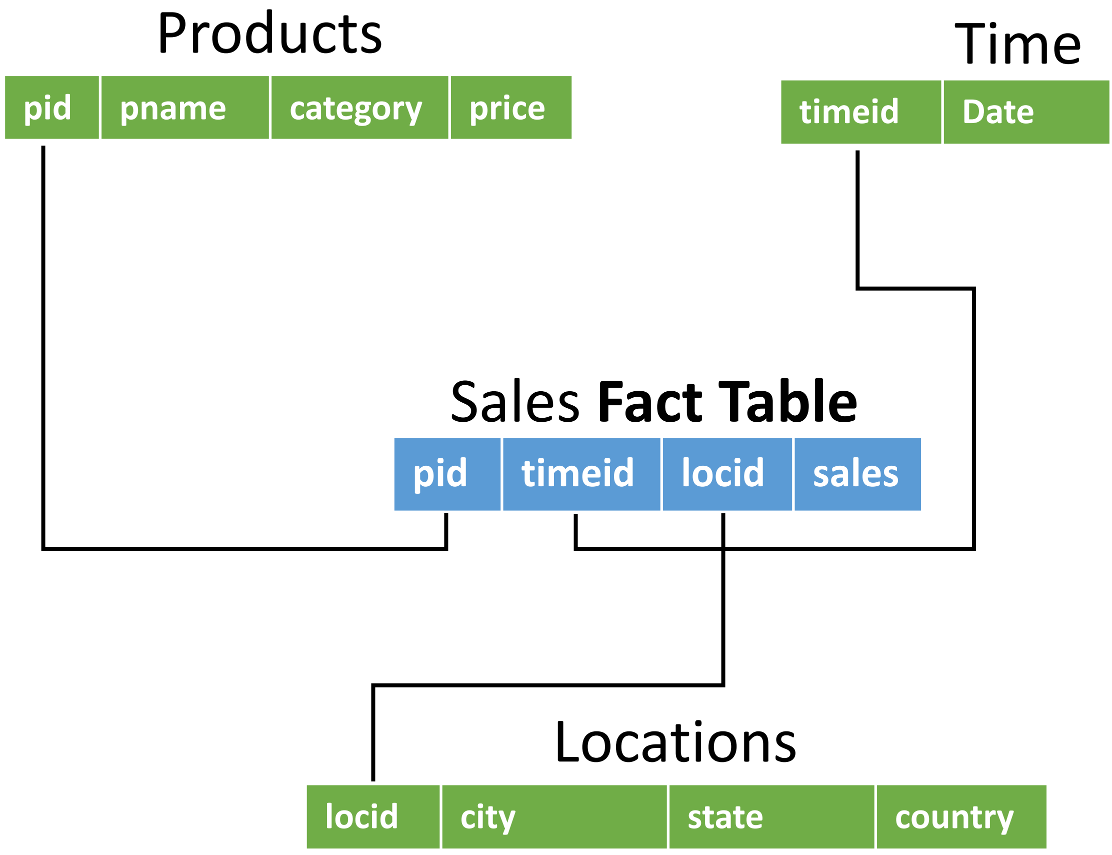

Products

pid	pname	category	price

Time

timeid	Date	Day

Sales **Fact Table**

pid	timeid	locid	sales

Locations

locid	city	state	country

Dimension Tables

← This looks like a star …

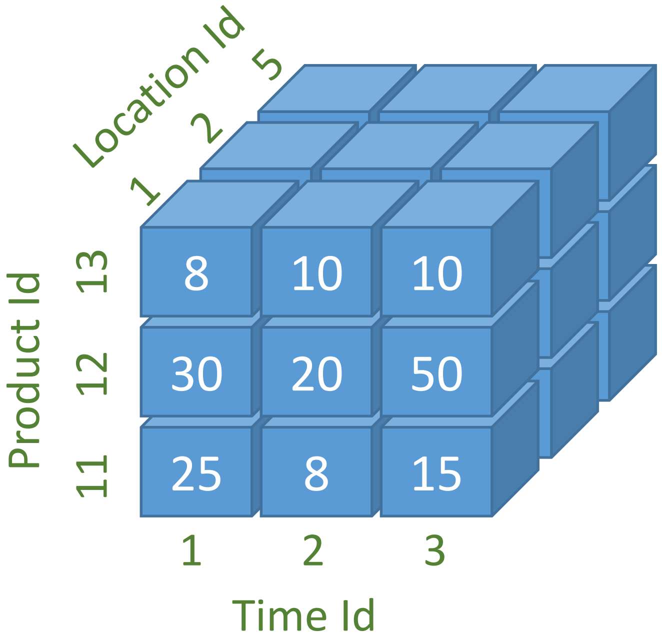

Product Id	1	2	3
13	8	10	10
12	30	20	50
11	25	8	15

Data Warehouse

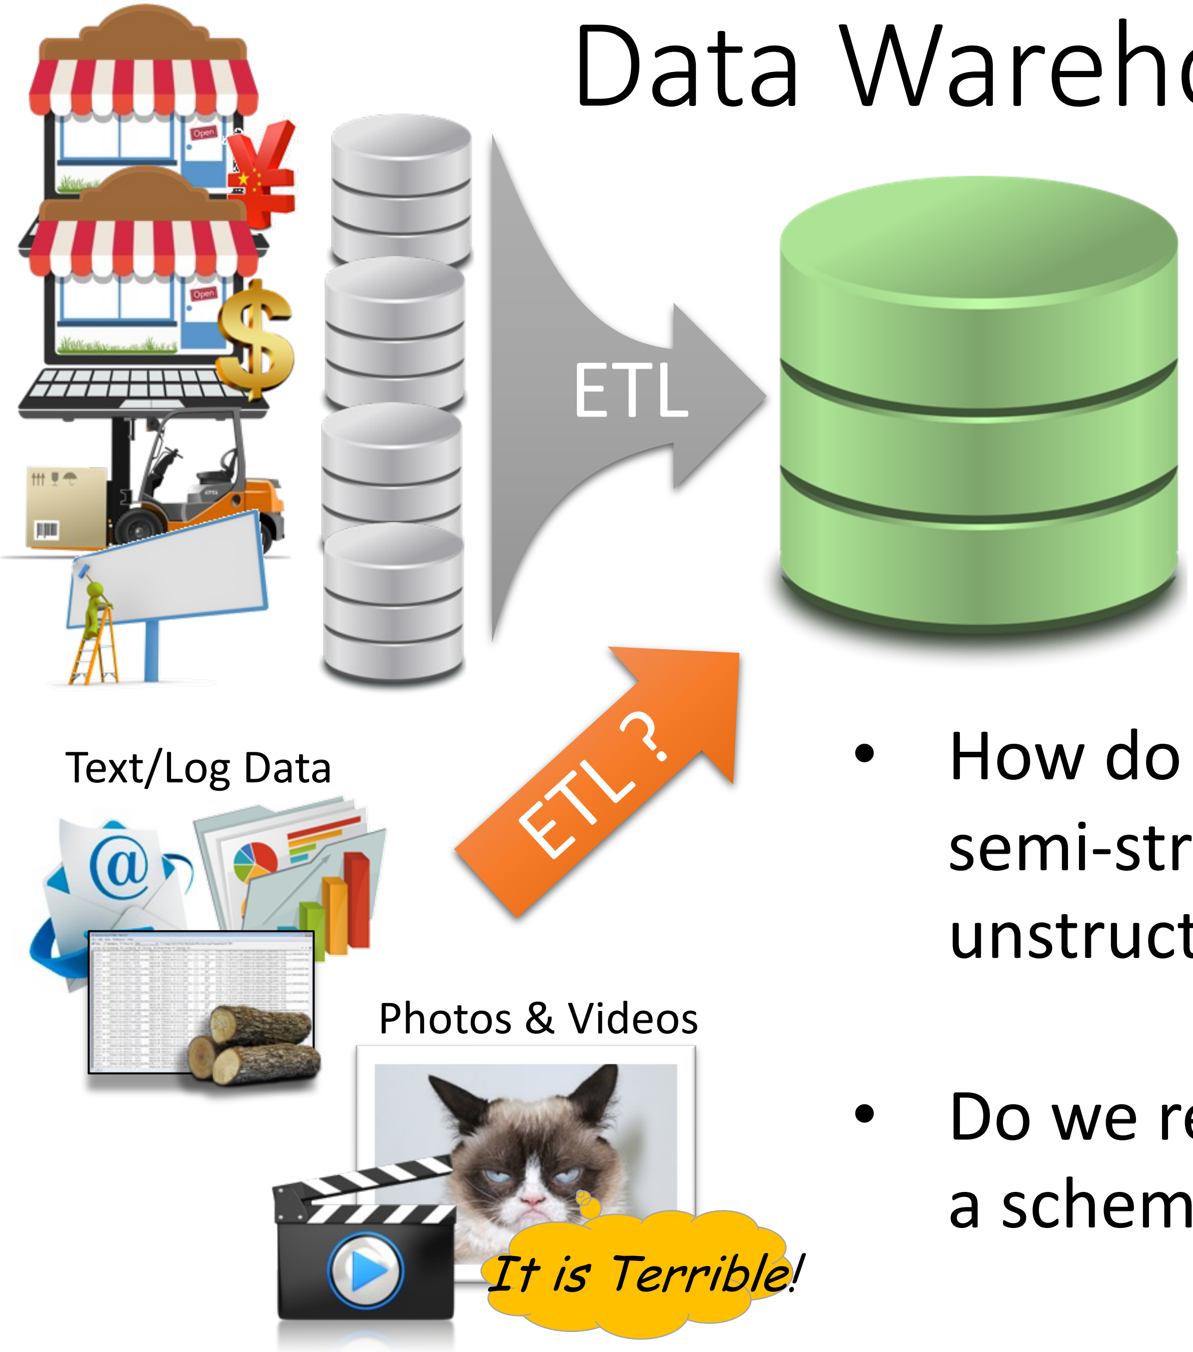

- How do we deal with semi-structured and unstructured data?
- Do we really want to force a schema on load?

Data Warehouse

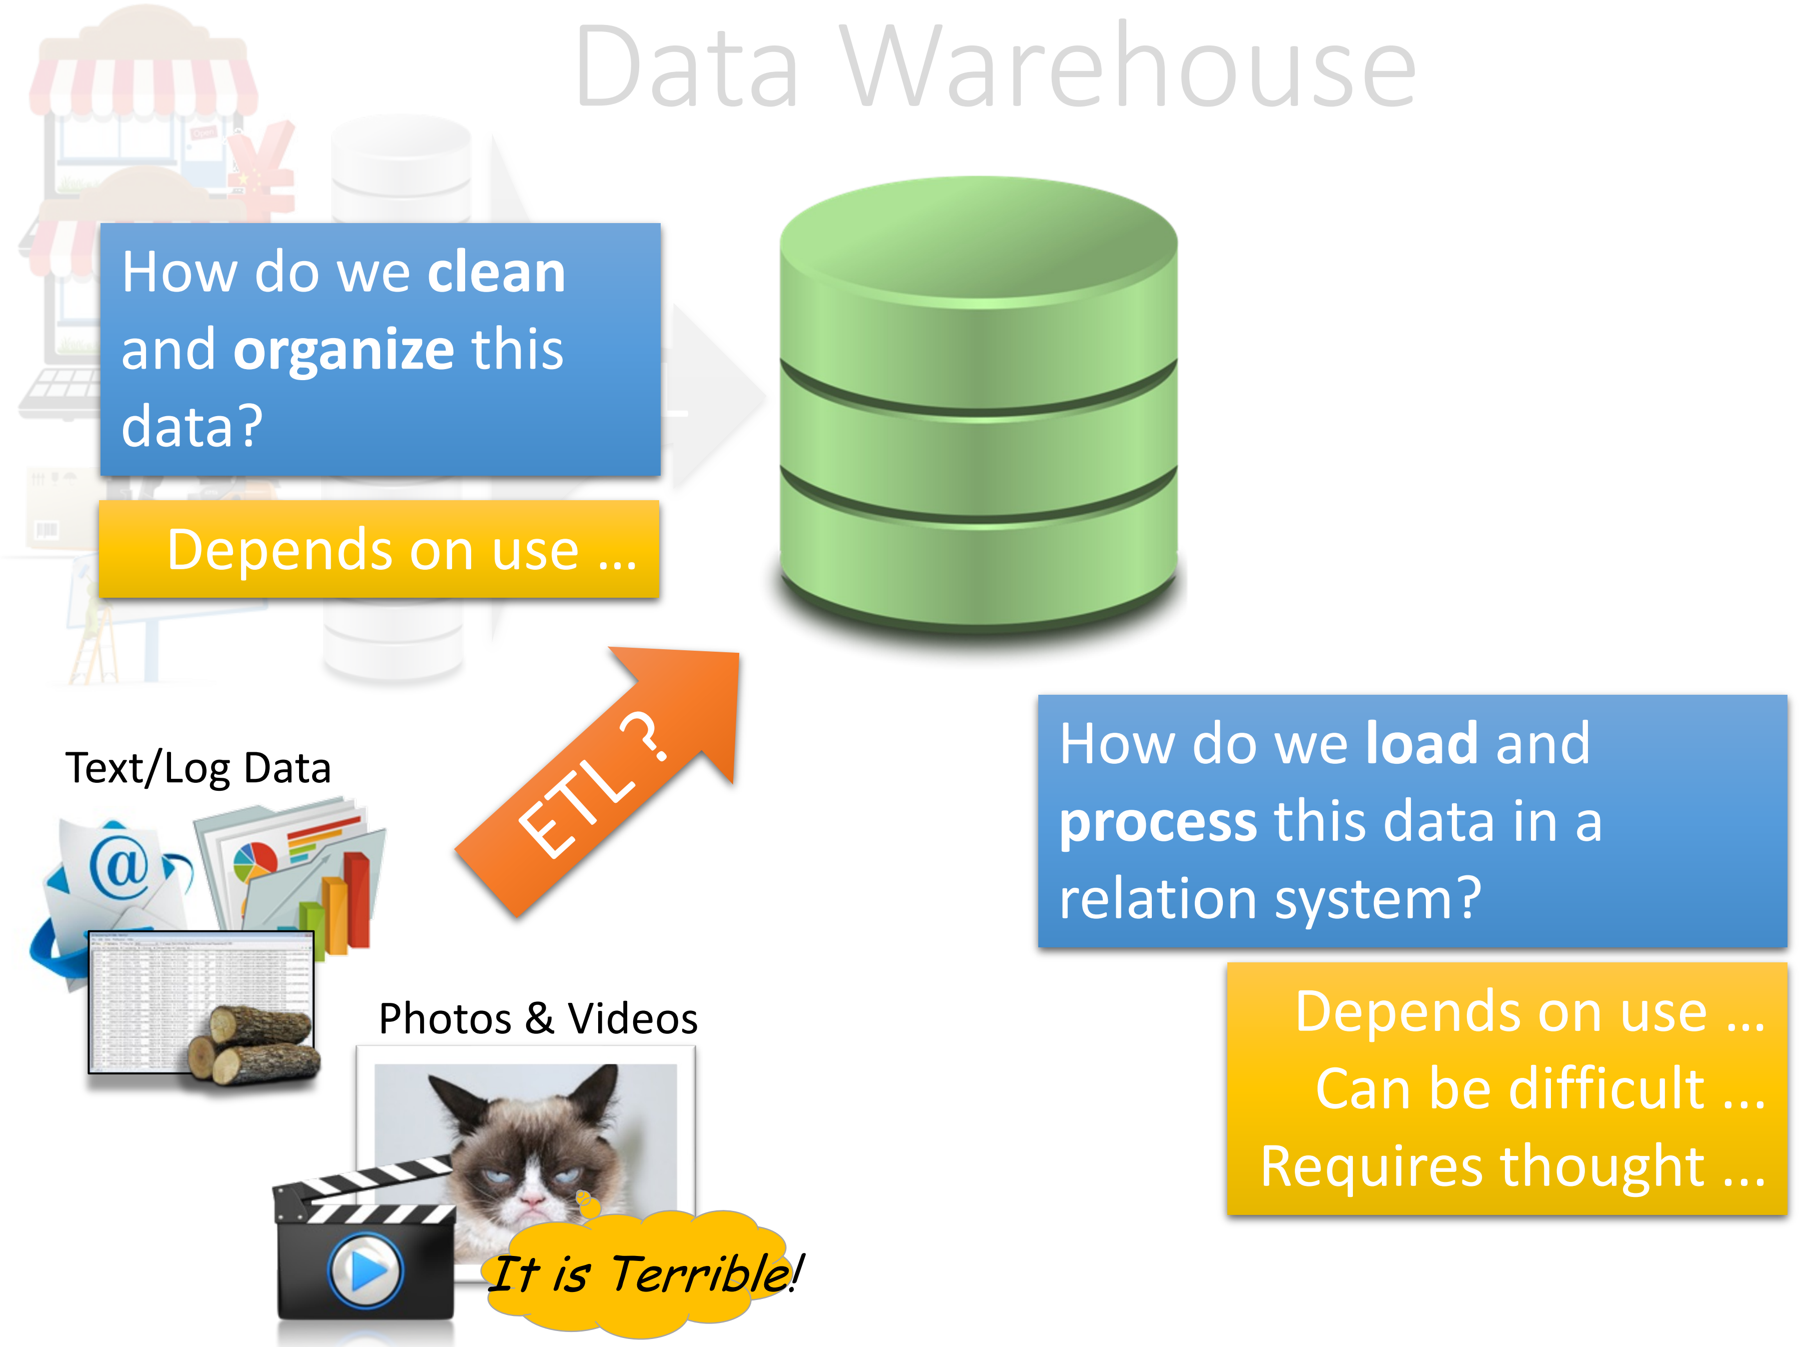

How do we **clean** and **organize** this data?

Depends on use …

How do we **load** and **process** this data in a relation system?

Depends on use …
Can be difficult ...
Requires thought ...

Data Lake*

*Still being defined…

[Buzzword Disclaimer]

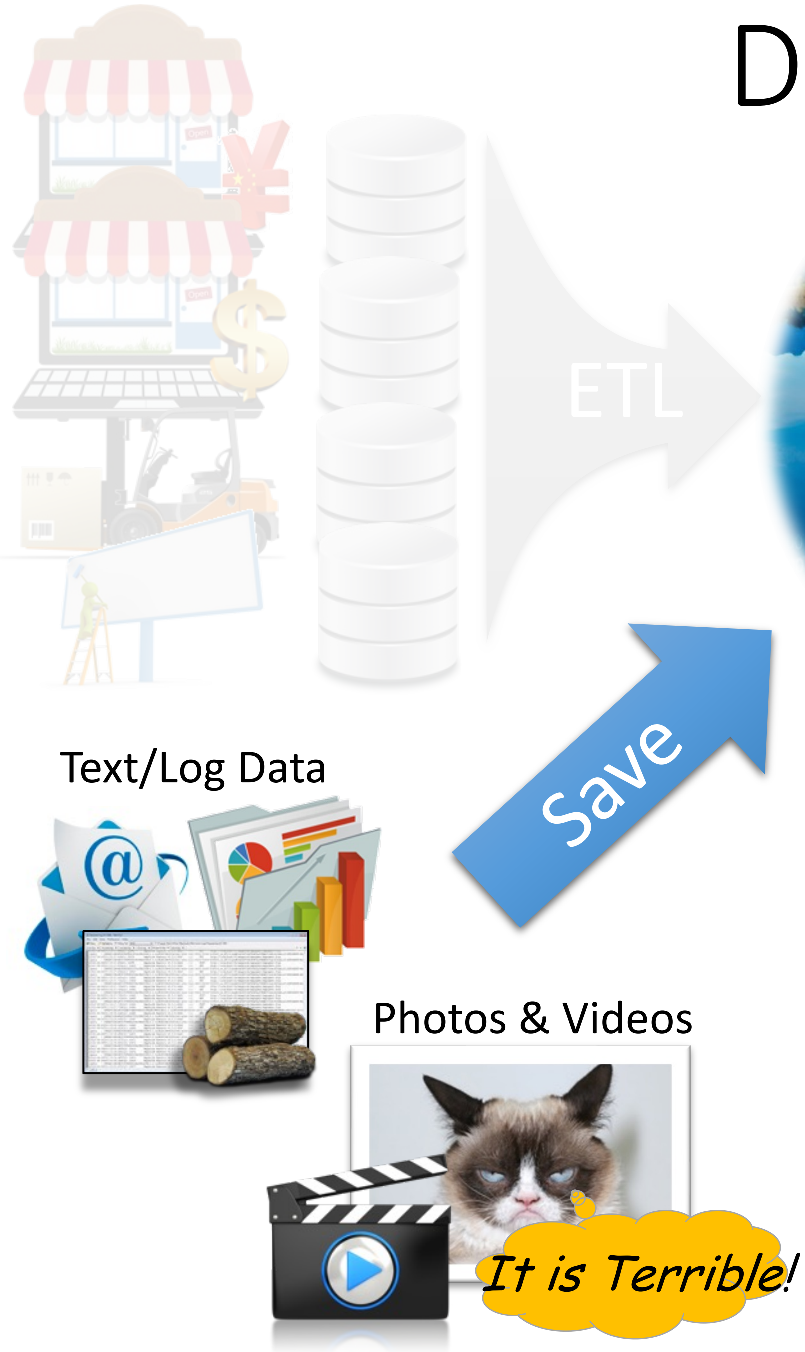

Big Idea:

Maintain a copy of all the data in one place and *free** data consumers to choose how to transform and use it.

*free to solve all the problems themselves

Origin of the Data Lake

- Attributed to James Dixon, CTO of Pentaho, 2010

“If you think of a **datamart** as a store of bottled water – **cleansed** and **packaged** and **structured** for **easy consumption** – the **data lake** is a **large body of water** in a more **natural state**.

The contents of the data lake **stream in** from a source to fill the lake, and various users of the lake can come to examine, dive in, or **take samples**.”

https://jamesdixon.wordpress.com/2010/10/14/pentaho-hadoop-and-data-lakes/

Data Lake

- Store unstructured data in **raw form**
  - Schema-on-**Read**: *determine the best organization **when data is used***
  - **Contrast:** Data Warehouses are Schema-on-Load (ET***L***)
    - Plan ahead (Fact tables and Dimensions)
- Often much **larger** than data warehouses
- Technologies
  - **Storage:** Large distributed file systems (e.g., HDFS)
    - Semi-structured formats (JSON, Parquet)
  - **Computation:** Map-Reduce
    - Recent trend to add SQL (or SQL like) functionality
- More Agile (?):
  - Don’t worry about schema & verification when loading
  - Disaggregated compute and storage → BYOF
    - bring your own compute frameworks …

- **What could go wrong?**

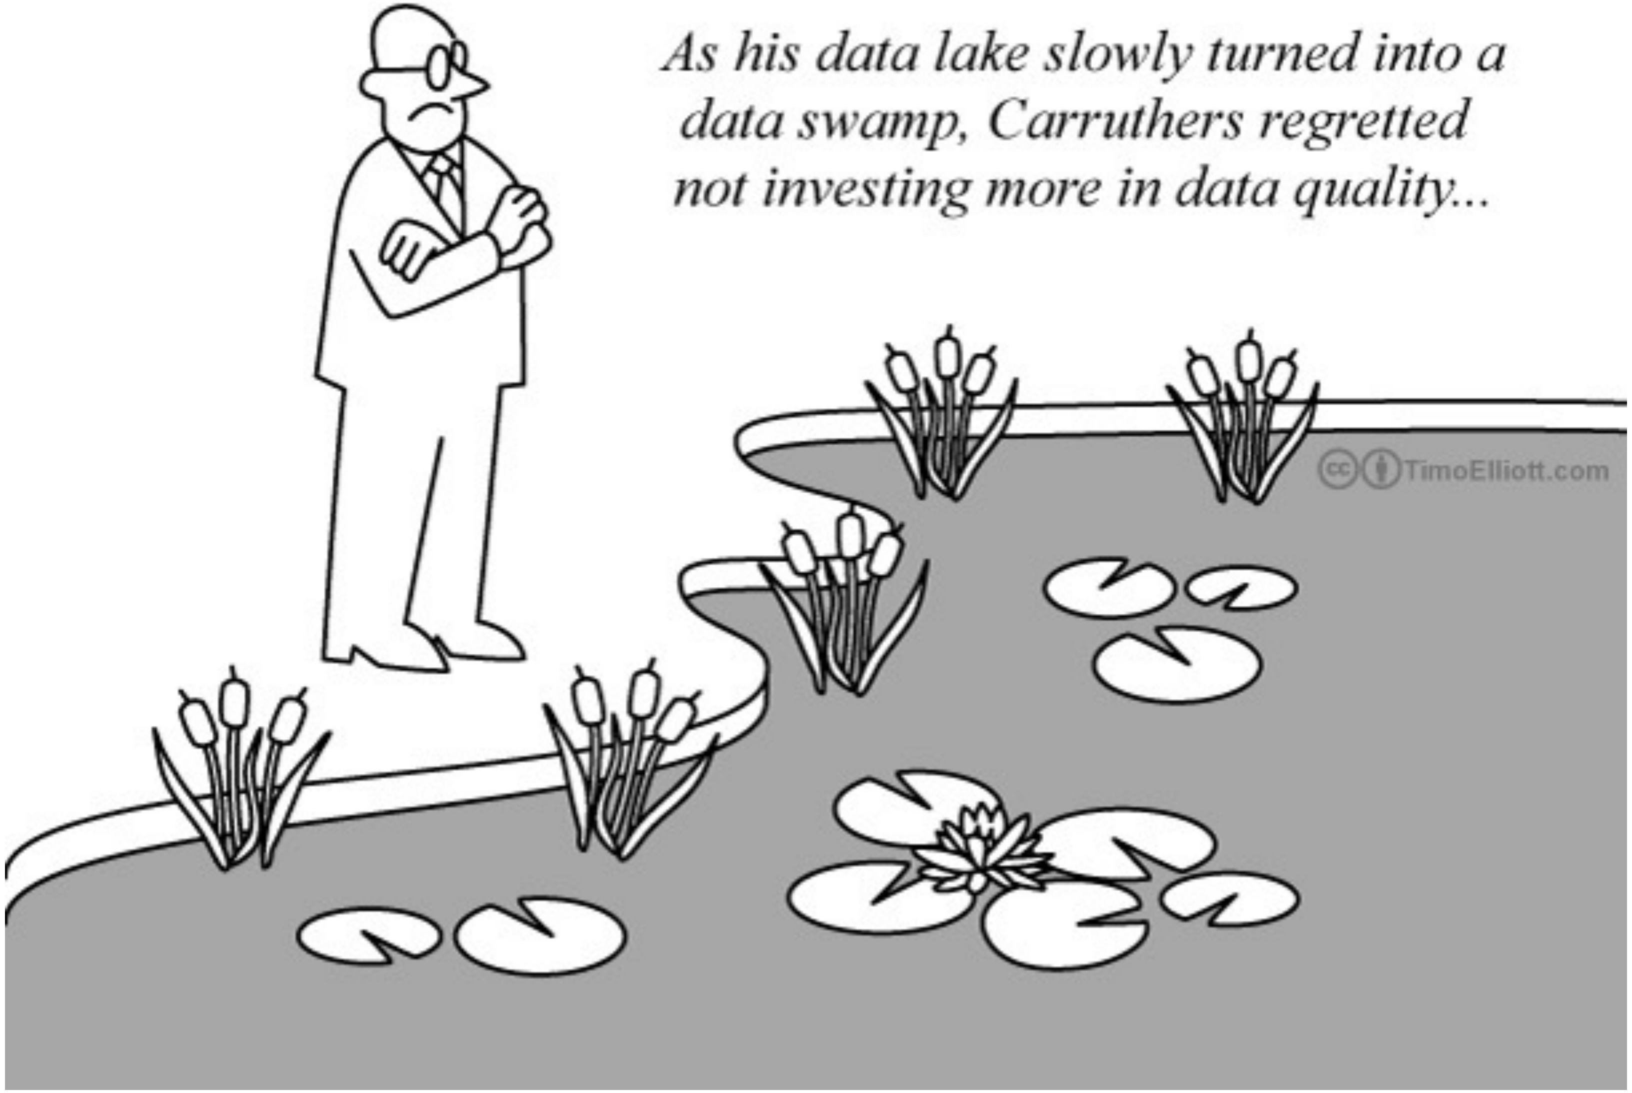
As his data lake slowly turned into a data swamp, Carruthers regretted not investing more in data quality...
TimoElliott.com

http://timoelliott.com/blog/2014/12/from-data-lakes-to-data-swamps.html

Data Lake → Data Swamp

- Cultural shift: *Curate* → Save **Everything**!
  - Signal to Noise ratio drops …
- Limited data governance → more agile →
  - **What** does it contain? **What** are all the **“fields”**
  - **When** and **how** and **from where** was it created
- Without cleaning and verification we begin to collect a rich history of **dirty data**
- Limited compatible with traditional tools

Data Lakes *Appear* to be Maturing

- Relational data-models + SQL:
    - **Hive:** SQL on top of Hadoop Map-Reduce
    - **SparkSQL:** SQL on top of Spark
- Tools are Improving:
    - Better data cleaning
    - Catalog Managers
    - Improved semi-structured “raw” data formats
- Improved data governance
    - Organization are recognizing the issues

Data Lake / Warehouse

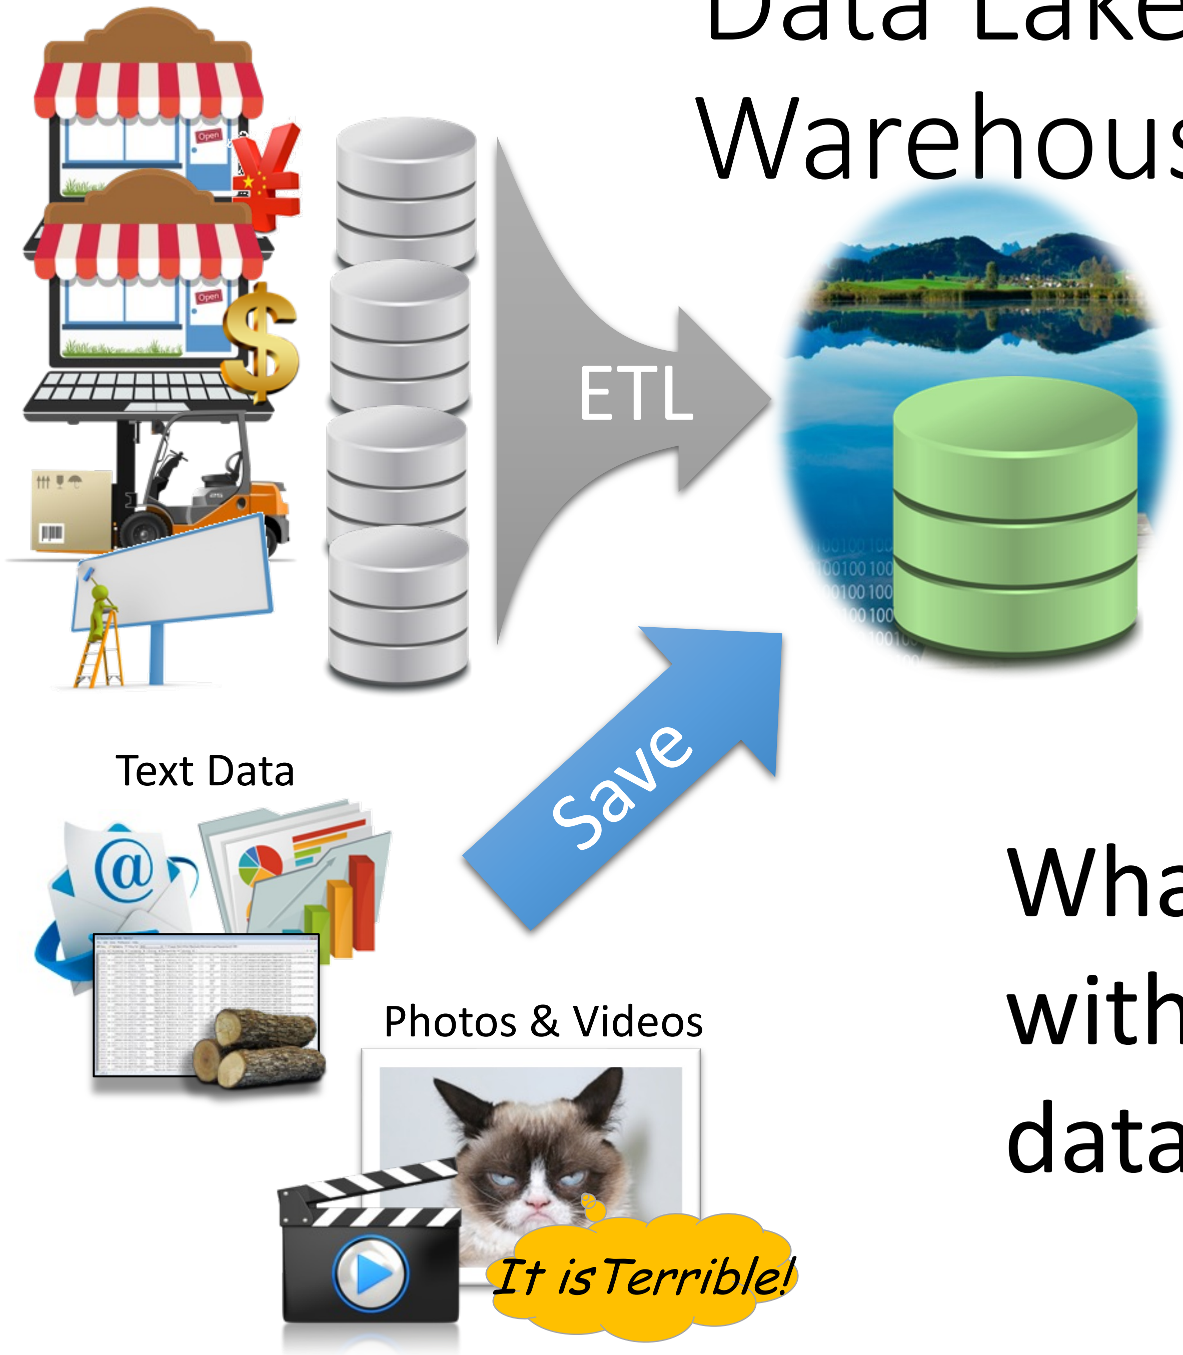

What do we do with all this data?

Data Lake / Warehouse

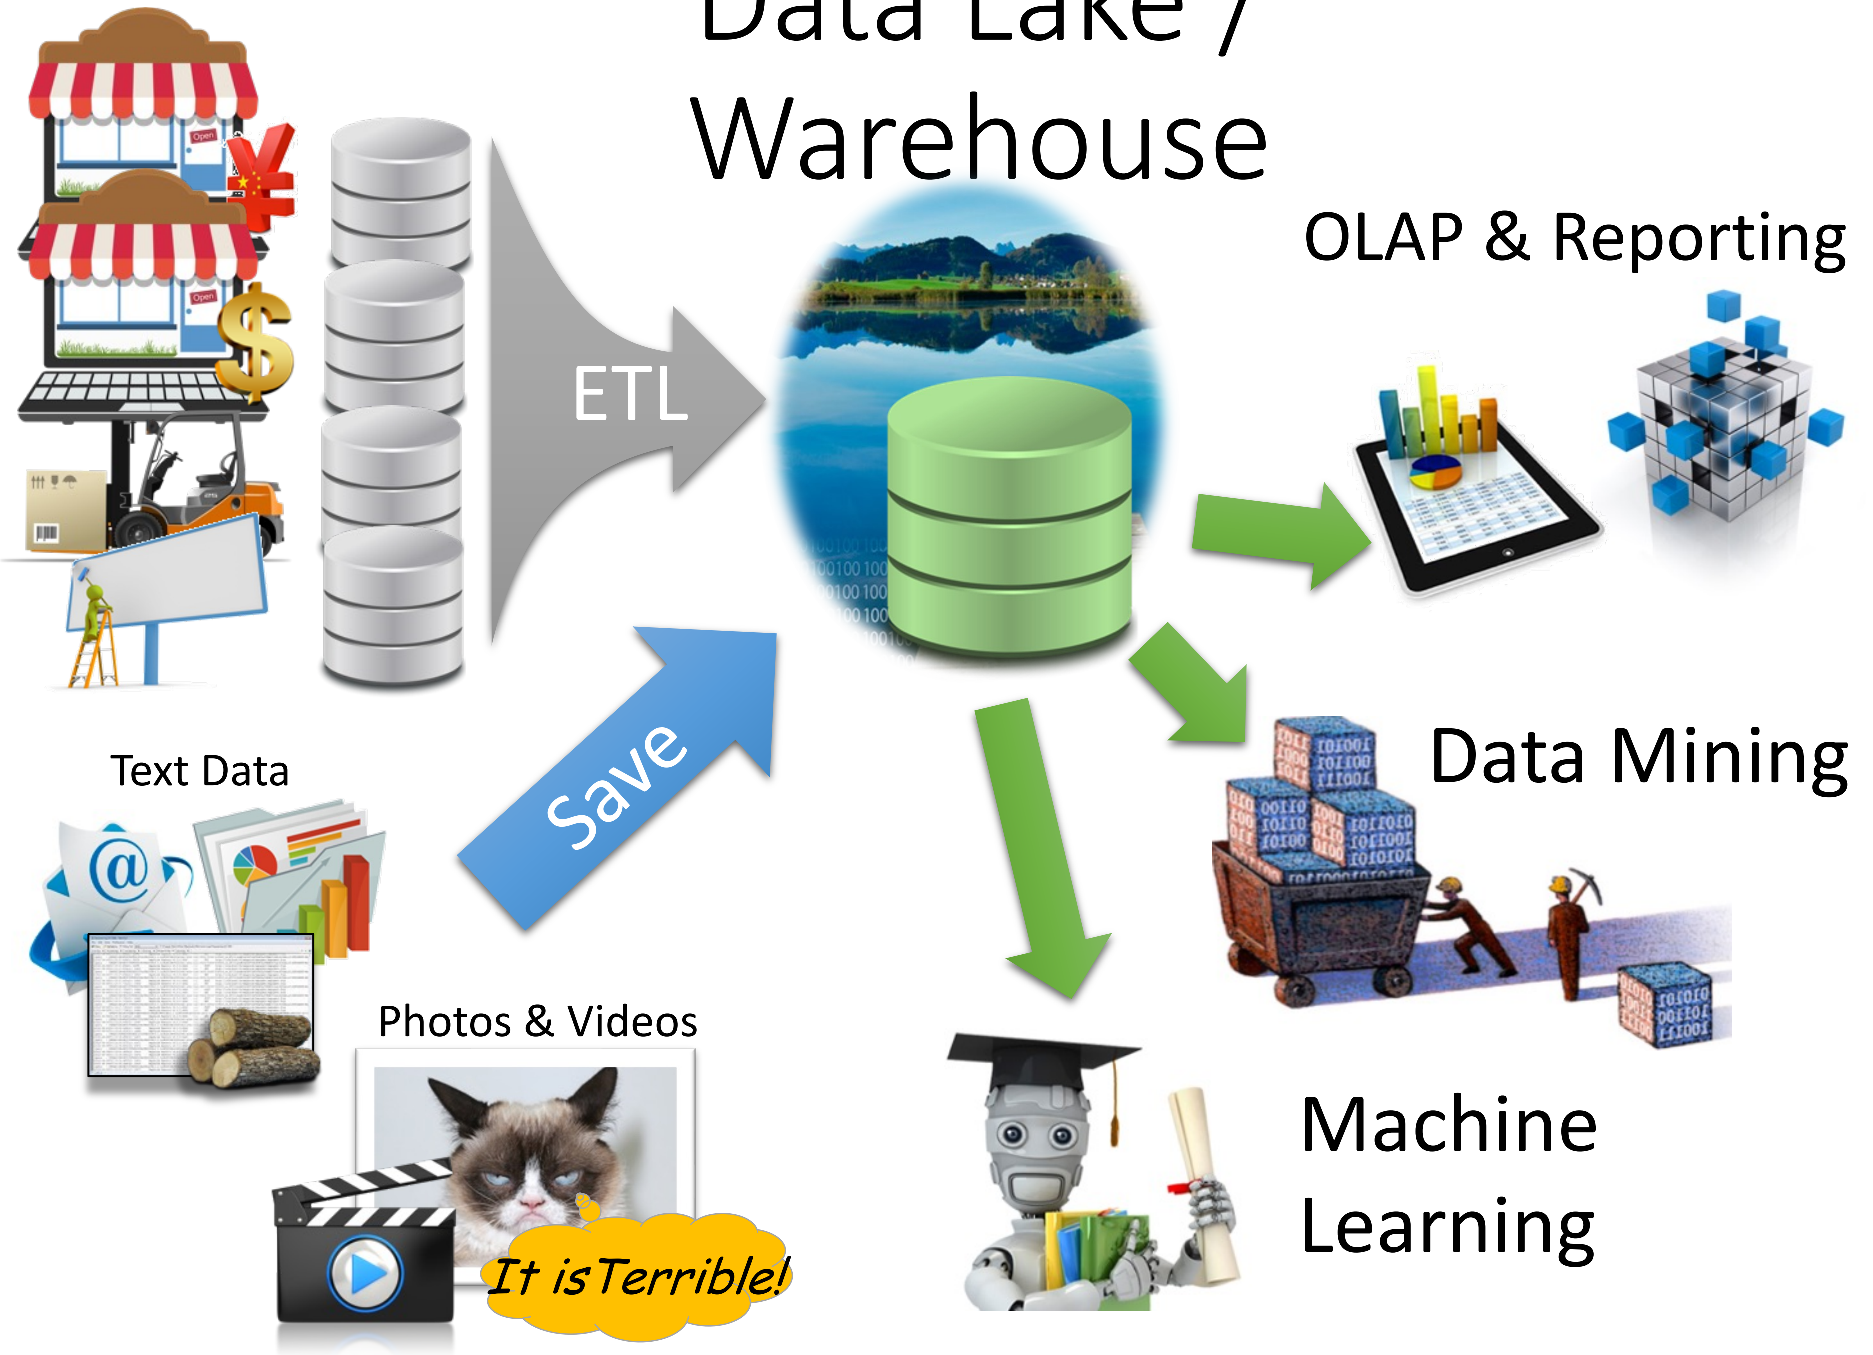

ETL

Save

Text Data

Photos & Videos

OLAP & Reporting

Data Mining

Machine Learning

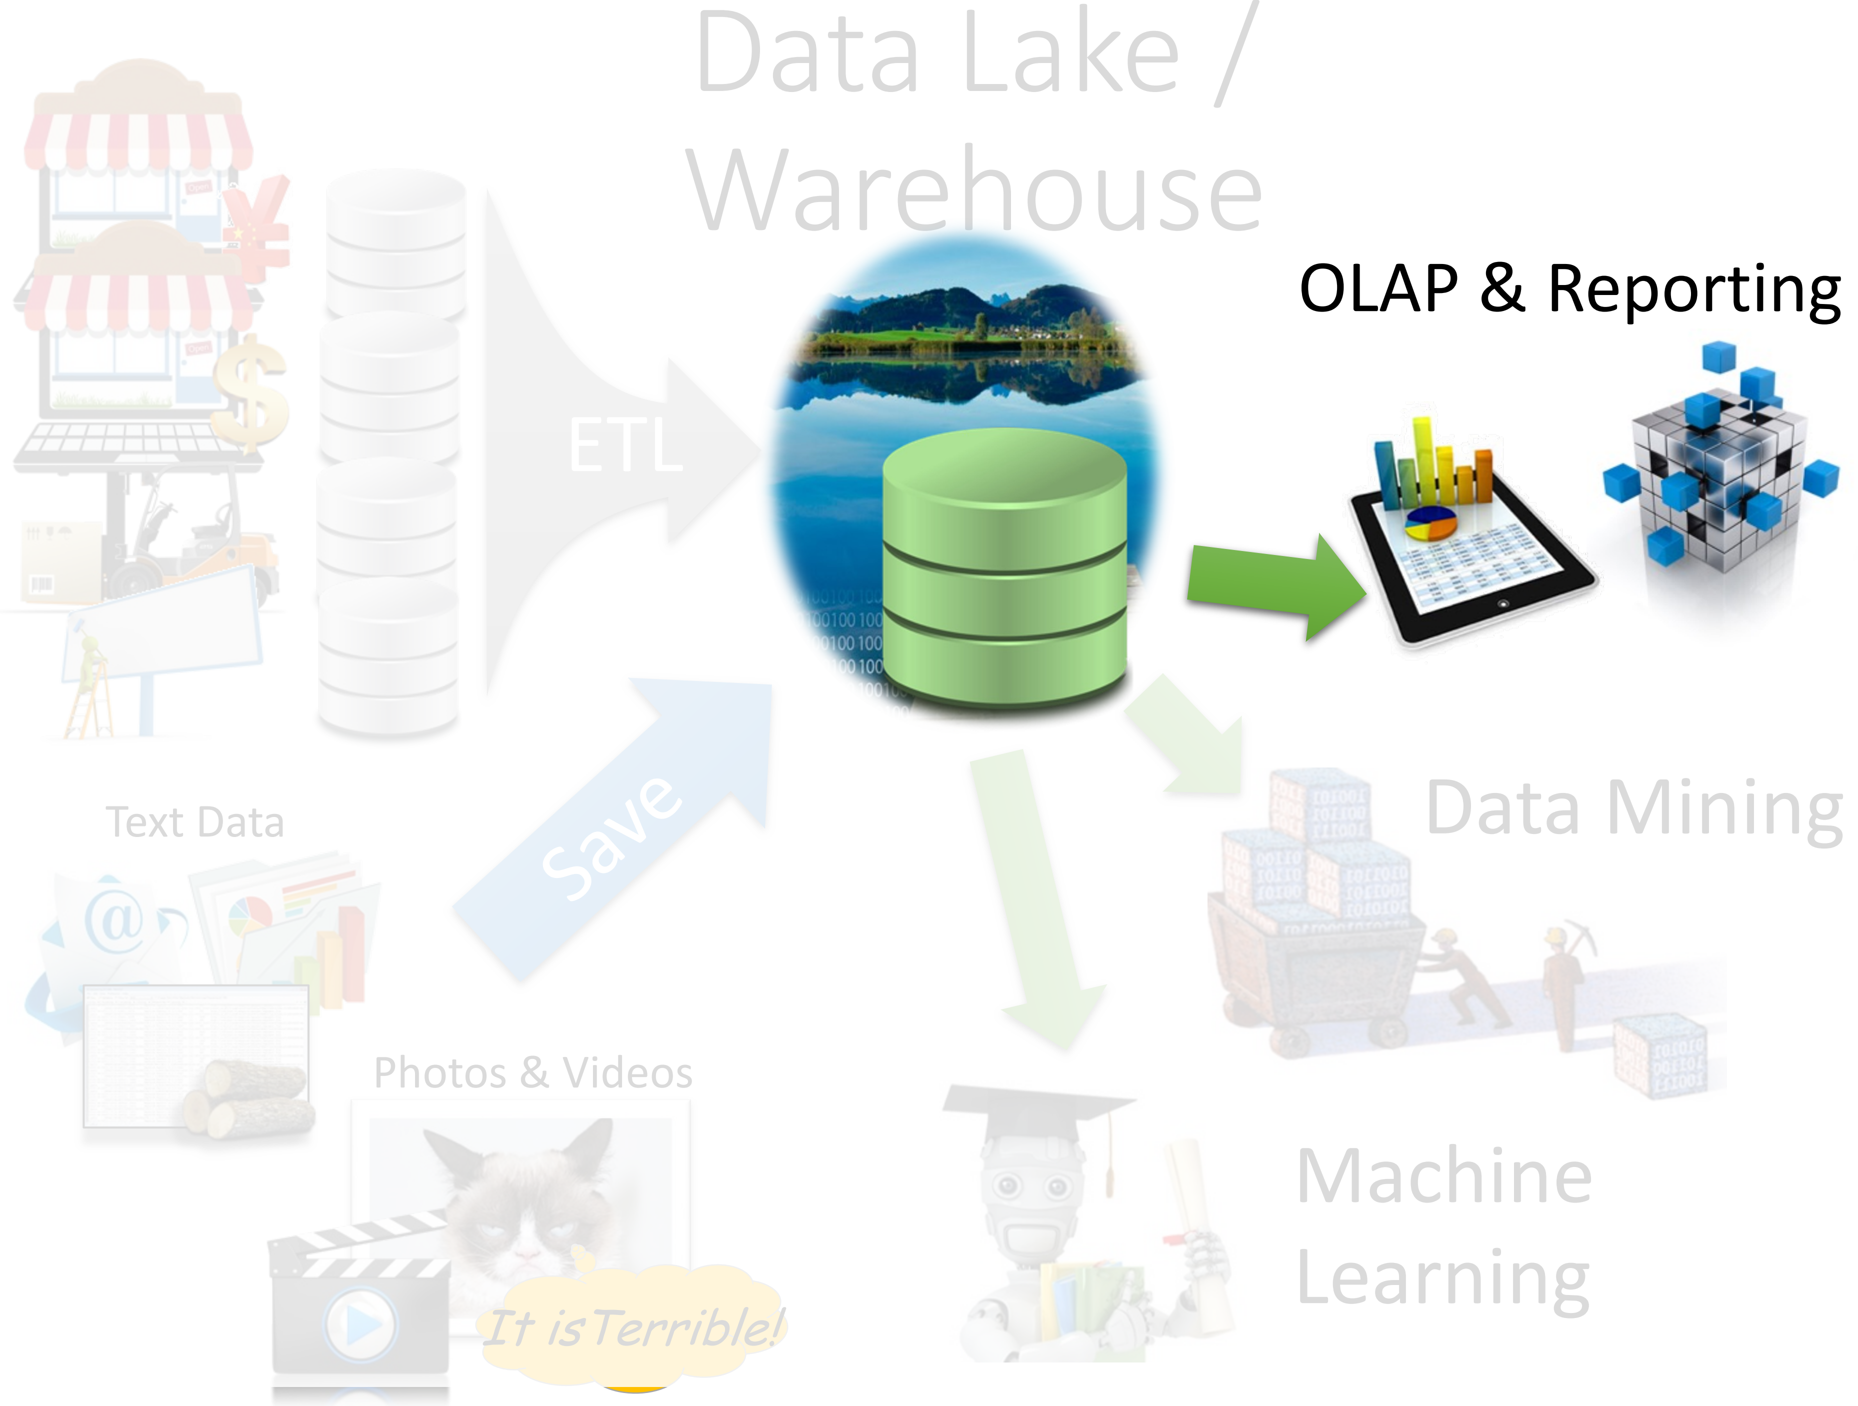

Data Lake / Warehouse

OLAP & Reporting

ETL

Save

Text Data

Photos & Videos

Data Mining

Machine Learning

Online Analytics Processing (OLAP)

Users interact with multidimensional data:

- Constructing ad-hoc and often complex SQL queries
- Using graphical tools that to construct queries
- Sharing views that summarize data across important dimensions

Cross Tabulation (Pivot Tables)

Item	Color	Quantity
Desk	Blue	2
Desk	Red	3
Sofa	Blue	4
Sofa	Red	5

		Item		
		Desk	Sofa	***Sum***
Color	Blue	2	4	***6***
	Red	3	5	***8***
	Sum	***5***	***9***	***14***

- Aggregate data across pairs of dimensions
  - **Pivot Tables:** *graphical interface* to select dimensions and aggregation function (e.g., SUM, MAX, MEAN)
  - **GROUP BY** queries

➢Related to contingency tables and marginalization in stats.

- What about many dimensions?

Cube Operator

- Generalizes cross-tabulation to higher dimensions.

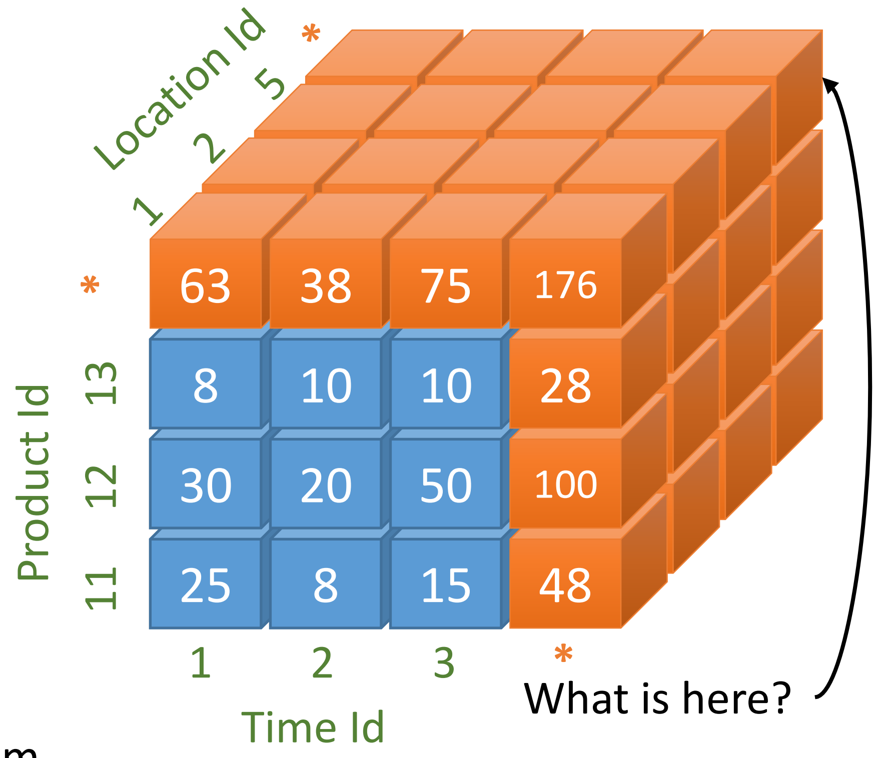

➢In SQL:

SELECT Item, Color, **SUM**(Quantity) **AS** QtySum
FROM Furniture
GROUP BY ***CUBE*** (Item, Color);

Item	Color	Quantity
Desk	Blue	2
Desk	Red	3
Sofa	Blue	4
Sofa	Red	5

Item	Color	QtySum
Desk	Blue	2
Desk	Red	3
Desk	*	5
Sofa	Blue	4
Sofa	Red	5
Sofa	*	9
*	*	14
*	Blue	6
*	Red	8

OLAP Queries

- **Slicing:** *selecting a value for a dimension*

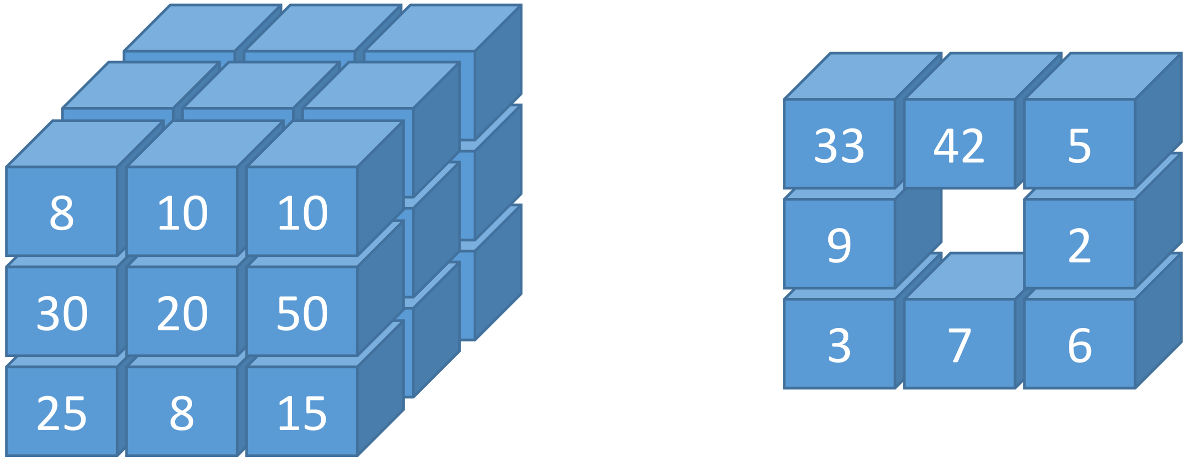

- **Dicing:** *selecting a range of values in multiple dimension*

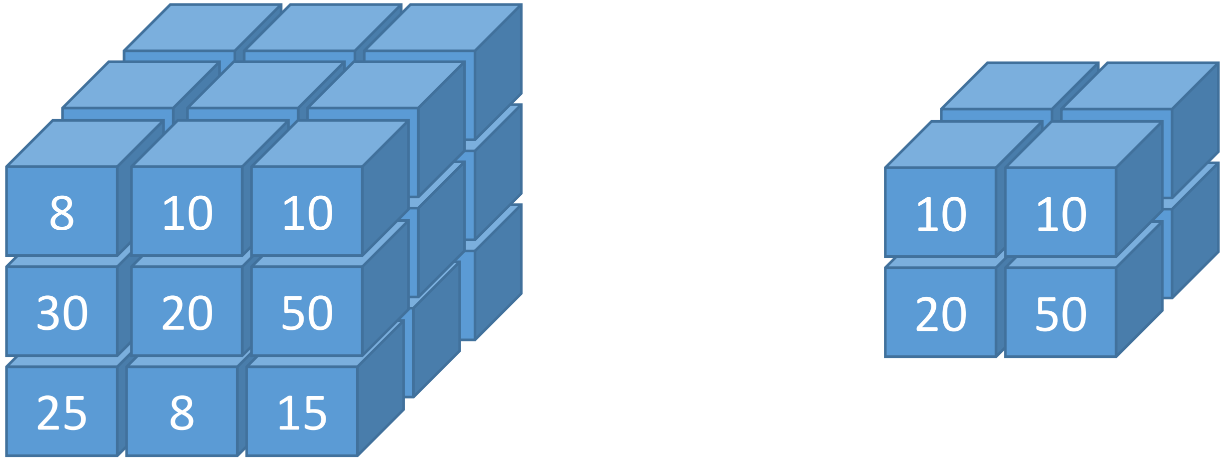

OLAP Queries

- **Rollup:** *Aggregating along a dimension*

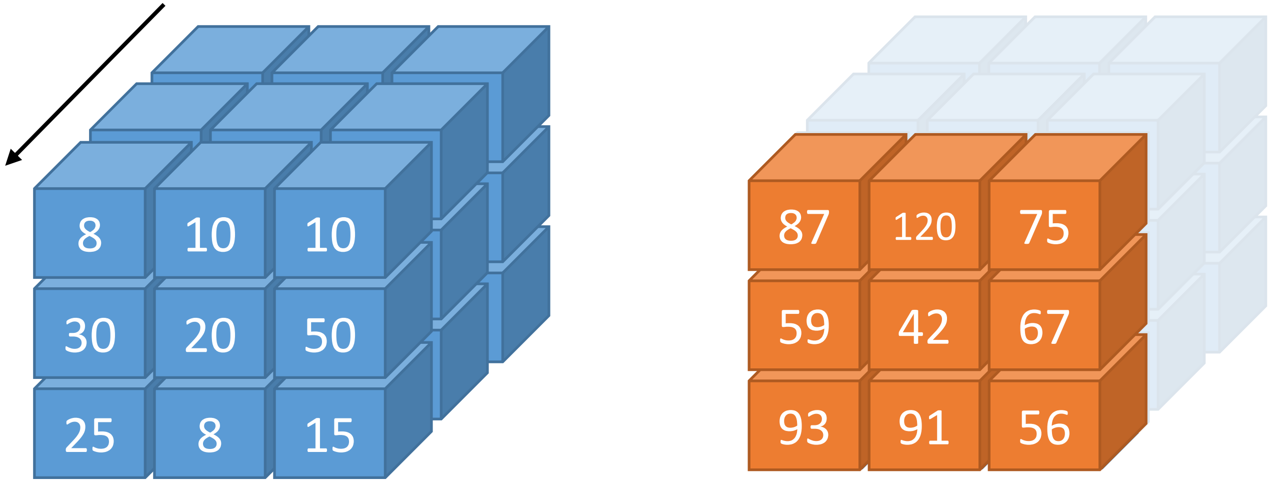

- **Drill-Down:** *de-aggregating along a dimension*

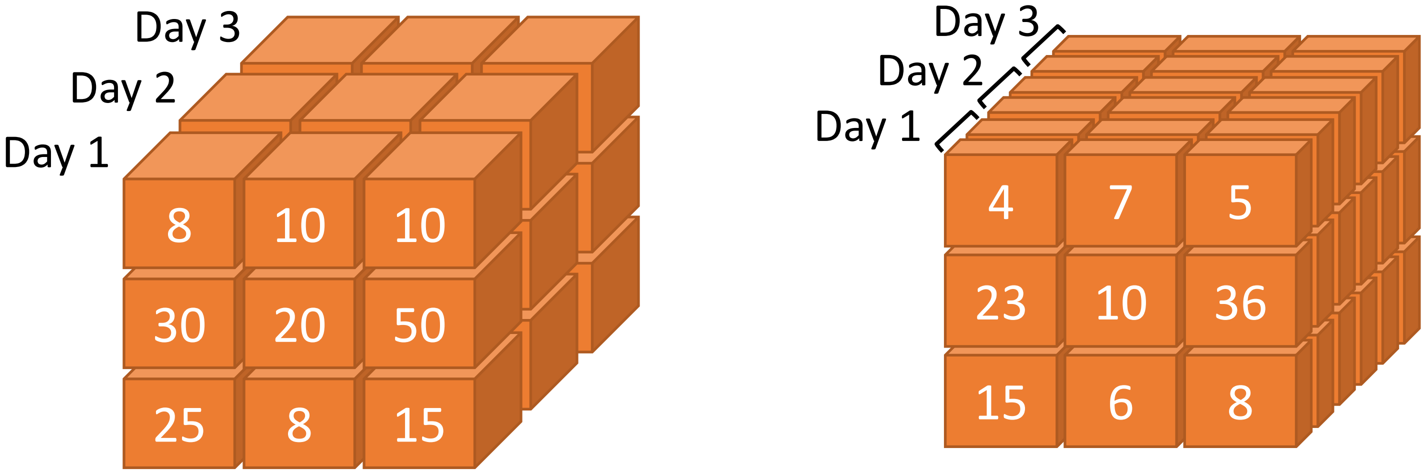

Reporting and Business Intelligence (BI)

- Use high-level tools to interact with their data:
  - Automatically generate SQL queries
    - Queries can get big!
- Common!

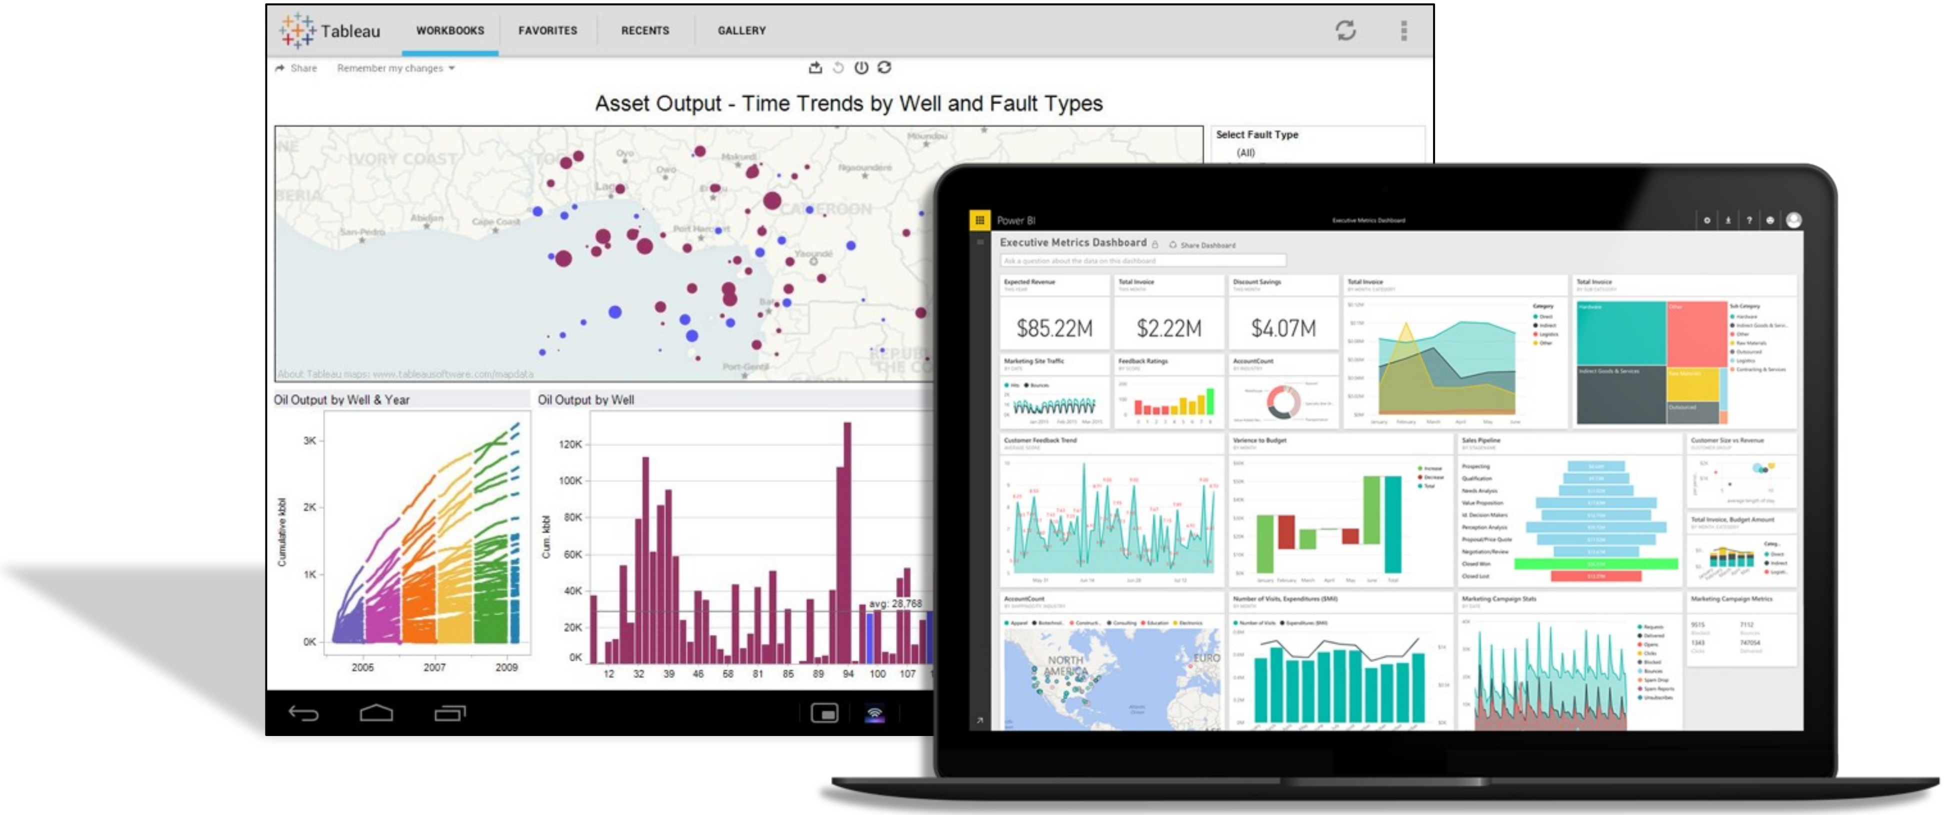

Data Lake / Warehouse

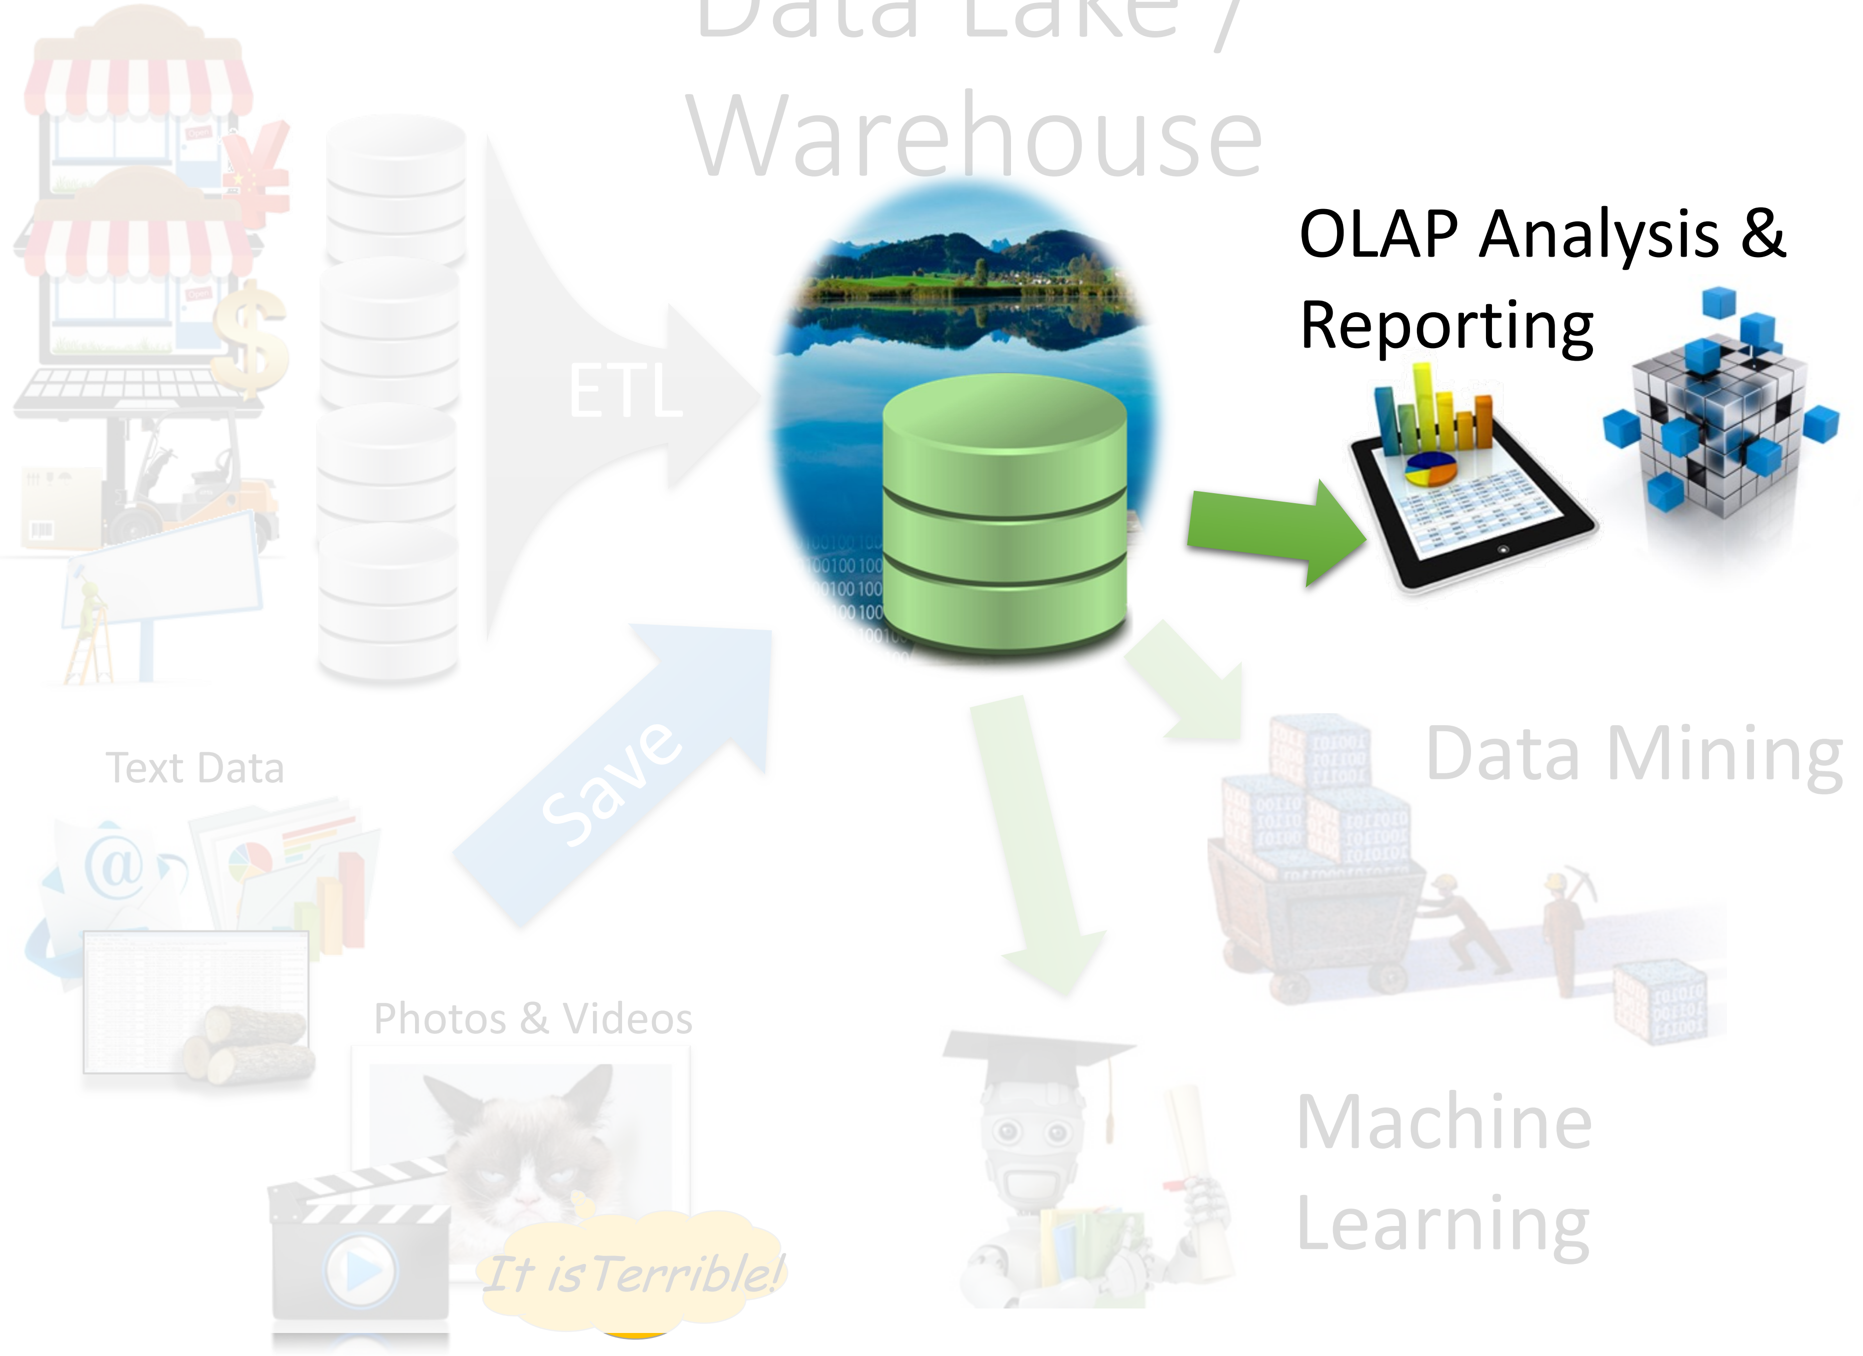

ETL

Save

Text Data

Photos & Videos

OLAP Analysis & Reporting

Data Mining

Machine Learning

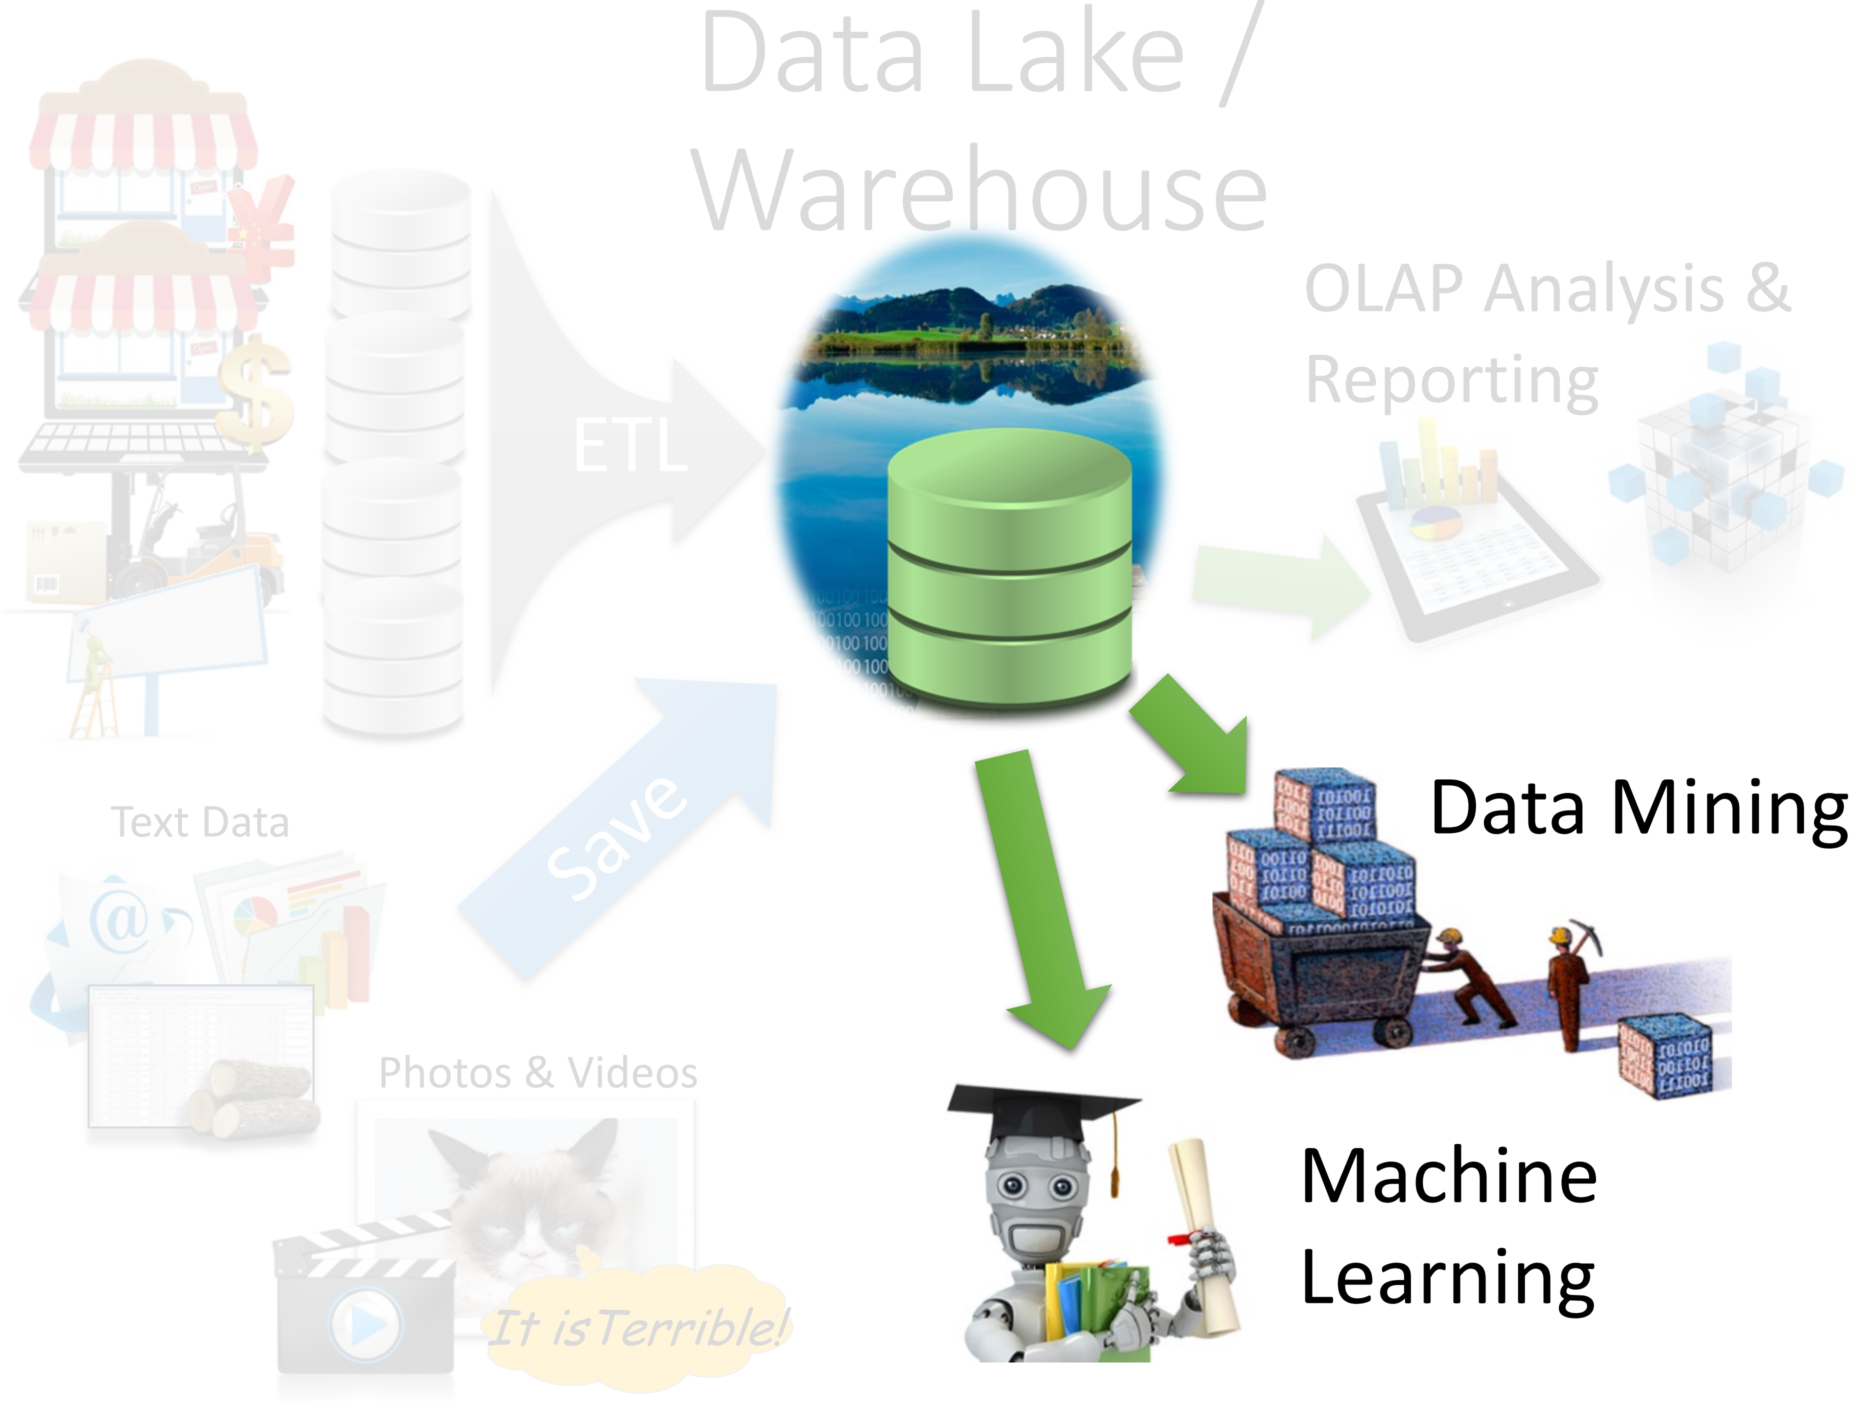

Data Lake / Warehouse

ETL

OLAP Analysis & Reporting

Text Data

Save

Photos & Videos

It is Terrible!

Data Mining

Machine Learning

Knowledge Discovery in Databases (KDD)

- Process of extracting ***knowledge*** from a ***data***
    - What does this mean?

Descriptive vs. Inferential Statistics

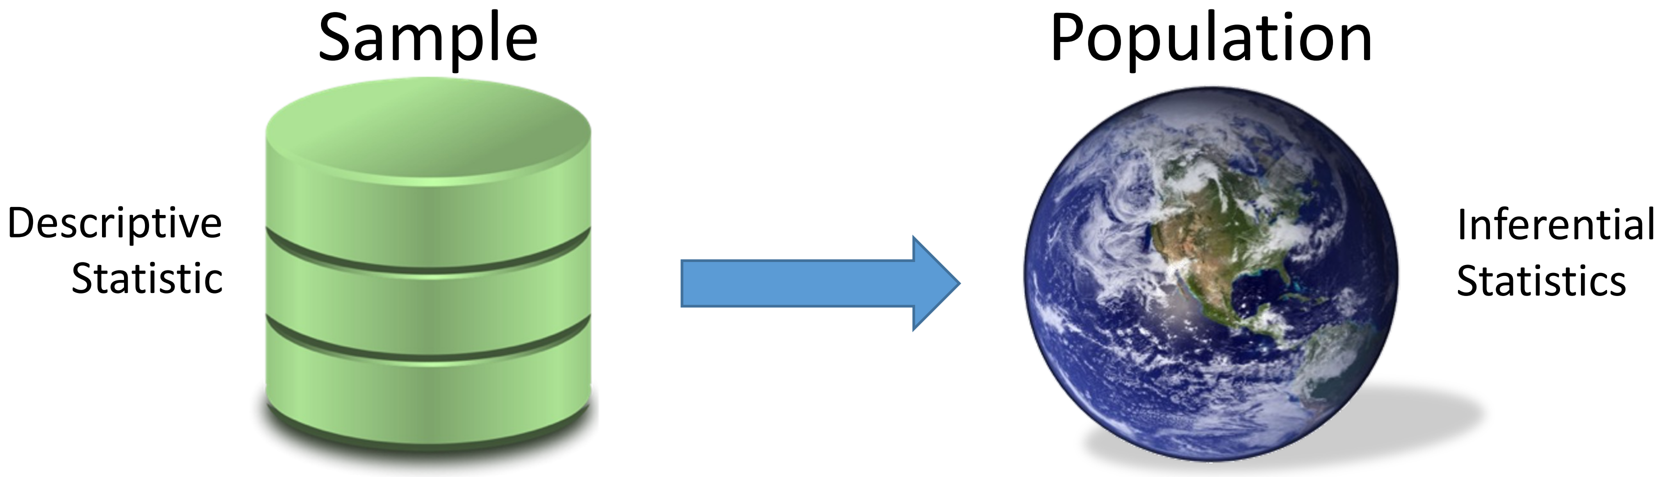

Sample

Population

Descriptive Statistic

Inferential Statistics

- **Descriptive Statistics:** *describe* the sample data
  - <u>Example:</u> *Average* sales last quarter
  - Can be **measured directly** from the database
- **Inferential Statistics:** *estimate* the population
  - <u>Example:</u> *Expected* sales next quarter
  - May be **estimated** using descriptive statistics

The Basic KDD Process

- **Data Selection:** *What data do I need for a given task?*
    - If data was already collected, how was the data collected?

- **Data Cleaning:** *Preparing the data for a given task*
    - Typically most challenging (time consuming) part.
    - Why might ETL not be enough?

- **Data Mining & ML:** *Running algorithms to infer patterns*
    - The fun part! Many tools, many options, complex tradeoffs.

- **Evaluation:** *Verifying that patterns are significant*
    - Algorithms will typically find patterns especially when none exist.

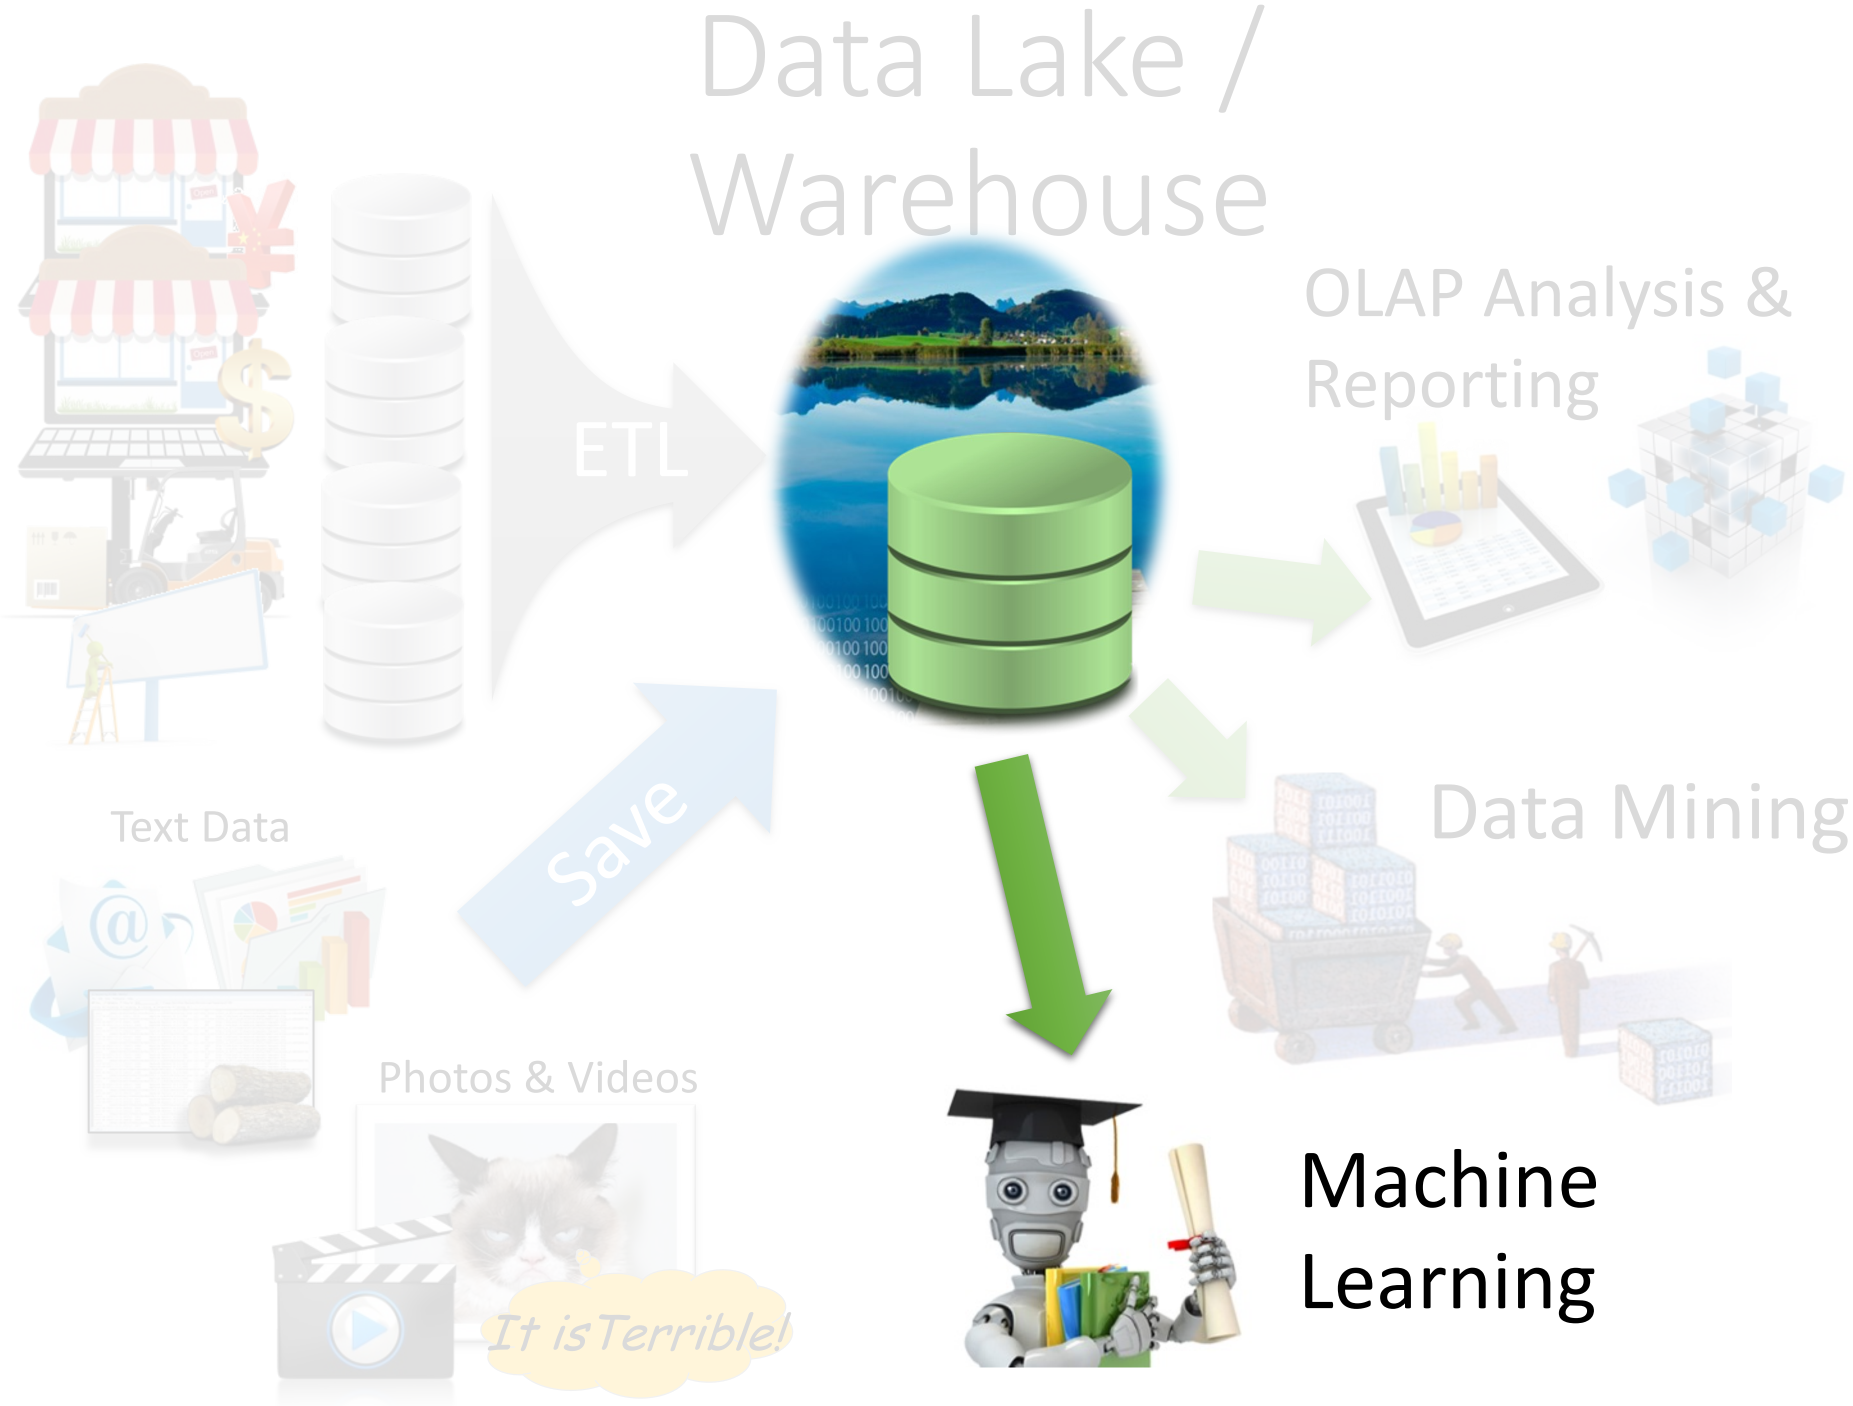

Data Lake / Warehouse

ETL

OLAP Analysis & Reporting

Data Mining

Text Data

Photos & Videos

Save

Machine Learning

What is Machine Learning?

Study of algorithms that:

- That improve their **performance**
  - Ability to understand what you are saying
- at some **task**
  - Voice recognition
- through **experience**
  - Transcribed speech data

-- Prof. Tom Mitchell, *CMU*

*“Machine Learning is the **second best** solution to any problem. The **first best** is of course to **solve the problem** directly.”*

-- Prof. Yaser S. Abu-Mostafa, *Caltech*

How would you write a program to recognize human speech?

You use ML every day!

What machine learning do you use every day?

- Spam detection
- Voice recognition
- Face tagging on Facebook
- Ad Targeting
- Credit card fraud detection
- Others? …

Machine Learning Lifecycle

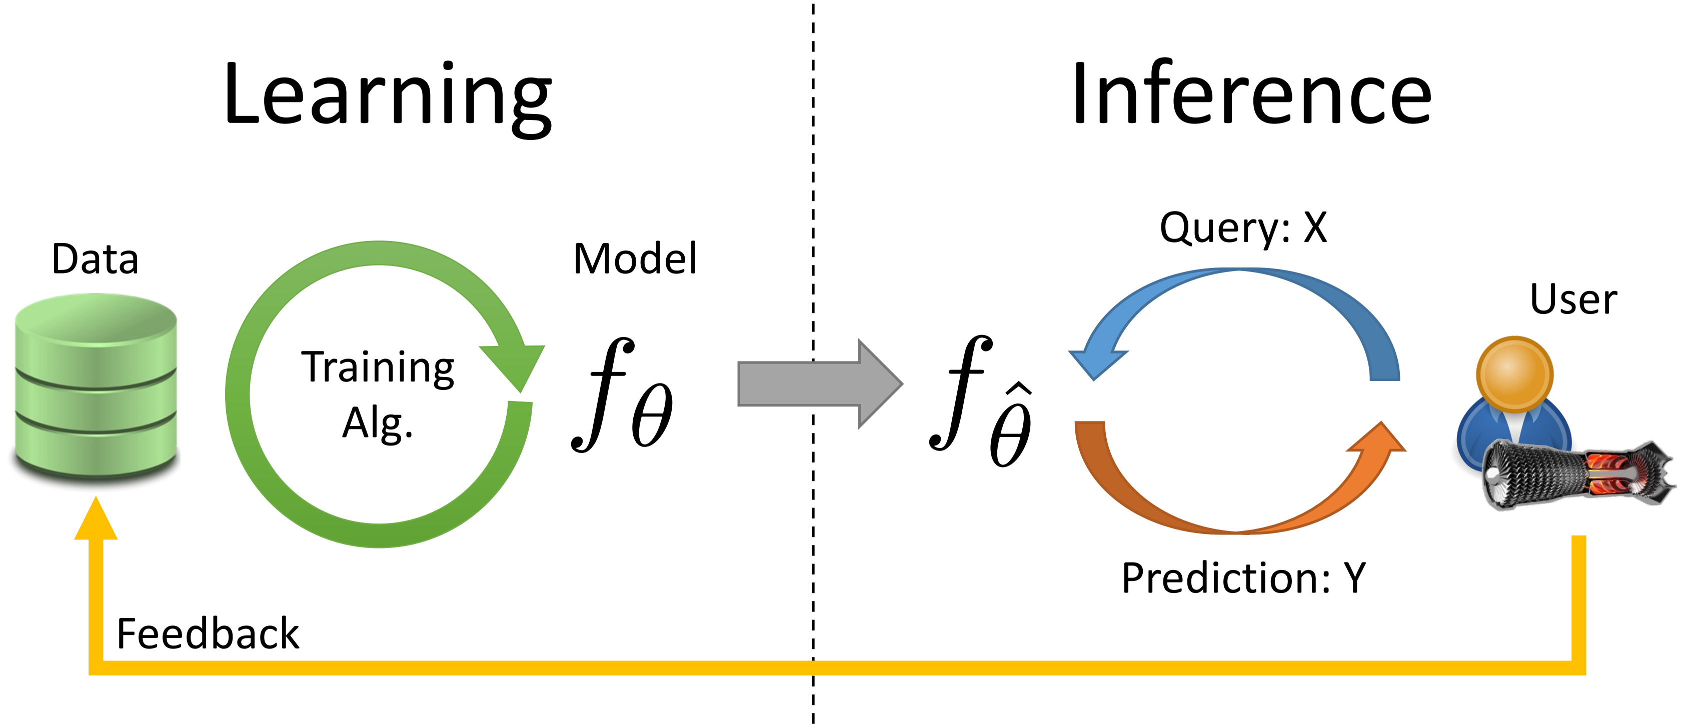

- Typically a time consuming iterative batch process
  - Feature engineering
  - Validation

➢ Focus is on making fast robust predictions

- Monitoring and tracking feedback
- Materialization + fast model inference

Learning: *Fitting the Model*

- **Training Data**
  - **X**: Features
  - **Y**: Label/Obs.
- Learn a function that **generalizes** the relationship between *X* and *Y*

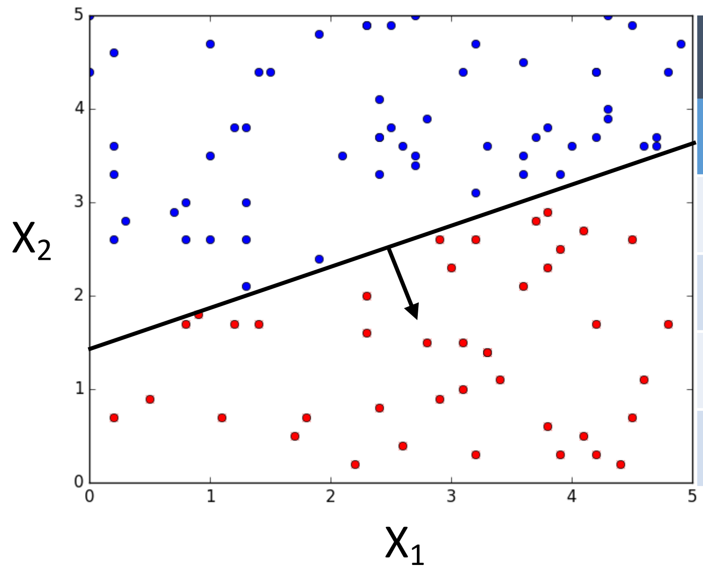

X		Y
x_1	x_2	
3.9	3.1	T
0.2	0.1	F
4.2	0.9	F
1.1	2.9	T

$$\min_{\theta} \|f_\theta(x) - Y\|_2^2$$

Features

Function class / Model Family

$$f_\theta(X) \rightarrow Y$$

Labels / Observations

Model Parameters

Finding the Best Parameters

$$f_\theta(X) \rightarrow Y$$

- Define some **objective** (e.g., prediction error)
- Search for best θ with respect to the objective

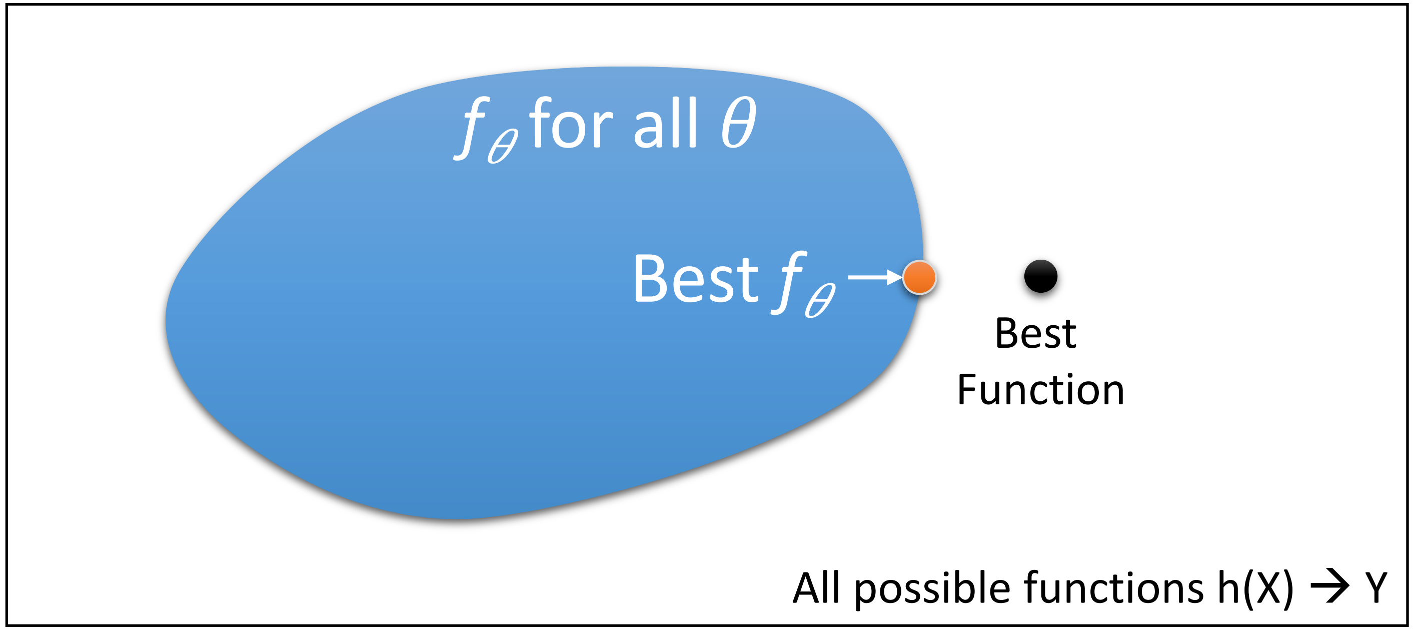

Generalization …

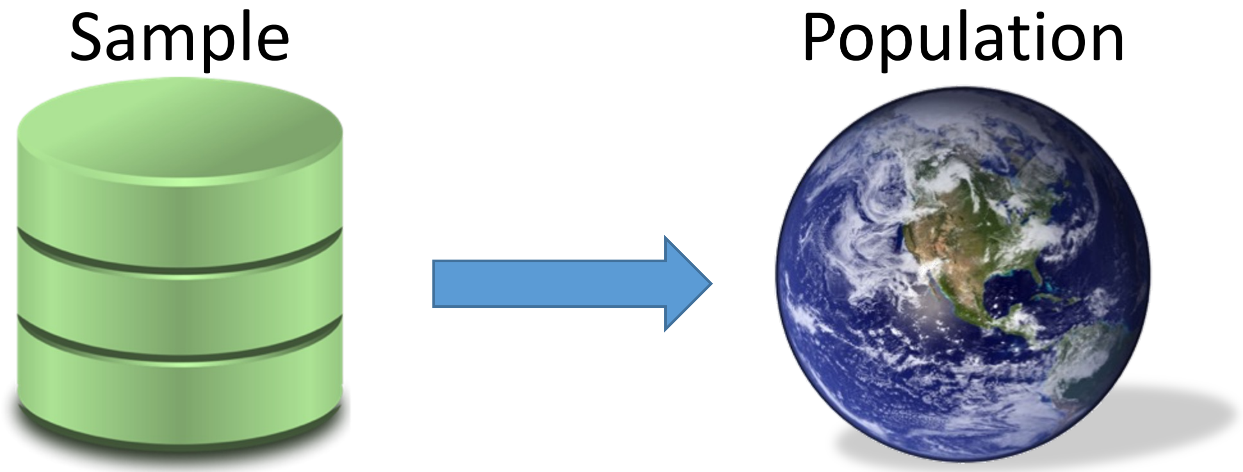

Sample

Population

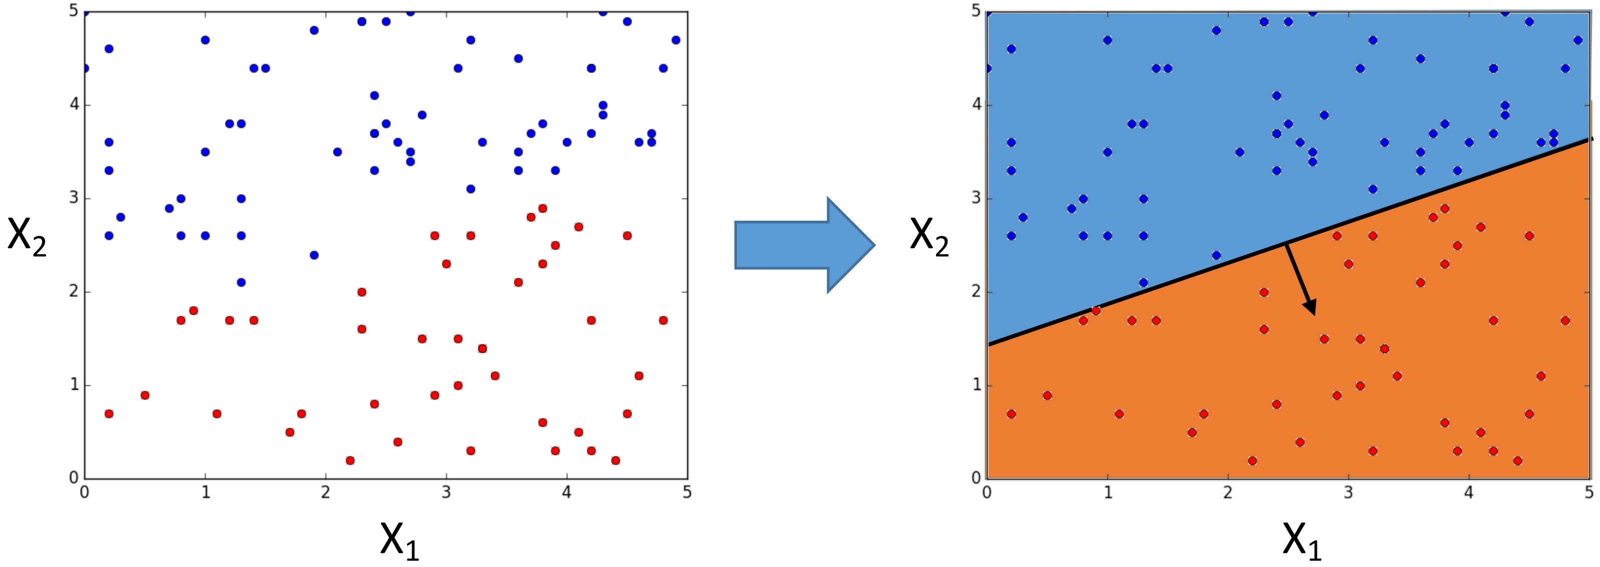

Inference: *Rendering Predictions*

- Evaluating the model on input queries:

$$f_{\hat{\theta}}(X) \to Y$$

- Online vs Offline:
  - Pre-computed **offline**: *movie rankings*
  - Computed **online** with each query: *speech recognition*
- May want to track confidence in prediction
- May require additional pre and post-processing
  - Feature lookup, content ranking, etc…

Feedback: *Incorporating New Data*

- After rendering a prediction we may get feedback on the results of the prediction:
  - **Explicit**: the *correct value* was “cat”
  - **Implicit**: the predicted animal was *incorrect*
  - Can be **noisy** …

- Watch out for **sample bias:**
  - Model affects the data is uses for training in the future
  - **Example**: only play top40 songs ...

Taxonomy

of Machine Learning

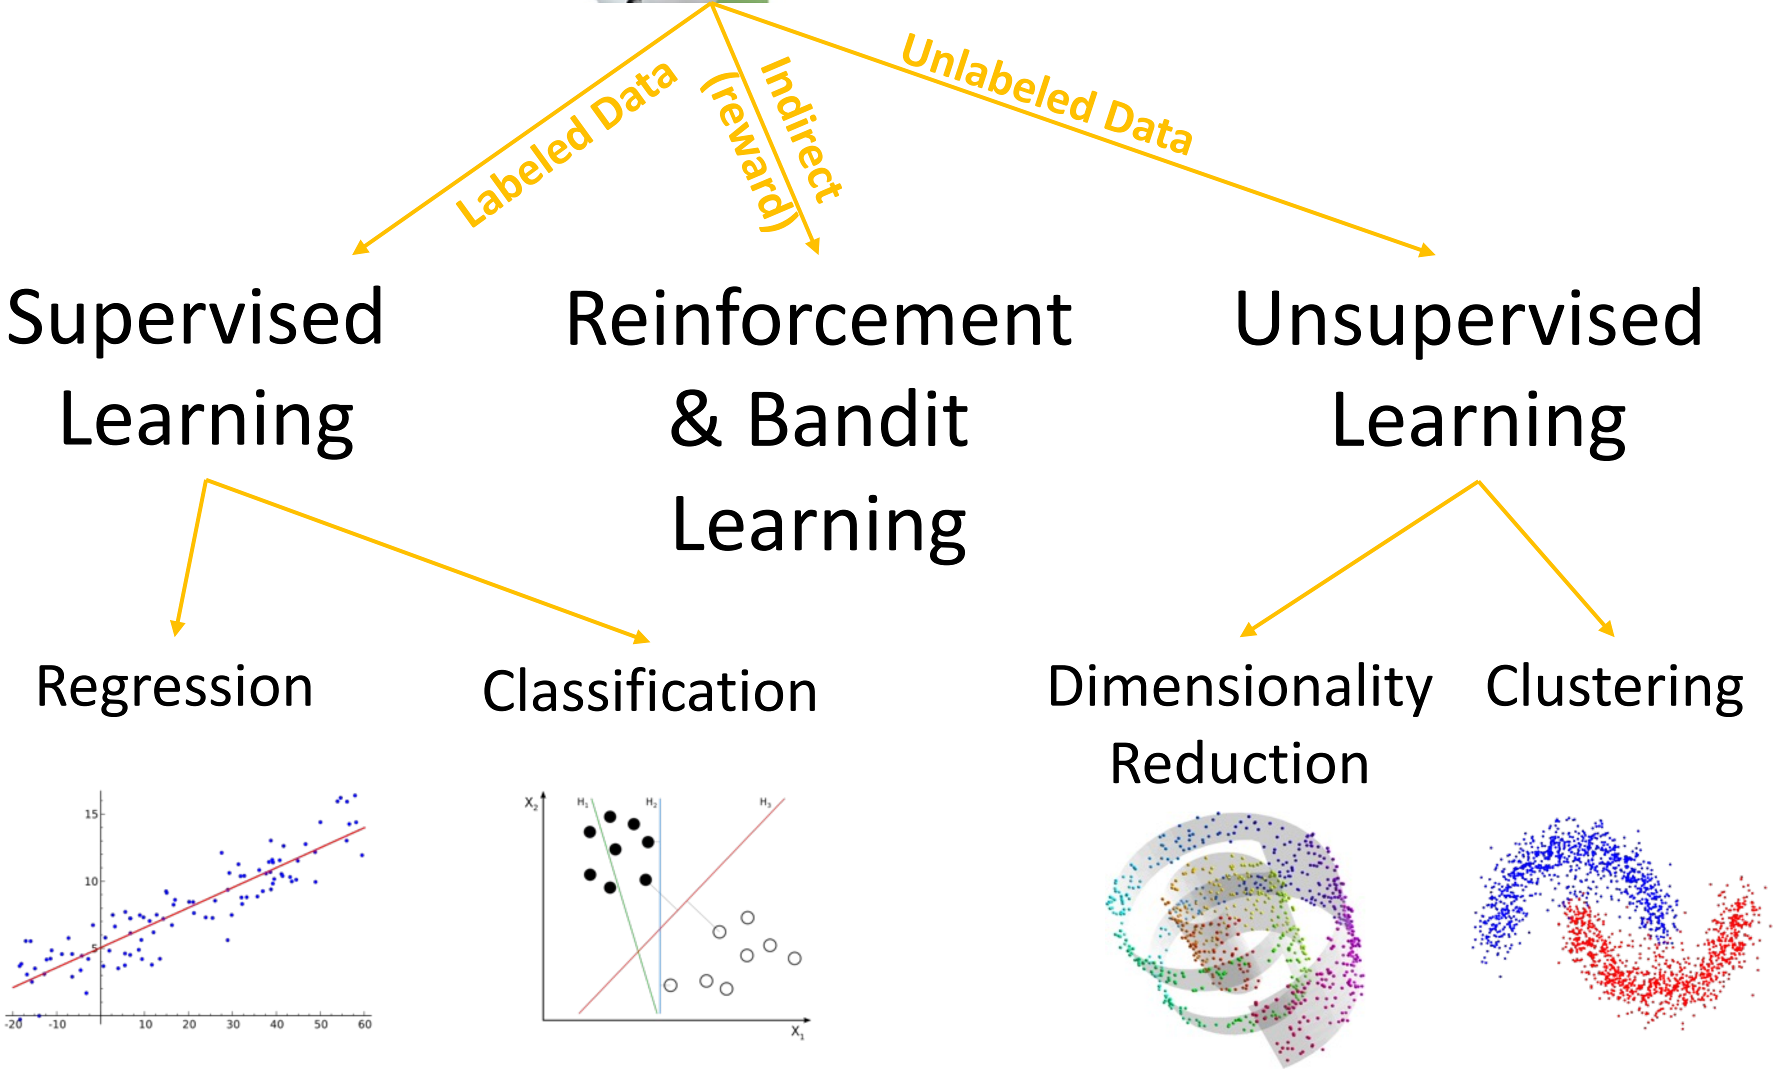

Labeled Data

Indirect (reward)

Unlabeled Data

Supervised Learning

Reinforcement & Bandit Learning

Unsupervised Learning

Regression

Classification

Dimensionality Reduction

Clustering