

CS240 Algorithm Design and Analysis
Fall 2024
Problem Set 4

Due: 23:59, Jan. 5, 2024

1. Submit your solutions to the course Gradescope.
2. If you want to submit a handwritten version, scan it clearly.
3. Late homeworks submitted within 24 hours of the due date will be marked down 25%. Homeworks submitted more than 24 hours after the due date will not be accepted unless there is a valid reason, such as a medical or family emergency.
4. You are required to follow ShanghaiTech's academic honesty policies. You are allowed to discuss problems with other students, but you must write up your solutions by yourselves. You are not allowed to copy materials from other students or from online or published resources. Violating academic honesty can result in serious penalties.

Problem 1:

Once upon a time, in a magical kingdom, there was a data structure that needed to be maintained by a group of skilled wizards. They performed a sequence of n operations on this data structure. The cost of each operation was determined by a mysterious rule: if the operation number was an exact power of 2 (like 1, 2, 4, 8, ...), the cost was equal to the operation number itself; otherwise, the cost was just 1.

The king of the kingdom asked the wizards to calculate the average cost of performing these operations, also known as the amortized cost per operation. You as the wizards of the kingdom, please use aggregate analysis or accounting methods to determine the amortized cost per operation.

Problem 2:

Suppose a shipping company needs to process packages in a queue-like manner, but they only have access to stack-based storage units. One way to simulate a queue is to use two stacks S_1 and S_2 . When a new package arrives, it is added to stack S_1 . To process a package (dequeue), the system first checks if stack S_2 is empty. If it is, all packages from S_1 are transferred to S_2 one by one (this is called a "dump"). Then, pop from S_2 .

For example, if the company processes the following operations: $\text{INSERT}(\mathbf{a})$, $\text{INSERT}(\mathbf{b})$, $\text{DELETE}()$, the results are:

Operation	Stacks State	
	$S_1 = []$	$S_2 = []$
$\text{INSERT}(\mathbf{a})$	$S_1 = [\mathbf{a}]$	$S_2 = []$
$\text{INSERT}(\mathbf{b})$	$S_1 = [\mathbf{b}, \mathbf{a}]$	$S_2 = []$
$\text{DELETE}()$	$S_1 = []$	$S_2 = [\mathbf{a}, \mathbf{b}]$ ("dump")
	$S_1 = []$	$S_2 = [\mathbf{b}]$ ("pop" returns \mathbf{a})

Suppose each stack operation (push or pop) takes 1 unit of time. Transferring n packages from S_1 to S_2 (a "dump") requires $2n$ units of work: n pops and n pushes.

- Suppose that (starting with an empty system) the company performs 3 arrivals (insertions), 2 removals (deletions), 3 more arrivals, and 2 additional removals. What is the total cost of these 10 operations, and how many packages remain in each stack at the end?
- If a total of n arrivals and n removals are performed in some order, what is the maximum possible time for a single operation? Provide an example sequence of operations that achieves this maximum time, and specify which operation causes it.
- Suppose the company processes an arbitrary sequence of arrivals and removals, starting with an empty system. What is the amortized cost of each operation? Provide the tightest possible (non-asymptotic) bounds. Use the accounting method to justify your answer. In particular, determine how much to charge (x) for an arrival and (y) for a removal. Prove your answer.

Problem 3:

Suppose you have $2n$ balls and 2 bins. For each ball, you throw it randomly into one of the bins. Let X_1 and X_2 denote the number of balls in the two bins after this process. Prove that for any $\varepsilon > 0$, there is a constant $c > 0$ such that the probability $\Pr[X_1 - X_2 > c\sqrt{n}] \leq \varepsilon$.

Problem 4:

Suppose there is a basic random function `rand50()` that returns 0 or 1 with equal probability (i.e., each outcome has a 50% chance of occurring). We want to construct a new random function `Rand75()` using calls to `rand50()` only—and without using any additional library functions or floating-point arithmetic—such that `Rand75()` returns 1 with probability 75% and 0 with probability 25%. Moreover, our goal is to **minimize the number of calls** to `rand50()`.

1. Provide a specific randomized algorithm (i.e., show how to call `rand50()` and how to process the returned values) to achieve this goal.

2. Prove the correctness of your proposed algorithm by showing:

- The probability that the algorithm returns 1 is indeed 75%, and the probability that it returns 0 is 25%.
- The algorithm uses as few calls to `rand50()` as possible (or give a justification that the number of calls is near-optimal under these constraints).

Hint: You may consider operating on two independent random bits from `rand50()` using an appropriate bitwise operation (e.g., bitwise OR) along with a suitable procedure to achieve the desired probabilities.


Problem 5:

A delivery company needs to load n packages into trucks. Each truck has a maximum load capacity of C , and the weight of the n packages is denoted as $\{w_1, w_2, \dots, w_n\}$, where $w_i \leq C$ ($1 \leq i \leq n$).

Design a polynomial-time 2-approximation algorithm to solve this problem and prove the correctness of your algorithm.

Problem 6:

Consider an $n \times n$ grid graph G , as shown below.



Each node v in G has a weight $w(v) > 0$. You want to choose an independent set of nodes with maximum total weight. That is, you want to choose a set of nodes S with maximum total weight such that for any $v \in S$, none of v 's neighbors are in S . To do this, consider the following greedy algorithm. Let V be the set of all nodes in G . Choose the node in V with the largest weight (breaking ties arbitrarily), add it to the independent set, then remove the node and all its neighbors from V . Repeat this process until V is empty. Let S be the output of this algorithm. Solve the following problems.

1. Let T be any independent set in G . Show that for each node $v \in T$, either $v \in S$, or there is a neighbor v' of v with $v' \in S$ and $w(v) \leq w(v')$.
2. Show that the greedy algorithm is a 4-approximation.