



# CS240 Algorithm Design and Analysis

## Lecture 22

### Randomized algorithms (Cont.)

Quan Li  
Fall 2024  
2024.12.12



# Hash Tables





# Hash Tables




- A hash table is a randomized data structure to efficiently implement a dictionary.
- Supports find, insert, and delete operations all in expected  $O(1)$  time.
  - But in the worst case, all operations are  $O(n)$ .
  - The worst case is provably very unlikely to occur.
- A hash table does not support efficient min / max or predecessor / successor functions.
  - All these take  $O(n)$  time on average.
- A practical, efficient alternative to binary search trees if only find, insert and delete needed.



# Direct addressing



- Suppose we want to store (key, value) pairs, where keys come from a finite universe  $U = \{0, 1, \dots, m-1\}$ .
- Use an array of size  $m$ .
  - **insert**( $k, v$ ) Store  $v$  in array position  $k$ .
  - **find**( $k$ ) Return the value in array position  $k$ .
  - **delete**( $k$ ) Clear the value in array position  $k$ .
- All operations take  $O(1)$  time.
- The problem is, if  $m$  is large, then we need to use a lot of memory.
  - Uses  $O(|U|)$  space.
  - **Ex** For 32 bit keys, need 4 GB memory. For 64 bit keys, more memory than in world.
- If only need to store few values, lots of space wasted.



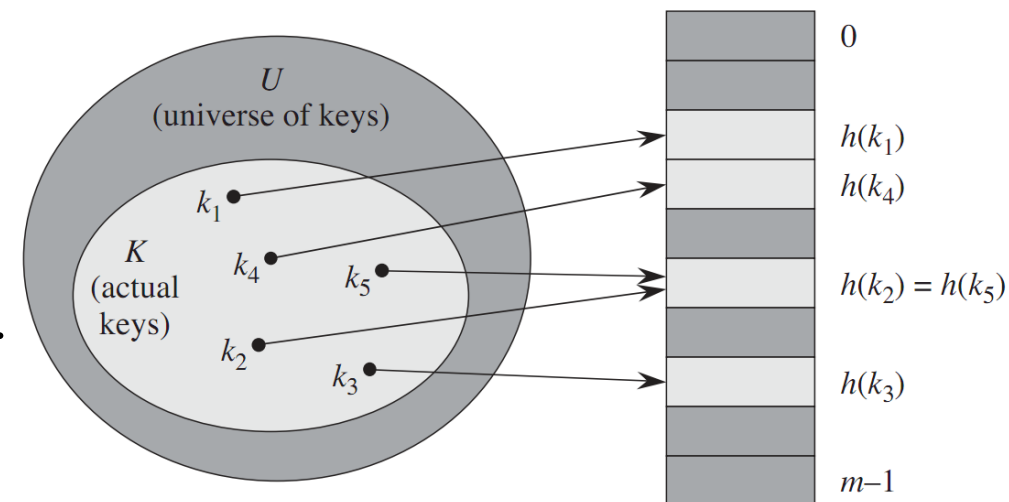
Source: Introduction to Algorithms, Cormen et al



# Hash Table



- Similar to direct addressing but uses much less space.
- **Idea** Instead of storing directly at key's location, convert key to much smaller value, and store at this location.
- A hash table consists of the following.
  - A universe  $U$  of keys.
  - An array of  $T$  of size  $m$ .
  - A hashing function  $h:U \rightarrow \{0,1,\dots,m-1\}$ .
- We'll talk later about how to pick good hash functions.
- **insert( $k, v$ )** Hash key to  $h(k)$ . Store  $v$  in  $T[h(k)]$ .
- **find( $k$ )** Return the value in  $T[h(k)]$
- **delete( $k$ )** Delete the value in  $T[h(k)]$
- Assuming  $h(k)$  takes  $O(1)$  time to compute, all ops still take  $O(1)$  time. Uses  $O(m)$  space.
- If  $m \ll |U|$ , then hashing uses much less space than direct addressing.
- However, our current scheme doesn't quite work, due to collisions.






# Collisions



- We store a key at array position  $h(k)$ .
- But what if two keys hash to the same location, i.e.,  $k_1 \neq k_2$ , but  $h(k_1) = h(k_2)$ ?
  - This is called a collision.
- Collisions are unavoidable when  $|U| > m$ .
  - By Pigeonhole Principle, must exist at least two different keys in  $U$  that hash to same value.






# Closed Addressing



- In closed addressing, every entry in hash table points to a linked list.
  - Keys that hash to the same location get added to the linked list.
  - For simplicity, we'll ignore values from now on and only focus on keys.
- **insert(k)** Add  $k$  to the linked list in  $T[h(k)]$ .
- **find(k)** Search the linked list in  $T[h(k)]$  for  $k$ .
- **delete(k)** Delete  $k$  from the linked list in  $T[h(k)]$ .
- Suppose the longest list has length  $\hat{n}$ , and average length list is  $\bar{n}$ .
  - Each operation takes worst case  $O(\hat{n})$  time.
  - An operation on a random key takes  $O(\bar{n})$  time.






# Load Factor



- The key to making closed addressing hashing fast is to make sure list lengths aren't too long.
- For this, we want the hash function to appear random.
  - Assume that any key is uniformly likely to be hashed to any table location.
- Suppose the hash table contains  $n$  items, and has size  $m$ .
- Then under the uniform hashing assumption, each table location has on average  $n/m$  keys.
  - Call  $\alpha = n/m$  the load factor.
- So the average time for each operation is  $O(\alpha)$ .
- However, even with uniform hashing, in the worst case, all keys can hash to the same location. So, the worst-case performance is  $O(n)$ .



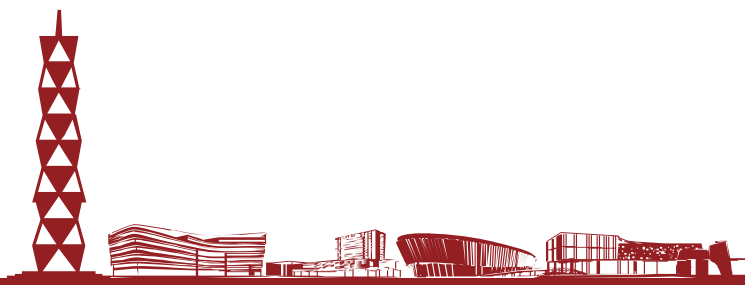




# Picking a hash function



- We saw that we want hash functions to hash keys to “random” locations.
  - However, note that each hash function is itself a deterministic function, i.e.  $h(k)$  always has the same value.
    - If  $h(k)$  can produce different values, we can't find key  $k$  in the hash table anymore.
- It's hard to find such random hash functions, since we don't assume anything about the distribution of input keys.
- In practice, we use a number of heuristic functions.





# Heuristic hash functions



- Assume the keys are natural numbers.
  - Convert other data types to numbers.
  - **Ex** To convert ASCII string to natural number, treat the string as a radix 128 number.  
E.g. "pt"  $\rightarrow (112 \times 128) + 116 = 14452$ .
- **Division method**  $h(k) = k \bmod m$ 
  - Often choose  $m$  a prime number not too close to a power of 2.
- **Multiplication method**  $h(k) = \lfloor m (k A \bmod 1) \rfloor$ , where  $A$  is some constant.
  - Knuth's suggestion is  $A = \frac{\sqrt{5}-1}{2} \approx 0.618034 \dots$





# Universal hashing



- As we said, regardless of the hash function, an adversary can choose a set of  $n$  inputs to make all operations  $O(n)$  time.
- Universal hashing overcomes this using randomization.
  - No matter what the  $n$  input keys are, every operation takes  $O(n/m)$  time in expectation, for a size  $m$  hash table.
  - Note  $O(n/m)$  time is optimal.
- Instead of using a fixed hash function, universal hashing uses a random hash function, chosen from some set of functions  $H$ .
- Say  $H$  is a universal hash family if for any keys  $x \neq y$

$$\Pr_{h \in H} [h(x) = h(y)] = 1/m$$

- So if we randomly choose a hash function from  $H$  and use it to hash any keys  $x, y$ , they have  $1/m$  probability of colliding.
- Note the hash functions in  $H$  are not random. However, we choose which function to use from  $H$  randomly.






# Universal hashing



- **Thm** Let  $H$  be a universal hash family. Let  $S$  be a set of  $n$  keys, and let  $x \in S$ . If  $h \in H$  is chosen at random, then the expected number of  $y \in S$  s.t.  $h(x) = h(y)$  is  $n/m$ .
- **Proof** Say  $S = \{x_1, \dots, x_n\}$ .
  - Let  $X$  be a random variable equal to the number of  $y \in S$  s.t.  $h(x) = h(y)$ .
  - Let  $X_i = 1$  if  $h(x_i) = h(x)$  and 0 otherwise.
  - $E[X_i] = \Pr_{h \in H}[h(x_i) = h(x)] \times 1 + \Pr_{h \in H}[h(x_i) \neq h(x)] \times 0 = 1/m$ .
    - First equality follows by universal hashing property.
  - $E[X] = E[X_1] + \dots + E[X_n] = n/m$ .





# Constructing universal hash family 1

- Choose a prime number  $p$  such that  $p > m$ , and  $p >$  all keys.
- Let  $h_{ab}(k) = ((ak + b) \bmod p) \bmod m$ .
- Let  $H_{pm} = \{h_{ab} \mid a \in \{1, 2, \dots, p-1\}, b \in \{0, 1, \dots, p-1\}\}$ .
- **Thm**  $H_{pm}$  is a universal hash family.
- **Proof** Let  $x, y < p$  be two different keys. For a given  $h_{ab}$  let
$$r = (ax + b) \bmod p, \quad s = (ay + b) \bmod p$$
- We have  $r \neq s$ , because  $r - s \equiv a(x - y) \bmod p \neq 0$ , since neither  $a$  nor  $x - y$  divide  $p$ .
- Also, each pair  $(a, b)$  leads to a different pair  $(r, s)$ , since
$$a = ((r - s)(x - y)^{-1} \bmod p), \quad b = (r - ax) \bmod p$$
  - Here,  $(x - y)^{-1} \bmod p$  is the unique multiplicative inverse of  $x - y$  in  $\mathbb{Z}_p^*$ .





# Constructing universal hash family 2



- Since there are  $p(p-1)$  pairs  $(a, b)$  and  $p(p-1)$  pairs  $(r, s)$  with  $r \neq s$ , then a random  $(a, b)$  produces a random  $(r, s)$ .
- The probability  $x$  and  $y$  collide equals the probability  $r \equiv s \pmod m$ .
- For fixed  $r$ , number of  $s \neq r$  s.t.  $r \equiv s \pmod m$  is  $(p-1)/m$ .
- So for each  $r$  and random  $s \neq r$ , probability that  $r \equiv s \pmod m$  is  $((p-1)/m)/(p-1) = 1/m$ .
- So  $\Pr_{h_{ab} \in H_{pm}} [h_{ab}(x) = h_{ab}(y)] = 1/m$  and  $H_{pm}$  is universal.






# Perfect hashing



- The hashing methods we've seen can ensure  $O(1)$  expected performance but are  $O(n)$  in the worst case due to collisions.
- However, if we have a fixed set of keys, perfect hashing can ensure no collisions at all.
  - Perfect hashing maintains a static set and allows  $\text{find}(k)$  and  $\text{delete}(k)$  in  $O(1)$  time.
  - It doesn't support  $\text{insert}(k)$ .
- **Ex** The fixed set of keys may represent the file names on a non-writable DVD.



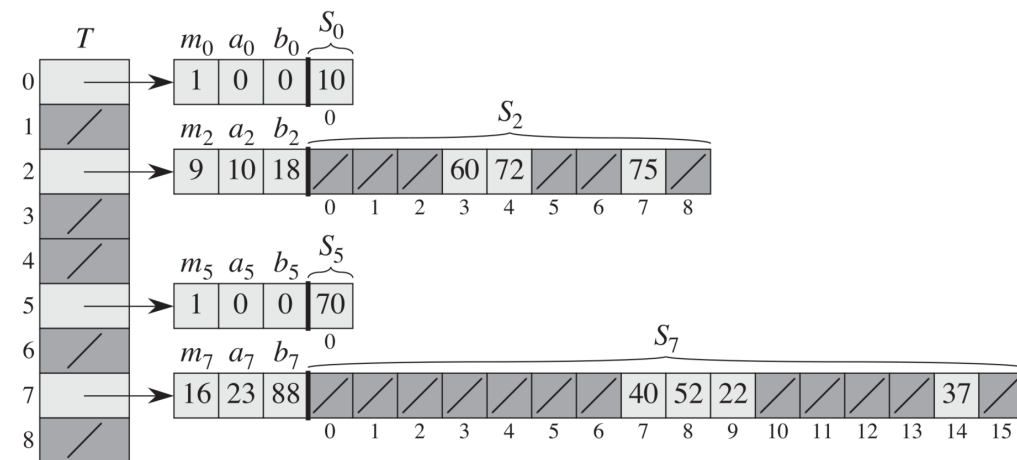


# Perfect hashing



- Suppose we want to store  $n$  items with no collisions.
- Perfect hashing uses two levels of universal hashing.
  - The first layer hash table has size  $m = n$ .
  - Use first layer hash function  $h$  to hash key to a location in  $T$ .
  - Each location  $j$  in  $T$  points to a hash table  $S_j$  with hash function  $h_j$ .
  - If  $n_j$  keys hash to location  $j$ , the size of  $S_j$  is  $m_j = n_j^2$ .
- We'll ensure there are no collisions in the secondary hash tables  $S_1, \dots, S_m$ .
  - So all operations take worst case  $O(1)$  time.
- Overall the space use is  $O(m + \sum_{j=1}^m n_j^2)$ .
  - We'll show this is  $O(n) = O(m)$ .
  - So perfect hashing uses same amount of space as normal hashing.

- $h(k) = ((3k + 42) \bmod 101) \bmod 9$
- $h_j(k) = ((a_j k + b_j) \bmod 101) \bmod m_j$








# Avoiding collisions



- **Lemma** Suppose we store  $n$  keys in a hash table of size  $m = n^2$  using universal hashing. Then with probability  $\geq 1/2$  there are no collision.
- **Proof** There are  $\binom{n}{2}$  pairs of keys that can collide.
  - Each collision occurs with probability  $1/m = 1/n^2$ , by universal hashing.
  - So the expected number of collisions is  $\frac{\binom{n}{2}}{n^2} \leq \frac{1}{2}$ .
  - By Markov's inequality the  $\Pr[\# \text{ collisions} \geq 1] \leq E[\# \text{ collisions}] \leq 1/2$ .
- When building each hash table  $S_j$ , there's  $< 1/2$  probability of having any collisions.
  - If collisions occur, pick another random hash function from the universal family and try again.
  - In expectation, we try twice before finding a hash function causing no collisions.

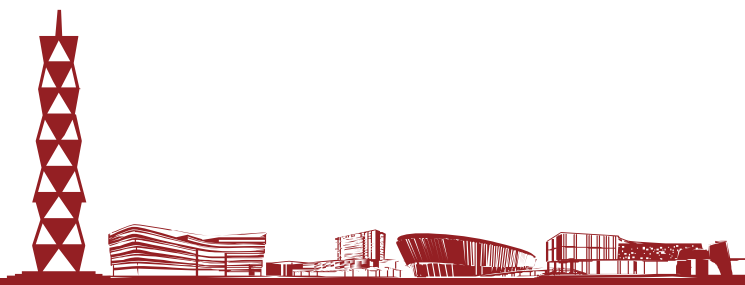


- **Lemma** Suppose we store  $n$  keys in a hash table of size  $m=n$ . Then the secondary hash tables use space  $E[\sum_{j=0}^{m-1} n_j^2] < 2n$ , where  $n_j$  is the number of keys hashing to location  $j$ .
- **Proof**  $E[\sum_{j=0}^{m-1} n_j^2] = E[\sum_{j=0}^{m-1} (n_j + 2 \binom{n_j}{2})] = E[\sum_{j=0}^{m-1} n_j] + 2 E[\sum_{j=0}^{m-1} \binom{n_j}{2}]$
- $\sum_{j=0}^{m-1} \binom{n_j}{2}$  is the total number of pairs of hash keys which collide in the first level hash table.
  - By universal hashing, this equals  $\binom{n}{2} \frac{1}{m} = \frac{n-1}{2}$ .
- $E[\sum_{j=0}^{m-1} n_j] = n$ .
- So  $E[\sum_{j=0}^{m-1} n_j^2] = n + \frac{2(n-1)}{2} < 2n$ .






# Bloom Filters



# Approximate Sets



- A Bloom filter is a data structure that can implement a set.
  - It only keeps track of which keys are present, not any values associated to keys.
  - It supports insert and find operations.
  - It doesn't support delete operations.
- Bloom filters use less memory than hash tables or other ways of implementing sets.
- However, Bloom filters are approximate.
  - It can produce false positives: it says an element is present even though it's not.
    - We can bound the probability of false positives.
  - But it doesn't produce false negatives: if it says an element isn't present, then it's not.





# Bloom Filter Applications



- Suppose we have a big database and querying it to check if an item is present is expensive.
- We store the set of items in the database using a Bloom filter.
  - This tells us whether an item is in database or not.
- If filter says an item's not present, it's definitely not in the database.
  - So, no need to do an expensive query.
- If filter says an item is present, then either item is present, or there's false positive.
  - When we query the database, there's a small probability we waste time querying for a nonexistent item.
- Overall, we save time by checking Bloom filter first before querying database.






# Bloom Filter Applications



- Consider a P2P network, where each node stores some files.
- If you want to get a file, you need to know which nodes have it.
- Keeping a list of all items stored at each node is too expensive.
- Instead, for every other node, keep a Bloom filter of its files.
- If filter says no for a node, it definitely doesn't have the file.
- If filter says yes, then either node has the file, or there's false positive and we make a useless request.
- Overall, we save space, and also won't waste much communication because we rarely make useless requests.






# Bloom Filters



- A Bloom filter consists of
  - An array  $A$  of size  $m$ , initially all 0's.
  - $k$  independent hash functions  $h_1, \dots, h_k$ , each mapping from keys to  $\{1, \dots, m\}$ .
- To store key  $x$ 
  - Set  $A[h_1(x)], A[h_2(x)], \dots, A[h_k(x)]$  all to 1.
  - Some locations can get set to 1 multiple times; that's fine.
- To check if key  $x$  is in the set
  - Read array locations  $A[h_1(x)], A[h_2(x)], \dots, A[h_k(x)]$ .
  - If all the values are 1, output "x is in set".
  - Otherwise, output "x is not in set".



A Bloom filter with  $k=3$  hash functions storing 2 items.






# Correctness



- Let's look at the correctness of the search function.
- If search for  $x$  returns no, then at least one of  $A[h_1(x)], \dots, A[h_k(x)]$  equals 0.
  - So  $x$  cannot be in the set, because if  $x$  had been inserted into the set, then we would have  $A[h_1(x)] = \dots = A[h_k(x)] = 1$ .
  - So there are no false negatives.
- If search for  $x$  returns yes, then  $A[h_1(x)] = \dots = A[h_k(x)] = 1$ .
  - So either  $x$  was inserted into the set.
  - Or we inserted some keys that hashed to the same  $k$  locations as  $x$ .
    - So it looks as if  $x$  was inserted, even though it wasn't.
    - This is a false positive. We'll bound the probability this happens.







# False Positive Probability 1



- False positive probability depends on  $k$  (number of hash functions),  $m$  (size of table) and  $n$  (number of keys inserted).
- Assume hash functions hash keys to random locations.
- When inserting one key, we set  $k$  random locations to 1.
- Fix any position  $i$ . Probability  $i$  is set to 1 by a hash function is  $1/m$ , so probability  $i$  stays 0 is  $1-1/m$ .
  - After  $k$  hashes, probability  $i$  still 0 is  $(1 - 1/m)^k$ .
  - To insert  $n$  items, we used  $nk$  hashes. So, probability  $i$  still 0 after all these is  $p = (1 - 1/m)^{nk}$ .
- We now use an approximation  $\left(1 - \frac{1}{m}\right)^{nk} \approx e^{-\frac{nk}{m}}$ .





# False Positive Probability 2



- So, probability any position  $i$  is 1 after  $n$  keys inserted is  $1 - p \approx 1 - e^{-\frac{nk}{m}}$ .
- Since there are  $m$  positions in the array, assume there are  $(1-p)m$  positions that are 1.
  - This isn't quite correct. The actual number of 1's in the array is a random variable, whose expectation is  $(1-p)m$ .
  - However, we can make the argument rigorous by showing that the actual number of 1's is  $(1 - p)m \pm \sqrt{m \log m}$  with high probability.
- We only get a false positive if when we check  $k$  random locations, they're all 1.
  - Probability is  $f = (1 - p)^k \approx \left(1 - e^{-\frac{nk}{m}}\right)^k$ .

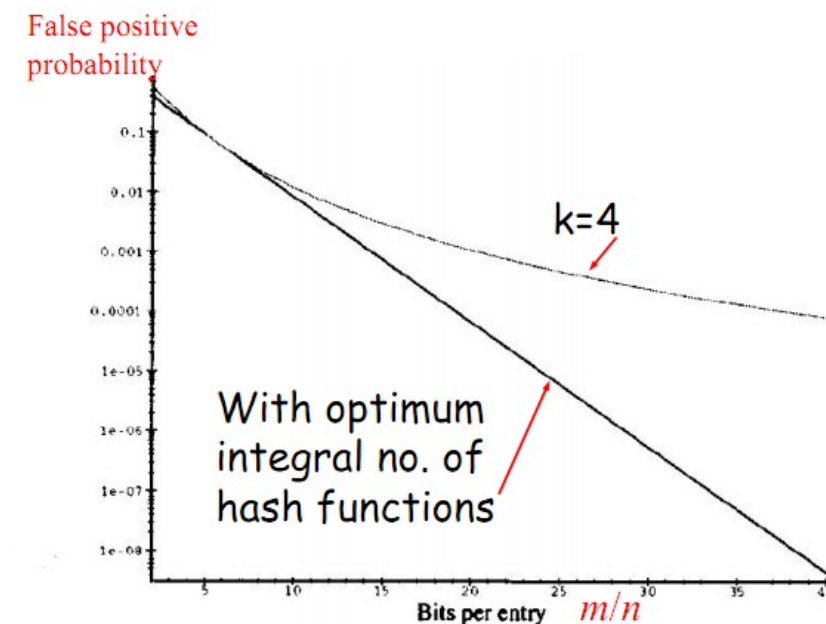




# False Positive Probability 3



- Notice the false prob.  $(1 - e^{-\frac{nk}{m}})^k$  is a function of  $k$ , the number of hash functions we use.
- We find  $k$  to minimize the false positive prob. by differentiating  $f$  wrt  $k$  and solving.
- The optimum  $k$  is  $\frac{m \ln(2)}{n}$ , which leads to  $f = \left(\frac{1}{2}\right)^k \approx 0.6185^{\frac{m}{n}}$ .
  - Notice that  $m/n$  is the average number of bits per item. So error rate decreases exponentially in space usage.





# Improvements



- Right now, Bloom filters can't handle deletes.
  - Say keys  $k_1$ ,  $k_2$  hash to two overlapping sets of locations. If you delete  $k_1$  by setting some of its locations to 0, you could also delete  $k_2$ .
- Deletes can be done by storing a count of how many keys hashed to that location, and inc / dec the counts when inserting or deleting.
  - But this uses more memory.
  - Also, what if the counts overflow?

