



# CS240 Algorithm Design and Analysis

## Lecture 25

### Approximation Algorithms

Quan Li  
Fall 2024  
2024.12.24



# Approximation Algorithms




- Up to now, most of our algorithms have been exact, i.e., they find an optimal solution.
- But there are many problems for which we don't know how to find an optimal solution.
  - A key example is NP-complete problems. We don't know efficient algorithms for any NPC problem.
- Many such problems are important in practice. What do we do?
- If we can't get find the best answer, let's try for good enough.
- Approximation algorithms find an approximately optimal answer.



# Approximation Ratio



- Let  $X$  be a maximization problem. Let  $A$  be an algorithm for  $X$ .
- Let  $a > 1$  be a constant.
- $A$  is an  $a$ -approximation algorithm for  $X$  if  $A$  always returns an answer that's at least  $1/a$  times the optimal.
  - **Ex** If  $X$  is max-flow,  $A$  is a 2-approx algorithm if it always returns a flow that's at least  $\frac{1}{2}$  the optimal.
  - The closer  $a$  is to 1, the better the approximation.
- If  $X$  is a minimization problem,  $A$  is an  $a$ -approximation algorithm for  $X$  if it always returns an answer that's at most  $a$  times larger than the optimal.
  - **Ex** If  $X$  is min-cut,  $A$  is a 2-approx algorithm if it always returns a cut that's at most 2 times the size of the optimal.
  - Again, the closer  $a$  is to 1, the better the approximation.







# Coverings



- Suppose there's a set of teachers, and each can teach a certain set of classes.
  - Let  $S_i$  be the set of classes teacher  $i$  can teach.
- The entire set of classes is  $X$ .
- We want to pick the minimum set of teachers to teach all the classes.
  - Let  $T$  be set of teachers we pick.
  - We want  $\bigcup_{i \in T} S_i = X$ , and  $T$  to be the smallest possible.






# Set Covering



- **Input** A collection  $F$  of sets. Each set has a cost. The union of all the sets is  $X$ .
- **Output** A subset  $G$  of  $F$ , whose union is  $X$ .
- **Goal** Minimize the total cost of the sets in  $G$ .



*If all sets have same cost,  $S_3$ ,  $S_4$  and  $S_5$  is a min cost set cover of  $X$ .*

- Minimum cost set cover is NP-complete.
- We'll see a  $\ln(n)$ -approximation algorithm, where  $n=|X|$ .





# A Greedy Approximation Algorithm



- A natural greedy heuristic is to choose sets which cover points most cheaply.
  - For each set, let  $c$  be its cost, and  $m$  be the number of points it covers.
  - We want to use the set with the smallest  $c/m$  value, because this is the cheapest way to cover some new points.
- After we pick this set, remove all the points it covers. Then we consider the per unit cost of the remaining sets and again choose the cheapest.





# A Greedy Approximation Algorithm



- $F$  is the entire collection of sets. The union of these sets is  $X$ .
  - Each set  $S$  in  $F$  has a cost  $\text{cost}(S)$ .
  - $U$  is the set of elements of  $X$  we haven't covered yet.
  - $C$  is the set cover we eventually output.
  - $U = X$
  - $C = \emptyset$
  - while  $U \neq \emptyset$ 
    - choose  $S \in F - C$  with  $\min |\text{cost}(S)|/|S \cap U|$
    - $C = C \cup \{S\}$
    - $U = U - S$
  - output  $C$
- Per unit cost to cover new elements.






# Proof of Correctness



- We always output a set cover, because the while loop continues till  $X$  is covered.
- We'll prove the approximation ratio is at most  $1 + 1/2 + 1/3 + \dots + 1/n \approx \ln(n)$ .
  - If the min cost of a set cover is  $V$ , our set cover costs at most  $\ln(n) * V$ .
- The basic plan is to bound the cost of the set cover the algorithm outputs using the "average cost" per element.



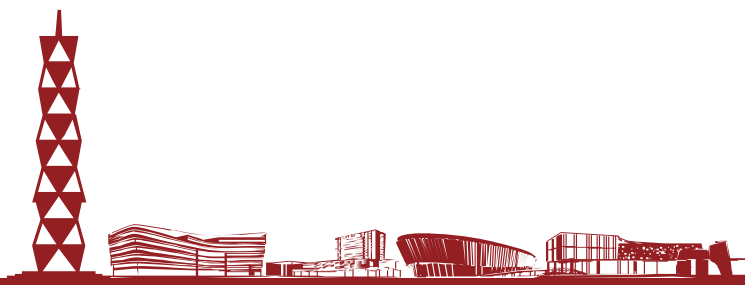




# Proof of Correctness



- Order the sets in  $C$  by when they're added to  $C$ , earliest set first.
  - Let the order be  $S_1, S_2, \dots, S_m$ .
- Cost of the set cover is  $L = \sum_i \text{cost}(S_i)$ .
- Order the elements in  $X$  by when they're added, earliest element first.
  - Let the order be  $e_1, e_2, \dots, e_n$ .
  - So, the first few  $e$ 's are added by  $S_1$ , the next few added by  $S_2$ , etc.
  - Every element in  $X$  is in the list, because  $C$  covers  $X$ .





# Proof of Correctness



- Let  $n_i$  be the number of new elements  $S_i$  covers.
    - So,  $n_i$  is the number of elements in  $S_i$ , but not in  $S_1, \dots, S_{i-1}$ .
  - Divide the cost of  $S_i$  evenly among the new elements it covers.
    - If  $e$  is newly covered by  $S_i$ , then  $\text{cost}(e) = \text{cost}(S_i)/n_i$ .
- $$\sum_k \text{cost}(e_k) = \sum_i n_i * \frac{\text{cost}(S_i)}{n_i} = \sum_i \text{cost}(S_i) = L$$
- Every element is covered by some  $S_i$ , and  $S_i$  covers  $n_i$  new elements.
- **We'll prove  $\text{cost}(e_k) \leq \text{OPT}/(n-k+1)$ , for any  $k$ .**
  - Suppose this is true, then
$$L = \sum_k \text{cost}(e_k) \leq \sum_k \text{OPT}/(n - k + 1) \approx \ln(n) * \text{OPT}$$





# The Per Element Cost



- Let's focus on some element  $e_k$ , and let  $S_j$  be the set which covers  $e_k$  for the first time.
- Let  $C_1, \dots, C_m$  be the sets in an optimal cover, each of which covers some elements of  $U = \{e_k, e_{k+1}, e_{k+2}, \dots, e_n\}$ .
  - Let  $n'_1, \dots, n'_m$  be the number of elements of  $U$  which  $C_1, \dots, C_m$  cover.
- **Obs 1:  $\sum_i n'_i \geq n - k + 1$ .**
  - Because  $C_1, \dots, C_m$  cover  $U$ .
- **Obs 2:  $\sum_i \text{cost}(C_i) \leq \text{OPT}$ .**
  - Because  $C_1, \dots, C_m$  are a subset of an optimal cover, which has cost  $\text{OPT}$ .





# The Per Element Cost



- **Obs 3** None of  $C_1, \dots, C_m$  are among  $S_1, \dots, S_{j-1}$ .
  - If some  $C_i$  is among  $S_1, \dots, S_{j-1}$ , then since  $C_i$  covers some  $e$  in  $U$ ,  $e$  would be covered by  $\{S_1, \dots, S_{j-1}\}$ . So,  $e$  would be among the first  $k-1$  elements covered. Contradiction.
- **Obs 4** There exists some  $C_i$  among  $C_1, \dots, C_m$  with  $\frac{\text{cost}(C_i)}{n'_i} \leq \text{OPT}/(n - k + 1)$ .
  - If every  $C_i$  in  $C_1, \dots, C_m$  has  $\frac{\text{cost}(C_i)}{n'_i} \geq \text{OPT}/(n - k + 1)$ , then
$$\begin{aligned} \text{OPT} &\geq \sum_i \text{cost}(C_i) = \sum_i n'_i * \frac{\text{cost}(C_i)}{n'_i} > \sum_i n'_i * \frac{\text{OPT}}{n - k + 1} \geq \text{OPT}/(n - k + 1) \sum_i n'_i \geq \frac{\text{OPT}}{n - k + 1} * (n - k + 1) = \text{OPT} \end{aligned}$$
Contradiction.





# Proof of Approximation Ratio




- **Lemma**  $cost(S_j)/n_j \leq OPT/(n - k + 1)$
- **Proof** When choosing  $S_j$ , the only sets the algorithm is not allowed to choose are  $S_1, \dots, S_{j-1}$ .
  - By obs 3,  $C_1, \dots, C_m$  aren't in  $S_1, \dots, S_{j-1}$ .
  - By obs 4, there's some  $C_i$  in  $C_1, \dots, C_m$ , with  $\frac{cost(C_j)}{n'_i} \leq OPT/(n - k + 1)$ .
  - $S_j$  was chosen so that  $cost(S_j)/n_j$  is min among all sets not in  $S_1, \dots, S_{j-1}$ .
  - So  $\frac{cost(S_j)}{n_j} \leq \frac{cost(C_i)}{n'_i} \leq OPT/(n - k + 1)$ .
- Since  $\frac{cost(S_j)}{n_j} = cost(e_k)$ , we have  $cost(e_k) \leq OPT/(n - k + 1)$ .
- The approximation ratio follows because

$$L = \sum_k cost(e_k) = \sum_k \frac{OPT}{n-k+1} \approx \ln(n) * OPT$$





# Scheduling







# Parallel Computing and Scheduling



- Computers today are parallel.
  - Multiple processors in a system.
  - Multiple tasks for the processors to run.
- Multiprocessor scheduling is the problem of deciding which tasks to run on which processors at what time.
- Many possible objectives.
  - Throughput, fairness, energy usage.
  - Latency, i.e. finishing all jobs as fast as possible.






# Makespan Scheduling



- $n$  independent jobs.
  - Jobs have different sizes, i.e. time needed to perform job.
  - Jobs can be done in any order.
  - Any job can be done on any machine.
- $m$  processors.
  - All have the same speed.
  - Each processors can do one job at a time.
- Assign the jobs to the processors.
- Makespan is when the last processor finishes all its jobs.
- Minimize the makespan.
  - i.e., finish all the jobs as fast as possible.







# Minimizing makespan is NPC



- The decision version of scheduling is obviously in NP.
- SUBSET-SUM: given a set of numbers  $S$  and target  $t$ , is there a subset of  $S$  summing to  $t$ ?
  - **Ex**  $S=\{1,3,8,9\}$ .  $t=9$ , yes.  $t=14$ , no.
  - This is NP-complete. We reduce SUBSET-SUM to scheduling.
- Let  $(S,t)$  be an instance of SUBSET-SUM.
  - Let  $s$  be sum of all elements in  $S$ .
- Make a set of jobs  $J = S \cup \{s-2t\}$ , and schedule them on 2 processors.

## • Recipe to establish NP-completeness of problem $Y$

- Step 1. Show that  $Y$  is in NP
- Step 2. Choose an NP-complete problem  $X$
- Step 3. Prove that  $X \leq_p Y$





# Minimizing makespan is NPC



- **Claim** If some subset of  $S$  sums to  $t$ , then min makespan is  $s-t$ .
- **Proof** Say  $S' \subseteq S$  sums to  $t$ . Schedule the jobs in  $S'$  and job  $s-2t$  on processor 1. So proc 1 finishes at time  $t+s-2t=s-t$ . Proc 2 does the jobs in  $S-S'$ , so it finishes at time  $s-t$  as well.
- **Claim** If the min makespan is  $s-t$ , there exists a subset of  $S$  that sums to  $t$ .
- **Proof** Suppose WLOG proc 1 does the  $s-2t$  job. Since makespan is  $s-t$ , the other jobs proc 1 does must have total size  $s-t-(s-2t)=t$ .
- So  $(S,t)$  is yes instance of SUBSET-SUM iff makespan =  $s-t$ .
  - So SUBSET-SUM  $\leq p$  scheduling, and scheduling is NP-complete.





# Graham's List Scheduling



- Since scheduling is NPC, it's unlikely we can find the min makespan in polytime.
- List scheduling is a simple greedy algorithm.
  - Finds a schedule with makespan at most twice the minimum.
  - A 2-approximation.
- If there are  $n$  tasks and  $m$  processors, list scheduling only takes  $O(n \log n + m)$  time.
  - Compare this to  $n!$   $C(n+m-1, m-1)$  time to try all possible schedules and pick the best.






# Graham's List Scheduling



- List the jobs in any order.
- As long as there are unfinished jobs.
  - If any processor doesn't have a job now, give it the next job in the list.
- Example
  - 3 processors. The jobs have length 2, 3, 3, 4, 5, 6, 8.
  - List them in any order. Say 4, 5, 3, 2, 6, 8, 3.
  - Initially, no proc has a job. Give first 3 jobs to the 3 procs.
  - At time 3, proc 3 is done. Give it next job in list, 2.
  - At time 4, proc 2 is done. Give it next job in list, 6.
  - At time 5, both 1, 3 are done. Give them next jobs in list, 8, 3.
  - Everybody finishes by time 13.
    - The makespan of this schedule is 13.







# The Worst Case for LS



- How badly can list scheduling do compared to optimal?
- Say there are  $m^2$  jobs with length 1, and one job with length  $m$ .
  - Suppose they're listed in the order  $1, 1, 1, \dots, 1, m$ .
  - LS has makespan  $2m$ . Optimal makespan is  $m+1$ .
  - $\text{makespan}(\text{LS}) / \text{makespan}(\text{opt}) = 2m/(m+1) \approx 2$ .
- This is worst possible case for list scheduling.



$\text{makespan}(\text{LS}) = 2m$



$\text{makespan}(\text{opt}) = m+1$





# LS is a 2-approximation




- Next, we prove LS always gives a schedule at most twice the optimal.
- Suppose LS gives makespan of  $M$ .
- Let the optimal schedule have makespan  $M^*$ .
- We prove that  $M \leq 2M^*$ .





# LS is a 2-approximation




- The picture above is the schedule produced by list scheduling.
- Consider task  $X$  that finishes last.
  - Say  $X$  starts at time  $T$ , and has length  $t$ .
- **Claim 1**  $M^* \geq t$ .
  - In any schedule,  $X$  has to run on some process.
  - Since  $X$  takes  $t$  time, every schedule, including the opt, takes  $\geq t$  time.





# LS is a 2-approximation



## ■ Claim 2 $M^* \geq T$ .


- Up to time  $T$ , no processor is ever idle.
  - Up to  $T$ , there's always some unfinished job.
  - As soon as a processor finishes one job, it's assigned another one.
- So at time  $T$ , each processor completed  $T$  units of work.
- So total amount of work in all the jobs is  $\geq mT$ . Up to  $T$ :  $mT$
- In the opt schedule,  $m$  processors complete at most  $m$  units of work per time unit.
- So length of opt schedule is  $\geq (\text{total work})/m \geq mT/m = T$ .







# LS is a 2-approximation



- From Claims 1 and 2, we have  $M^* \geq t$  and  $M^* \geq T$ .
- So  $M^* \geq \max(T, t)$ .
- $M = T + t$ , because X is last job to finish.
- So  $M/M^* \leq (T+t)/\max(T, t) \leq 2$ .





# LPT Scheduling



- Worst case for LS occurred when longest job was scheduled last.
  - Large jobs are “dangerous” at end.
- Let's try to schedule longest jobs first.
- Longest processing time (LPT) schedule is just like list scheduling, except it first sorts tasks by nonincreasing order of size.
- **Ex** For three processors and tasks with sizes 2, 3, 3, 4, 5, 6, 8, LPT first sorts the jobs as 8,6,5,4,3,3,2. Then it assigns  $p_1$  tasks 8,3,  $p_2$  tasks 6,3,  $p_3$  tasks 5,4,2, for a makespan of 11.
- LPT has an approximation ratio of  $4/3$ .





# LPT is a $4/3$ -approximation



- **Thm** Suppose the optimal makespan is  $M^*$ , and LPT produces a schedule with makespan  $M$ . Then  $M \leq 4/3 M^*$ .
- Let  $X$  be the last job to finish. Assume it starts at time  $T$  and has size  $t$ .
- Assume WLOG that  $X$  is the last job to start.
  - If not, then say  $Y$  starts after  $T$ .
  - $Y$  finishes before  $T+t$ . So we can remove  $Y$  without increasing the makespan.
- **Cor 1**  $X$  is the smallest job.
  - $X$  is the last job to start, so due to LPT scheduling it's the smallest.






# LPT is a 4/3-approximation



- **Claim 1** LPT's makespan =  $T + t \leq M^* + t$ .
  - As in LS, no processor is idle up to time  $T$ , so  $M^* \geq T$ .
- **Case 1**  $t \leq M^*/3$ .
  - Then LPT's makespan  $\leq M^* + t \leq M^* + M^*/3 = 4/3 M^*$ .
- **Case 2**  $t > M^*/3$ .
  - Since  $X$  is the smallest task, all tasks have size  $> M^*/3$ .
  - So the optimal schedule has at most 2 tasks per processor. So  $n \leq 2m$ .
  - If  $1 \leq n \leq m$ , then LPT and optimal schedule both put one task per processor.
  - If  $m < n \leq 2m$ , then optimal schedule is to put tasks in nonincreasing order on processors  $1, \dots, m$ , then on  $m, \dots, 1$ .
    - LPT also schedules tasks this way, so it's optimal.





# LS VS. LPT



- LPT gives better approximation ratio, has same running time. Why bother with LS?
- LS is online.
  - Imagine the jobs are coming one by one.
  - LS just puts them on any idle computer.
- LPT is offline
  - It needs to know all the jobs that will ever arrive, in order to sort them.
- In a realistic parallel computation, you get jobs on the fly.
  - Online is more realistic.
  - LS is usually more useful.

