



Lecture 2: Basic Artificial Neural Networks


Lan Xu
SIST, ShanghaiTech
Fall, 2022

Outline

- Single layer neural networks
 - Network models
 - Example: Logistic Regression
- Multi-layer neural networks
 - Limitations of single layer networks
 - Networks with single hidden layer

Acknowledgement: Hugo Larochelle's, Mehryar Mohri@NYU's & Yingyu Liang@Princeton's course notes

Mathematical model of a neuron



artificial neuron: *weighted sum and non-linearity*

$$s = \underset{\substack{\text{bias} \\ \uparrow}}{b} + \underset{\substack{\text{weights} \\ \uparrow}}{w_1}x_1 + w_2x_2 + \cdots + w_Mx_M = \mathbf{w}^T \mathbf{x}$$

input features \rightarrow (pointing to x_1, x_2, \dots, x_M)

sum \rightarrow (pointing to s)

$$h = g(s)$$

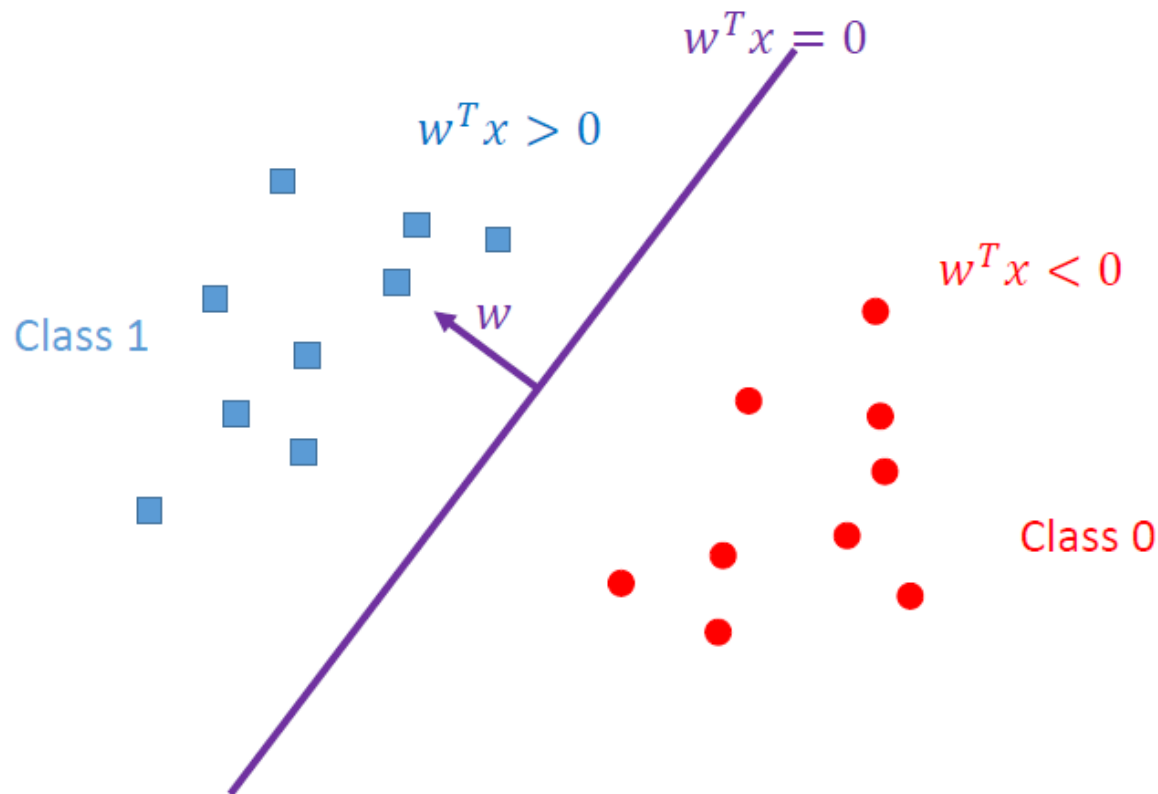
output \rightarrow (pointing to h)

non-linearity \rightarrow (pointing to g)

sum \rightarrow (pointing to s)

Single neuron as a linear classifier

- Binary classification




Outline


- Artificial neuron
 - Perceptron algorithm
- Single layer neural networks
 - Network models
 - Example: Logistic Regression
- Multi-layer neural networks
 - Limitations of single layer networks
 - Networks with single hidden layer

Acknowledgement: Hugo Larochelle's, Mehryar Mohri@NYU's & Yingyu Liang@Princeton's course notes

Single layer neural network



Single layer neural network




layer: *parallelized weighted sum and non-linearity*


$$\begin{array}{c} \text{one sum} \\ \text{per weight vector} \end{array} s_j = \mathbf{w}_j^T \mathbf{x} \longrightarrow \mathbf{s} = \mathbf{W}^T \mathbf{x} \begin{array}{c} \text{vector of sums} \\ \text{from weight matrix} \end{array}$$

$$\mathbf{h} = \sigma(\mathbf{s})$$

Single layer neural network

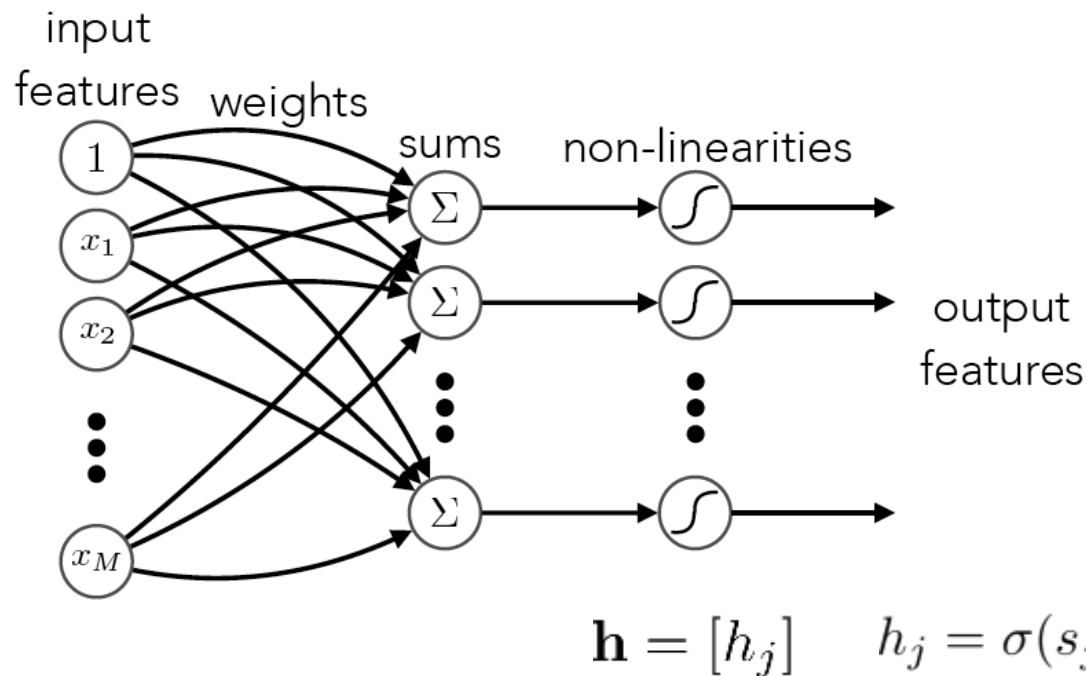


layer: *parallelized weighted sum and non-linearity*




What is the output?

- Element-wise nonlinear functions
 - Independent feature/attribute detectors



What is the output?

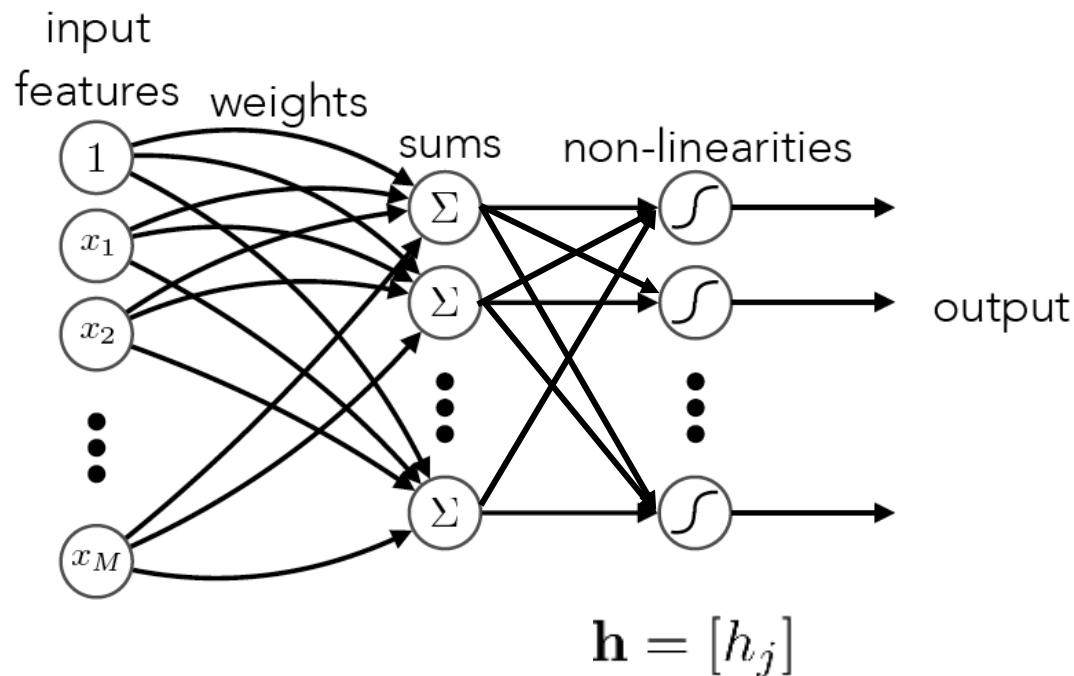
- Nonlinear functions with vector input
 - Competition between neurons



$$h_j = g(\mathbf{s}) = g(\mathbf{w}_1^\top \mathbf{x}, \dots, \mathbf{w}_m^\top \mathbf{x})$$

What is the output?


- Nonlinear functions with vector input
 - Example: Winner-Take-All (WTA)



$$h_j = g(\mathbf{s}) = \begin{cases} 1 & \text{if } j = \arg \max_i \mathbf{w}_i^\top \mathbf{x} \\ 0 & \text{if otherwise} \end{cases}$$

A probabilistic perspective

- Change the output nonlinearity



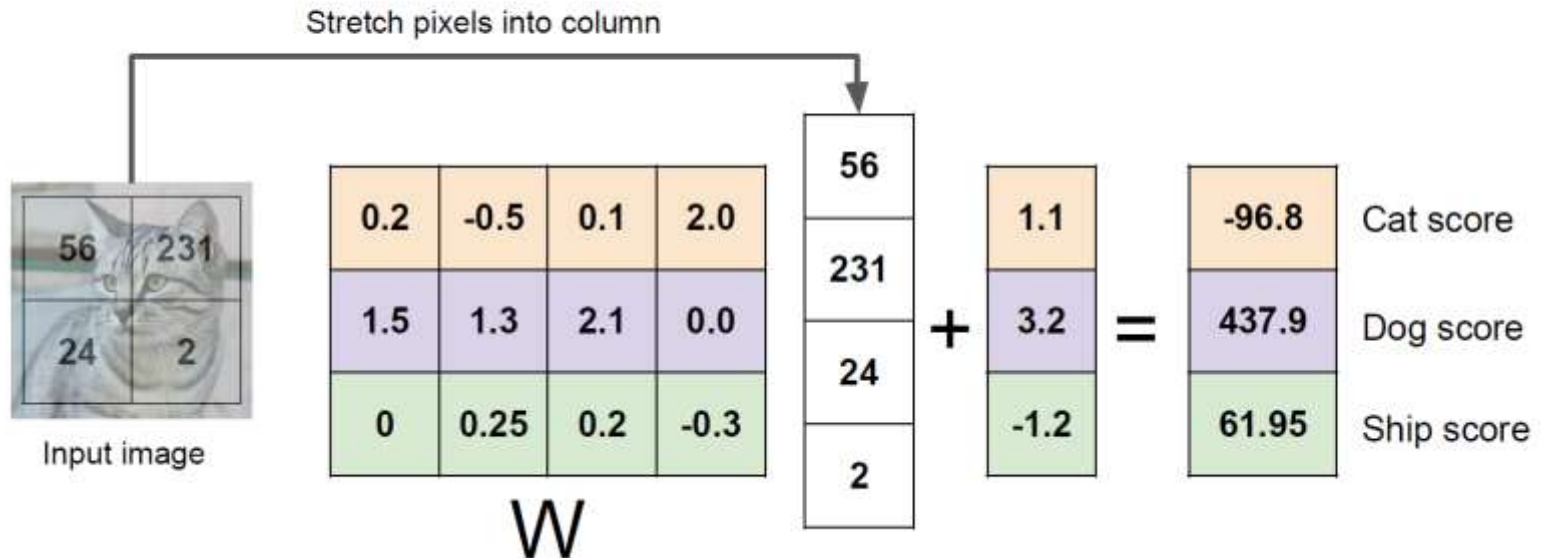
- From WTA to Softmax function

scores = unnormalized log probabilities of the classes.

$$P(Y = k | X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}} \quad \text{where} \quad s = f(x_i; W)$$

Multiclass linear classifiers

- Example with an image with 4 pixels, and 3 classes (cat/dog/ship)



- The WTA prediction: one-hot encoding of its predicted label

$$y = 1 \Leftrightarrow y = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} \quad y = 2 \Leftrightarrow y = \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix} \quad y = 3 \Leftrightarrow y = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$$

Probabilistic outputs

scores = unnormalized log probabilities of the classes.




$$P(Y = k|X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}}$$

where

$$s = f(x_i; W)$$

unnormalized probabilities



How to learn a multiclass classifier?

■ Define a loss function and do minimization

- Given training data $\{(x_i, y_i): 1 \leq i \leq n\}$ i.i.d. from distribution D
- Find $y = f(x) \in \mathcal{H}$ that minimizes $\hat{L}(f) = \frac{1}{n} \sum_{i=1}^n l(f, x_i, y_i)$
- s.t. the expected loss is small

$$L(f) = \mathbb{E}_{(x,y) \sim D}[l(f, x, y)]$$



Empirical loss

Learning a multiclass linear classifier

- Design a loss function for multiclass classifiers
 - Perceptron?
 - Yes, see homework
 - Hinge loss
 - The SVM and max-margin (see CS231n)
 - Probabilistic formulation
 - Log loss and logistic regression
- Generalization issue
 - Avoid overfitting by regularization

Example: Logistic Regression

- Learning loss: negative log likelihood

scores = unnormalized log probabilities of the classes.


$$P(Y = k|X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}} \quad \text{where} \quad s = f(x_i; W)$$

Want to maximize the log likelihood, or (for a loss function) to minimize the negative log likelihood of the correct class:

$$L_i = -\log P(Y = y_i|X = x_i)$$

Logistic Regression

- Learning loss: example



Logistic Regression

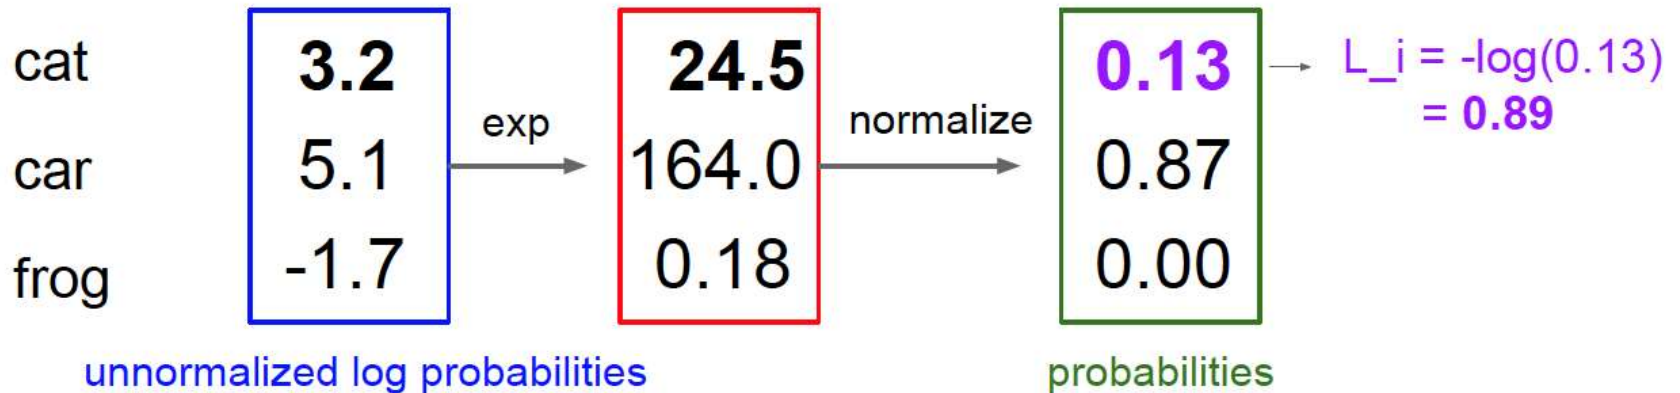
- Learning loss: questions



$$L_i = -\log\left(\frac{e^{sy_i}}{\sum_j e^{s_j}}\right)$$

unnormalized probabilities

Q: What is the min/max possible loss L_i ?



Logistic Regression

- Learning loss: questions



$$L_i = -\log\left(\frac{e^{sy_i}}{\sum_j e^{s_j}}\right)$$

unnormalized probabilities

Q2: Usually at initialization W is small so all $s \approx 0$. What is the loss?

cat
car
frog

3.2
5.1
-1.7

exp

24.5
164.0
0.18

normalize

0.13
0.87
0.00

unnormalized log probabilities

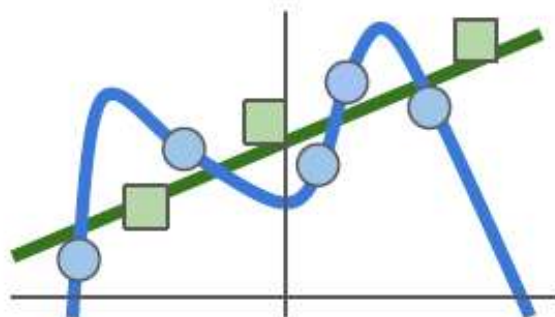
probabilities

$$L_i = -\log(0.13) = 0.89$$

Learning with regularization

- Constraints on hypothesis space
 - Similar to Linear Regression

$$L(W) = \underbrace{\frac{1}{N} \sum_{i=1}^N L_i(f(x_i, W), y_i)}_{\text{Data loss: Model predictions should match training data}} + \underbrace{\lambda R(W)}_{\text{Regularization: Model should be "simple", so it works on test data}}$$



Learning with regularization

■ Regularization terms

In common use:

L2 regularization $R(W) = \sum_k \sum_l W_{k,l}^2$

L1 regularization $R(W) = \sum_k \sum_l |W_{k,l}|$

Elastic net (L1 + L2) $R(W) = \sum_k \sum_l \beta W_{k,l}^2 + |W_{k,l}|$

Max norm regularization (might see later)

■ Priors on the weights

- Bayesian: integrating out weights
- Empirical: computing MAP estimate of W

L1 vs L2 regularization



<https://www.youtube.com/watch?v=jEVh0uheCPk>

L1 vs L2 regularization

■ Sparsity

$$x = [1, 1, 1, 1]$$

$$w_1 = [1, 0, 0, 0]$$

$$w_2 = [0.25, 0.25, 0.25, 0.25]$$

$$w_3 = [0.5, 0.5, 0, 0]$$

$$f(x) = w^\top x$$


$$w_1^\top x = w_2^\top x = w_3^\top x$$

$$\|w_1\|^2 = |w_1| = 1$$


$$\|w_2\|^2 = 4/16 = 1/4, |w_2| = 1$$

$$\|w_3\|^2 = 2/4 = 1/2, |w_3| = 1$$

A L1 regularization



B L2 regularization



Optimization: gradient descent

- Gradient descent

```
# Vanilla Gradient Descent
```


```
while True:
```

```
    weights_grad = evaluate_gradient(loss_fun, data, weights)
```


```
    weights += - step_size * weights_grad # perform parameter update
```

- Learning rate matters

$\eta_t = t$, it is too big



too small η_t , after 100 iterations



Optimization: gradient descent

■ Stochastic gradient descent

$$L(W) = \frac{1}{N} \sum_{i=1}^N L_i(x_i, y_i, W) + \lambda R(W)$$

$$\nabla_W L(W) = \frac{1}{N} \sum_{i=1}^N \nabla_W L_i(x_i, y_i, W) + \lambda \nabla_W R(W)$$

Full sum expensive
when N is large!

Approximate sum
using a **minibatch** of
examples
32 / 64 / 128 common

```
# Vanilla Minibatch Gradient Descent
```

```
while True:
```

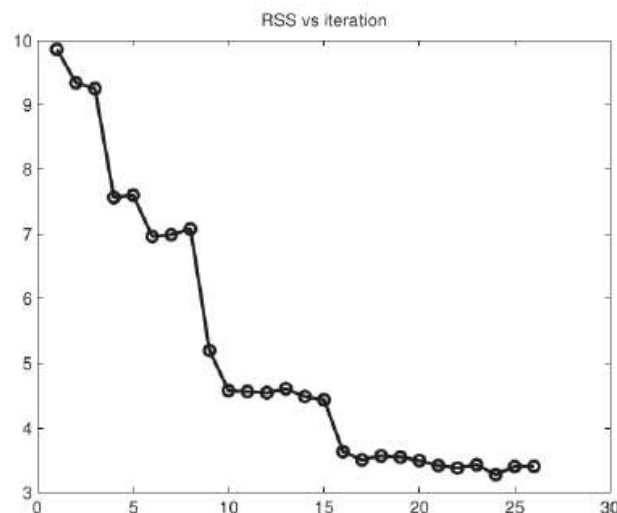

```
    data_batch = sample_training_data(data, 256) # sample 256 examples
```

```
    weights_grad = evaluate_gradient(loss_fun, data_batch, weights)
```

```
    weights += - step_size * weights_grad # perform parameter update
```

Optimization: gradient descent

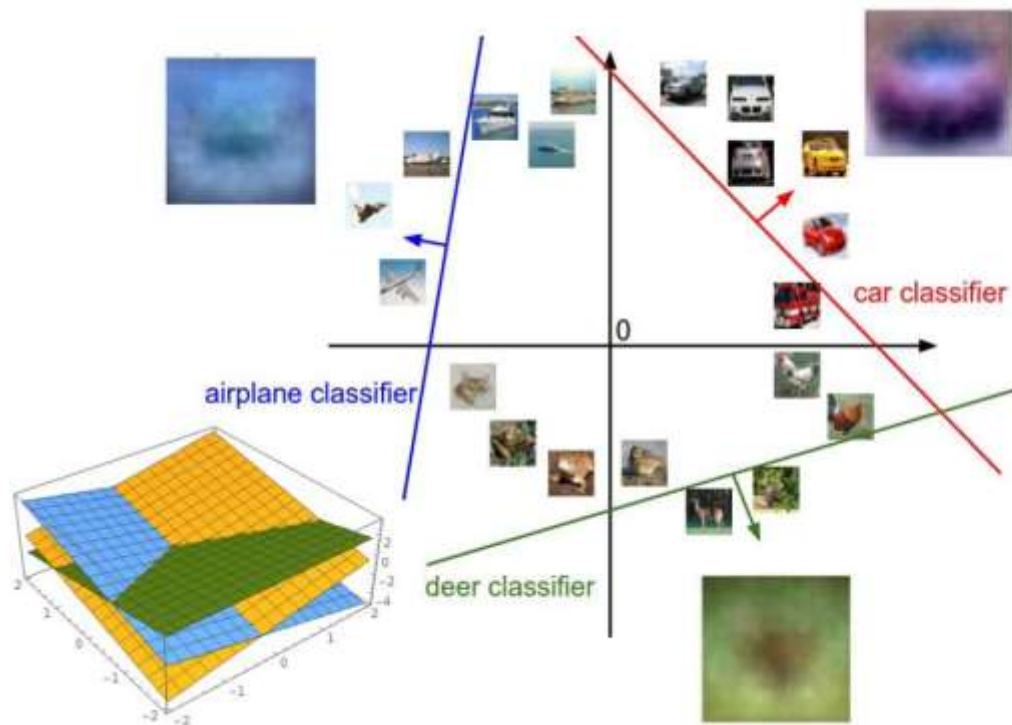
■ Stochastic gradient descent



- ▶ the objective does not always decrease for each step
- ▶ comparing to GD, SGD needs more steps, but each step is cheaper
- ▶ mini-batch, say pick up 100 samples and do average, may accelerate the convergence

Interpreting network weights

- What are those weights?



$$f(x, W) = Wx + b$$



Array of **32x32x3** numbers
(3072 numbers total)

Outline


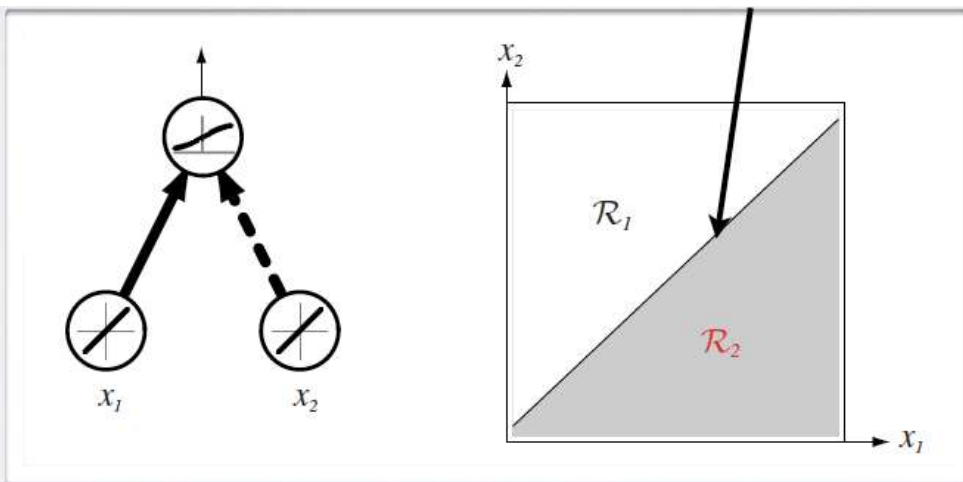
- Single layer neural networks
 - Network models
 - Example: Logistic Regression
- Multi-layer neural networks
 - Limitations of single layer networks
 - Networks with single hidden layer

Acknowledgement: Hugo Larochelle's, Mehryar Mohri@NYU's & Yingyu Liang@Princeton's course notes

Capacity of single neuron

■ Binary classification

- A neuron estimates $P(y = 1|\mathbf{x}) = \sigma(\mathbf{w}^\top \mathbf{x})$
- Its decision boundary is linear, determined by its weights



Capacity of single neuron

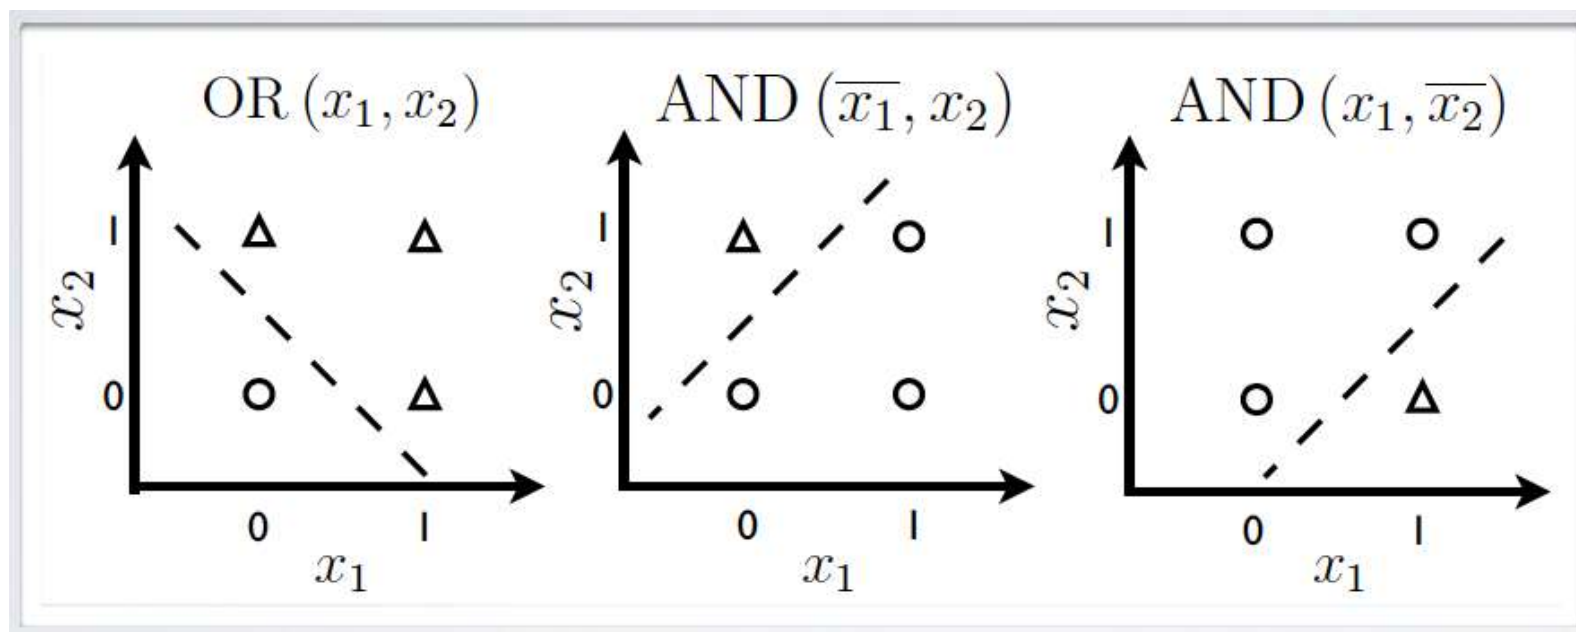
- Can solve linearly separable problems

$$\mathcal{D} = \mathcal{D}^+ \cup \mathcal{D}^-$$

$$\exists \mathbf{w}^*, \mathbf{w}^{*\top} \mathbf{x} > 0, \forall \mathbf{x} \in \mathcal{D}^+$$


$$\mathbf{w}^{*\top} \mathbf{x} < 0, \forall \mathbf{x} \in \mathcal{D}^-$$

- Examples



Capacity of single neuron


- Can't solve non linearly separable problems



- Can we use multiple neurons to achieve this?

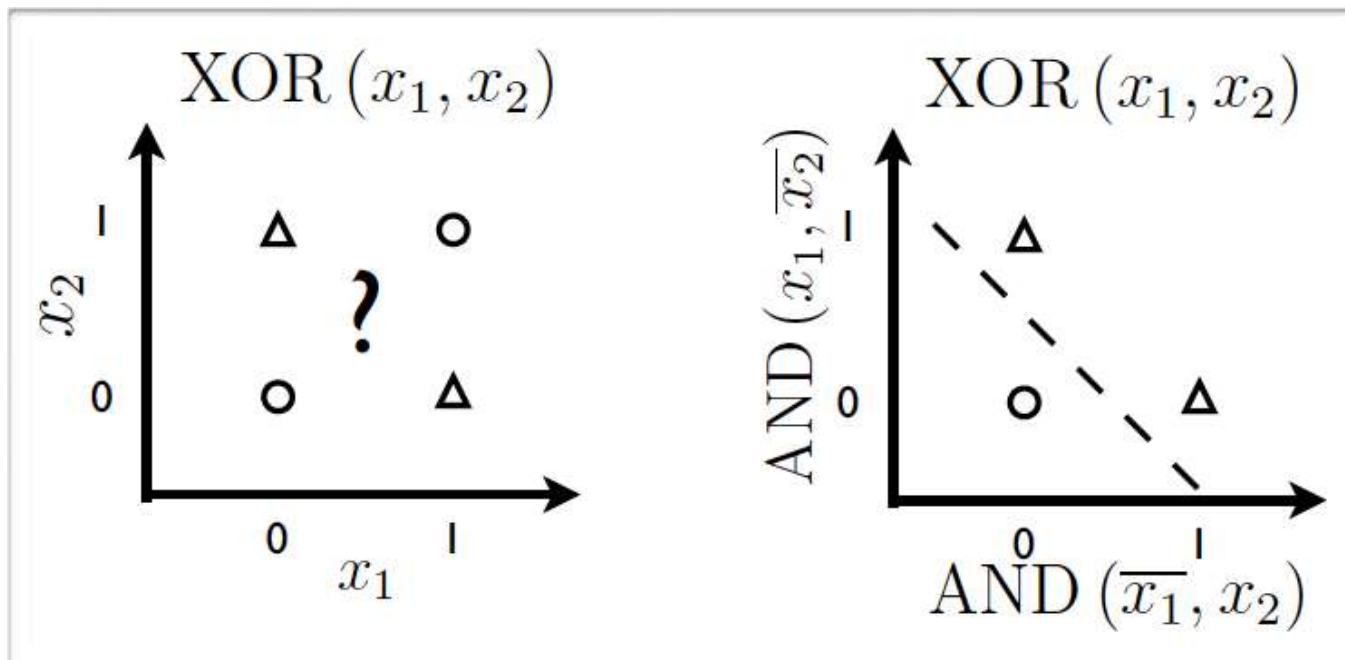
Capacity of single neuron

- Can't solve non linearly separable problems
- Unless the input is transformed in a better representation



Capacity of single neuron

- Can't solve non linearly separable problems



- Unless the input is transformed in a better representation

Adding one more layer

- Single hidden layer neural network
 - 2-layer neural network: ignoring input units

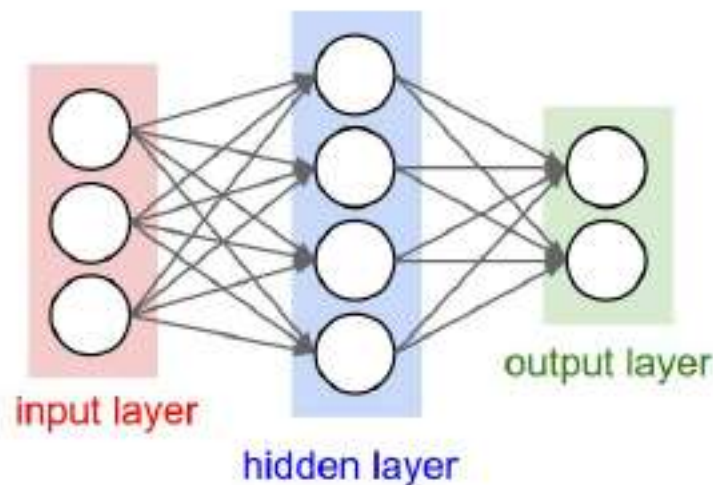
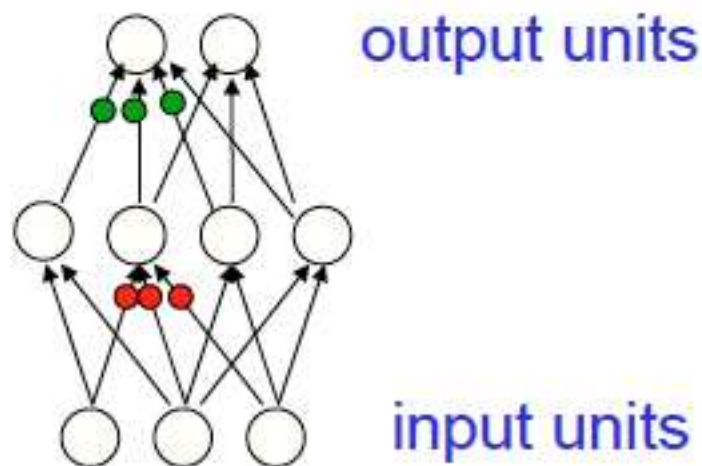
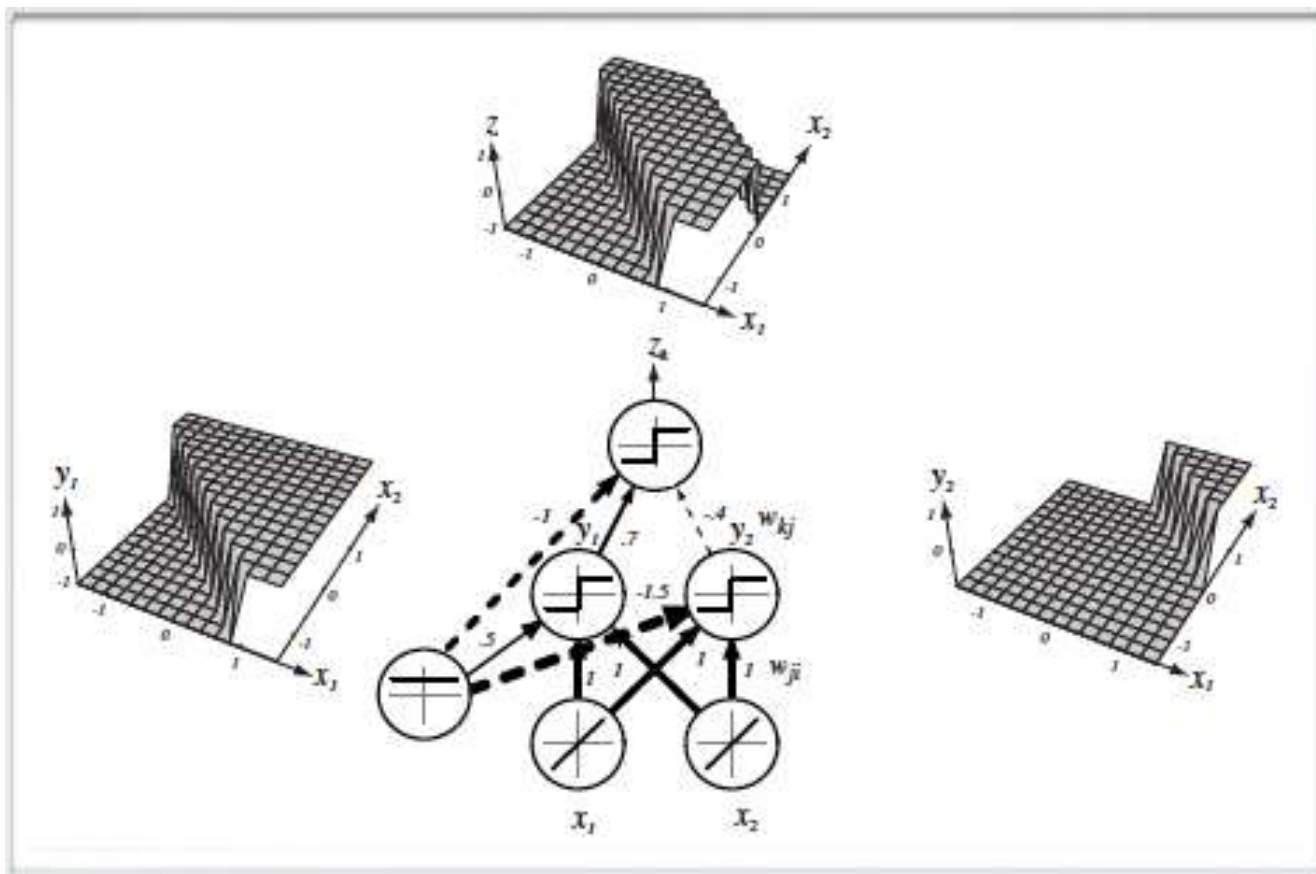


Figure : Two different visualizations of a 2-layer neural network. In this example: 3 input units, 4 hidden units and 2 output units

- Q: What if using linear activation in hidden layer?


Capacity of neural network

- Single hidden layer neural network
 - Partition the input space into regions




Capacity of neural network

- Single hidden layer neural network
 - Form a stump/delta function



Capacity of neural network

- Single hidden layer neural network




Multi-layer perceptron

■ Boolean case

- Multilayer perceptrons (MLPs) can compute more complex Boolean functions
- MLPs can compute **any** Boolean function
 - Since they can emulate individual gates
- MLPs are *universal Boolean functions*

$$((A \& \bar{X} \& Z) | (A \& \bar{Y})) \& ((X \& Y) | \overline{(X \& Z)})$$



Capacity of neural network

- Universal approximation

- Theorem (Hornik, 1991)

- A single hidden layer neural network with a linear output unit can approximate any continuous function arbitrarily well, **given enough hidden units**.

- The result applies for sigmoid, tanh and many other hidden layer activation functions

- Caveat: good result but not useful in practice

- How many hidden units?
 - How to find the parameters by a learning algorithm?

General neural network

- Multi-layer neural network





Figure : A 3-layer neural net with 3 input units, 4 hidden units in the first and second hidden layer and 1 output unit

- Naming conventions; a N -layer neural network:
 - ▶ $N - 1$ layers of hidden units
 - ▶ One output layer

Multilayer networks



network: sequence of parallelized weighted sums and non-linearities

DEFINE $\mathbf{x}^{(0)} \equiv \mathbf{x}$, $\mathbf{x}^{(1)} \equiv \mathbf{h}$, ETC.

1st layer

$$\mathbf{s}^{(1)} = \mathbf{W}^{(1)} \mathbf{x}^{(0)}$$

$$\mathbf{x}^{(1)} = \sigma(\mathbf{s}^{(1)})$$

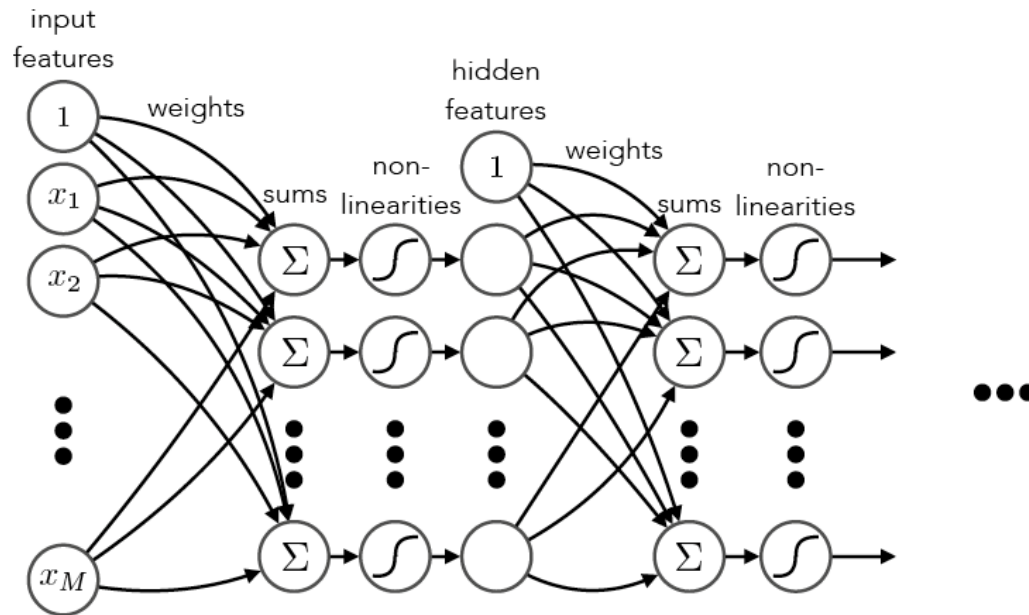
2nd layer

$$\mathbf{s}^{(2)} = \mathbf{W}^{(2)} \mathbf{x}^{(1)}$$

$$\mathbf{x}^{(2)} = \sigma(\mathbf{s}^{(2)})$$

...

Multilayer networks



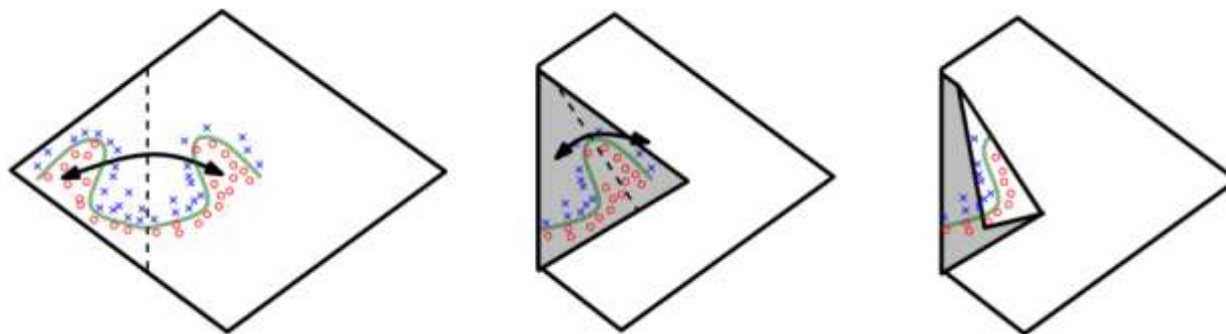
network: *sequence of parallelized weighted sums and non-linearities*

$$\begin{array}{c} \text{output} \end{array} = \sigma \left(\dots \sigma \left(\begin{array}{c} \text{2nd weights} \end{array} \sigma \left(\begin{array}{c} \text{1st weights} \end{array} \begin{array}{c} \text{input} \end{array} \right) \right) \dots \right)$$

The equation shows the mathematical representation of the network. The output is a vertical blue bar. It is equal to a sequence of non-linear functions σ applied to weighted sums. The weights are represented by blue squares: '2nd weights' and '1st weights'. The input is a vertical blue bar. Arrows point from the text labels to their corresponding symbols in the equation.

Why more layers (deeper)?

- A deep architecture can represent certain functions more compactly
 - (Montufar et al., NIPS'14)
 - Functions representable with a **deep rectifier net** can require an exponential number of hidden units with a shallow one.



Why more layers (deeper)?

- A deep architecture can represent certain functions more compactly
 - Example: Boolean functions
 - There are Boolean functions which require an exponential number of hidden units in the single layer case
 - require a **polynomial number of hidden units** if we can adapt the number of layers
 - Example: multivariate polynomials (Rolnick & Tegmark, ICLR'18)
 - Total number of neurons m required to approximate natural classes of multivariate polynomials of n variables
 - grows **only linearly with n** for deep neural networks, but grows exponentially when merely a single hidden layer is allowed.

Why more layers (deeper)?



<https://youtu.be/aircAruvnKk?list=PLZHQObOWTQDN>
U6R1_67000Dx_ZCJB-3pi

Other network connectivity


sequential connectivity: *information must flow through the entire sequence to reach the output*



information may not be able to propagate easily


→ *make shorter paths to output*

residual & highway
connections



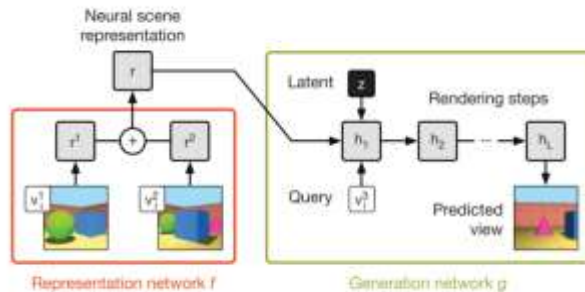
Deep residual learning for image recognition, He et al., 2016
Highway networks, Srivastava et al., 2015

dense (concatenated)
connections



Densely connected convolutional networks, Huang et al., 2017

Modern MLP as Implicit Representation




Generative Query Networks
[Eslami et al. 2018]



[Flynn et al., 2016; Zhou et al., 2018b;
Mildenhall et al. 2019]


Multiplane Images (MPIs)



Voxel Grids + CNN decoder



DeepVoxels
[Sitzmann et al. 2019]




Neural Volumes
[Lombardi et al. 2019]

Voxel Grids + Ray Marching



SRN
[Sitzmann et al. 2019b]



NeRF
[Mildenhall et al. 2020]




IDR
[Yariv et al. 2020]

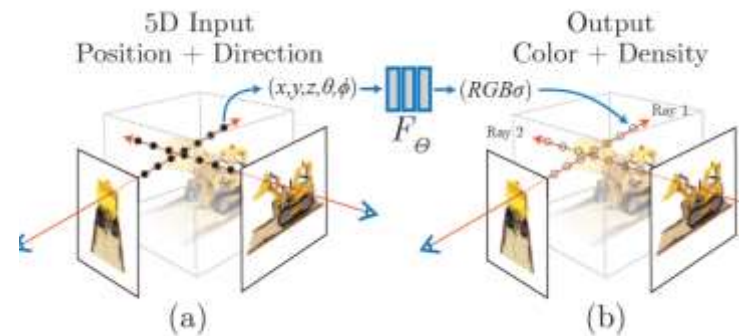
Implicit Fields

Modern MLP in NeRF

- - Color + Density
- Positional Encoding
- Volume Rendering



Representing Scenes as Neural Radiance Fields for View Synthesis, Mildenhall et al., *ECCV 2020 Oral - Best Paper Honorable Mention*



Modern MLP in NeRF

NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis

Ben Mildenhall*
UC Berkeley

Pratul P. Srinivasan*
UC Berkeley


Matthew Tancik*
UC Berkeley

Jonathan T. Barron
Google Research

Ravi Ramamoorthi
UC San Diego

Ren Ng
UC Berkeley

* Denotes Equal Contribution



<https://youtu.be/JuH79E8rdKc>

Outline

- Single layer neural networks
 - Network models; Example: Logistic Regression
- Multi-layer neural networks
 - Limitations of single layer networks
 - Neural networks with single hidden layer
 - Sequential network architecture and variants
- Inference and learning
 - Forward and Backpropagation
 - Examples: one-layer network
 - General BP algorithm

Computation in neural network

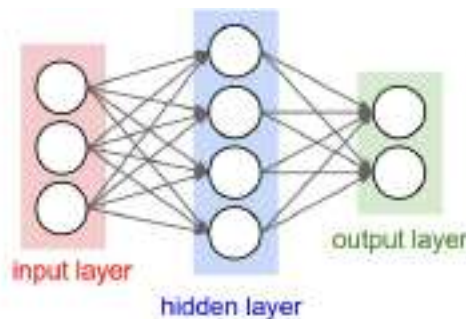
- We only need to know two algorithms
 - Inference/prediction: simply forward pass
 - Parameter learning: needs backward pass
- Basic fact:
 - A neural network is a function of composed operations

$$f_L(\mathbf{w}_L, f_{L-1}(\mathbf{w}_{L-1}, \dots f_1(\mathbf{w}_1, \mathbf{x}) \dots))$$

- All the f functions are linear + (simple) nonlinear (differentiable a.e.) operators

Inference example: Forward Pass

- What does the network compute?




- Output of the network can be written as:

$$h_j(\mathbf{x}) = f(v_{j0} + \sum_{i=1}^D x_i v_{ji})$$
$$o_k(\mathbf{x}) = g(w_{k0} + \sum_{j=1}^J h_j(\mathbf{x}) w_{kj})$$

(j indexing hidden units, k indexing the output units, D number of inputs)

Forward Pass in Python

- Example code for a forward pass for a 3-layer network in Python:



```
# forward-pass of a 3-layer neural network:  
f = lambda x: 1.0/(1.0 + np.exp(-x)) # activation function (use sigmoid)  
x = np.random.randn(3, 1) # random input vector of three numbers (3x1)  
h1 = f(np.dot(W1, x) + b1) # calculate first hidden layer activations (4x1)  
h2 = f(np.dot(W2, h1) + b2) # calculate second hidden layer activations (4x1)  
out = np.dot(W3, h2) + b3 # output neuron (1x1)
```

- Can be implemented efficiently using matrix operations

Parameter learning: Backward Pass

■ Supervised learning framework

- Find weights:

$$\mathbf{w}^* = \operatorname{argmin}_{\mathbf{w}} \sum_{n=1}^N \operatorname{loss}(\mathbf{o}^{(n)}, \mathbf{t}^{(n)})$$

where $\mathbf{o} = f(\mathbf{x}; \mathbf{w})$ is the output of a neural network

- Define a loss function, eg:

- ▶ Squared loss: $\sum_k \frac{1}{2} (o_k^{(n)} - t_k^{(n)})^2$
- ▶ Cross-entropy loss: $-\sum_k t_k^{(n)} \log o_k^{(n)}$

- Gradient descent:

$$\mathbf{w}^{t+1} = \mathbf{w}^t - \eta \frac{\partial E}{\partial \mathbf{w}^t}$$

where η is the learning rate (and E is error/loss)

Backward pass

■ Backpropagation


- An efficient method for computing gradients in NNs
- A neural network as a function of composed operations

$$f_L(\mathbf{w}_L, f_{L-1}(\mathbf{w}_{L-1}, \dots f_1(\mathbf{w}_1, \mathbf{x}) \dots))$$

and the loss \mathcal{L} is a function of the network output

→ use chain rule to calculate gradients

Backward pass



<https://www.youtube.com/watch?v=llg3gGewQ5U>

Gradient descent iteration

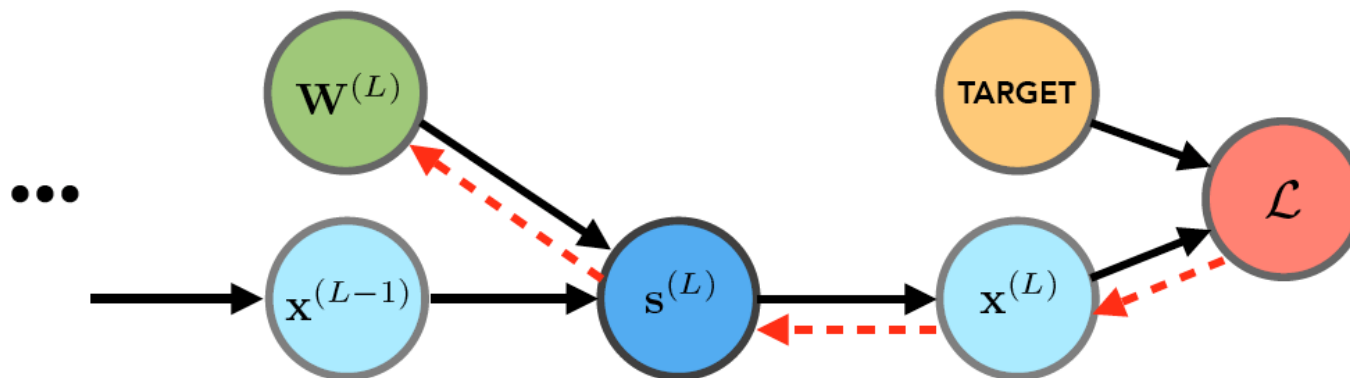
■ Forward pass

1st layer	2nd layer	...	Loss
$\mathbf{s}^{(1)} = \mathbf{W}^{(1)} \mathbf{x}^{(0)}$	$\mathbf{s}^{(2)} = \mathbf{W}^{(2)} \mathbf{x}^{(1)}$		\mathcal{L}
$\mathbf{x}^{(1)} = \sigma(\mathbf{s}^{(1)})$	$\mathbf{x}^{(2)} = \sigma(\mathbf{s}^{(2)})$		

■ Backward pass


calculate $\nabla_{W^{(1)}} \mathcal{L}, \nabla_{W^{(2)}} \mathcal{L}, \dots$ let's start with the final layer: $\nabla_{W^{(L)}} \mathcal{L}$

to determine the chain rule ordering, we'll draw the dependency graph



Gradient descent iteration

■ Backward pass



$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}^{(L)}} = \frac{\partial \mathcal{L}}{\partial \mathbf{x}^{(L)}} \frac{\partial \mathbf{x}^{(L)}}{\partial \mathbf{s}^{(L)}} \frac{\partial \mathbf{s}^{(L)}}{\partial \mathbf{W}^{(L)}}$$

depends on the
form of the loss

derivative of the
non-linearity

$$\frac{\partial}{\partial \mathbf{W}^{(L)}} (\mathbf{W}^{(L)\top} \mathbf{x}^{(L-1)}) \\ = \mathbf{x}^{(L-1)\top}$$

note $\nabla_{\mathbf{W}^{(L)}} \mathcal{L} \equiv \frac{\partial \mathcal{L}}{\partial \mathbf{W}^{(L)}}$ is notational convention

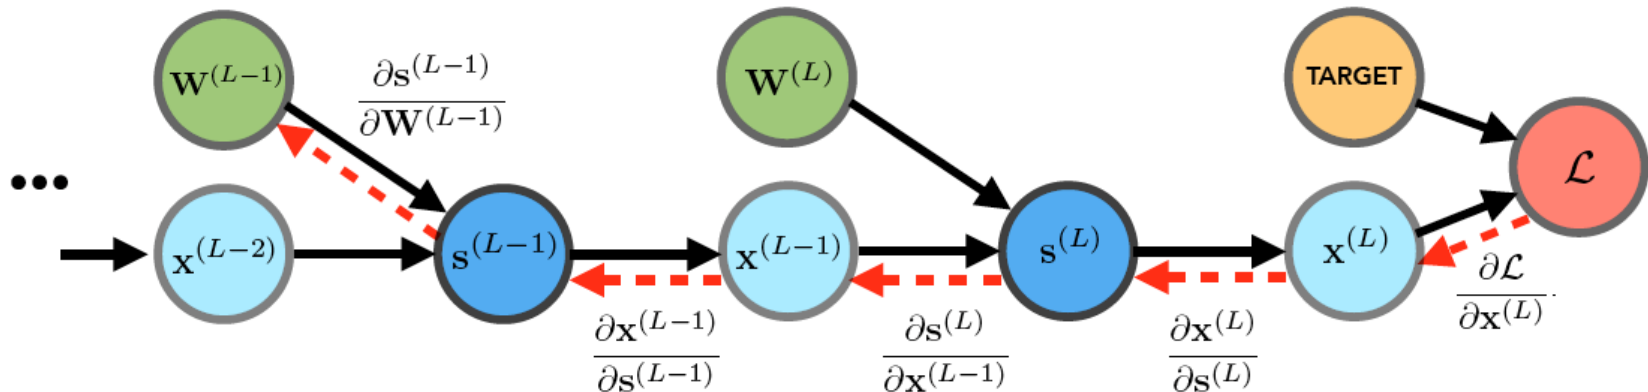
The order needs
to be reversed for
Jacobians! See
LN04 for details

Gradient descent iteration

■ Backward pass

now let's go back one more layer...

again we'll draw the dependency graph:

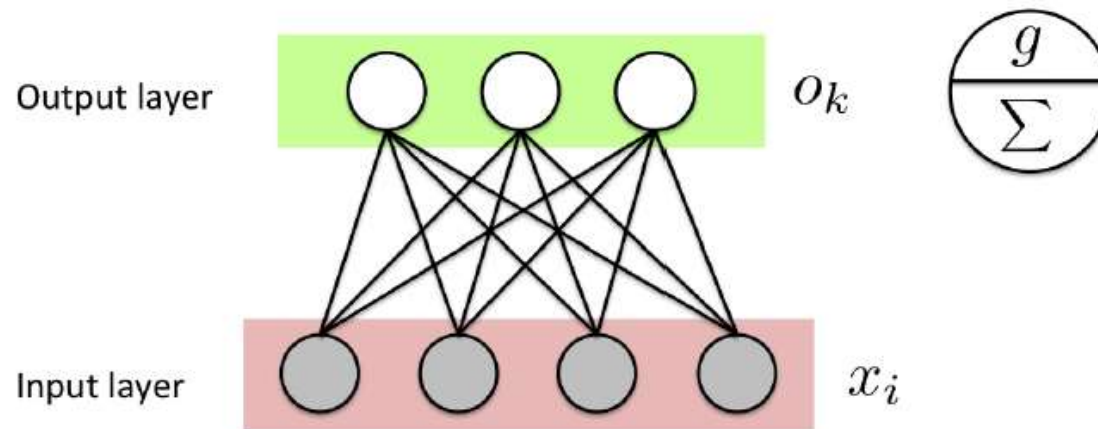


$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}^{(L)}} = \frac{\partial \mathcal{L}}{\partial \mathbf{x}^{(L)}} \frac{\partial \mathbf{x}^{(L)}}{\partial \mathbf{s}^{(L)}} \frac{\partial \mathbf{s}^{(L)}}{\partial \mathbf{x}^{(L-1)}} \frac{\partial \mathbf{x}^{(L-1)}}{\partial \mathbf{s}^{(L-1)}} \frac{\partial \mathbf{s}^{(L-1)}}{\partial \mathbf{W}^{(L-1)}}$$

The order needs to be reversed for Jacobians! See LN04 for details



Example: Single Layer Network

- Let's take a single layer network




Example: Single Layer Network

- Let's take a single layer network and draw it a bit differently




Example: Single Layer Network



- Error gradients for single layer network:

$$\frac{\partial E}{\partial w_{ki}} =$$

Example: Single Layer Network




- Error gradients for single layer network:

$$\frac{\partial E}{\partial w_{ki}} = \frac{\partial E}{\partial o_k} \frac{\partial o_k}{\partial z_k} \frac{\partial z_k}{\partial w_{ki}}$$

- Error gradient is computable for any continuous activation function $g()$, and any continuous error function


Example: Single Layer Network



- Error gradients for single layer network:

$$\frac{\partial E}{\partial w_{ki}} = \underbrace{\frac{\partial E}{\partial o_k}}_{\delta_k^o} \frac{\partial o_k}{\partial z_k} \frac{\partial z_k}{\partial w_{ki}}$$

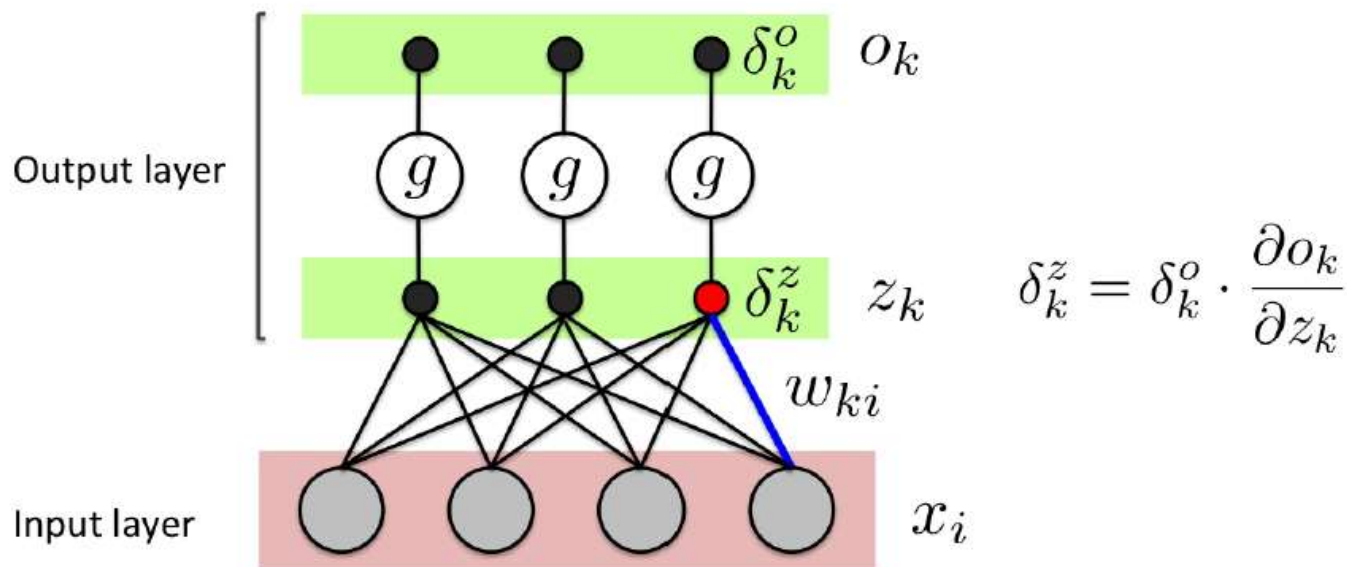
Example: Single Layer Network



- Error gradients for single layer network:

$$\frac{\partial E}{\partial w_{ki}} = \frac{\partial E}{\partial o_k} \frac{\partial o_k}{\partial z_k} \frac{\partial z_k}{\partial w_{ki}} = \delta_k^o \frac{\partial o_k}{\partial z_k} \frac{\partial z_k}{\partial w_{ki}}$$


Example: Single Layer Network



- Error gradients for single layer network:

$$\frac{\partial E}{\partial w_{ki}} = \frac{\partial E}{\partial o_k} \frac{\partial o_k}{\partial z_k} \frac{\partial z_k}{\partial w_{ki}} = \underbrace{\delta_k^o \cdot \frac{\partial o_k}{\partial z_k}}_{\delta_k^z} \frac{\partial z_k}{\partial w_{ki}}$$

Example: Single Layer Network



- Error gradients for single layer network:

$$\frac{\partial E}{\partial w_{ki}} = \frac{\partial E}{\partial o_k} \frac{\partial o_k}{\partial z_k} \frac{\partial z_k}{\partial w_{ki}} = \delta_k^z \frac{\partial z_k}{\partial w_{ki}} = \delta_k^z \cdot x_i$$

Outline

- Multi-layer neural networks
 - Limitations of single layer networks
 - Neural networks with single hidden layer
 - Sequential network architecture and variants
- Inference and learning
 - Forward and Backpropagation
 - Examples: one-layer network
 - General BP algorithm


An implementation perspective

- Example: Univariate logistic least square model

$$s = wx + b$$

$$y = \sigma(s)$$

$$\mathcal{L} = \frac{1}{2}(y - t)^2$$



Univariate chain rule

- A structured way to implement it
 - The goal is to write [a program](#) that efficiently computes the derivatives

Computing the loss:

$$s = wx + b$$

$$y = \sigma(s)$$

$$\mathcal{L} = \frac{1}{2}(y - t)^2$$

Computing the derivatives:

$$\frac{d\mathcal{L}}{dy} = y - t$$


$$\frac{d\mathcal{L}}{ds} = \frac{d\mathcal{L}}{dy} \sigma'(s)$$

$$\frac{d\mathcal{L}}{dw} = \frac{d\mathcal{L}}{ds} x$$

$$\frac{d\mathcal{L}}{db} = \frac{d\mathcal{L}}{ds}$$

Computation graph

- Represent the computations using a **computation graph**
 - Nodes: inputs & computed quantities
 - Edges: which nodes are computed directly as function of which other nodes



Univariate chain rule

■ A shorthand notation

- Use $\delta_y := d\mathcal{L}/dy$, called the error signal
- Note that the error signals are values computed by the program

Computing the loss:

$$s = wx + b$$

$$y = \sigma(s)$$

$$\mathcal{L} = \frac{1}{2}(y - t)^2$$


Computing the derivatives:

$$\delta_y = y - t$$

$$\delta_s = \delta_y \sigma'(s)$$

$$\delta_w = \delta_s x$$


$$\delta_b = \delta_s$$



Multivariate chain rule

- The computation graph has fan-out > 1

L_2 -Regularized regression



$$z = wx + b$$


$$y = \sigma(z)$$

$$\mathcal{L} = \frac{1}{2}(y - t)^2$$

$$\mathcal{R} = \frac{1}{2}w^2$$

$$\mathcal{L}_{\text{reg}} = \mathcal{L} + \lambda \mathcal{R}$$

Multiclass logistic regression



$$z_\ell = \sum_j w_{\ell j} x_j + b_\ell$$

$$y_k = \frac{e^{z_k}}{\sum_\ell e^{z_\ell}}$$

$$\mathcal{L} = - \sum_k t_k \log y_k$$

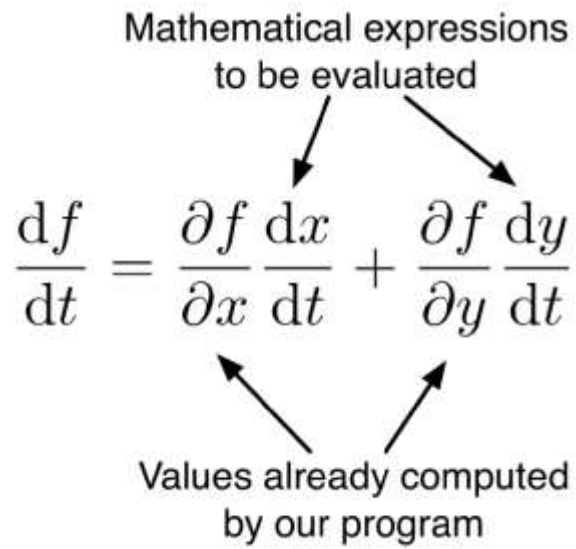
Multivariable chain rule


- Recall the distributed chain rule

Mathematical expressions
to be evaluated

$$\frac{df}{dt} = \frac{\partial f}{\partial x} \frac{dx}{dt} + \frac{\partial f}{\partial y} \frac{dy}{dt}$$

Values already computed
by our program





- The shorthand notation:

$$\delta_t = \delta_x \frac{dx}{dt} + \delta_y \frac{dy}{dt}$$

General Backpropagation

- Given a computation graph

Let v_1, \dots, v_N be a **topological ordering** of the computation graph (i.e. parents come before children.)

v_N denotes the variable we're trying to compute derivatives of (e.g. loss)

forward pass


For $i = 1, \dots, N$
Compute v_i as a function of $\text{Pa}(v_i)$

backward pass

$\delta_{v_N} = 1$
For $i = N - 1, \dots, 1$
$$\delta_{v_i} = \sum_{j \in \text{Ch}(v_i)} \delta_{v_j} \frac{\partial v_j}{\partial v_i}$$

General Backpropagation

- Example: univariate logistic least square regression



Forward pass:

$$z = wx + b$$

$$y = \sigma(z)$$

$$\mathcal{L} = \frac{1}{2}(y - t)^2$$

$$\mathcal{R} = \frac{1}{2}w^2$$

$$\mathcal{L}_{\text{reg}} = \mathcal{L} + \lambda\mathcal{R}$$

Backward pass:

$$\delta_{\mathcal{L}_{\text{reg}}} =$$

$$\delta_z =$$

$$\delta_{\mathcal{R}} =$$

$$=$$

$$=$$

$$\delta_w =$$

$$\delta_{\mathcal{L}} =$$

$$=$$

$$=$$

$$\delta y =$$


$$\delta_b =$$

$$=$$

$$=$$

General Backpropagation

■ Example: Multilayer Perceptron (multiple outputs)



Forward pass:

$$z_i = \sum_j w_{ij}^{(1)} x_j + b_i^{(1)}$$

$$h_i = \sigma(z_i)$$

$$y_k = \sum_i w_{ki}^{(2)} h_i + b_k^{(2)}$$

$$\mathcal{L} = \frac{1}{2} \sum_k (y_k - t_k)^2$$

Backward pass:

$$\bar{\mathcal{L}} = 1$$

$$\bar{y}_k = \bar{\mathcal{L}} (y_k - t_k)$$

$$\bar{w}_{ki}^{(2)} = \bar{y}_k h_i$$

$$\bar{b}_k^{(2)} = \bar{y}_k$$

$$\bar{h}_i = \sum_k \bar{y}_k w_{ki}^{(2)}$$


$$\bar{z}_i = \bar{h}_i \sigma'(z_i)$$

$$\bar{w}_{ij}^{(1)} = \bar{z}_i x_j$$

$$\bar{b}_i^{(1)} = \bar{z}_i$$

General Backpropagation


- Backprop as message passing:



- Each node receives a set of messages from its children, which are aggregated into its error signal, then it passes messages to its parents
- **Modularity:** each node only has to know how to compute derivatives w.r.t. its arguments – **local computation in the graph**

Patterns in backward flow

- Multiplicative node

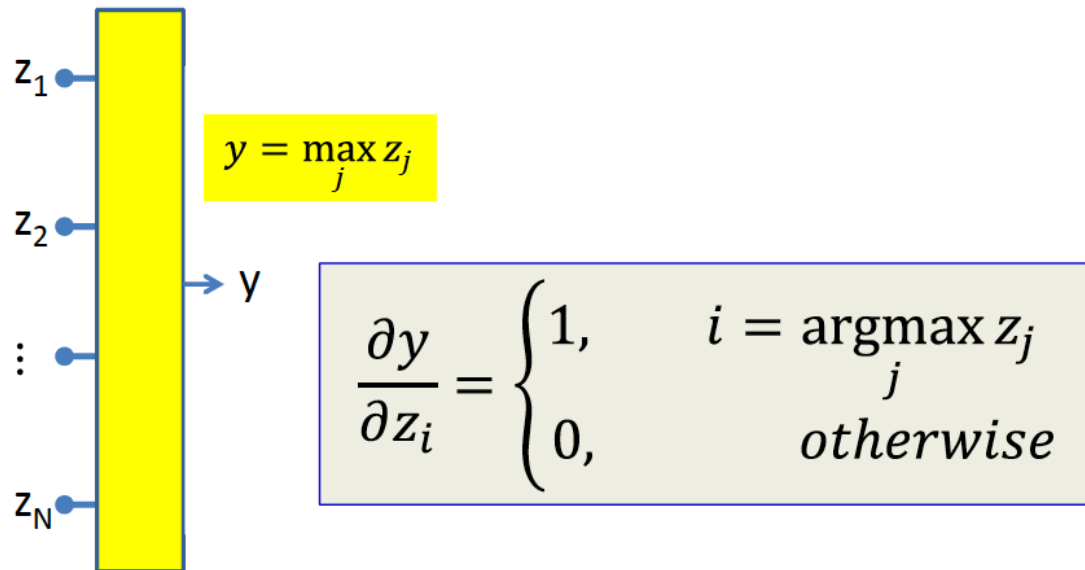


Forward: $o_i^{(k)} = y_j^{(k-1)} y_l^{(k-1)}$

$$\frac{\partial L}{\partial y_j^{(k-1)}} = \frac{\partial L}{\partial o_i^{(k)}} \frac{\partial o_i^{(k)}}{\partial y_j^{(k-1)}} = y_l^{(k-1)} \frac{\partial L}{\partial o_i^{(k)}}$$

Patterns in backward flow


■ Max node



- Vector equivalent of subgradient
 - 1 w.r.t. the largest incoming input
 - Incremental changes in this input will change the output
 - 0 for the rest
 - Incremental changes to these inputs will not change the output

Computation cost

- Forward pass: one add-multiply operation per weight
- Backward pass: two add-multiply operations per weight



$$z_i = \sum_j w_{ij}^{(1)} x_j + b_i^{(1)}$$

$$\begin{aligned} \bar{w}_{ki}^{(2)} &= \bar{y}_k h_i \\ \bar{h}_i &= \sum_k \bar{y}_k w_{ki}^{(2)} \end{aligned}$$

- For a multilayer network, the cost is linear in the number of layers, quadratic in the number of units per layer

Backpropagation

- Backprop is used to train the majority of neural nets
 - Even generative network learning, or advanced optimization algorithms (second-order) use backprop to compute the update of weights
- However, backprop seems biologically implausible
 - No evidence for biological signals analogous to error derivatives
 - All the existing biologically plausible alternatives learn much more slowly on computers.
 - So how on earth does the brain learn???

Coding examples

- Getting familiar with Pytorch

- ☐ Python Tutorial: <https://cs231n.github.io/python-numpy-tutorial/>
- ☐ PyTorch in 60 mins:
https://pytorch.org/tutorials/beginner/deep_learning_60min_blitz.html

- Predicting house prices

- ☐ https://d2l.ai/chapter_multilayer-perceptrons/kaggle-house-price.html

Summary

- Artificial neurons, Single-layer network
- Multi-layer neural networks
- Inference and learning
 - Forward and Backpropagation
- Next time ...
 - Modern topics about MLP, CNN
- **Reference:**
 - [d2l.ai: 4.1-4.3, 4.7](#)
 - [DLBook: Chapter 6](#)