



Lecture 04: CNNs II – Network Regularization & Architecture

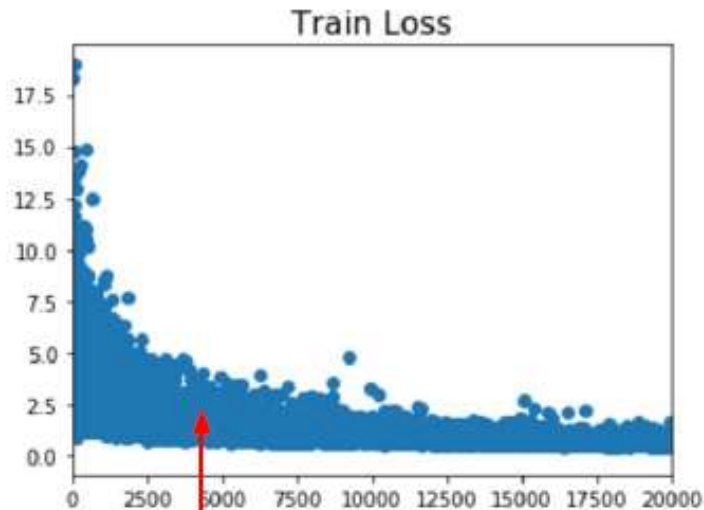
Lan Xu
SIST, ShanghaiTech
Fall, 2023

Training overview

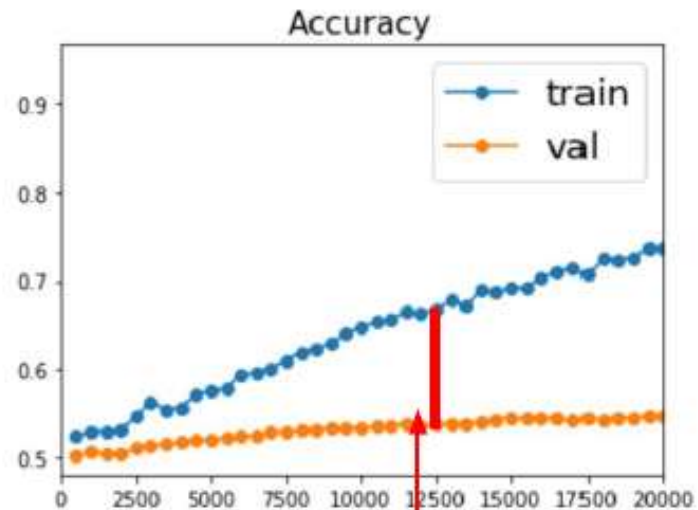
- Two aspects of training networks
 - Optimization
 - How do we minimize the loss function effectively?
 - Generalization
 - How do we avoid overfitting?
- CNN training pipeline
 - Data processing
 - Weight initialization
 - Parameter updates
 - Batch normalization
- Avoid overfitting: Regularization

Beyond Training Error

- How do we generalize to unseen data?
 - Well studied but still poorly understood



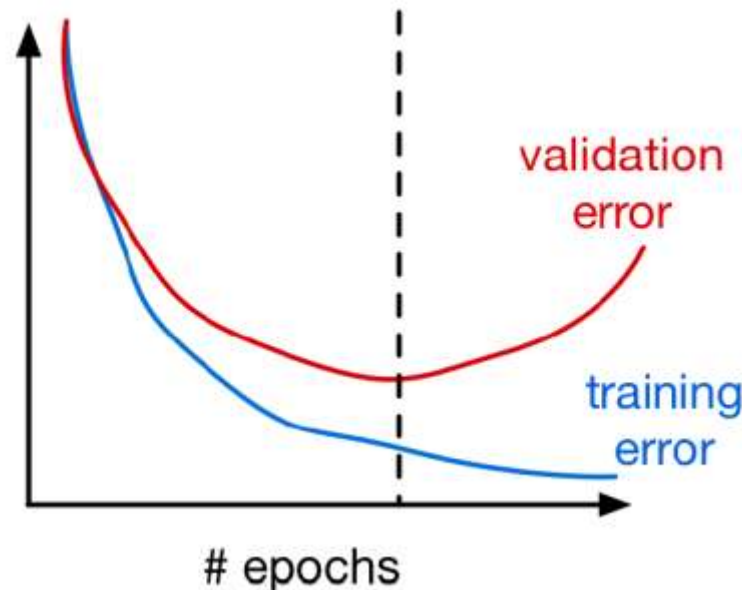
Better optimization algorithms
help reduce training loss



But we really care about error on new
data - how to reduce the gap?

Early Stopping

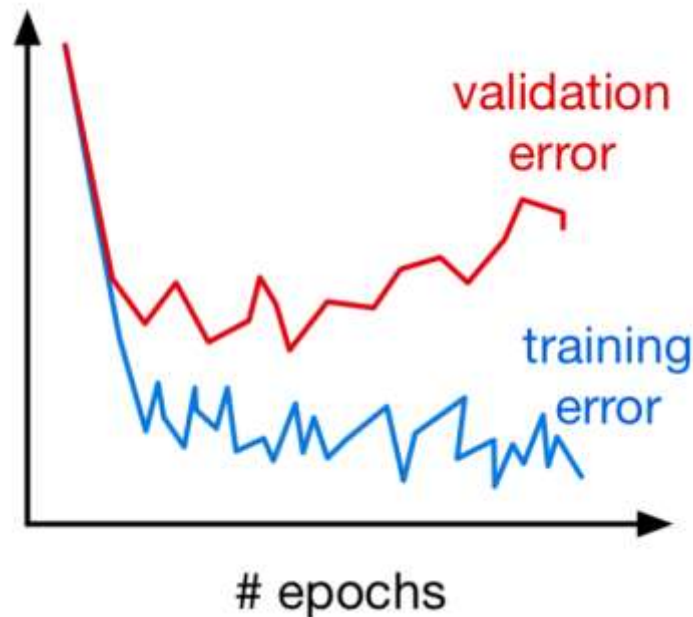
- Early stopping: monitor performance on a validation set, stop training when the validation error starts going up.
 - We don't always want to find a global (or even local) optimum of our cost function.



- Weights start out small, so it takes time for them to grow large. Therefore, it has a similar effect to weight decay.

Early Stopping

- A slight catch: validation error fluctuates because of stochasticity in the updates.
 - Determining when the validation error has actually leveled off can be tricky.
 - May use temporal smoothing

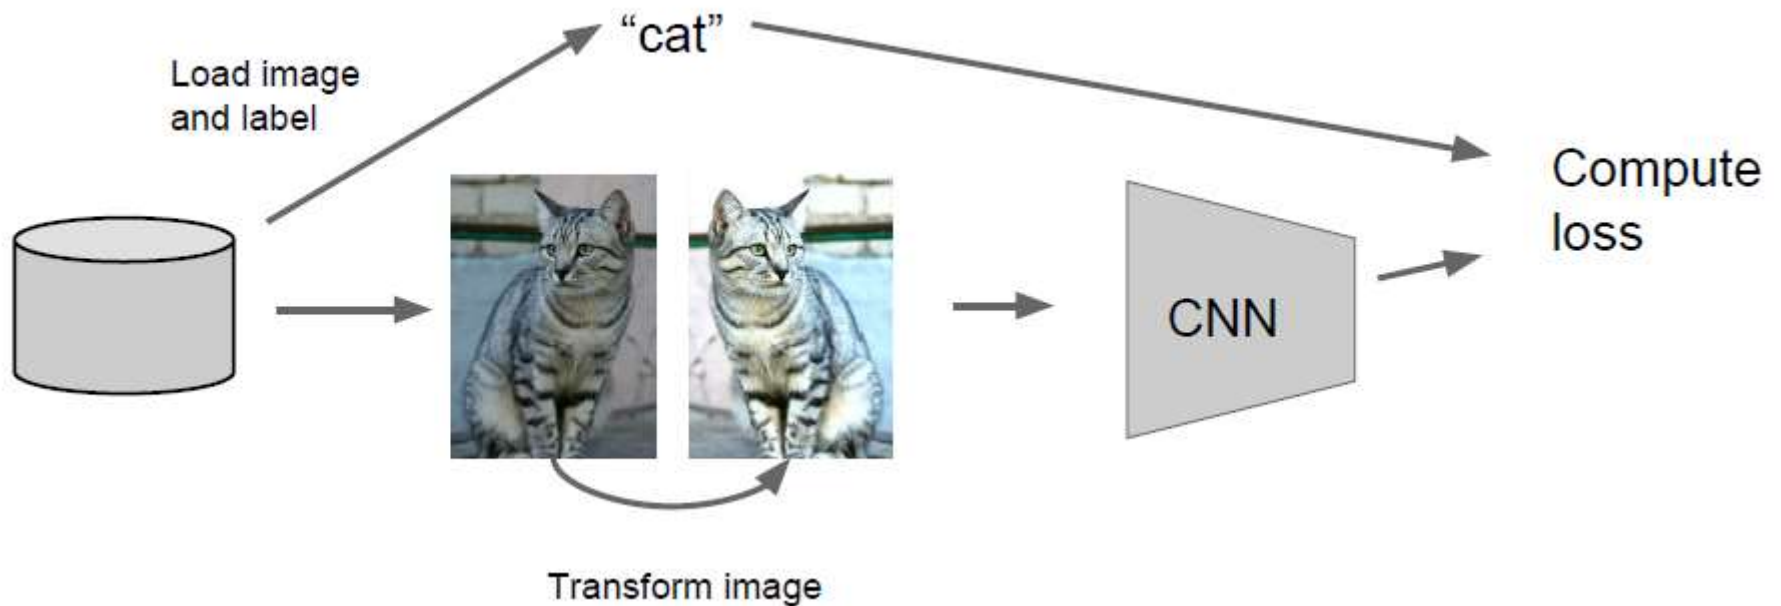


Outline

- Regularization in CNN training
 - Data Augmentation
 - Weight Regularization & Transfer Learning
 - Stochastic Regularization
 - Hyper-parameter optimization

Data Augmentation

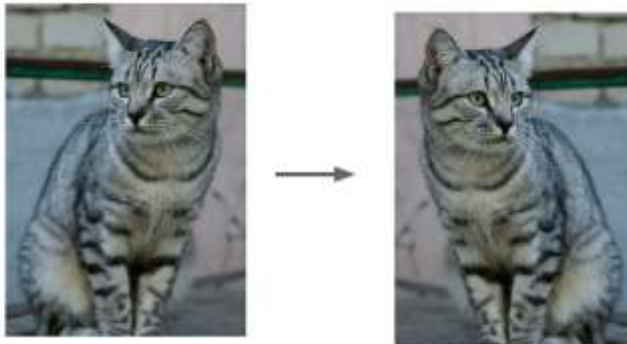
- Create more data for regularization



Data Augmentation

■ Create more data for regularization

Horizontal Flips

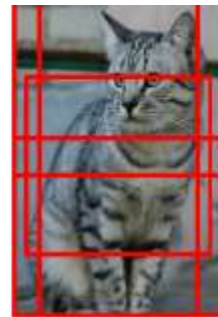


Random crops and scales

Training: sample random crops / scales

ResNet:

1. Pick random L in range $[256, 480]$
2. Resize training image, short side = L
3. Sample random 224×224 patch



Testing: average a fixed set of crops

ResNet:

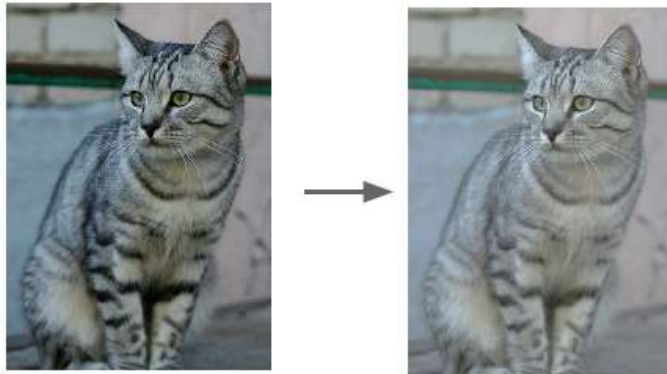
1. Resize image at 5 scales: $\{224, 256, 384, 480, 640\}$
2. For each size, use 10 224×224 crops: 4 corners + center, + flips

Data Augmentation

- Create more data for regularization

Color Jitter

Simple: Randomize
contrast and brightness



More Complex:

1. Apply PCA to all [R, G, B] pixels in training set
2. Sample a “color offset” along principal component directions
3. Add offset to all pixels of a training image

(As seen in [Krizhevsky et al. 2012], ResNet, etc)

Data Augmentation

- Create more data for regularization
- Examples (for visual recognition)
 - translation
 - horizontal or vertical
 - flip
 - rotation
 - smooth warping
 - noise (e.g. flip random pixels)
- The choice of transformations depends on the task.
 - E.g. horizontal flip for object recognition, but not handwritten digit recognition.

Data Augmentation

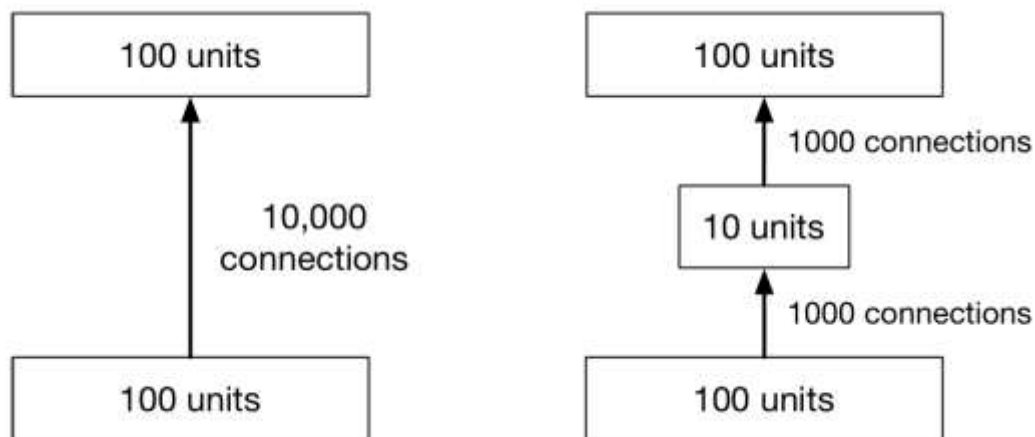
- AutoAugment (Cubuk et al, Arxiv 2018)
 - An automatic way to design custom data augmentation policies for computer vision datasets,
 - Selecting an optimal policy from a search space of 2.9×10^{32} image transformation possibilities.
 - E.g., guiding the selection of basic image transformation operations, such as flipping an image horizontally/vertically, rotating an image, changing the color of an image, etc.
 - Using reinforcement learning strategy (More later...)
- Results
 - New state of the art: ImageNet: 83.54% top1 accuracy; SVHN: error rate 1.02%.
 - AutoAugment policies are found to be transferable to other vision datasets.

Outline

- Regularization in CNN training
 - Data Augmentation
 - Weight Regularization & Transfer Learning
 - Stochastic Regularization
 - Hyper-parameter optimization

Reducing # of Parameters

- Reducing the number of layers or the number of parameters per layer.
- Adding a linear **bottleneck layer**:



- The first network is strictly more expressive than the second (i.e. it can represent a strictly larger class of functions). (Why?)
- Remember how linear layers don't make a network more expressive? They might still improve generalization.

Weight Regularization

- L_2 regularization / weight decay
 - Encouraging the weights to be small in magnitude

$$\mathcal{E}_{\text{reg}} = \mathcal{E} + \lambda \mathcal{R} = \mathcal{E} + \frac{\lambda}{2} \sum_j w_j^2$$

- The gradient update can be interpreted as weight decay

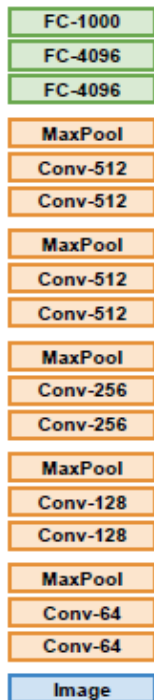
$$\begin{aligned} \mathbf{w} &\leftarrow \mathbf{w} - \alpha \left(\frac{\partial \mathcal{E}}{\partial \mathbf{w}} + \lambda \frac{\partial \mathcal{R}}{\partial \mathbf{w}} \right) \\ &= \mathbf{w} - \alpha \left(\frac{\partial \mathcal{E}}{\partial \mathbf{w}} + \lambda \mathbf{w} \right) \\ &= (1 - \alpha \lambda) \mathbf{w} - \alpha \frac{\partial \mathcal{E}}{\partial \mathbf{w}} \end{aligned}$$

Transfer Learning

Transfer Learning with CNNs

1. Train on Imagenet

Donahue et al, "DeCAF: A Deep Convolutional Activation Feature for Generic Visual Recognition", ICML 2014
Razavian et al, "CNN Features Off-the-Shelf: An Astounding Baseline for Recognition", CVPR Workshops 2014

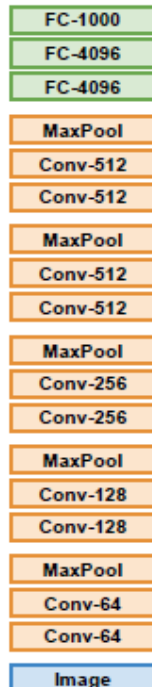


Transfer Learning

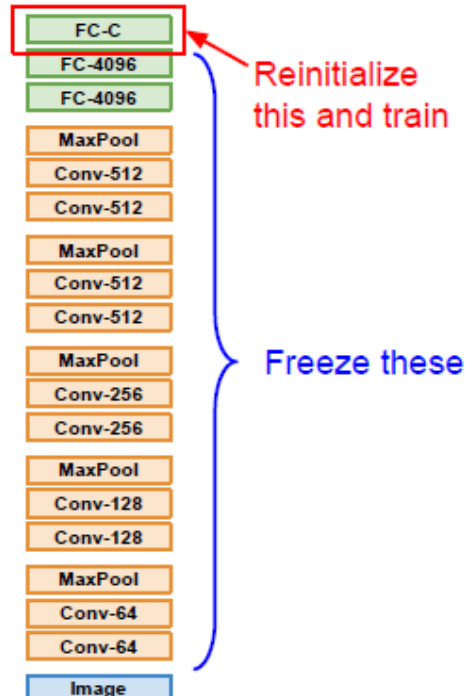
Transfer Learning with CNNs

Donahue et al, "DeCAF: A Deep Convolutional Activation Feature for Generic Visual Recognition", ICML 2014
Razavian et al, "CNN Features Off-the-Shelf: An Astounding Baseline for Recognition", CVPR Workshops 2014

1. Train on Imagenet



2. Small Dataset (C classes)

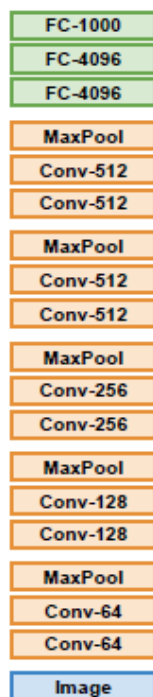


Transfer Learning

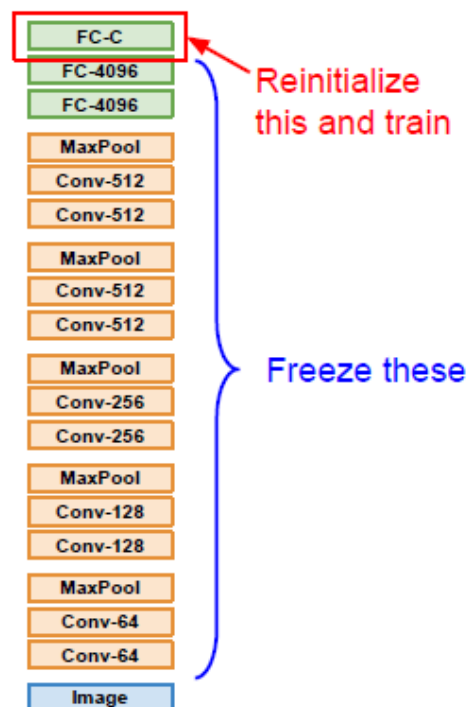
Transfer Learning with CNNs

Donahue et al, "DeCAF: A Deep Convolutional Activation Feature for Generic Visual Recognition", ICML 2014
Razavian et al, "CNN Features Off-the-Shelf: An Astounding Baseline for Recognition", CVPR Workshops 2014

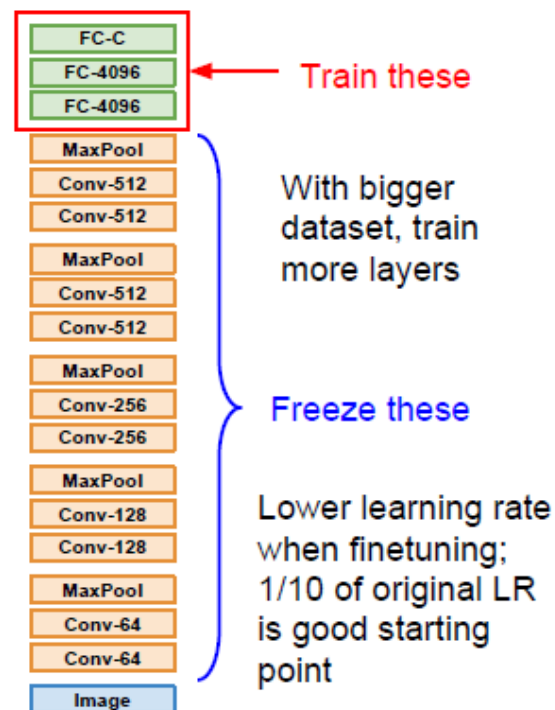
1. Train on Imagenet



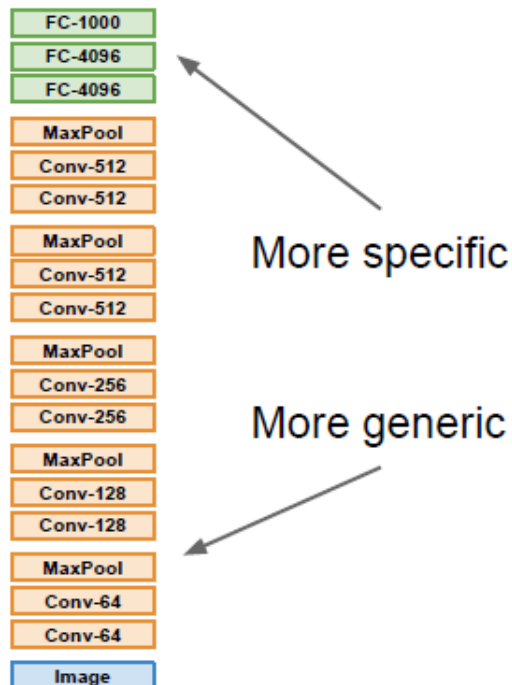
2. Small Dataset (C classes)



3. Bigger dataset



Transfer Learning



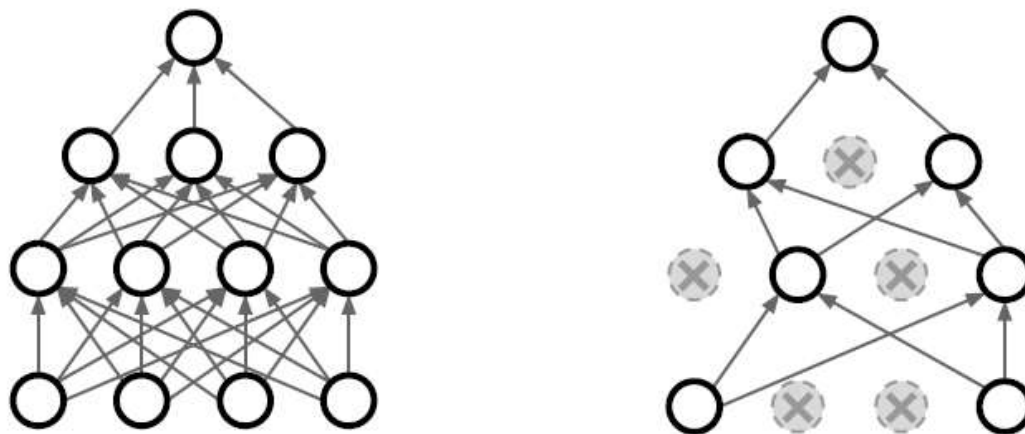
	very similar dataset	very different dataset
very little data	Use Linear Classifier on top layer	You're in trouble... Try linear classifier from different stages
quite a lot of data	Finetune a few layers	Finetune a larger number of layers

Outline

- Regularization in CNN training
 - Data Augmentation
 - Weight Regularization & Transfer Learning
 - Stochastic Regularization
 - Hyper-parameter optimization
- Network Architectures

Stochastic Regularization

- For a network to overfit, its computations need to be really precise. This suggests regularizing them by injecting noise into the computations, a strategy known as **stochastic regularization**.
- Dropout is a stochastic regularizer which randomly deactivates a subset of the units



Srivastava et al, "Dropout: A simple way to prevent neural networks from overfitting", JMLR 2014

Dropout

■ Operations

$$h_i = m_i \cdot \phi(z_i),$$

where m_i is a Bernoulli random variable, independent for each hidden unit.

Regularization: Dropout

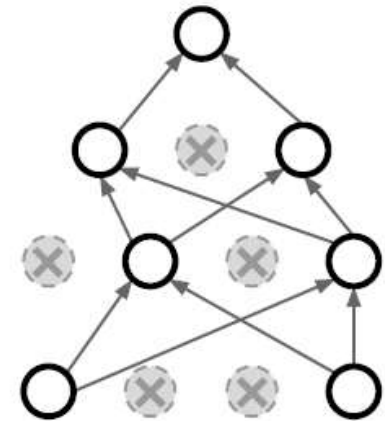
```
p = 0.5 # probability of keeping a unit active. higher = less dropout

def train_step(X):
    """ X contains the data """

    # forward pass for example 3-layer neural network
    H1 = np.maximum(0, np.dot(W1, X) + b1)
    U1 = np.random.rand(*H1.shape) < p # first dropout mask
    H1 *= U1 # drop!
    H2 = np.maximum(0, np.dot(W2, H1) + b2)
    U2 = np.random.rand(*H2.shape) < p # second dropout mask
    H2 *= U2 # drop!
    out = np.dot(W3, H2) + b3

    # backward pass: compute gradients... (not shown)
    # perform parameter update... (not shown)
```

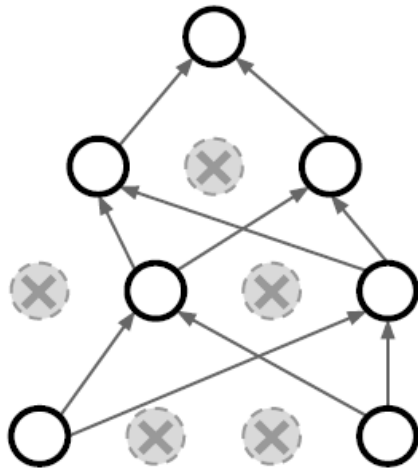
Example forward pass with a 3-layer network using dropout



Understanding Dropout

Regularization: Dropout

How can this possibly be a good idea?

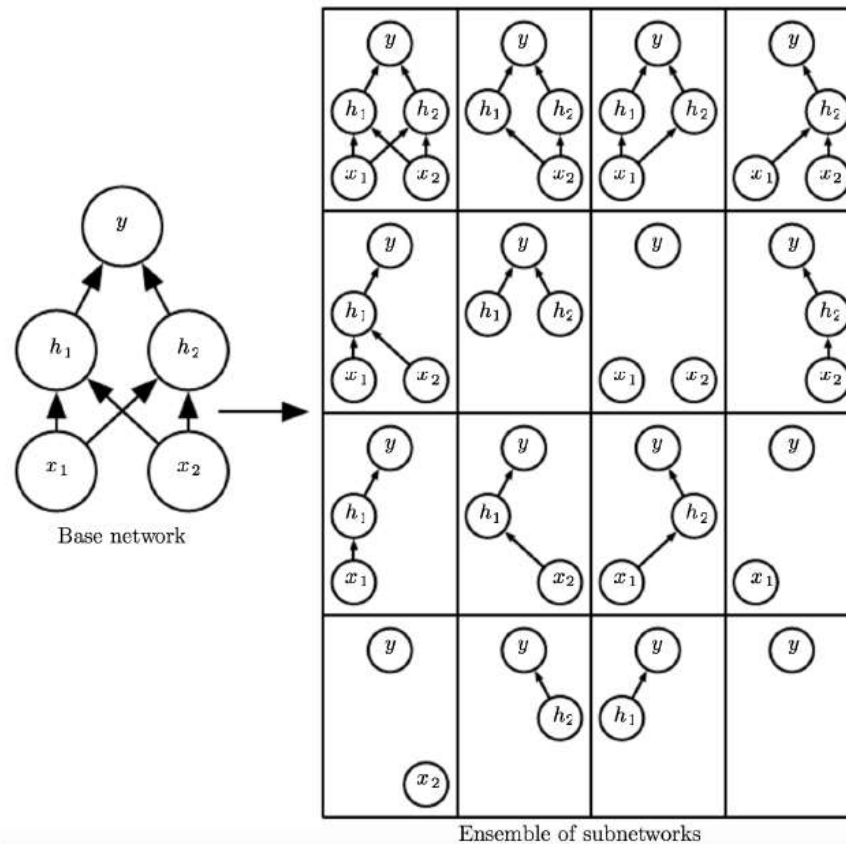


Forces the network to have a redundant representation;
Prevents co-adaptation of features



Understanding Dropout

- Dropout can be seen as training an ensemble of 2^D different architectures with shared weights (where D is the number of units):



— Goodfellow et al., *Deep Learning*

Dropout

- Dropout at test time

Dropout makes our output random!

$$\begin{array}{c} \text{Output} \\ \text{(label)} \end{array} \boxed{y} = f_W \left(\begin{array}{c} \text{Input} \\ \text{(image)} \end{array} \boxed{x}, \begin{array}{c} \text{Random} \\ \text{mask} \end{array} \boxed{z} \right)$$

Want to “average out” the randomness at test-time

$$y = f(x) = E_z [f(x, z)] = \int p(z) f(x, z) dz$$

But this integral seems hard ...

Dropout

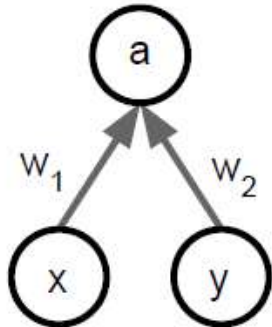
- Dropout at test time

Want to approximate
the integral

$$y = f(x) = E_z[f(x, z)] = \int p(z)f(x, z)dz$$

Consider a single neuron.

At test time we have: $E[a] = w_1x + w_2y$



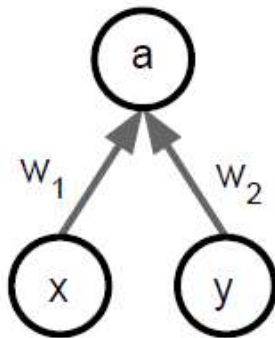
Dropout

- Dropout at test time

Want to approximate the integral

$$y = f(x) = E_z[f(x, z)] = \int p(z)f(x, z)dz$$

Consider a single neuron.



At test time we have: $E[a] = w_1x + w_2y$

During training we have:
$$E[a] = \frac{1}{4}(w_1x + w_2y) + \frac{1}{4}(w_1x + 0y) + \frac{1}{4}(0x + 0y) + \frac{1}{4}(0x + w_2y) = \frac{1}{2}(w_1x + w_2y)$$

At test time, **multiply**
by dropout probability

Dropout

■ Dropout at test time

```
def predict(X):  
    # ensembled forward pass  
    H1 = np.maximum(0, np.dot(W1, X) + b1) * p # NOTE: scale the activations  
    H2 = np.maximum(0, np.dot(W2, H1) + b2) * p # NOTE: scale the activations  
    out = np.dot(W3, H2) + b3
```

At test time all neurons are active always

=> We must scale the activations so that for each neuron:

output at test time = expected output at training time

Dropout

■ Implementation: Inverted dropout

```
p = 0.5 # probability of keeping a unit active. higher = less dropout

def train_step(X):
    # forward pass for example 3-layer neural network
    H1 = np.maximum(0, np.dot(W1, X) + b1)
    U1 = (np.random.rand(*H1.shape) < p) / p # first dropout mask. Notice /p!
    H1 *= U1 # drop!
    H2 = np.maximum(0, np.dot(W2, H1) + b2)
    U2 = (np.random.rand(*H2.shape) < p) / p # second dropout mask. Notice /p!
    H2 *= U2 # drop!
    out = np.dot(W3, H2) + b3

    # backward pass: compute gradients... (not shown)
    # perform parameter update... (not shown)

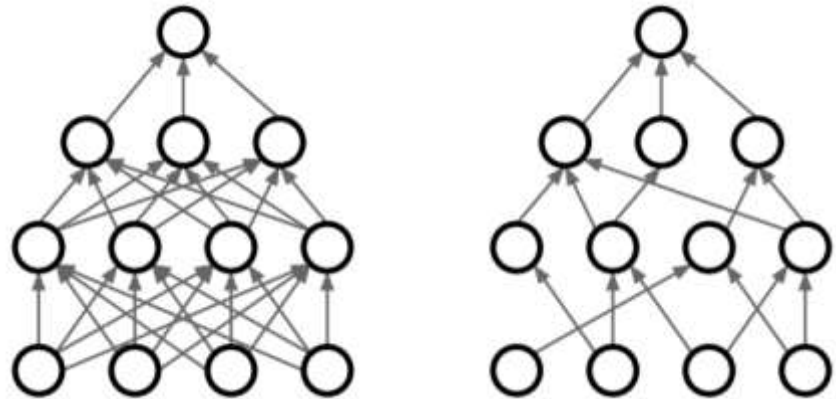
def predict(X):
    # ensembled forward pass
    H1 = np.maximum(0, np.dot(W1, X) + b1) # no scaling necessary
    H2 = np.maximum(0, np.dot(W2, H1) + b2)
    out = np.dot(W3, H2) + b3
```

test time is unchanged!

Stochastic Regularization

- Lots of other stochastic regularizers have been proposed:
 - **DropConnect** drops connections instead of activations.

- Training: Drop connections between neurons (set weights to 0)
- Testing: Use all the connections



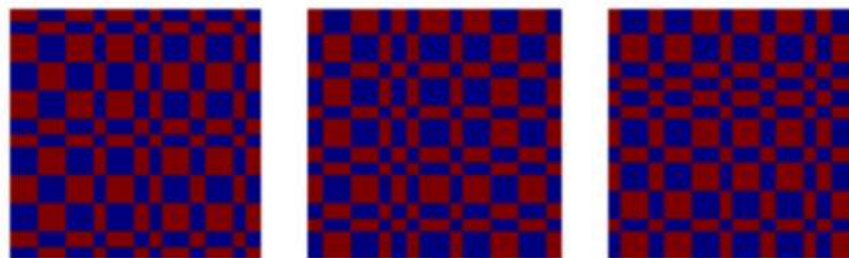
Wan et al, “Regularization of Neural Networks using DropConnect”, ICML 2013

Stochastic Regularization

- Lots of other stochastic regularizers have been proposed:

- Fractional Pooling

- Training: Use randomized pooling regions
 - Testing: Average predictions from several regions



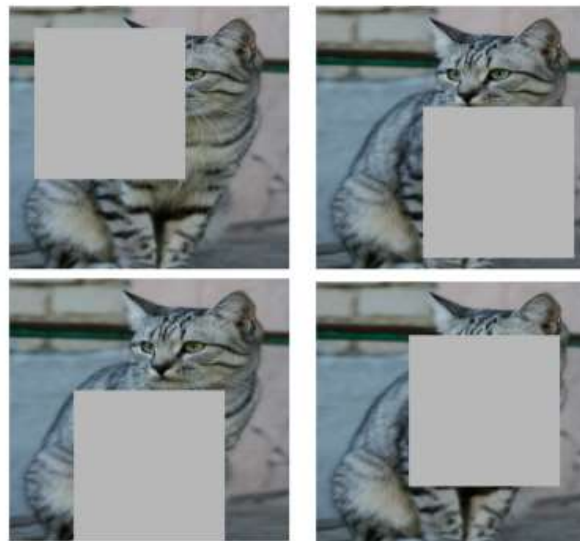
Graham, “Fractional Max Pooling”, arXiv 2014

Stochastic Regularization

- Lots of other stochastic regularizers have been proposed:

- Cutout

- Training: Set random image regions to zero
 - Testing: Use full image predictions from several regions



Works very well for small datasets like CIFAR,
less common for large datasets like ImageNet

DeVries and Taylor, “Improved Regularization of Convolutional Neural Networks with Cutout”,
arXiv 2017

Stochastic Regularization

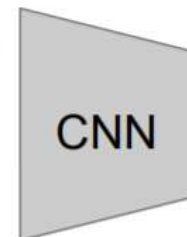
- Lots of other stochastic regularizers have been proposed:

- Mixup

- Training: Train on random blends of images
 - Testing: Use original images



Randomly blend the pixels of pairs of training images, e.g. 40% cat, 60% dog



Target label:
cat: 0.4
dog: 0.6

Zhang et al, “mixup: Beyond Empirical Risk Minimization”, ICLR 2018

Stochastic Regularization

- Lots of other stochastic regularizers have been proposed:
 - Training: Add random noise
 - Testing: Marginalize over the noise
- In practice
 - Consider **dropout** for large fully-connected layers
 - **Batch normalization** and **data augmentation** almost always a good idea
 - Try **cutout** and **mixup** especially for small classification datasets

Outline

- Regularization in CNN training
 - Data Augmentation
 - Weight Regularization & Transfer Learning
 - Stochastic Regularization
 - Hyper-parameter optimization

Hyperparameter optimization

- (Cross-)validation strategy

coarse -> **fine** cross-validation in stages

First stage: only a few epochs to get rough idea of what params work

Second stage: longer running time, finer search

... (repeat as necessary)

Tip for detecting explosions in the solver:

If the cost is ever $> 3 \times$ original cost, break out early

Hyperparameter optimization

For example: run coarse search for 5 epochs

```
max_count = 100
for count in xrange(max_count):
    reg = 10**uniform(-5, 5)
    lr = 10**uniform(-3, -6)

    trainer = ClassifierTrainer()
    model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
    trainer = ClassifierTrainer()
    best_model_local, stats = trainer.train(X_train, y_train, X_val, y_val,
                                           model, two_layer_net,
                                           num_epochs=5, reg=reg,
                                           update='momentum', learning_rate_decay=0.9,
                                           sample_batches = True, batch_size = 100,
                                           learning_rate=lr, verbose=False)
```

note it's best to optimize
in log space!

```
val_acc: 0.412000, lr: 1.405206e-04, reg: 4.793564e-01, (1 / 100)
val_acc: 0.214000, lr: 7.231888e-06, reg: 2.321281e-04, (2 / 100)
val_acc: 0.208000, lr: 2.119571e-06, reg: 8.011857e+01, (3 / 100)
val_acc: 0.196000, lr: 1.551131e-05, reg: 4.374936e-05, (4 / 100)
val_acc: 0.079000, lr: 1.753300e-05, reg: 1.200424e+03, (5 / 100)
val_acc: 0.223000, lr: 4.215128e-05, reg: 4.196174e+01, (6 / 100)
val_acc: 0.441000, lr: 1.750259e-04, reg: 2.110807e-04, (7 / 100)
val_acc: 0.241000, lr: 6.749231e-05, reg: 4.226413e+01, (8 / 100)
val_acc: 0.482000, lr: 4.296863e-04, reg: 6.642555e-01, (9 / 100)
val_acc: 0.079000, lr: 5.401602e-06, reg: 1.599828e+04, (10 / 100)
val_acc: 0.154000, lr: 1.618508e-06, reg: 4.925252e-01, (11 / 100)
```

nice

Hyperparameter optimization

Now run finer search...

```
max_count = 100
for count in xrange(max_count):
    reg = 10**uniform(-5, 5)
    lr = 10**uniform(-3, -6)
```

adjust range

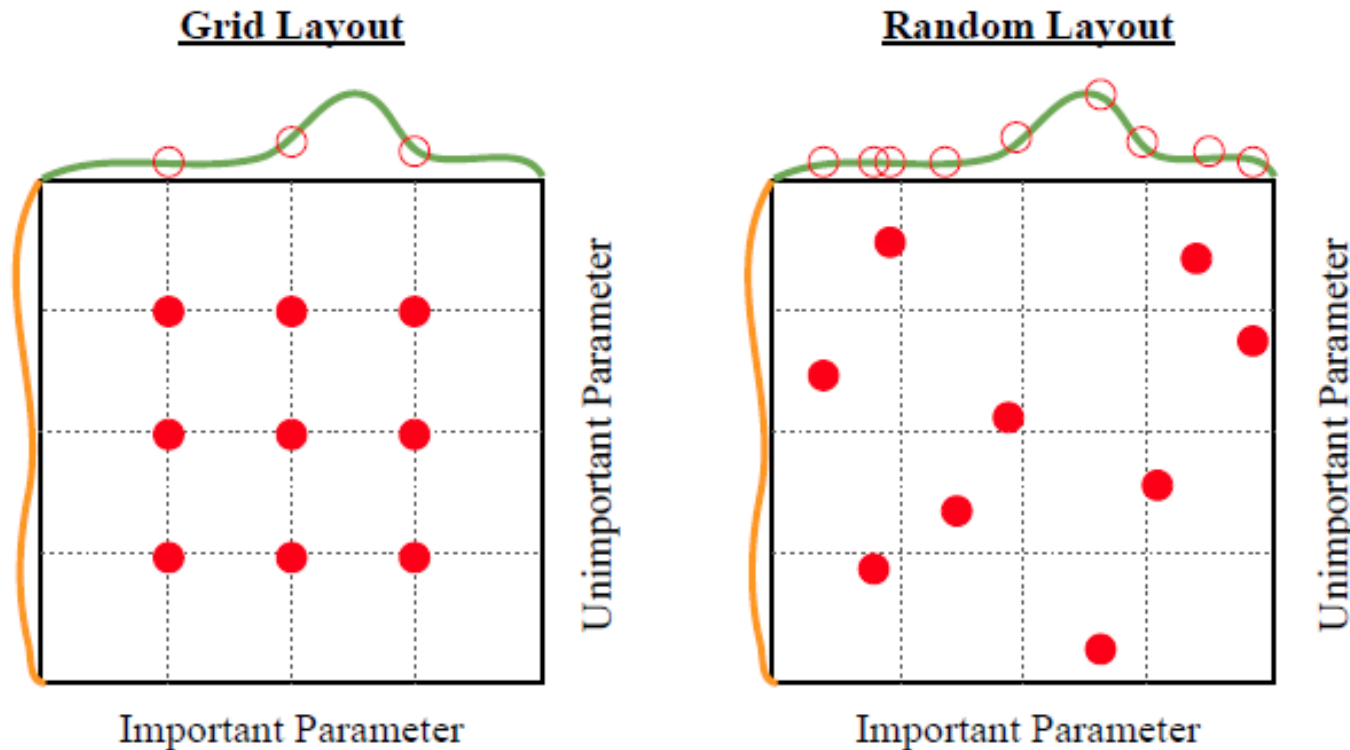
```
max_count = 100
for count in xrange(max_count):
    reg = 10**uniform(-4, 0)
    lr = 10**uniform(-3, -4)
```

val_acc: 0.527000, lr: 5.340517e-04, reg: 4.097824e-01, (0 / 100)
val_acc: 0.492000, lr: 2.279484e-04, reg: 9.991345e-04, (1 / 100)
val_acc: 0.512000, lr: 8.680827e-04, reg: 1.349727e-02, (2 / 100)
val_acc: 0.461000, lr: 1.028377e-04, reg: 1.220193e-02, (3 / 100)
val_acc: 0.460000, lr: 1.113730e-04, reg: 5.244309e-02, (4 / 100)
val_acc: 0.498000, lr: 9.477776e-04, reg: 2.001293e-03, (5 / 100)
val_acc: 0.469000, lr: 1.484369e-04, reg: 4.328313e-01, (6 / 100)
val_acc: 0.522000, lr: 5.586261e-04, reg: 2.312685e-04, (7 / 100)
val_acc: 0.530000, lr: 5.808183e-04, reg: 8.259964e-02, (8 / 100)
val_acc: 0.489000, lr: 1.979168e-04, reg: 1.010889e-04, (9 / 100)
val_acc: 0.490000, lr: 2.036031e-04, reg: 2.406271e-03, (10 / 100)
val_acc: 0.475000, lr: 2.021162e-04, reg: 2.287807e-01, (11 / 100)
val_acc: 0.460000, lr: 1.135527e-04, reg: 3.905040e-02, (12 / 100)
val_acc: 0.515000, lr: 6.947668e-04, reg: 1.562808e-02, (13 / 100)
val_acc: 0.531000, lr: 9.471549e-04, reg: 1.433895e-03, (14 / 100)
val_acc: 0.509000, lr: 3.140888e-04, reg: 2.857518e-01, (15 / 100)
val_acc: 0.514000, lr: 6.438349e-04, reg: 3.033781e-01, (16 / 100)
val_acc: 0.502000, lr: 3.921784e-04, reg: 2.707126e-04, (17 / 100)
val_acc: 0.509000, lr: 9.752279e-04, reg: 2.850865e-03, (18 / 100)
val_acc: 0.500000, lr: 2.412048e-04, reg: 4.997821e-04, (19 / 100)
val_acc: 0.466000, lr: 1.319314e-04, reg: 1.189915e-02, (20 / 100)
val_acc: 0.516000, lr: 8.039527e-04, reg: 1.528291e-02, (21 / 100)

53% - relatively good
for a 2-layer neural net
with 50 hidden neurons.

Hyperparameter optimization

- Random search vs. Grid search



Random Search for Hyper-Parameter Optimization, Bergstra and Bengio, 2012

Hyperparameter optimization

- **Hyperparameters to play with:**
 - network architecture
 - learning rate, its decay schedule, update type
 - regularization (L2/Dropout strength)
- **Other hyperparameter optimization methods**
 - Shahriari, et al. "Taking the human out of the loop: A review of Bayesian optimization." Proceedings of the IEEE 104.1 (2016): 148-175.

Outline

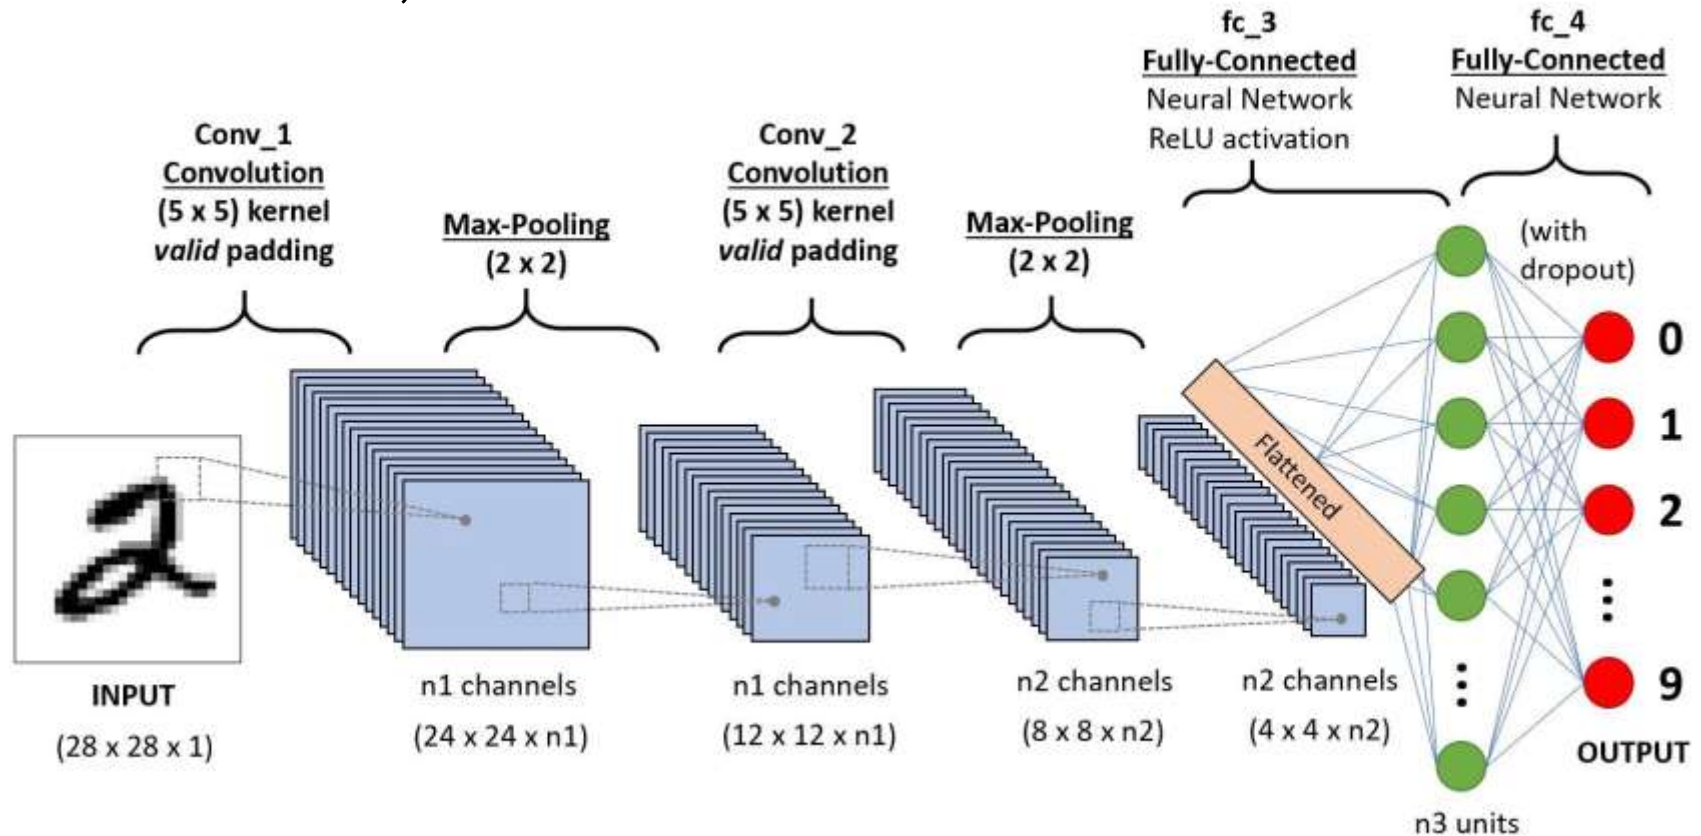
- CNN architectures

- ☐ Sequential structure: LeNet/AlexNet/VGGNet
- ☐ Parallel branches: GoogLeNet
- ☐ Residual structure: ResNet/DenseNet
- ☐ Network Architecture Search

Acknowledgement: Zemel et al's CSC411 and Feifei Li et al's cs231n notes

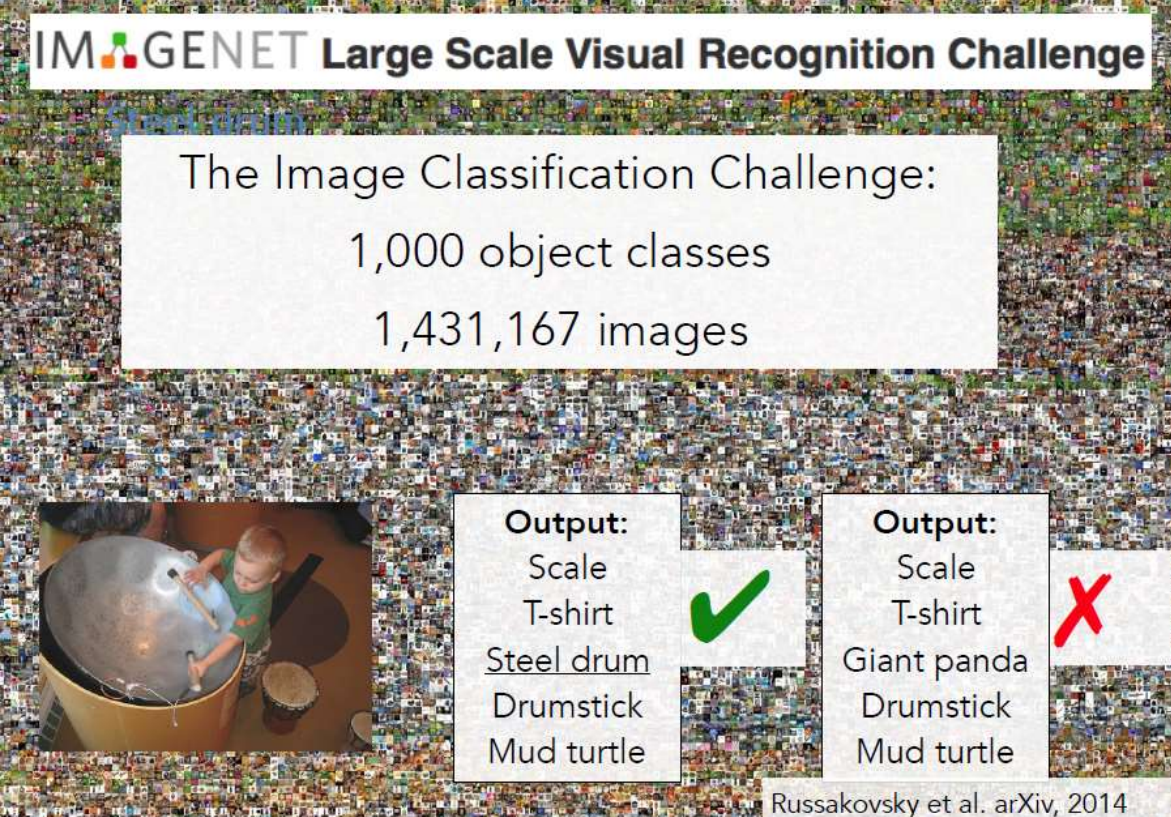
LeNet-5

- Handwritten digit recognition
- LeCun et al., 1998




Background: Image/Object Classification


- Problem Setup
 - Input: Image
 - Output: Object class



IMAGENET Large Scale Visual Recognition Challenge

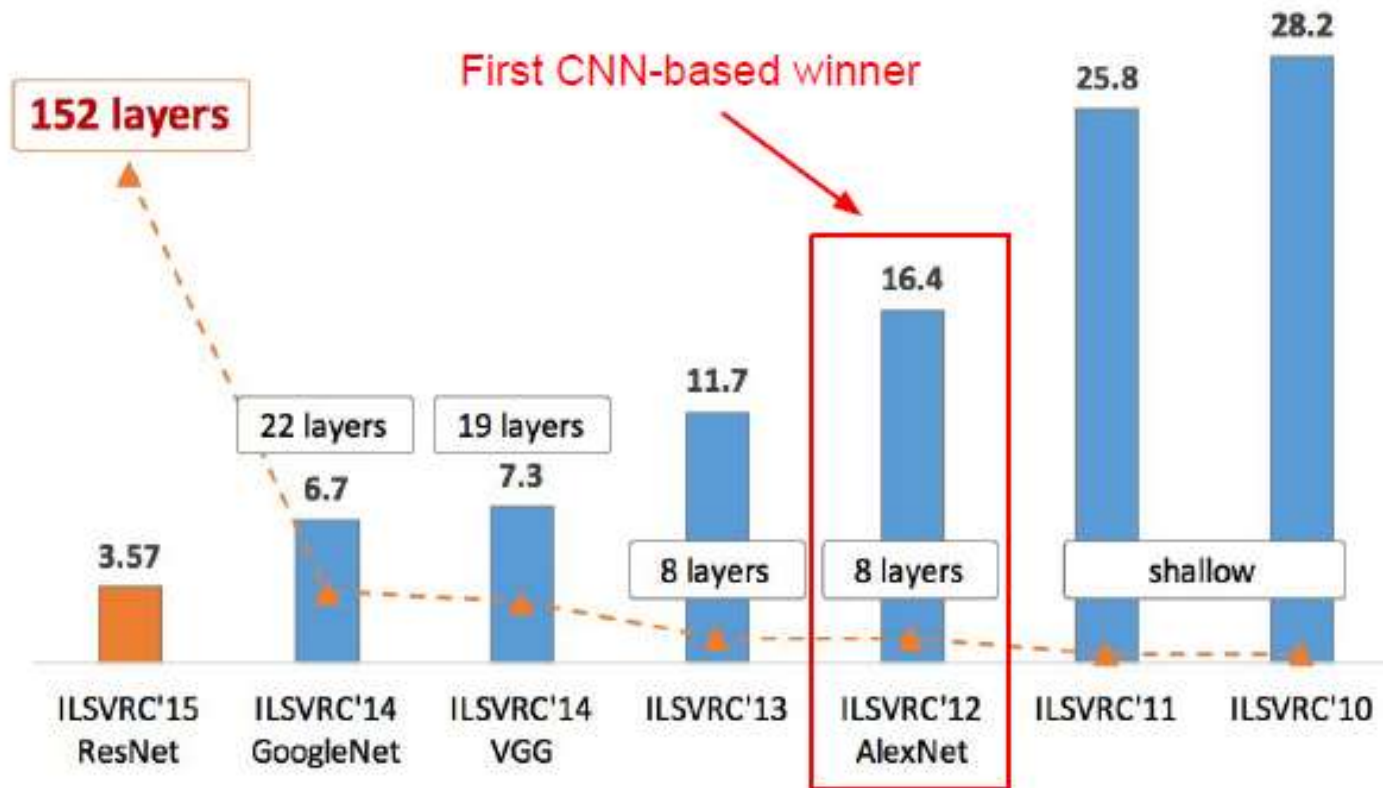
The Image Classification Challenge:
1,000 object classes
1,431,167 images



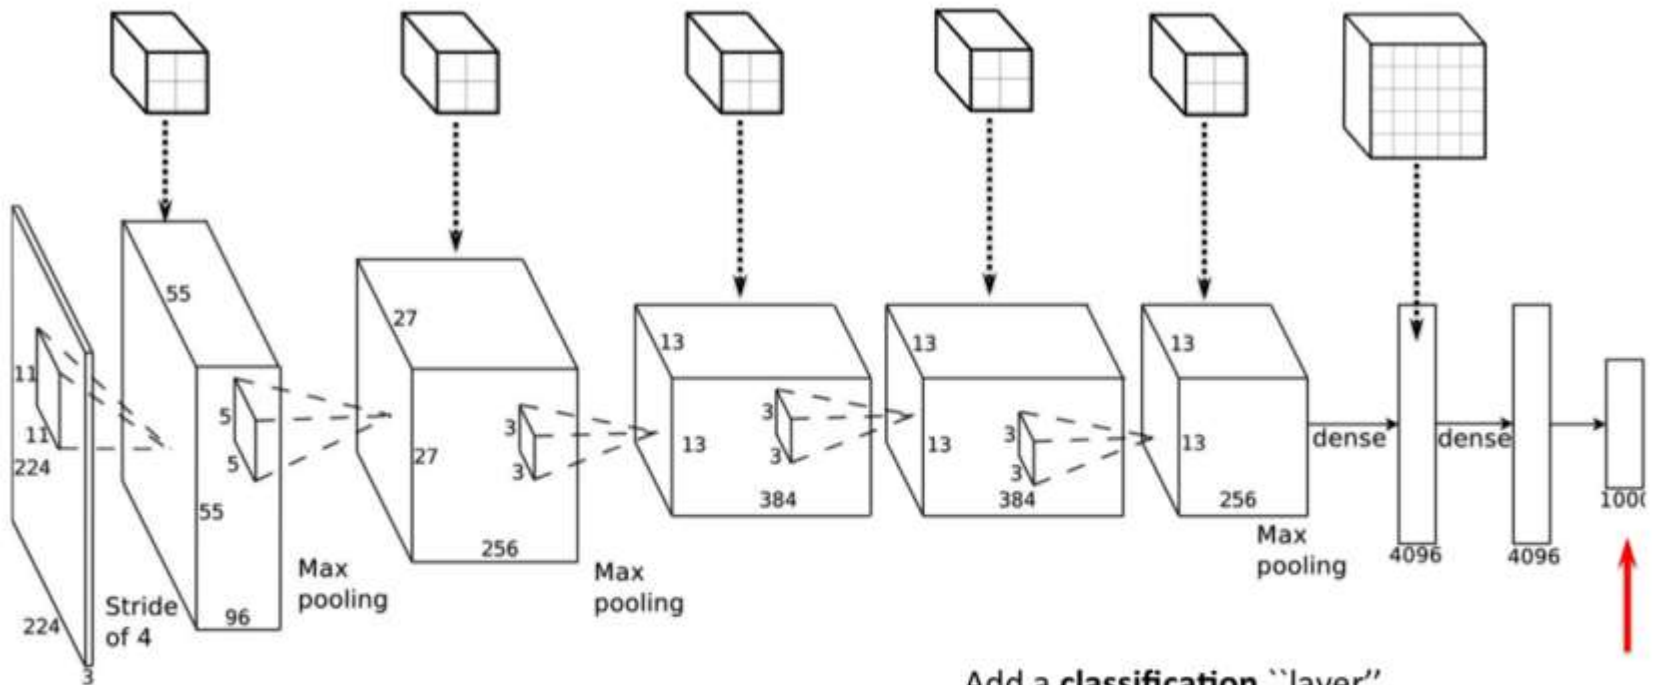
Output:		Output:
Scale		Scale
T-shirt		T-shirt
<u>Steel drum</u>		Giant panda
Drumstick		Drumstick
Mud turtle		Mud turtle

Russakovsky et al. arXiv, 2014

ImageNet (ILSVRC)



AlexNet



- heavy data augmentation
- dropout 0.5
- batch size 128
- SGD Momentum 0.9
- Learning rate $1e-2$, reduced by 10 manually when val accuracy plateaus
- L2 weight decay $5e-4$
- 7 CNN ensemble: 18.2% \rightarrow 15.4%

Add a **classification** "layer".

For an input image, the value in a particular dimension of this vector tells you the probability of the corresponding object class.

AlexNet

■ Deeper network structure

- More convolution layers
- Local contrast normalization
- ReLu instead of Tanh
- Dropout as regularization

[227x227x3] INPUT

[55x55x96] **CONV1**: 96 11x11 filters at stride 4, pad 0

[27x27x96] **MAX POOL1**: 3x3 filters at stride 2

[27x27x96] **NORM1**: Normalization layer

[27x27x256] **CONV2**: 256 5x5 filters at stride 1, pad 2

[13x13x256] **MAX POOL2**: 3x3 filters at stride 2

[13x13x256] **NORM2**: Normalization layer

[13x13x384] **CONV3**: 384 3x3 filters at stride 1, pad 1

[13x13x384] **CONV4**: 384 3x3 filters at stride 1, pad 1

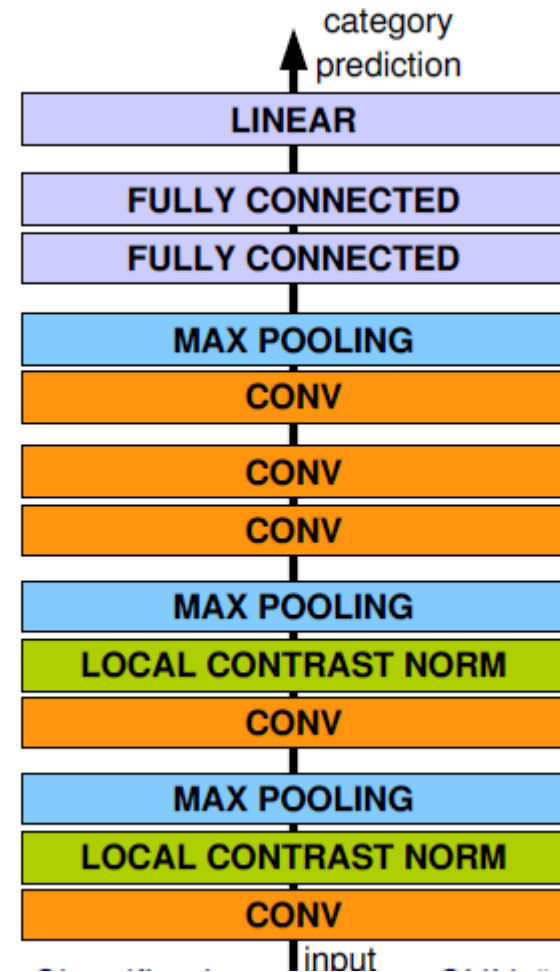
[13x13x256] **CONV5**: 256 3x3 filters at stride 1, pad 1

[6x6x256] **MAX POOL3**: 3x3 filters at stride 2

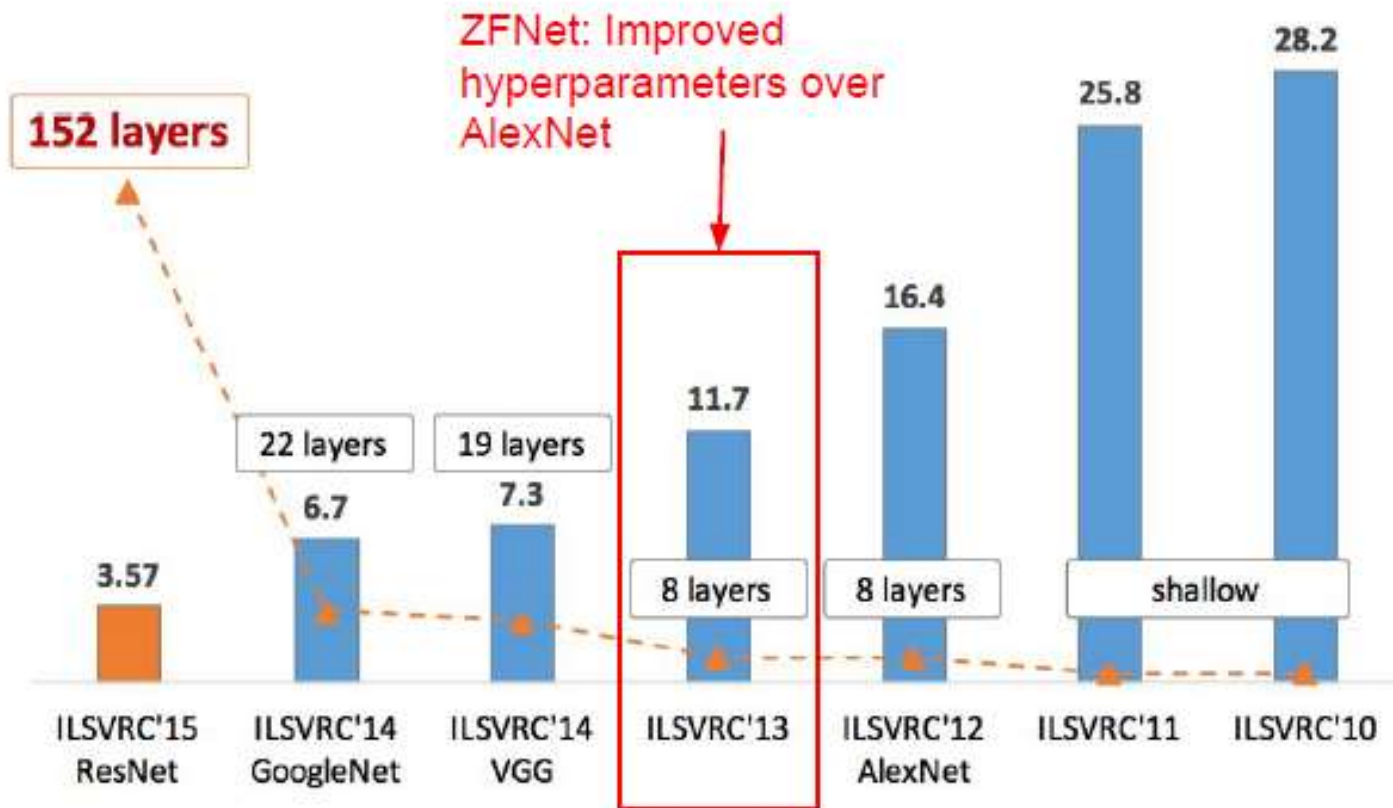
[4096] **FC6**: 4096 neurons

[4096] **FC7**: 4096 neurons

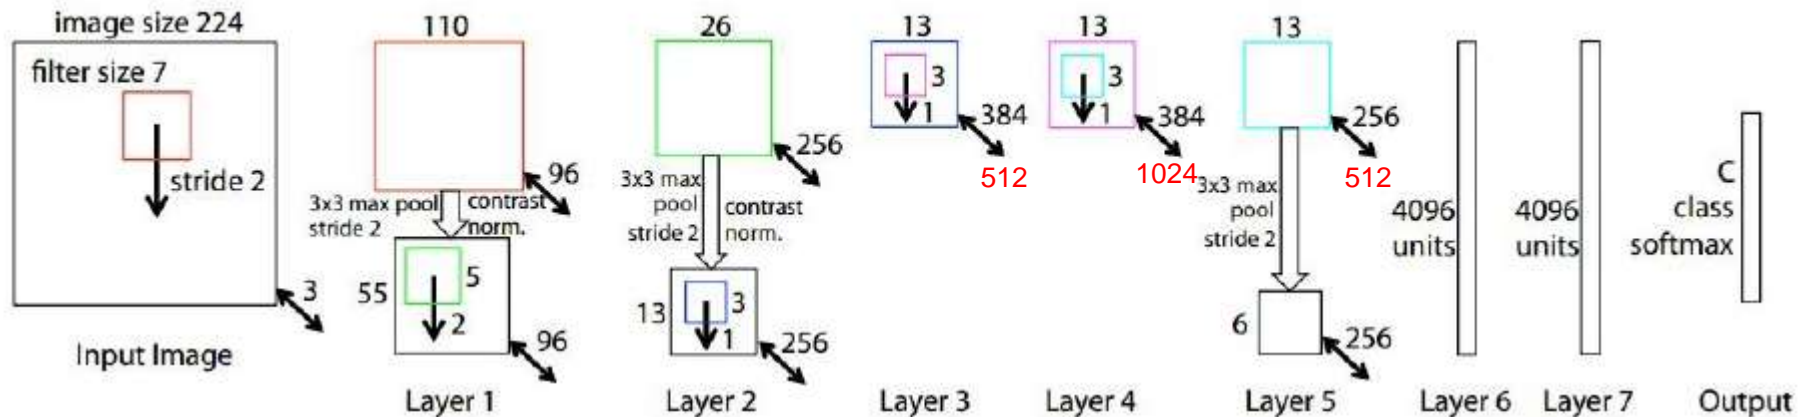
[1000] **FC8**: 1000 neurons (class scores)



ImageNet (ILSVRC)



ZFNet



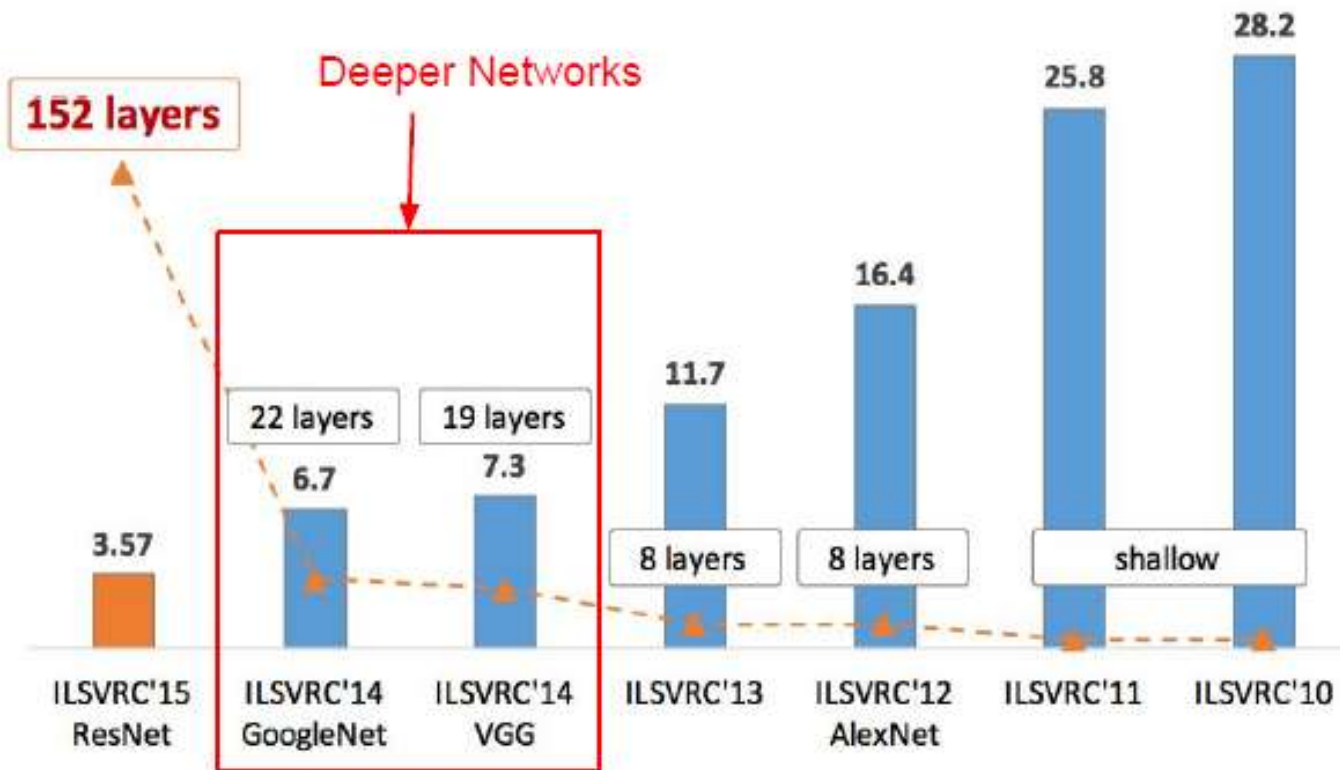
AlexNet but:

CONV1: change from (11x11 stride 4) to (7x7 stride 2)

CONV3,4,5: instead of 384, 384, 256 filters use 512, 1024, 512

ImageNet top 5 error: 16.4% -> 11.7%

ImageNet (ILSVRC)



VGGNet

Case Study: VGGNet

[Simonyan and Zisserman, 2014]

Small filters, Deeper networks

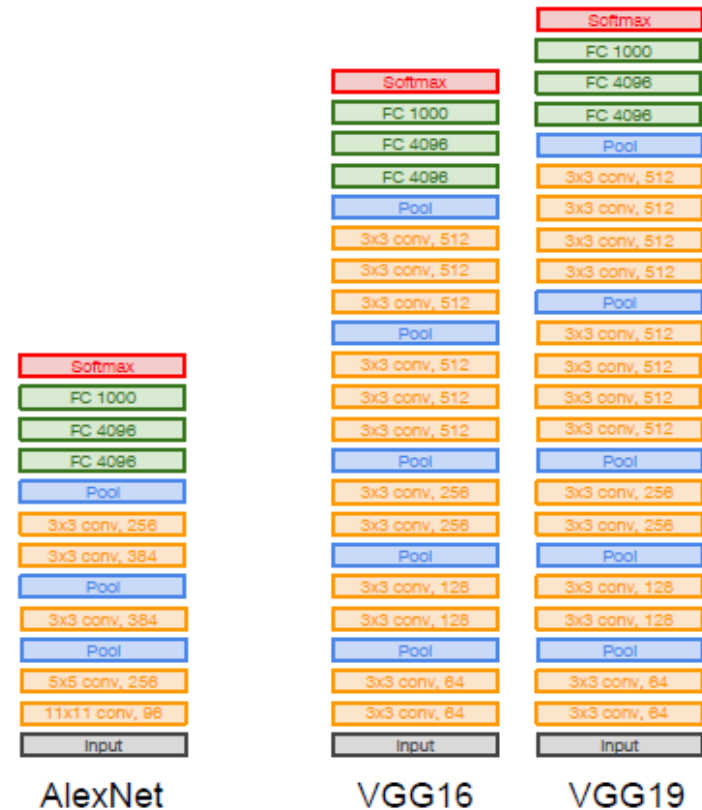
8 layers (AlexNet)

-> 16 - 19 layers (VGG16Net)

Only 3x3 CONV stride 1, pad 1
and 2x2 MAX POOL stride 2

11.7% top 5 error in ILSVRC'13
(ZFNet)

-> 7.3% top 5 error in ILSVRC'14



VGGNet

Case Study: VGGNet

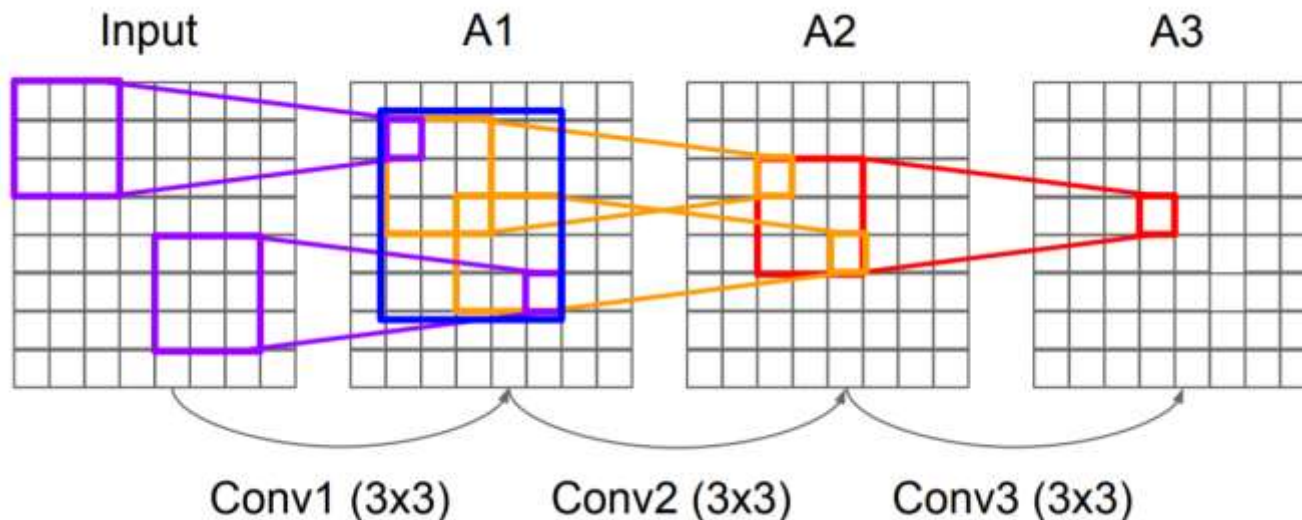
[Simonyan and Zisserman, 2014]

Stack of three 3x3 conv (stride 1) layers has same **effective receptive field** as one 7x7 conv layer

Q: Why use smaller filters? (3x3 conv)

But deeper, more non-linearities

And fewer parameters: $3 * (3^2 C^2)$ vs. $7^2 C^2$ for C channels per layer



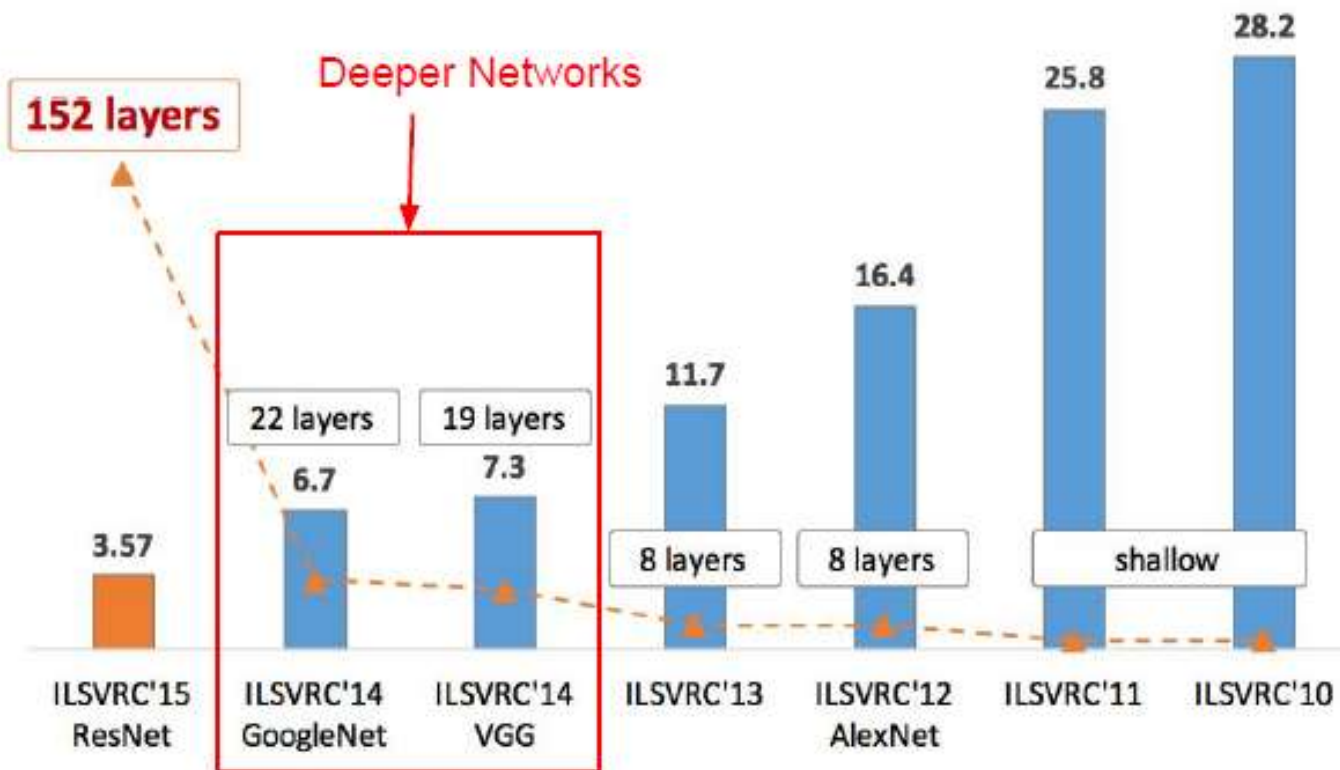
Outline

- CNN architectures

- Sequential structure: LeNet/AlexNet/VGGNet
- Parallel branches: GoogLeNet
- Residual structure: ResNet/DenseNet
- Network Architecture Search

Acknowledgement: Zemel et al's CSC411 and Feifei Li et al's cs231n notes

ImageNet (ILSVRC)



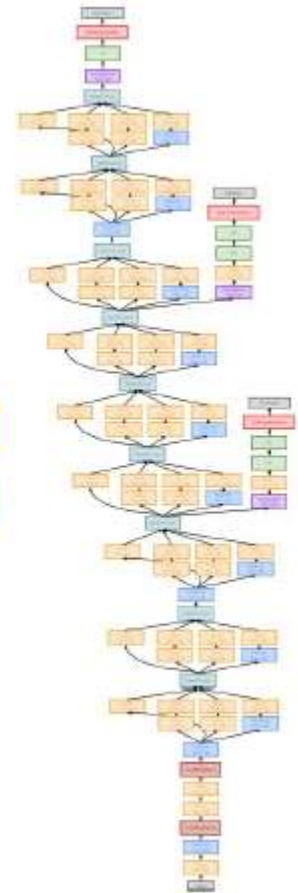
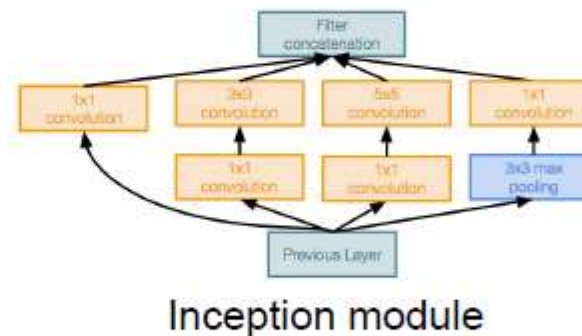
GoogLeNet

Case Study: GoogLeNet

[Szegedy et al., 2014]

Deeper networks, with computational efficiency

- 22 layers
- Efficient “Inception” module
- No FC layers
- Only 5 million parameters!
12x less than AlexNet
- ILSVRC’14 classification winner
(6.7% top 5 error)

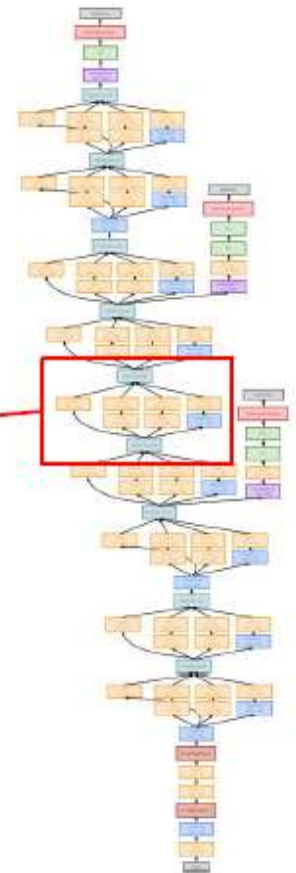
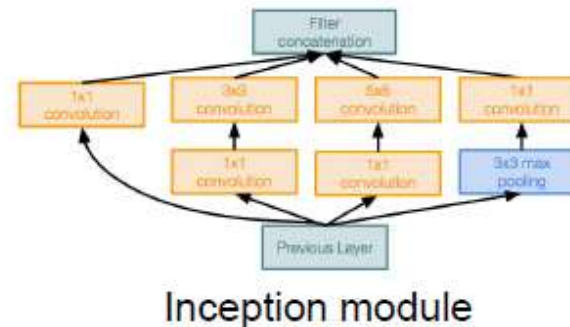


GoogLeNet

Case Study: GoogLeNet

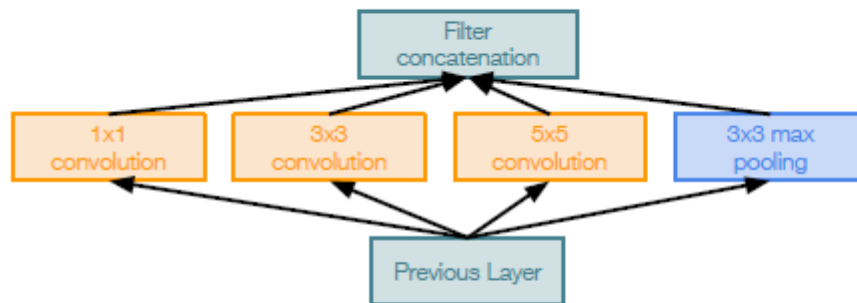
[Szegedy et al., 2014]

“Inception module”: design a good local network topology (network within a network) and then stack these modules on top of each other



GoogLeNet

■ Inception Module



Naive Inception module

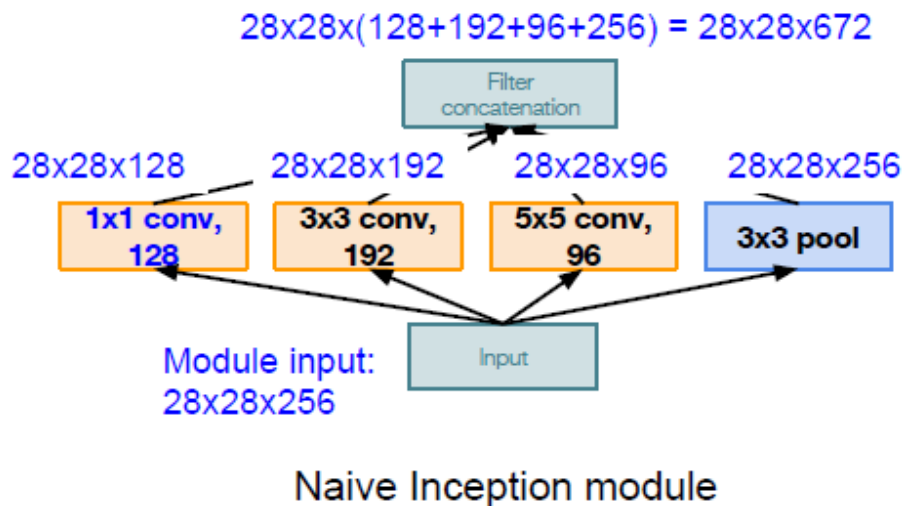
Apply parallel filter operations on the input from previous layer:

- Multiple receptive field sizes for convolution (1x1, 3x3, 5x5)
- Pooling operation (3x3)

Concatenate all filter outputs together depth-wise

GoogLeNet

■ Inception Module



Conv Ops:

[1x1 conv, 128] $28 \times 28 \times 128 \times 1 \times 1 \times 256$

[3x3 conv, 192] $28 \times 28 \times 192 \times 3 \times 3 \times 256$

[5x5 conv, 96] $28 \times 28 \times 96 \times 5 \times 5 \times 256$

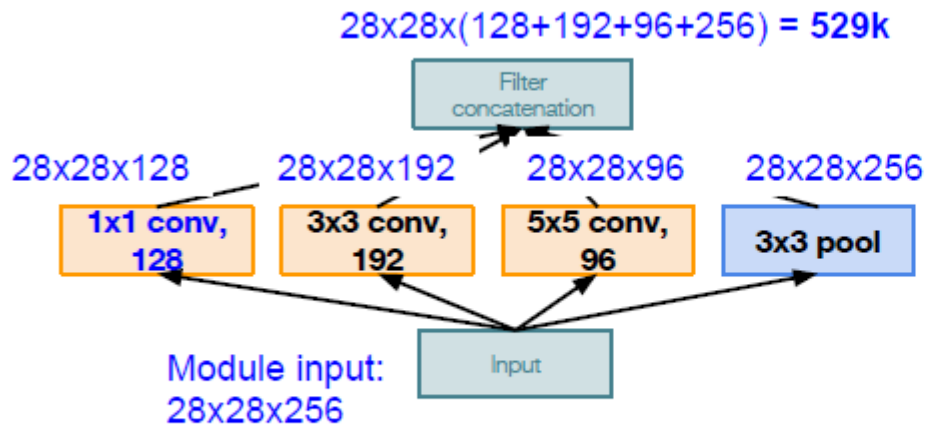
Total: 854M ops

Very expensive compute

Pooling layer also preserves feature depth, which means total depth after concatenation can only grow at every layer!

GoogLeNet

■ Inception Module

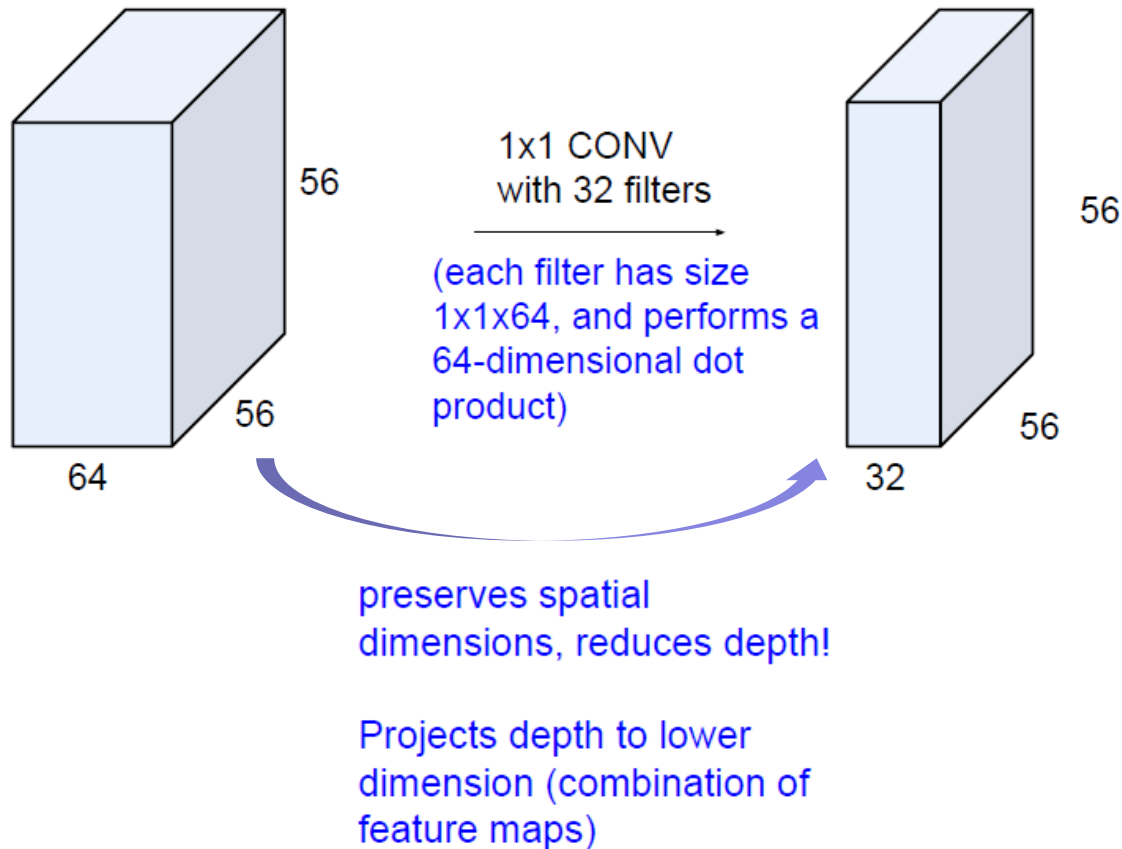


Naive Inception module

Solution: “bottleneck” layers that use 1x1 convolutions to reduce feature depth

GoogLeNet

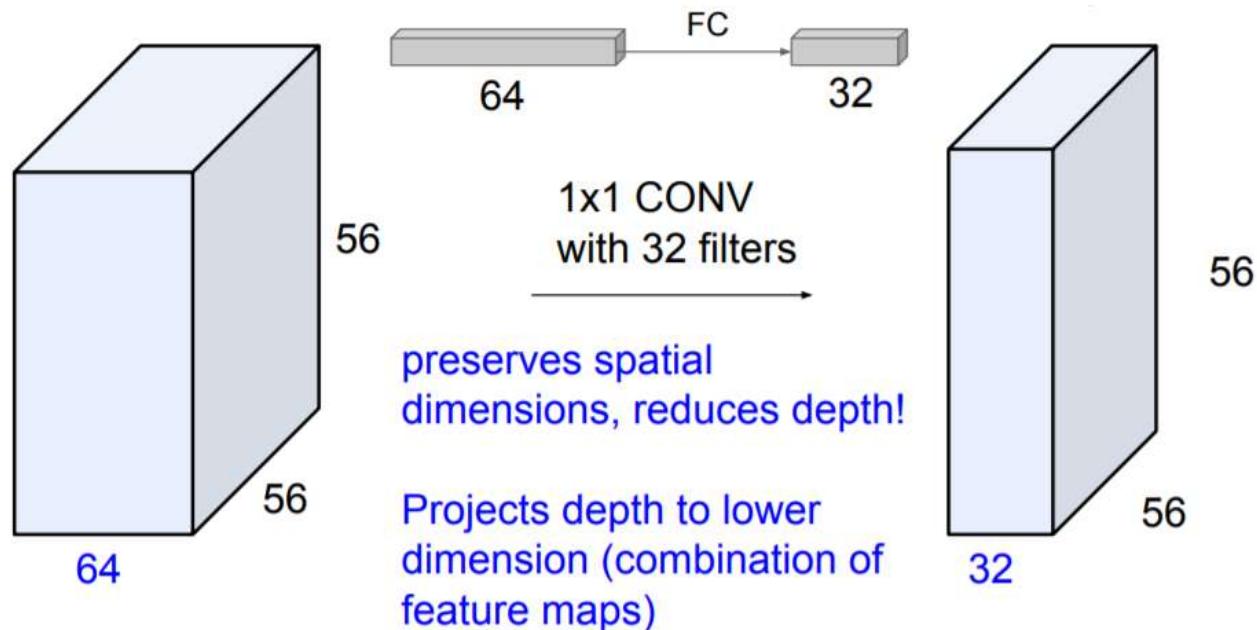
■ Bottleneck layer



GoogLeNet

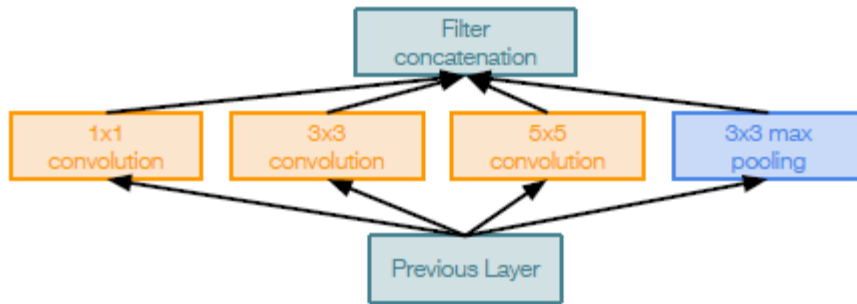
■ 1x1 Convolutions

- Alternatively, interpret it as applying the same FC layer on each input pixel

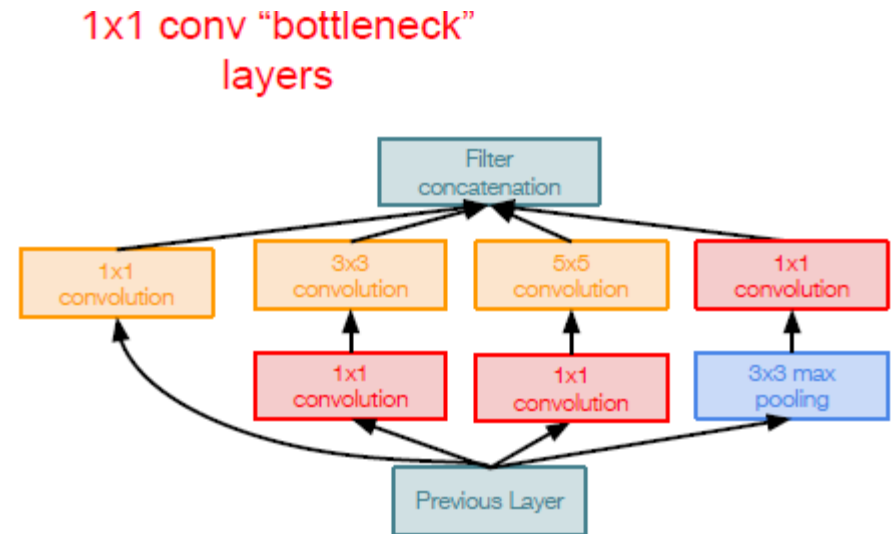


GoogLeNet

■ Inception Module



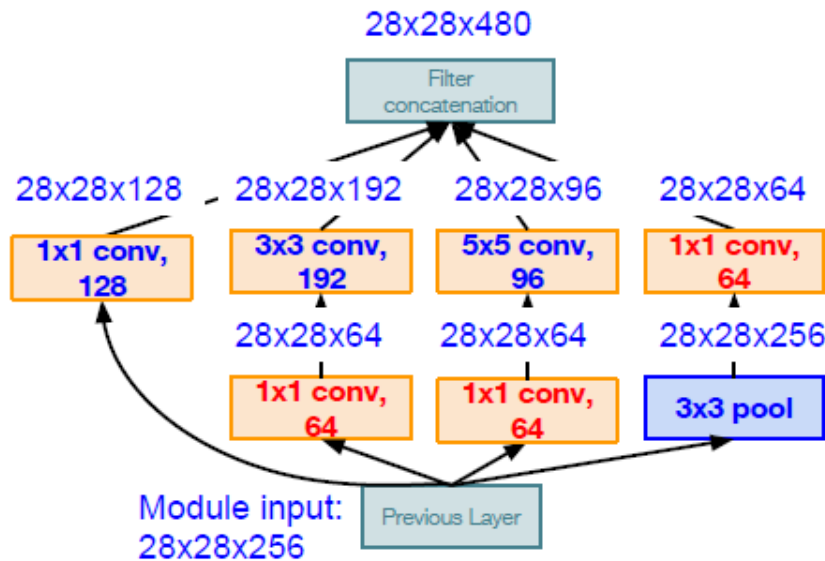
Naive Inception module



Inception module with dimension reduction

GoogLeNet

■ Inception Module



Inception module with dimension reduction

Conv Ops:

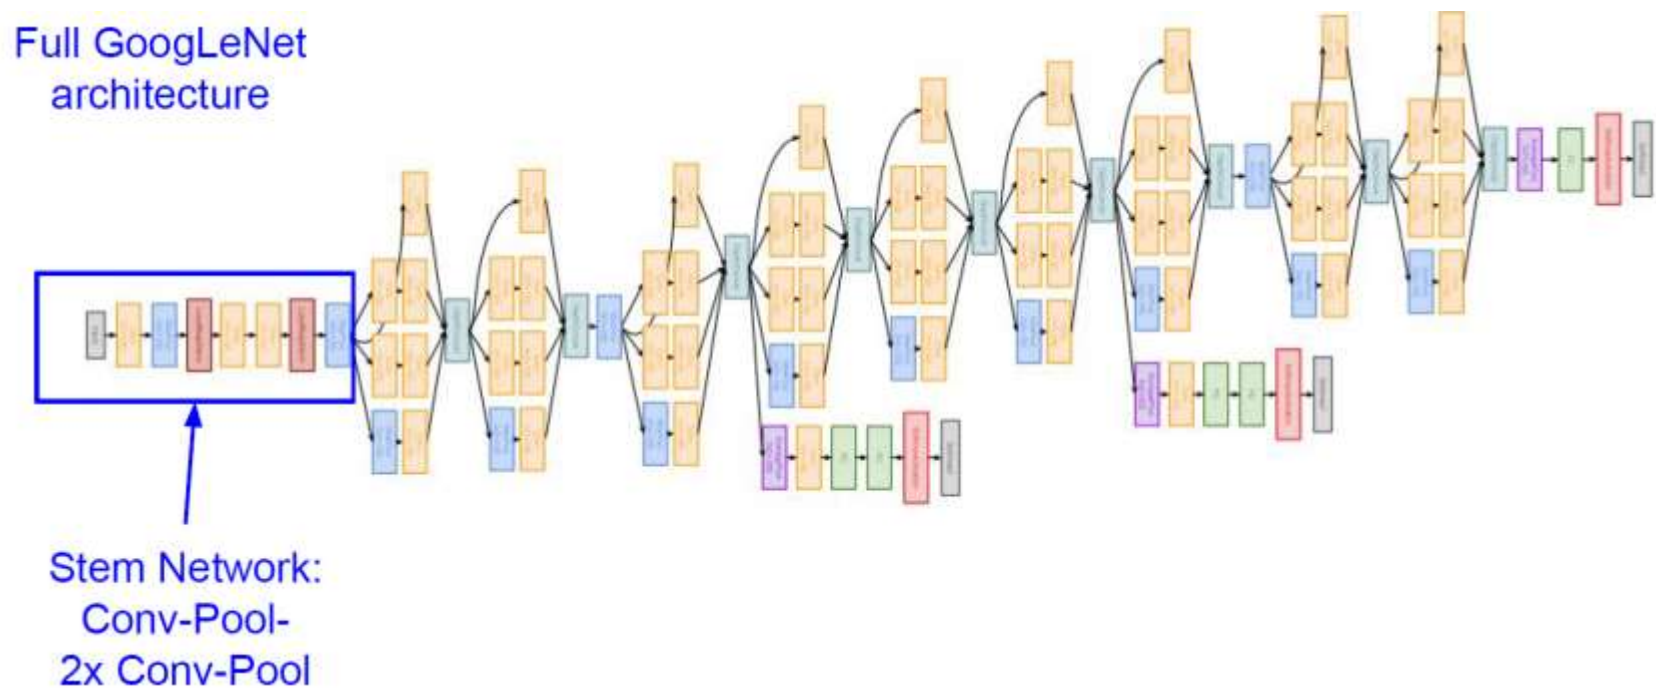
[1x1 conv, 64] 28x28x64x1x1x256
[1x1 conv, 64] 28x28x64x1x1x256
[1x1 conv, 128] 28x28x128x1x1x256
[3x3 conv, 192] 28x28x192x3x3x64
[5x5 conv, 96] 28x28x96x5x5x64
[1x1 conv, 64] 28x28x64x1x1x256

Total: 358M ops

Compared to 854M ops for naive version
Bottleneck can also reduce depth after
pooling layer

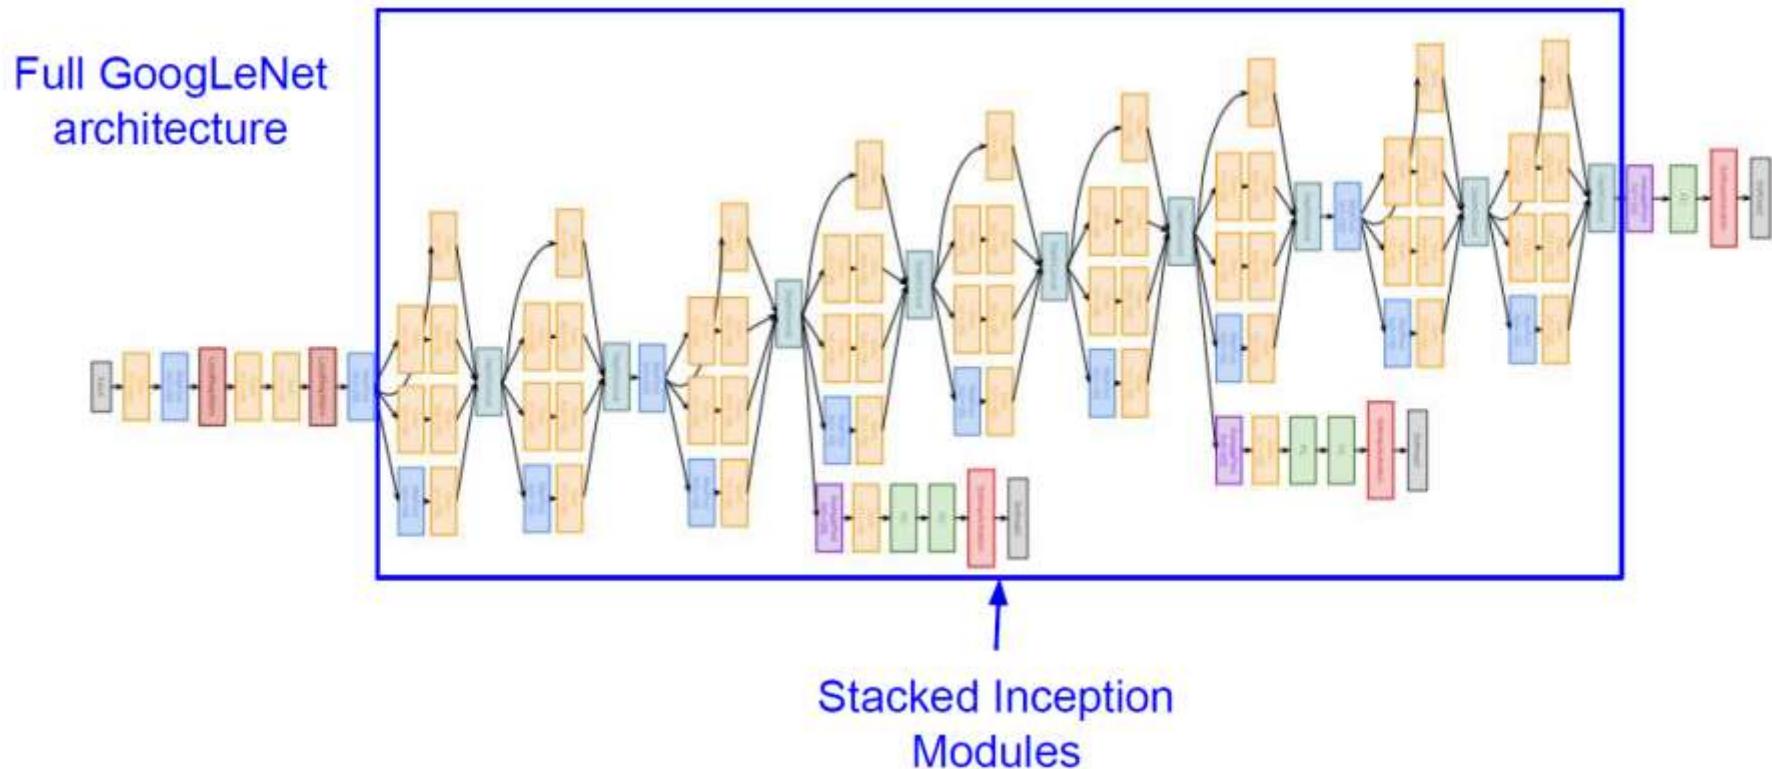
GoogLeNet

■ Overall network structure



GoogLeNet

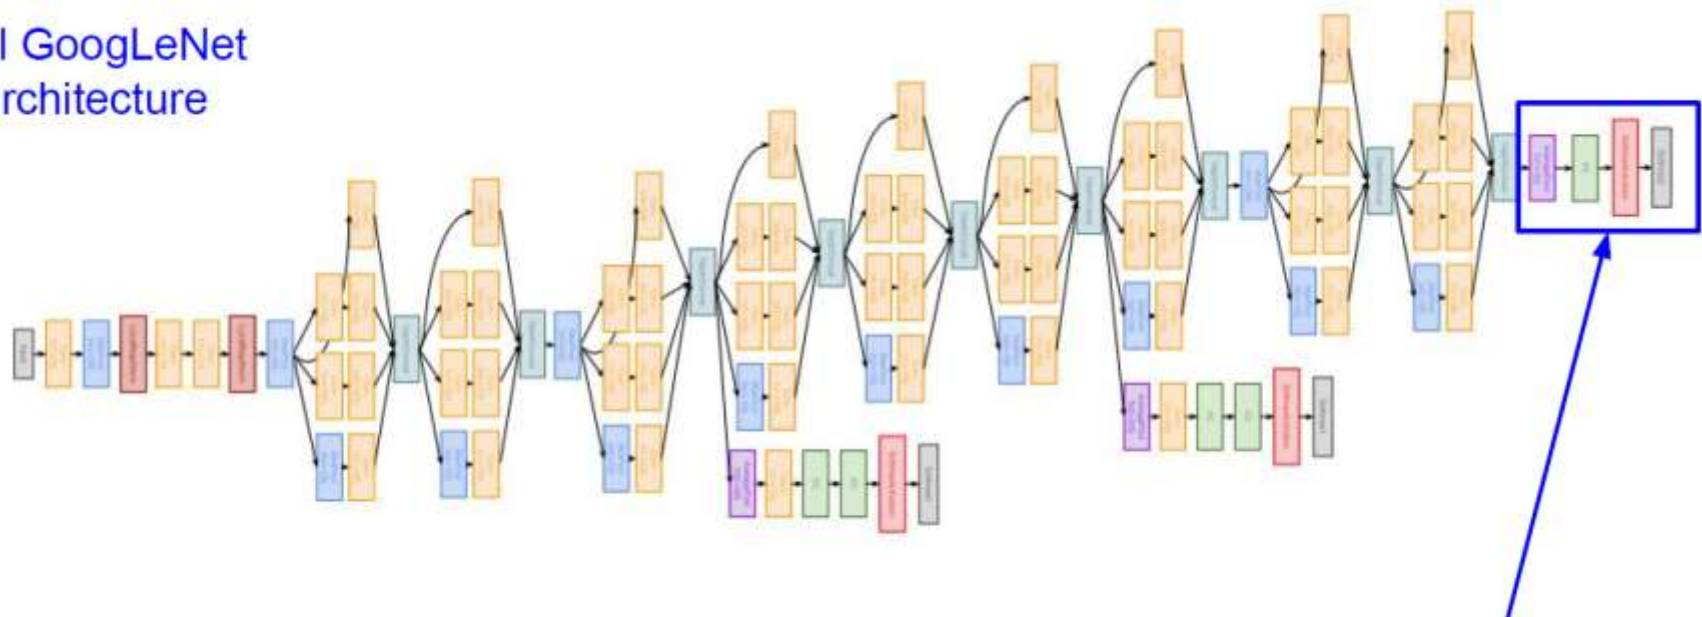
- Overall network structure



GoogLeNet

- Overall network structure

Full GoogLeNet architecture

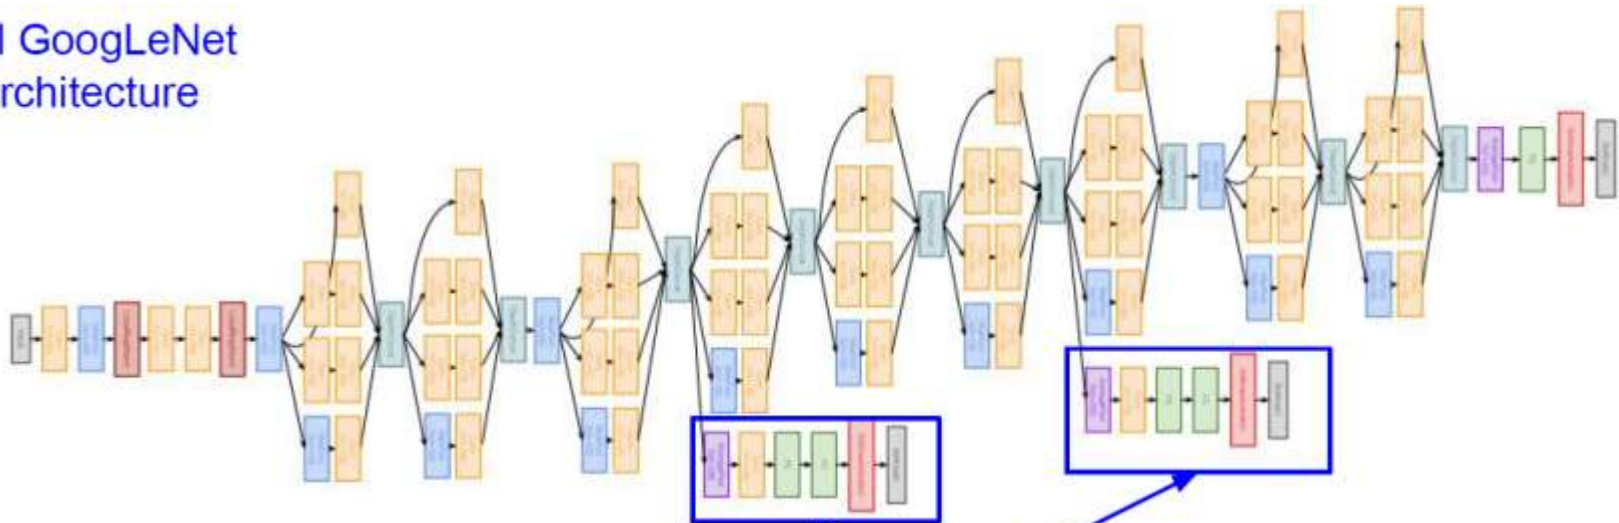


Classifier output
(removed expensive FC layers!)

GoogLeNet

■ Overall network structure

Full GoogLeNet architecture

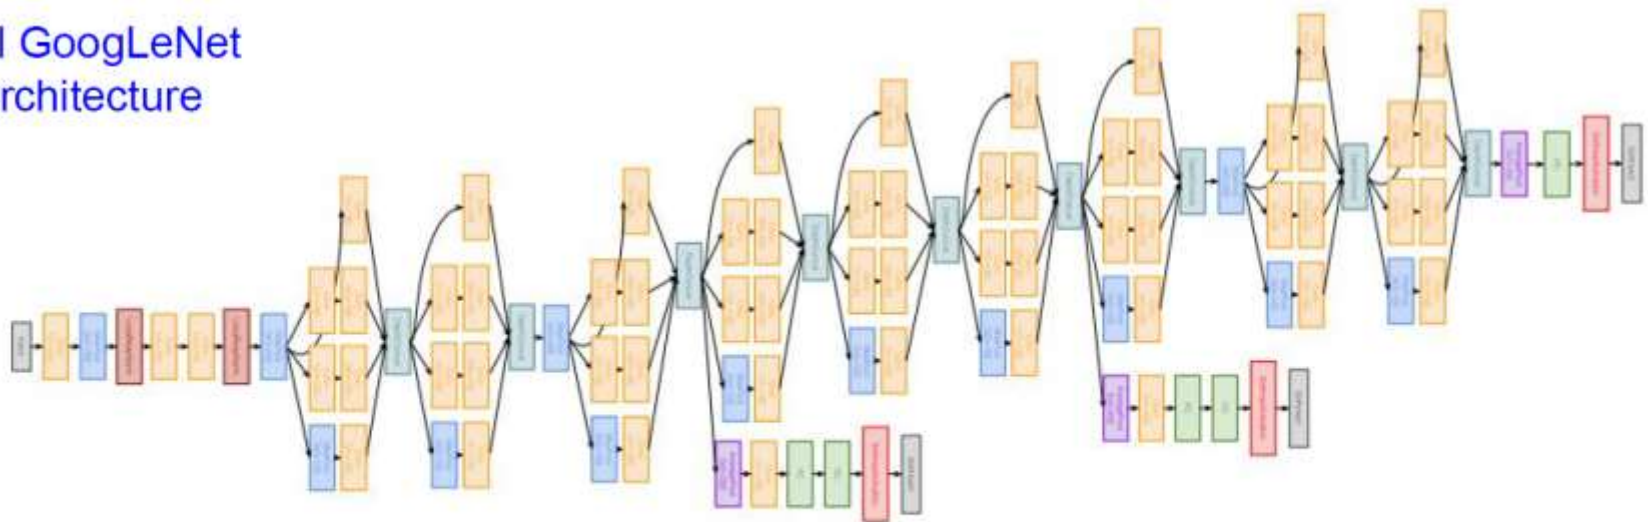


Auxiliary classification outputs to inject additional gradient at lower layers
(AvgPool-1x1Conv-FC-FC-Softmax)

GoogLeNet

- Overall network structure

Full GoogLeNet
architecture



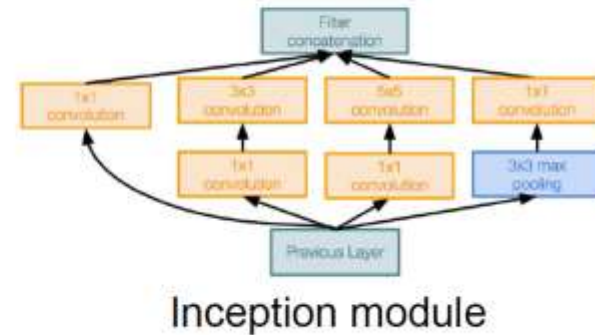
22 total layers with weights (including each parallel layer in an Inception module)

GoogLeNet

■ Summary

Deeper networks, with computational efficiency

- 22 layers
- Efficient “Inception” module
- No FC layers
- 12x less params than AlexNet
- ILSVRC’14 classification winner (6.7% top 5 error)



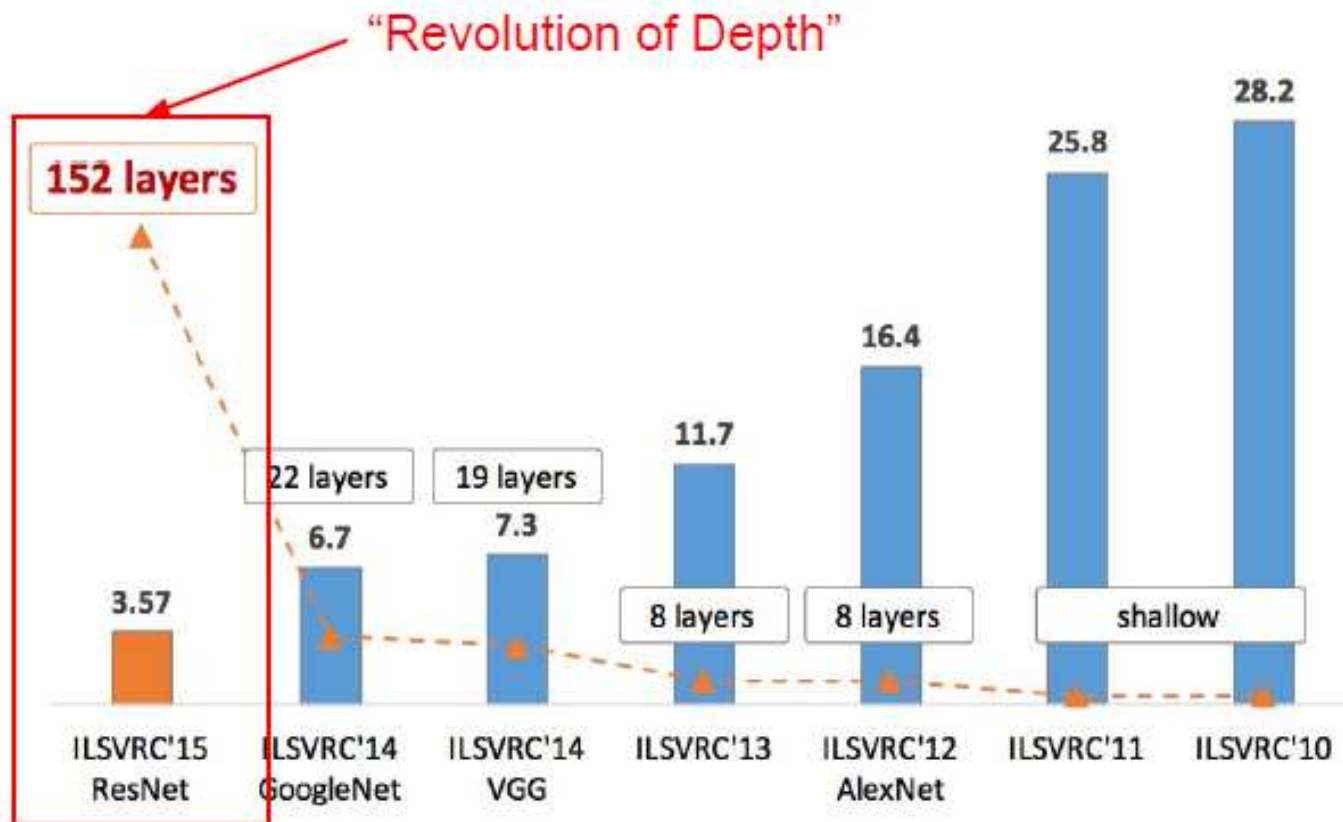
Outline

- CNN architectures

- Sequential structure: LeNet/AlexNet/VGGNet
- Parallel branches: GoogLeNet
- Residual structure: ResNet/DenseNet
- Network Architecture Search

Acknowledgement: Zemel et al's CSC411 and Feifei Li et al's cs231n notes

ImageNet (ILSVRC)



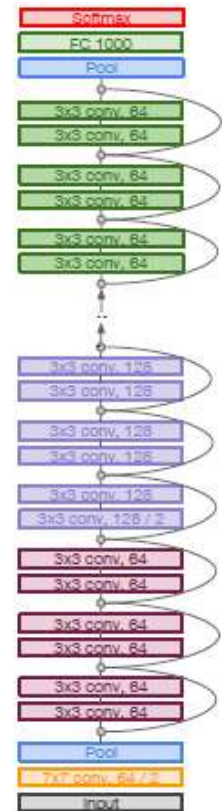
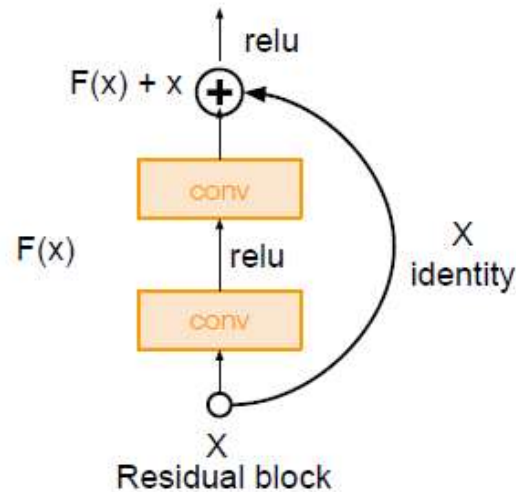
ResNet

Case Study: ResNet

[He et al., 2015]

Very deep networks using residual connections

- 152-layer model for ImageNet
- ILSVRC'15 classification winner (3.57% top 5 error)
- Swept all classification and detection competitions in ILSVRC'15 and COCO'15!



ResNet

- What happens when stacking deeper plain conv layers?



56-layer model performs worse on both training and test error
-> The deeper model performs worse, but it's not caused by overfitting!

ResNet

■ Hypothesis:

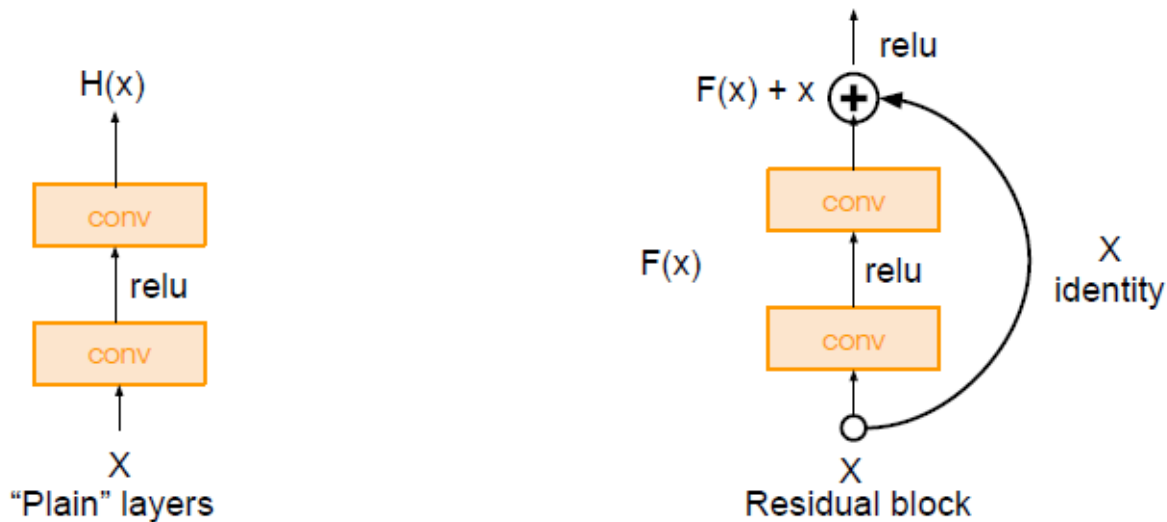
- The problem is an optimization problem, deeper models are harder to optimize

The deeper model should be able to perform at least as well as the shallower model.

A solution by construction is copying the learned layers from the shallower model and setting additional layers to identity mapping.

ResNet

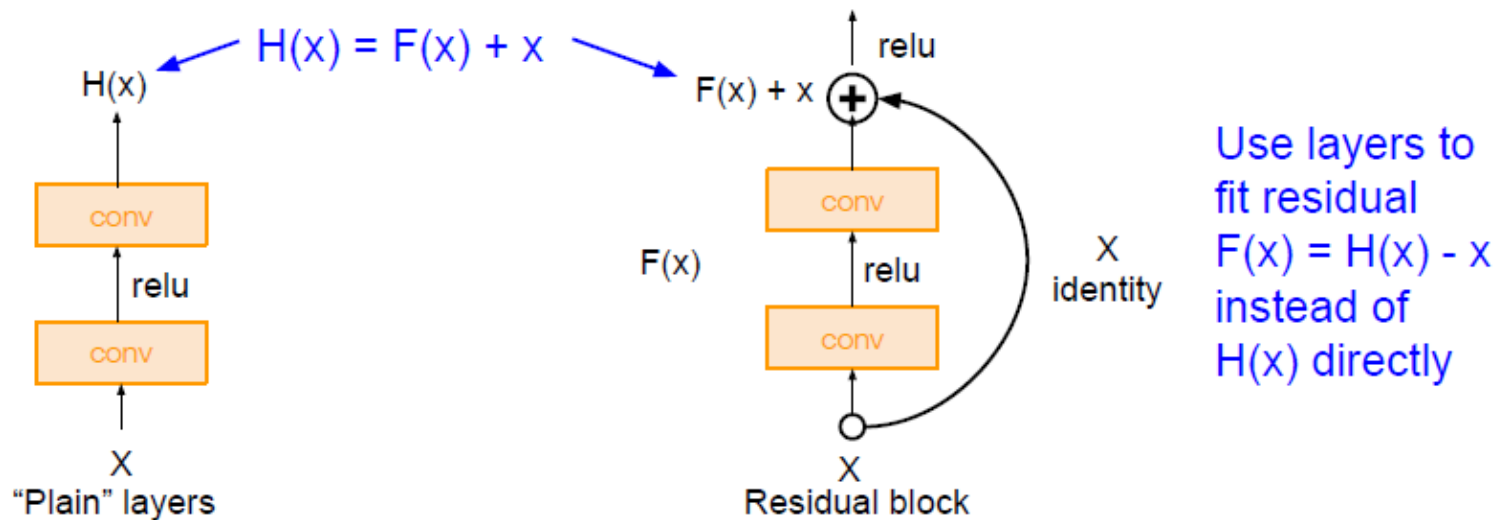
- Solution:
 - Use network layers to fit a residual mapping



He et al “Deep Residual Learning for Image Recognition”, CVPR 2016

ResNet

- Solution:
 - Use network layers to fit a residual mapping



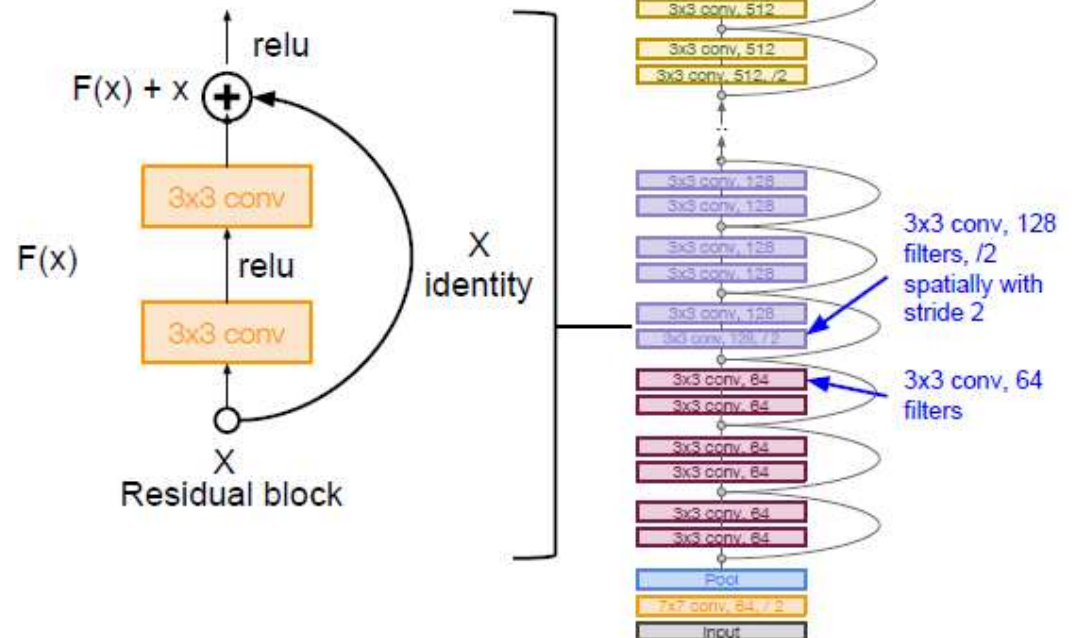
ResNet

Case Study: ResNet

[He et al., 2015]

Full ResNet architecture:

- Stack residual blocks
- Every residual block has two 3x3 conv layers
- Periodically, double # of filters and downsample spatially using stride 2 (/2 in each dimension)



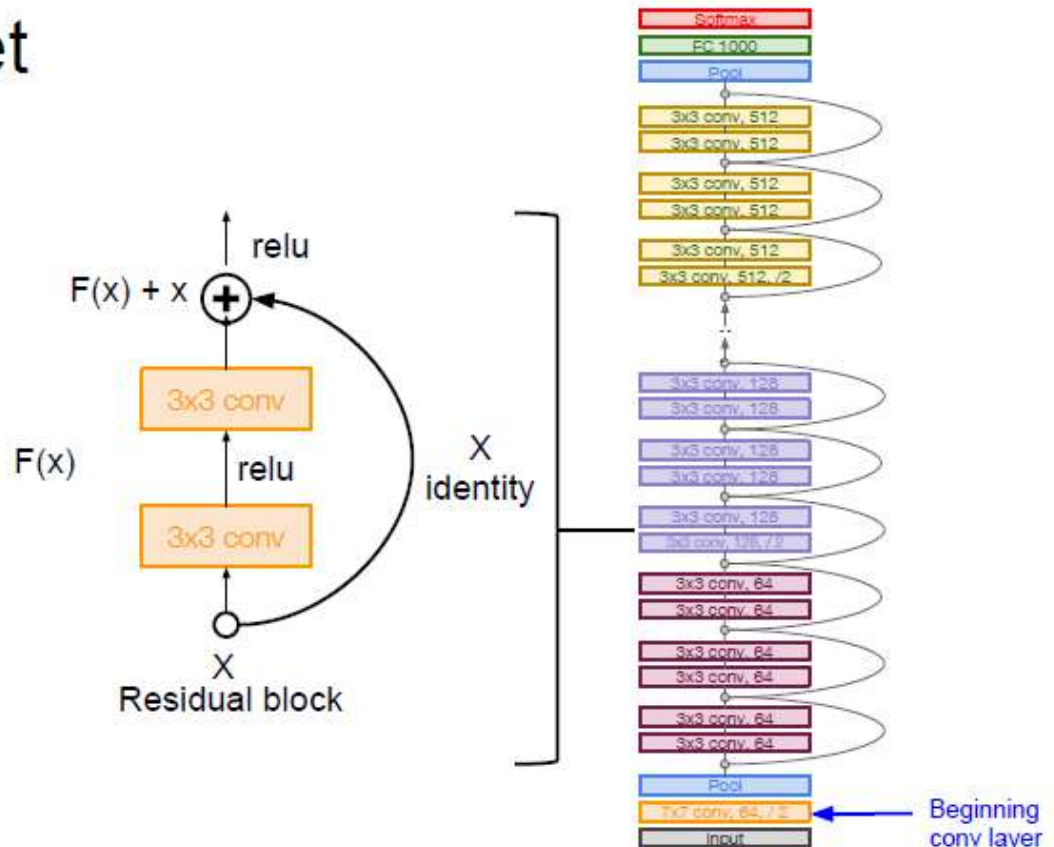
ResNet

Case Study: ResNet

[He et al., 2015]

Full ResNet architecture:

- Stack residual blocks
- Every residual block has two 3x3 conv layers
- Periodically, double # of filters and downsample spatially using stride 2 (/2 in each dimension)
- Additional conv layer at the beginning



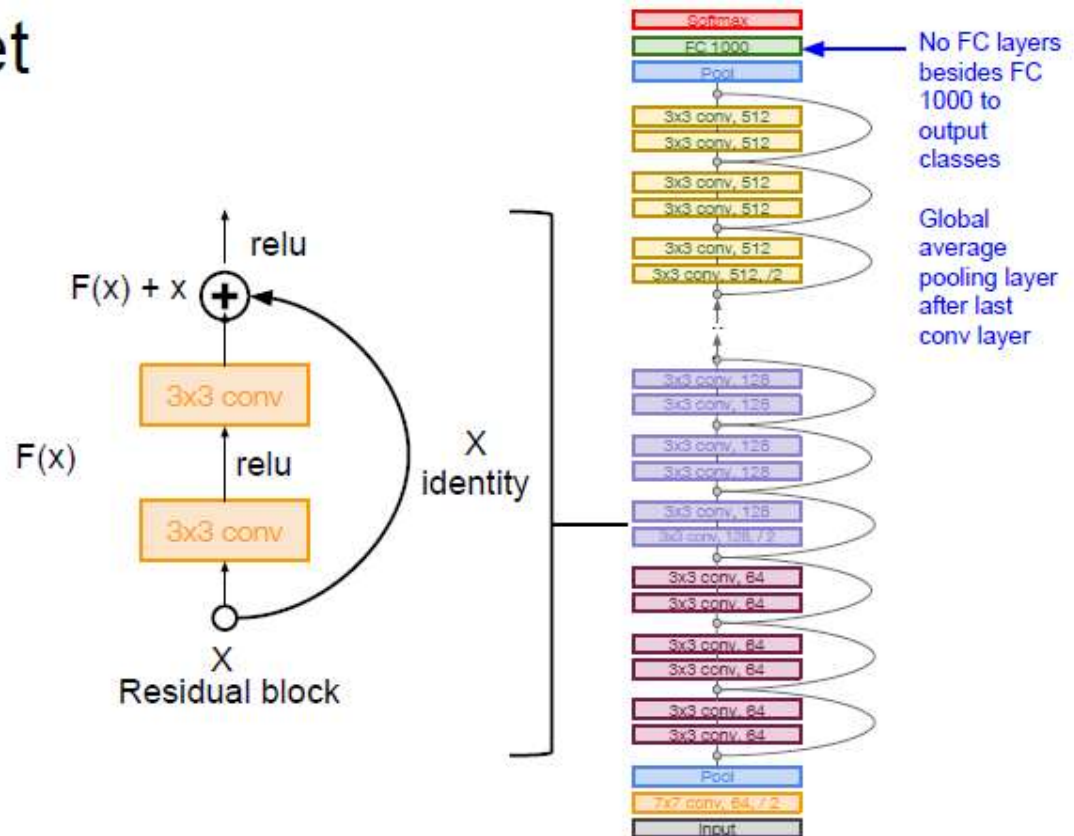
ResNet

Case Study: ResNet

[He et al., 2015]

Full ResNet architecture:

- Stack residual blocks
- Every residual block has two 3x3 conv layers
- Periodically, double # of filters and downsample spatially using stride 2 (/2 in each dimension)
- Additional conv layer at the beginning
- No FC layers at the end (only FC 1000 to output classes)

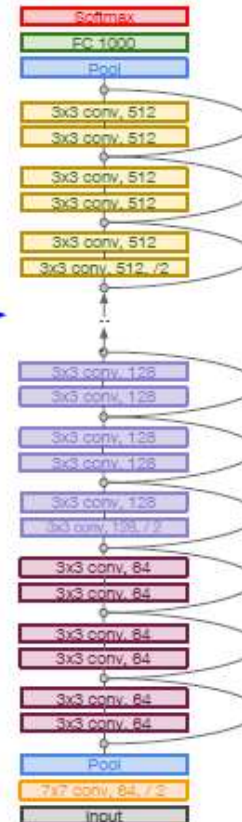


ResNet

Case Study: ResNet

[He et al., 2015]

Total depths of 34, 50, 101, or 152 layers for ImageNet

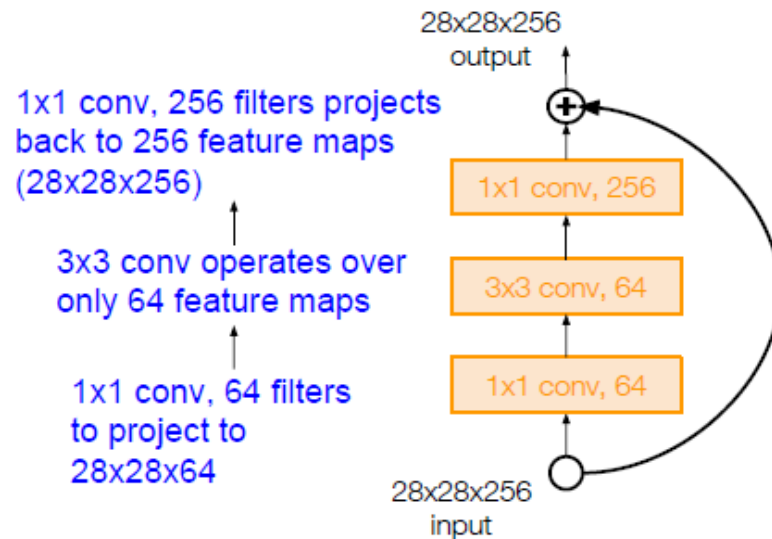


ResNet

Case Study: ResNet

[He et al., 2015]

For deeper networks
(ResNet-50+), use “bottleneck”
layer to improve efficiency
(similar to GoogLeNet)



ResNet

■ Training details

Training ResNet in practice:

- Batch Normalization after every CONV layer
- Xavier/2 initialization from He et al.
- SGD + Momentum (0.9)
- Learning rate: 0.1, divided by 10 when validation error plateaus
- Mini-batch size 256
- Weight decay of $1e-5$
- No dropout used

ResNet

■ Results

Experimental Results

- Able to train very deep networks without degrading (152 layers on ImageNet, 1202 on Cifar)
- Deeper networks now achieve lowering training error as expected
- Swept 1st place in all ILSVRC and COCO 2015 competitions

MSRA @ ILSVRC & COCO 2015 Competitions

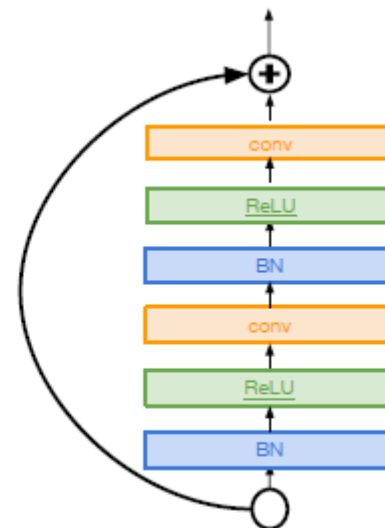
- **1st places in all five main tracks**

- ImageNet Classification: *"Ultra-deep"* (quote Yann) **152-layer** nets
- ImageNet Detection: **16%** better than 2nd
- ImageNet Localization: **27%** better than 2nd
- COCO Detection: **11%** better than 2nd
- COCO Segmentation: **12%** better than 2nd

ILSVRC 2015 classification winner (3.6% top 5 error) -- better than "human performance"! (Russakovsky 2014)

Other: Identity Mappings in ResNet

- Improved ResNet block design from creators of ResNet
- Creates a more direct path for propagating information throughout network (moves activation to residual mapping pathway)
- Gives better performance

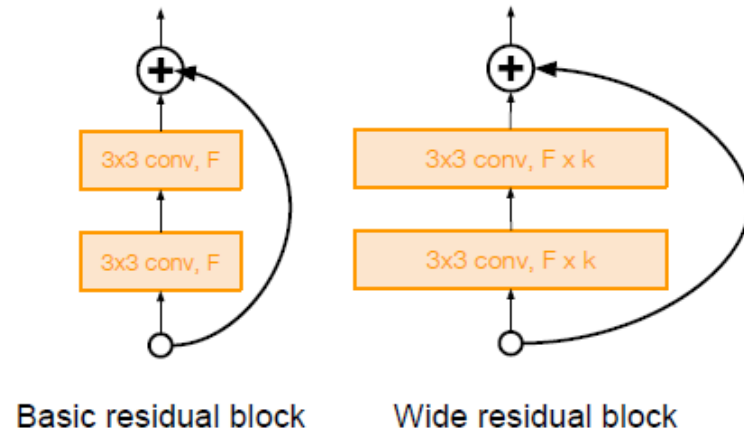


Other: Wide ResNets

Wide Residual Networks

[Zagoruyko et al. 2016]

- Argues that residuals are the important factor, not depth
- User wider residual blocks ($F \times k$ filters instead of F filters in each layer)
- 50-layer wide ResNet outperforms 152-layer original ResNet
- Increasing width instead of depth more computationally efficient (parallelizable)

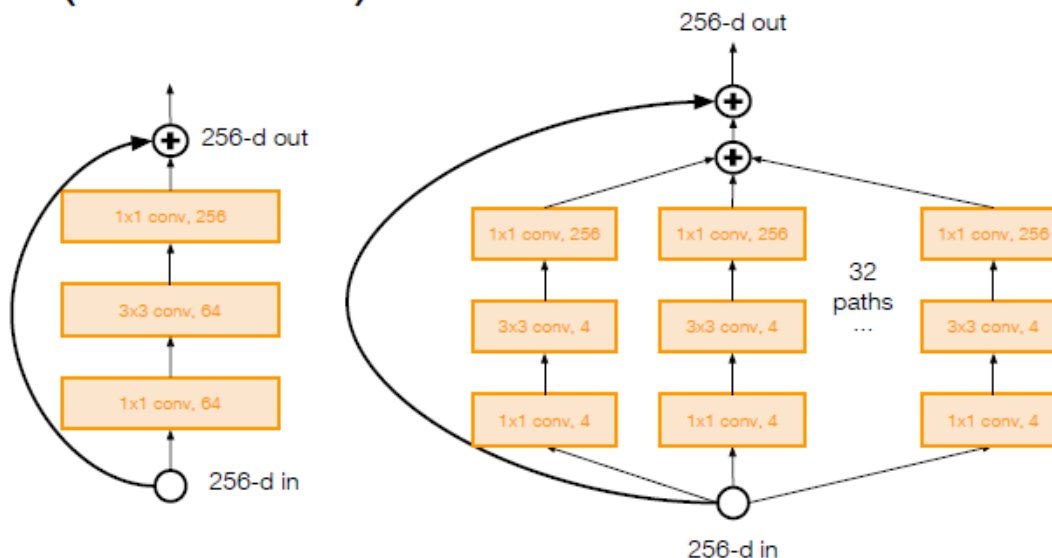


Other: ResNeXt

Aggregated Residual Transformations for Deep Neural Networks (ResNeXt)

[Xie et al. 2016]

- Also from creators of ResNet
- Increases width of residual block through multiple parallel pathways (“cardinality”)
- Parallel pathways similar in spirit to Inception module

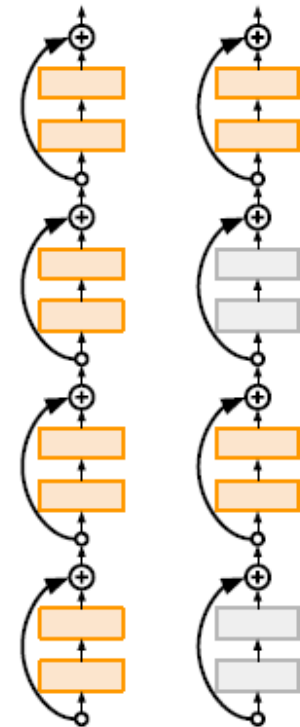


Other: ResNet with Stochastic Depth

Deep Networks with Stochastic Depth

[Huang et al. 2016]

- Motivation: reduce vanishing gradients and training time through short networks during training
- Randomly drop a subset of layers during each training pass
- Bypass with identity function
- Use full deep network at test time

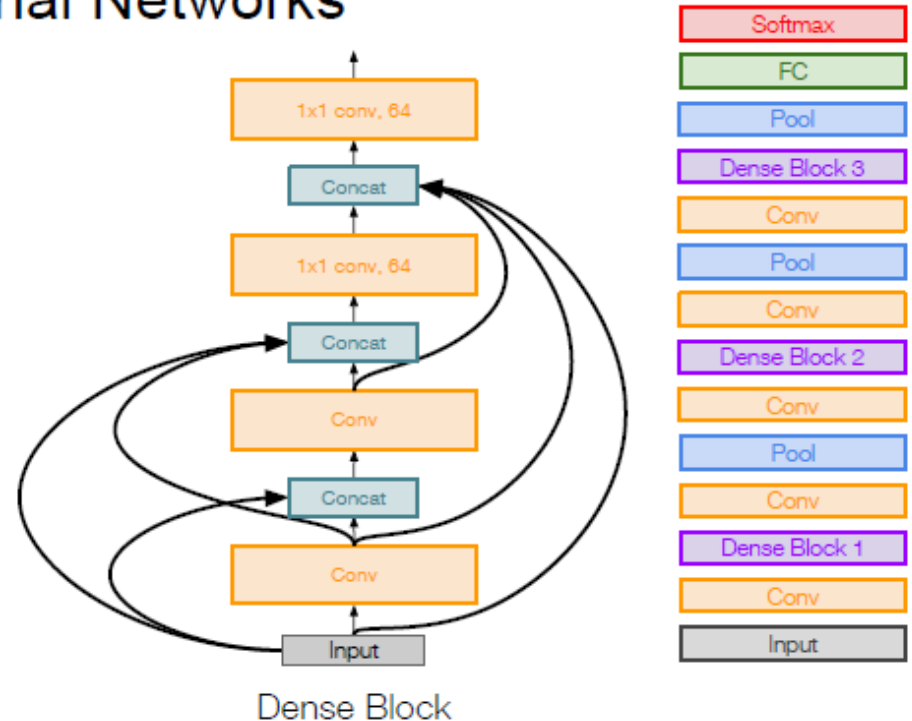


DenseNet

Densely Connected Convolutional Networks

[Huang et al. 2017]

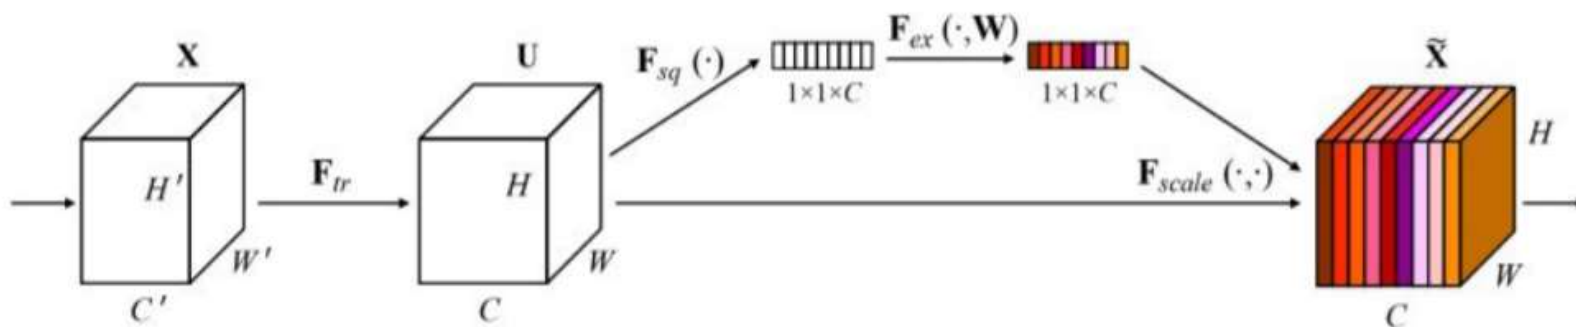
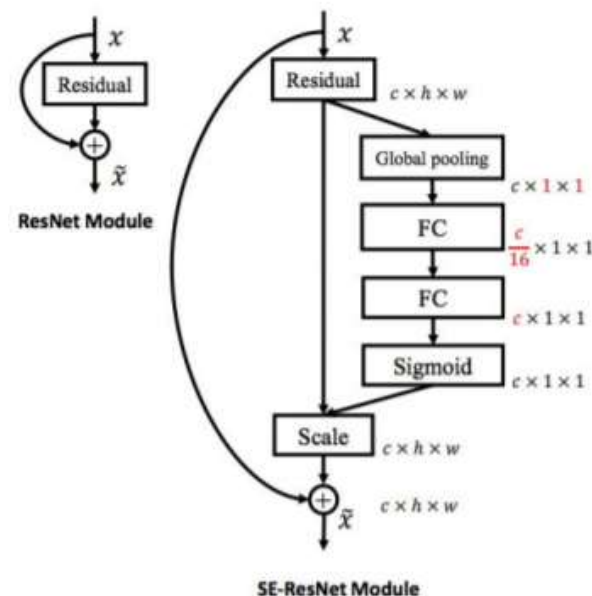
- Dense blocks where each layer is connected to every other layer in feedforward fashion
- Alleviates vanishing gradient, strengthens feature propagation, encourages feature reuse



Squeeze-and-Excitation Networks (SENet)

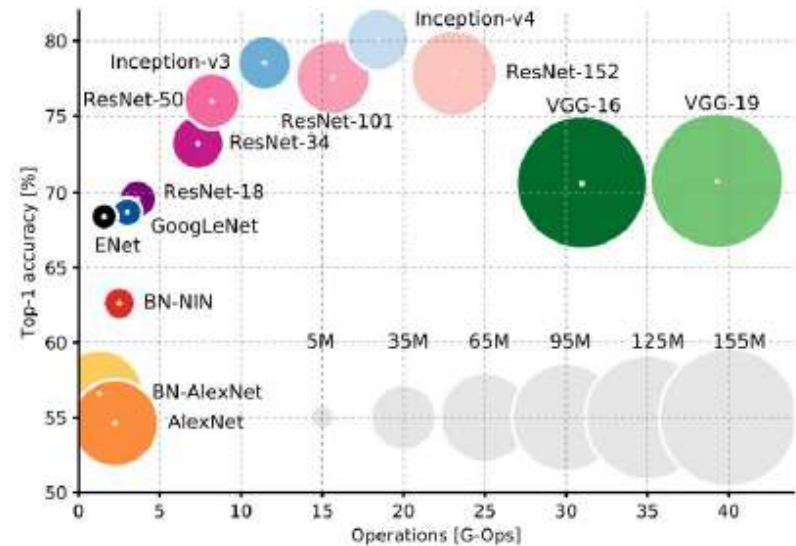
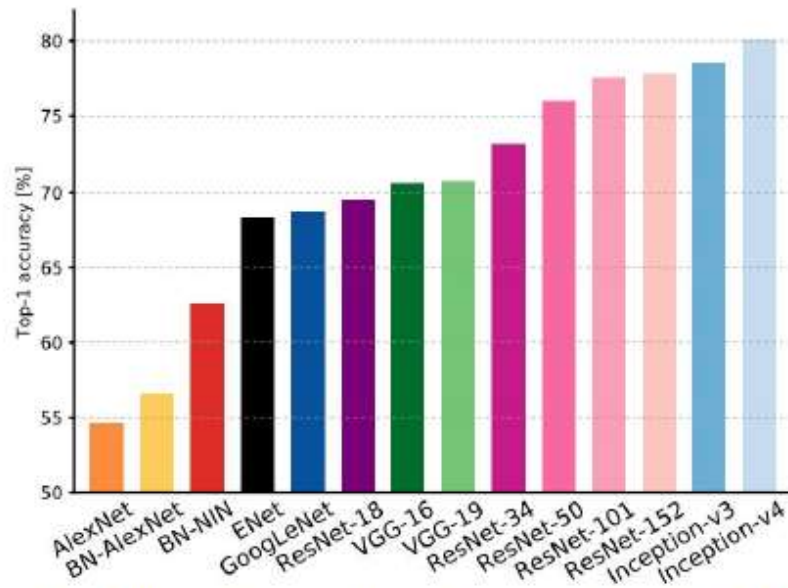
[Hu et al. 2017]

- Add a “feature recalibration” module that learns to adaptively reweight feature maps
- Global information (global avg. pooling layer) + 2 FC layers used to determine feature map weights
- ILSVRC'17 classification winner (using ResNeXt-152 as a base architecture)



Model complexity

Comparing complexity...



An Analysis of Deep Neural Network Models for Practical Applications, 2017.

Outline

■ CNN architectures

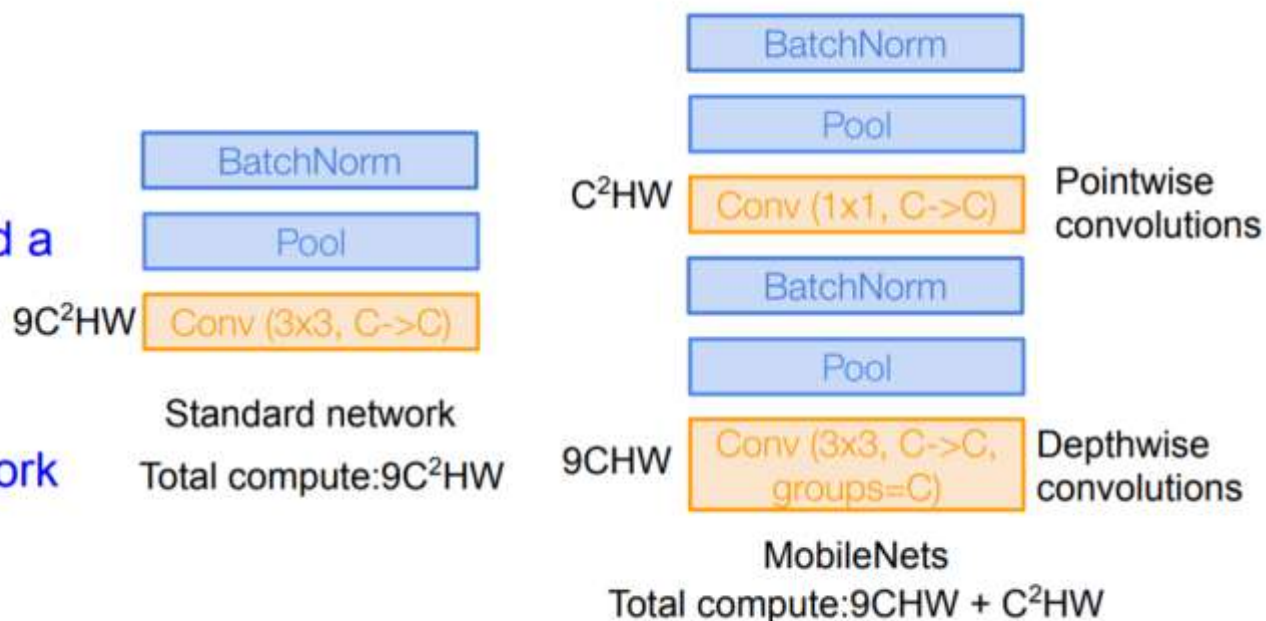
- Sequential structure: LeNet/AlexNet/VGGNet
- Parallel branches: GoogLeNet
- Residual structure: ResNet/DenseNet
- Network Architecture Search

Acknowledgement: Zemel et al's CSC411 and Feifei Li et al's cs231n notes

Efficient networks

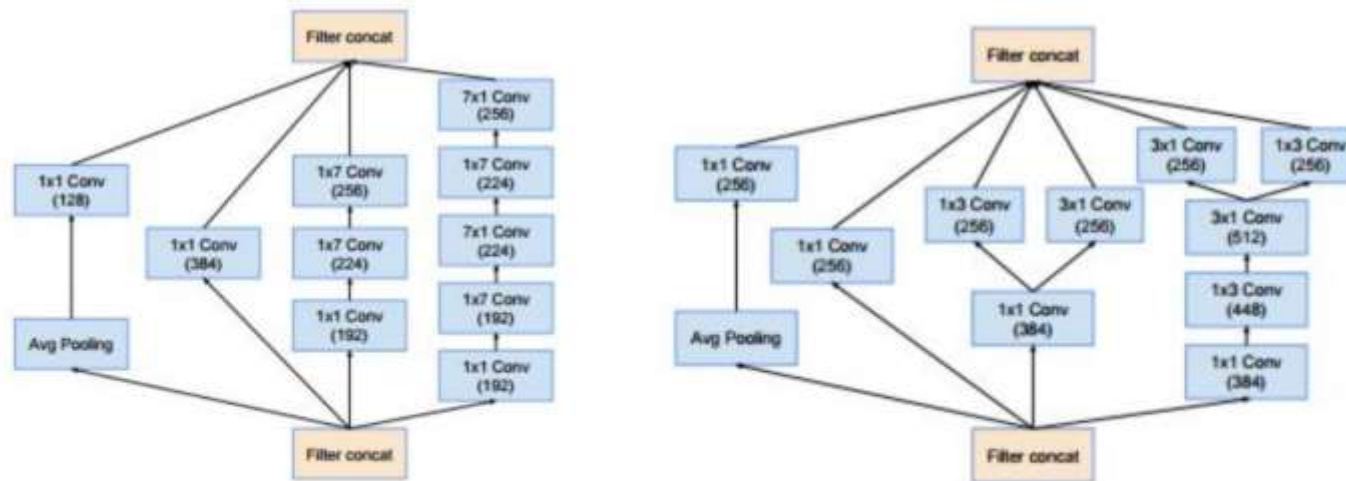
■ MobileNets: Efficient Convolutional Neural Networks for Mobile Applications [Howard et al. 2017]

- Depthwise separable convolutions replace standard convolutions by factorizing them into a depthwise convolution and a 1×1 convolution
- Much more efficient, with little loss in accuracy
- Follow-up MobileNetV2 work in 2018 (Sandler et al.)
- ShuffleNet: Zhang et al, CVPR 2018



Network Architecture

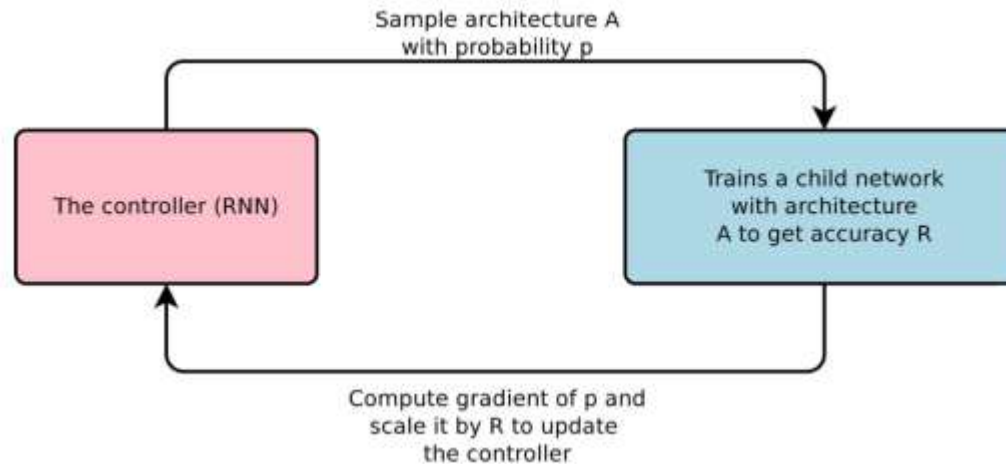
- Problems with network architecture
 - Designing NA is hard
 - Lots of human efforts go into tuning them
 - Not a lot of intuition into how to design them well
 - Can we learn good architectures automatically?



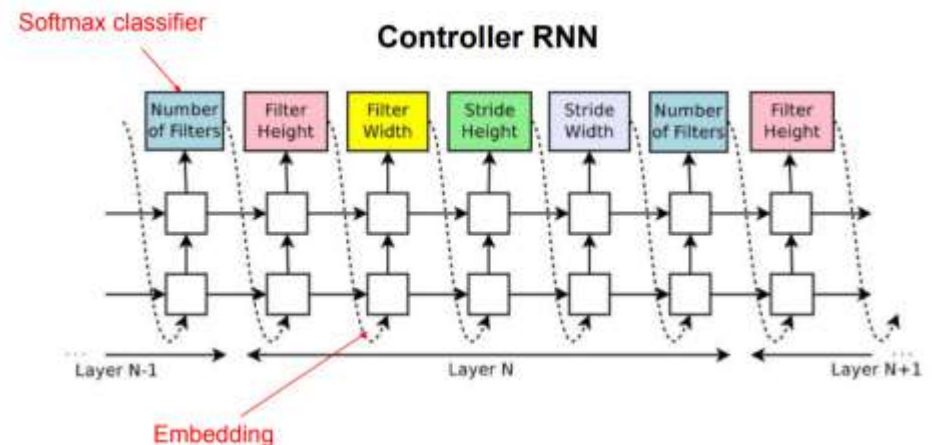
Two layers from the famous Inception V4 computer vision model.
Szegedy et al, 2017

Network Architecture

■ Neural architecture search (Zoph and Le, ICLR 2016)

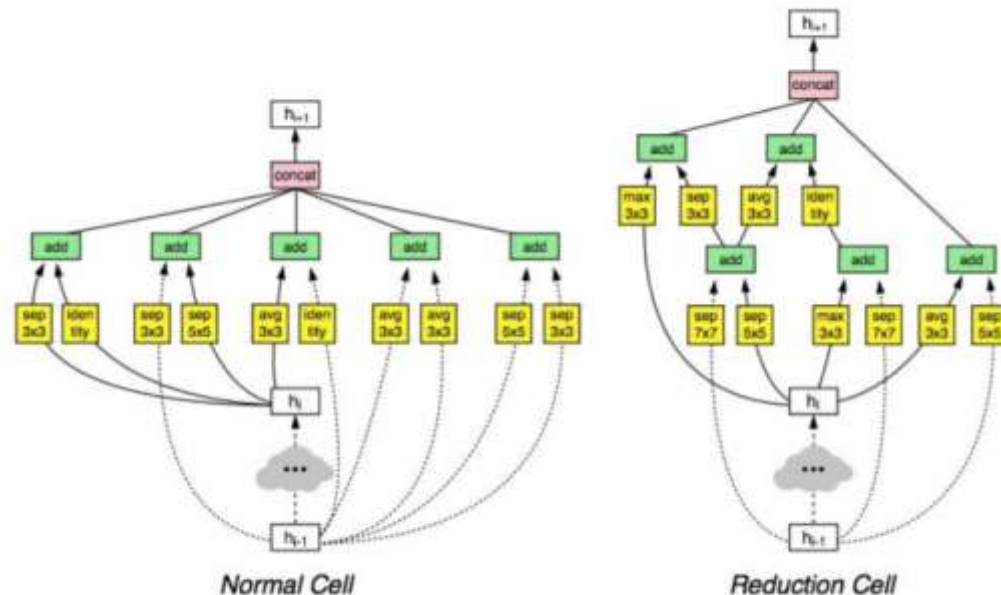


- “Controller” network that learns to design a good network architecture (output a string corresponding to network design)
- Iterate:
 - 1) Sample an architecture from search space
 - 2) Train the architecture to get a “reward” R corresponding to accuracy
 - 3) Compute gradient of sample probability, and scale by R to perform controller parameter update (i.e. increase likelihood of good architecture being sampled, decrease likelihood of bad architecture)



Network Architecture

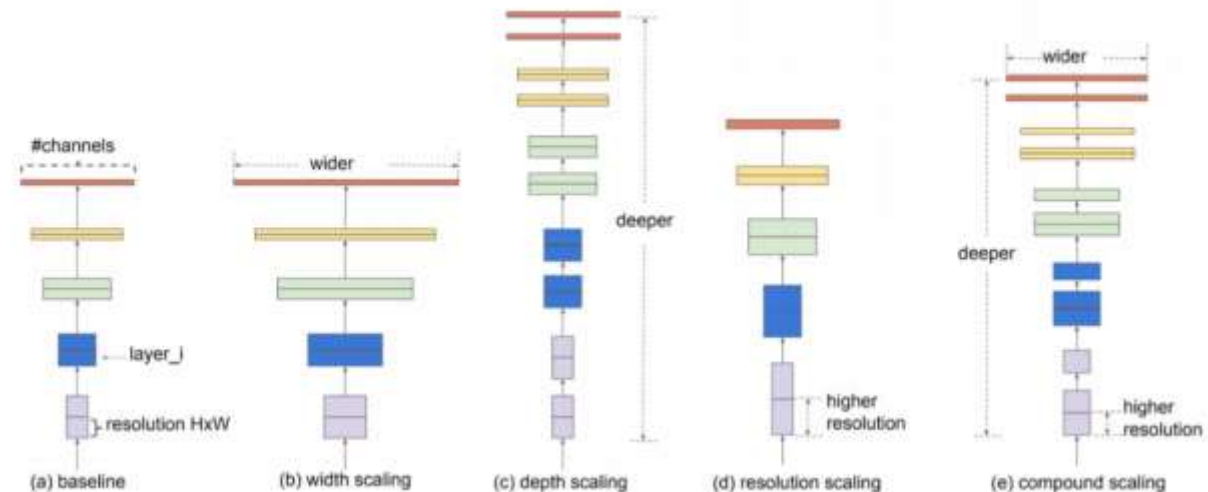
- Neural architecture search (Zoph et al. 2017)
 - Design a search space of building blocks (“cells”) that can be flexibly stacked
 - NASNet: Use NAS to find best cell structure on smaller CIFAR-10 dataset, then transfer architecture to ImageNet
 - Many follow-up works in this space e.g. AmoebaNet (Real et al. 2019) and ENAS (Pham, Guan et al. 2018)



Network Architecture

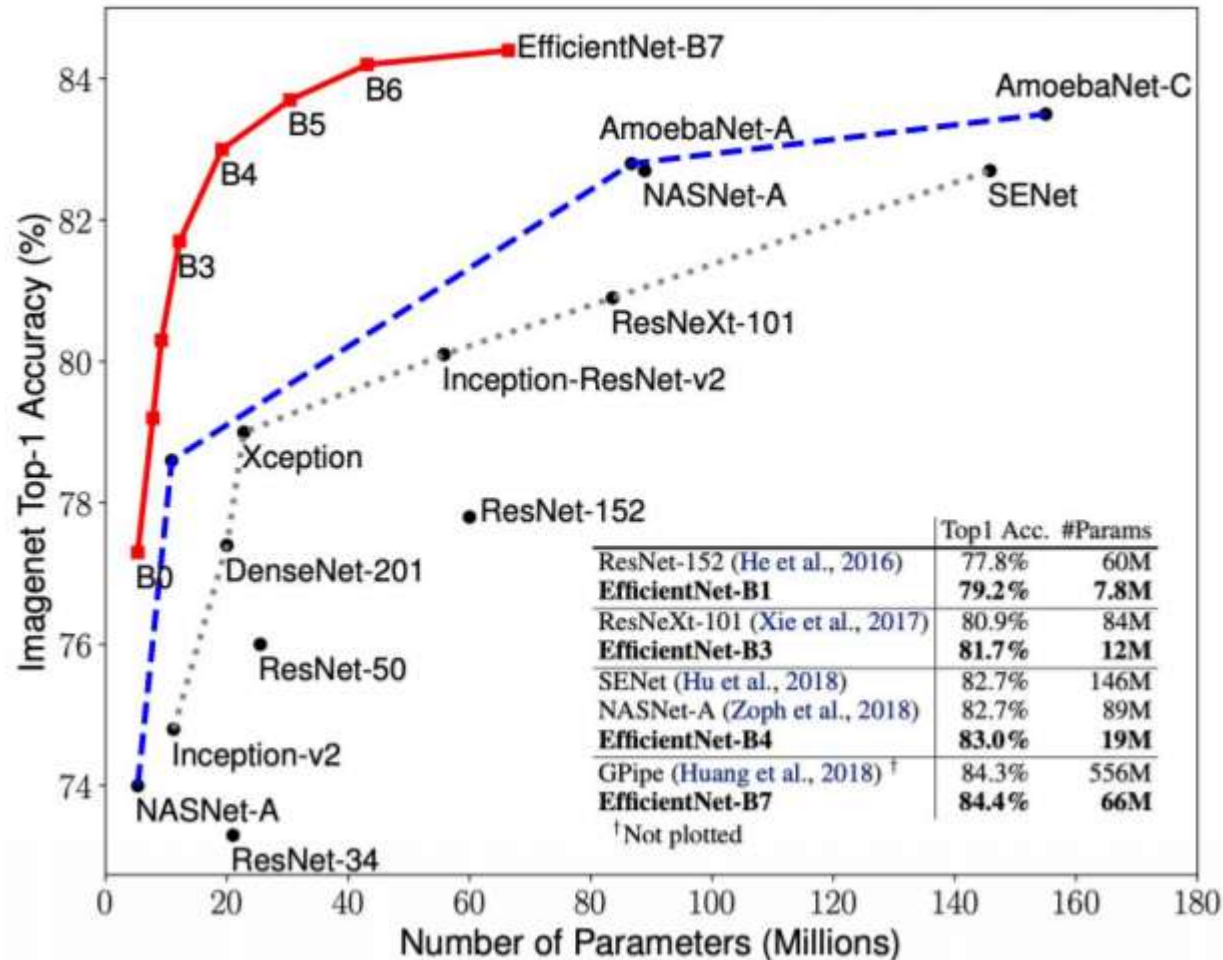
- EfficientNet: Smart Compound Scaling [Tan and Le. 2019]
 - Increase network capacity by scaling width, depth, and resolution, while balancing accuracy and efficiency.
 - Search for optimal set of compound scaling factors given a compute budget (target memory & flops).
 - Scale up using smart heuristic rules

depth: $d = \alpha^\phi$
width: $w = \beta^\phi$
resolution: $r = \gamma^\phi$
s.t. $\alpha \cdot \beta^2 \cdot \gamma^2 \approx 2$
 $\alpha \geq 1, \beta \geq 1, \gamma \geq 1$



Network Architecture

- EfficientNet: Smart Compound Scaling [Tan and Le. 2019]



Network structure summary

- AlexNet showed that you can use CNNs to train Computer Vision models.
- ZFNet, VGG shows that bigger networks work better
- GoogLeNet is one of the first to focus on efficiency using 1x1 bottleneck convolutions and global avg pool instead of FC layers
- ResNet showed us how to train extremely deep networks
 - Limited only by GPU & memory!
 - Showed diminishing returns as networks got bigger
- After ResNet: CNNs were *better than the human metric* and focus shifted to Efficient networks:
 - Lots of tiny networks aimed at mobile devices: MobileNet, ShuffleNet
- Neural Architecture Search can now automate architecture design

Summary

- Bag of tricks for improving generalization
 - Pros: you have a toolbox to use
 - Cons: many trial and error, tedious process
- Seeking fully automatic approaches to model selection
 - Bayesian optimization
 - Reinforcement learning
- Next time
 - CNN in Vision, RNN
- Reference
 - CS231n course notes