



# Lecture 18: Diffusion Model II: Score-based Methods

Lan Xu  
SIST, ShanghaiTech  
Fall, 2023



# Outline

- Diffusion: Denoising Diffusion Probabilistic Models
- Score-based Diffusion Models
- Accelerated Sampling
- Conditional Generation and Guidance


# The Landscape of Generative Models



# Diffusion: A Generative Learning Big Bang



# Diffusion: new trend in CVPR



# Recall DDPM

- Denoising Diffusion Probabilistic Models
- Target: understand the training and sampling phases!

---

## Algorithm 1 Training

---

```
1: repeat  
2:    $\mathbf{x}_0 \sim q(\mathbf{x}_0)$   
3:    $t \sim \text{Uniform}(\{1, \dots, T\})$   
4:    $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$   
5:   Take gradient descent step on  
        $\nabla_{\theta} \|\epsilon - \epsilon_{\theta}(\sqrt{\bar{\alpha}_t}\mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon, t)\|^2$   
6: until converged
```

---

## Algorithm 2 Sampling

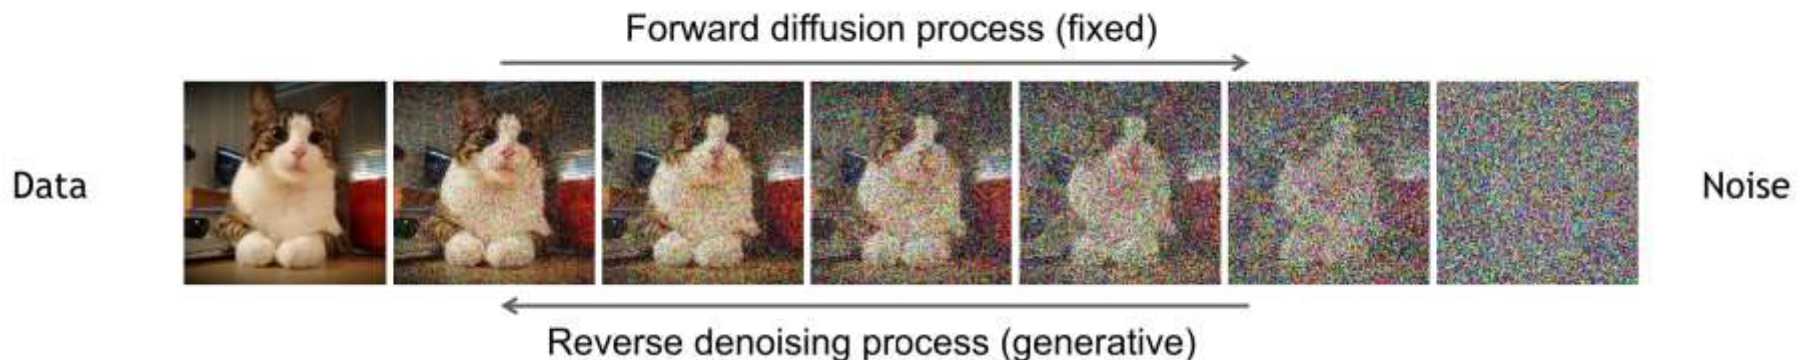
---

```
1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$   
2: for  $t = T, \dots, 1$  do  
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$   
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_{\theta}(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$   
5: end for  
6: return  $\mathbf{x}_0$ 
```

---

# Recall DDPM

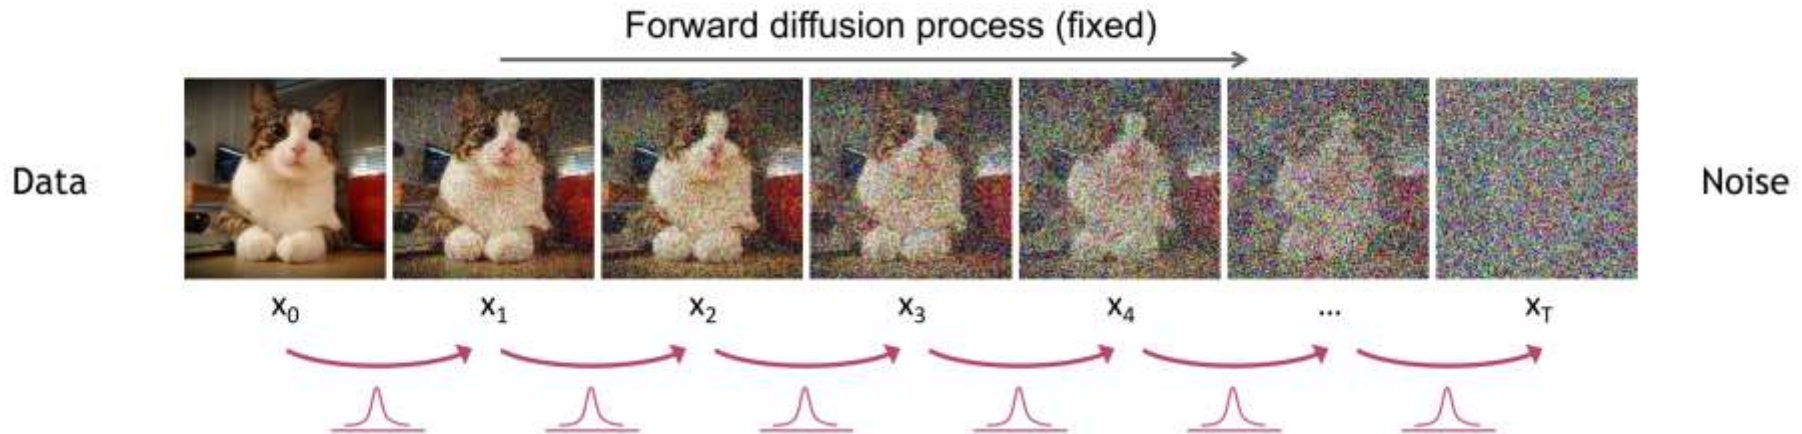
- Learning to generate by denoising
- **Forward diffusion** process that gradually adds noise to input
- **Reverse denoising** process that learns to generate data by denoising





# Recall DDPM

- Forward Diffusion Process in T steps



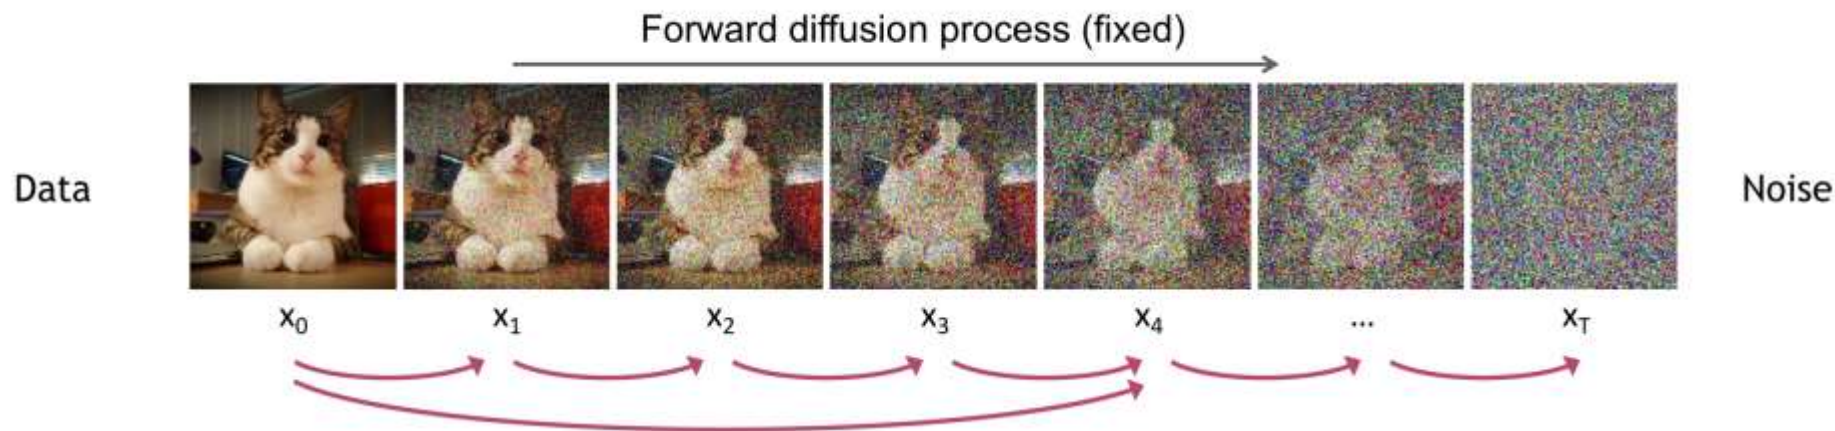
$$q(\mathbf{x}_t|\mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_t; \sqrt{1 - \beta_t}\mathbf{x}_{t-1}, \beta_t\mathbf{I}) \quad q(\mathbf{x}_{1:T}|\mathbf{x}_0) = \prod_{t=1}^T q(\mathbf{x}_t|\mathbf{x}_{t-1})$$

$$x_t = \sqrt{1 - \beta_t}x_{t-1} + \sqrt{\beta_t}\epsilon \quad \text{where} \quad \epsilon \sim \mathcal{N}(0, I)$$



# Recall DDPM

## ■ Diffusion Kernel and Reparametrization Trick



$$\alpha_t = 1 - \beta_t \quad \bar{\alpha}_t = \prod_{i=1}^t \alpha_i$$

$$q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I})$$


$$\mathbf{x}_t = \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \epsilon, \quad \epsilon \sim \mathcal{N}(0, \mathbf{I})$$

$$= \sqrt{\alpha_t} \mathbf{x}_{t-1} + \sqrt{1 - \alpha_t} \epsilon$$

$$= \sqrt{\alpha_t \alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{1 - \alpha_t \alpha_{t-1}} \epsilon$$

$$= \dots$$

$$= \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon$$



$$q(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I})$$


$$\mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon$$

The noise schedule) is designed such that  $\bar{\alpha}_T \rightarrow 0$

# Recall DDPM

- What happens to a distribution in the forward diffusion?
- The diffusion kernel  $q(\mathbf{x}_t|\mathbf{x}_0)$  is **Gaussian convolution**!

$$\underbrace{q(\mathbf{x}_t)}_{\text{Diffused data dist.}} = \int \underbrace{q(\mathbf{x}_0, \mathbf{x}_t)}_{\text{Joint dist.}} d\mathbf{x}_0 = \int \underbrace{q(\mathbf{x}_0)}_{\text{Input data dist.}} \underbrace{q(\mathbf{x}_t|\mathbf{x}_0)}_{\text{Diffusion kernel}} d\mathbf{x}_0$$



- We can sample  $\mathbf{x}_t \sim q(\mathbf{x}_t)$  by first sampling  $\mathbf{x}_0 \sim q(\mathbf{x}_0)$ , and then sampling  $\mathbf{x}_t \sim q(\mathbf{x}_t|\mathbf{x}_0)$  (ancestral sampling).


# Recall DDPM

- Reverse: Generative Learning by Denoising
- Diffusion parameters are designed:  $q(\mathbf{x}_T) \approx \mathcal{N}(\mathbf{x}_T; \mathbf{0}, \mathbf{I})$

Generation:

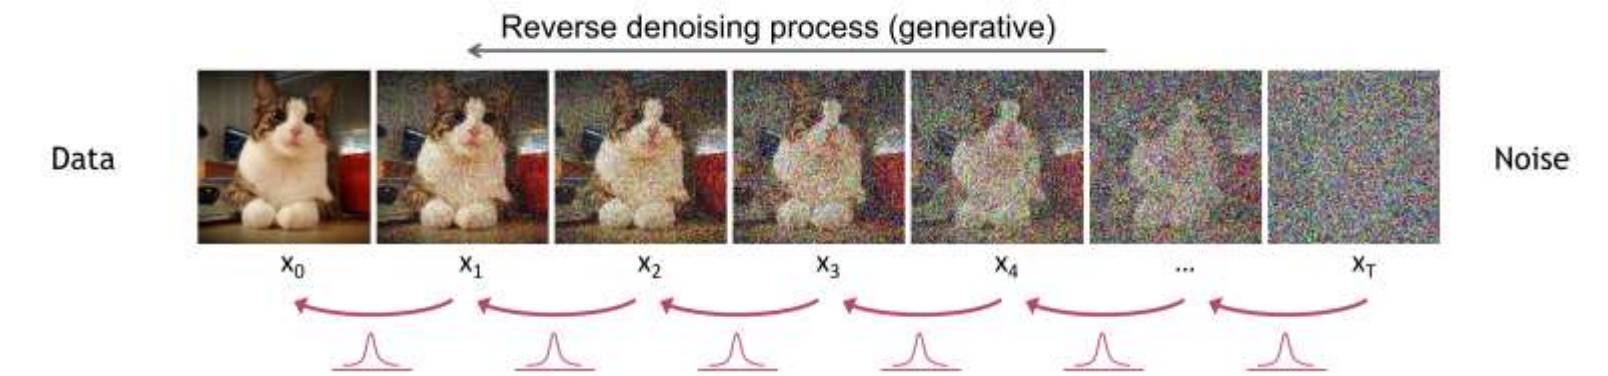
Sample  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{x}_T; \mathbf{0}, \mathbf{I})$

Iteratively sample  $\mathbf{x}_{t-1} \sim \underbrace{q(\mathbf{x}_{t-1}|\mathbf{x}_t)}_{\text{True Denoising Dist.}}$



# Recall DDPM

## ■ Reverse Denoising Process



$$p_{\theta}(\mathbf{x}_{0:T}) = p(\mathbf{x}_T) \prod_{t=1}^T p_{\theta}(\mathbf{x}_{t-1} | \mathbf{x}_t)$$

## ■ From $q(x_{t-1} | x_t)$ to $q(x_{t-1} | x_t, x_0)$

$$\begin{aligned} q(x_{t-1} | x_t, x_0) &= \frac{q(x_{t-1}, x_t, x_0)}{q(x_t, x_0)} = \frac{q(x_t | x_{t-1}, x_0) \cdot q(x_{t-1}, x_0)}{q(x_t, x_0)} = \frac{q(x_t | x_{t-1}, x_0) \cdot q(x_{t-1} | x_0)}{q(x_t | x_0)} \\ &= \frac{q(x_t | x_{t-1}) \cdot q(x_{t-1} | x_0)}{q(x_t | x_0)} \end{aligned}$$

# Recall DDPM

- Learning Denoising Model

$$\frac{q(x_t|x_{t-1}) \cdot q(x_{t-1}|x_0)}{q(x_t|x_0)}$$

$$q(x_{t-1}|x_t, x_0) = \mathcal{N}(x_{t-1}; \tilde{\mu}(x_t, x_0), \tilde{\beta}_t \mathbf{I})$$

where  $\tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0) := \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1-\bar{\alpha}_t}\mathbf{x}_0 + \frac{\sqrt{\alpha_t}(1-\bar{\alpha}_{t-1})}{1-\bar{\alpha}_t}\mathbf{x}_t$  and  $\tilde{\beta}_t := \frac{1-\bar{\alpha}_{t-1}}{1-\bar{\alpha}_t}\beta_t$

- Since  $x_0 = \frac{1}{\sqrt{\bar{\alpha}_t}}(x_t - \sqrt{1-\bar{\alpha}_t}\bar{\epsilon}_t)$ , thus:  $\tilde{\mu}_t = \frac{1}{\sqrt{\alpha_t}}(x_t - \frac{1-\alpha_t}{\sqrt{1-\bar{\alpha}_t}}\bar{\epsilon}_t)$
- Use NN to regress the noise:

$$\mu_\theta(\mathbf{x}_t, t) = \frac{1}{\sqrt{1-\beta_t}} \left( \mathbf{x}_t - \frac{\beta_t}{\sqrt{1-\bar{\alpha}_t}} \epsilon_\theta(\mathbf{x}_t, t) \right)$$

$$L = \mathbb{E}_{\mathbf{x}_0 \sim q(\mathbf{x}_0), t \sim \mathcal{U}\{1, T\}, \epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})} \left[ \lambda_t ||\epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1-\bar{\alpha}_t} \epsilon, t)||^2 \right]$$

# Recall DDPM

- The authors of DDPM say that it's fine to drop all that baggage in the front and instead just use

$$\begin{aligned} L &= \mathbb{E}_{\mathbf{x}_0, \epsilon} \left[ \|\epsilon - \epsilon_{\theta}(\mathbf{x}_t, t)\|_2^2 \right] \\ &= \mathbb{E}_{\mathbf{x}_0, \epsilon} \left[ \left\| \epsilon - \epsilon_{\theta} \left( \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t \right) \right\|_2^2 \right] \end{aligned}$$

- Note that this is not a variational lower bound on the log-likelihood anymore: in fact, you can view it as a **reweighted version of ELBO** that emphasizes reconstruction quality!

# Recall DDPM

- If we have the noise, sampling by using Gaussians:

$$q(x_{t-1}|x_t, x_0) = \mathcal{N}(x_{t-1}; \tilde{\mu}(x_t, x_0), \tilde{\beta}_t \mathbf{I})$$

- 1) sampling  $z_t$
- 2) sampling  $x_{t-1}$ , using the estimated noise

$$x_{t-1} = \tilde{\mu}_t + \tilde{\beta}_t \cdot z_t = \frac{1}{\sqrt{\alpha_t}} \left( x_t - \frac{1-\alpha_t}{\sqrt{1-\bar{\alpha}_t}} \bar{\epsilon}_t \right) + \frac{1-\bar{\alpha}_{t-1}}{1-\bar{\alpha}_t} \cdot \beta_t \cdot z_t$$

$$x_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( x_t - \frac{1-\alpha_t}{\sqrt{1-\bar{\alpha}_t}} \epsilon_{\theta}(x_t, t) \right) + \frac{1-\bar{\alpha}_{t-1}}{1-\bar{\alpha}_t} \cdot \beta_t \cdot z_t$$



# Recall DDPM

## ■ Rethinking the Training and Sampling processes.....

---

### Algorithm 1 Training

---

```
1: repeat  
2:    $\mathbf{x}_0 \sim q(\mathbf{x}_0)$   
3:    $t \sim \text{Uniform}(\{1, \dots, T\})$   
4:    $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$   
5:   Take gradient descent step on  
        $\nabla_{\theta} \|\epsilon - \epsilon_{\theta}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t)\|^2$   
6: until converged
```

---

---

### Algorithm 2 Sampling

---


```
1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$   
2: for  $t = T, \dots, 1$  do  
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$   
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_{\theta}(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$   
5: end for  
6: return  $\mathbf{x}_0$ 
```

---



- During training, add noise from 0 to t, then estimate it
- During sampling, note that  $\sigma_t = \frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t} \cdot \beta_t$
- As t increases,  $\bar{\alpha}_t$  decreases,  $\sqrt{1 - \bar{\alpha}_t}$  increases
- Thus,  $\epsilon_{\theta}(\mathbf{x}_t, t)$  works as denoise auto-encoder for various noise levels!

# Recall DDPM

## ■ Forward/Reverse process for Image Generation




Forward process:  
converting the image  
distribution to pure  
noise



Reverse process:  
sampling from the  
image distribution,  
starting with pure  
noise

# Implementation Considerations

## ■ UNet + Other Stuff




Diffusion models typically use a U-Net on steroids as the noise predictive model — you take the good ol' model that you are already familiar with and add:

- Positional Embeddings
- ResNet Blocks
- ConvNext Blocks
- Attention Modules
- Group Normalization
- Swish and GeLU

It's a massive kitchen sink of modern CV tricks

# Implementation Considerations

- Diffusion models often use **U-Net architectures** with ResNet blocks and self-attention layers



- Time representation: sinusoidal positional embeddings or random Fourier features.
- Time features are fed to the residual blocks using either simple spatial addition or using adaptive group normalization layers.

# Outline


- Diffusion: Denoising Diffusion Probabilistic Models
- **Score-based Diffusion Models**
- Accelerated Sampling
- Conditional Generation and Guidance

# Crash Course in Differential Equations

## ■ ODE and SDE

Ordinary Differential Equation (ODE):

$$\frac{dx}{dt} = f(x, t) \quad \text{or} \quad dx = f(x, t)dt$$



Analytical Solution:  $\mathbf{x}(t) = \mathbf{x}(0) + \int_0^t \mathbf{f}(\mathbf{x}, \tau) d\tau$

Iterative Numerical Solution:  $\mathbf{x}(t + \Delta t) \approx \mathbf{x}(t) + \mathbf{f}(\mathbf{x}(t), t)\Delta t$




# Crash Course in Differential Equations

## ■ ODE and SDE

Ordinary Differential Equation (ODE):

$$\frac{dx}{dt} = f(x, t) \quad \text{or} \quad dx = f(x, t)dt$$




Analytical Solution: 
$$\mathbf{x}(t) = \mathbf{x}(0) + \int_0^t \mathbf{f}(\mathbf{x}, \tau) d\tau$$

Iterative Numerical Solution: 
$$\mathbf{x}(t + \Delta t) \approx \mathbf{x}(t) + \mathbf{f}(\mathbf{x}(t), t)\Delta t$$

Stochastic Differential Equation (SDE):


$$\frac{dx}{dt} = \underbrace{f(x, t)}_{\text{drift coefficient}} + \underbrace{\sigma(x, t)\omega_t}_{\text{diffusion coefficient}}$$

$$(dx = f(x, t)dt + \sigma(x, t)d\omega_t)$$



$$\mathbf{x}(t + \Delta t) \approx \mathbf{x}(t) + \mathbf{f}(\mathbf{x}(t), t)\Delta t + \sigma(\mathbf{x}(t), t)\sqrt{\Delta t}\mathcal{N}(\mathbf{0}, \mathbf{I})$$

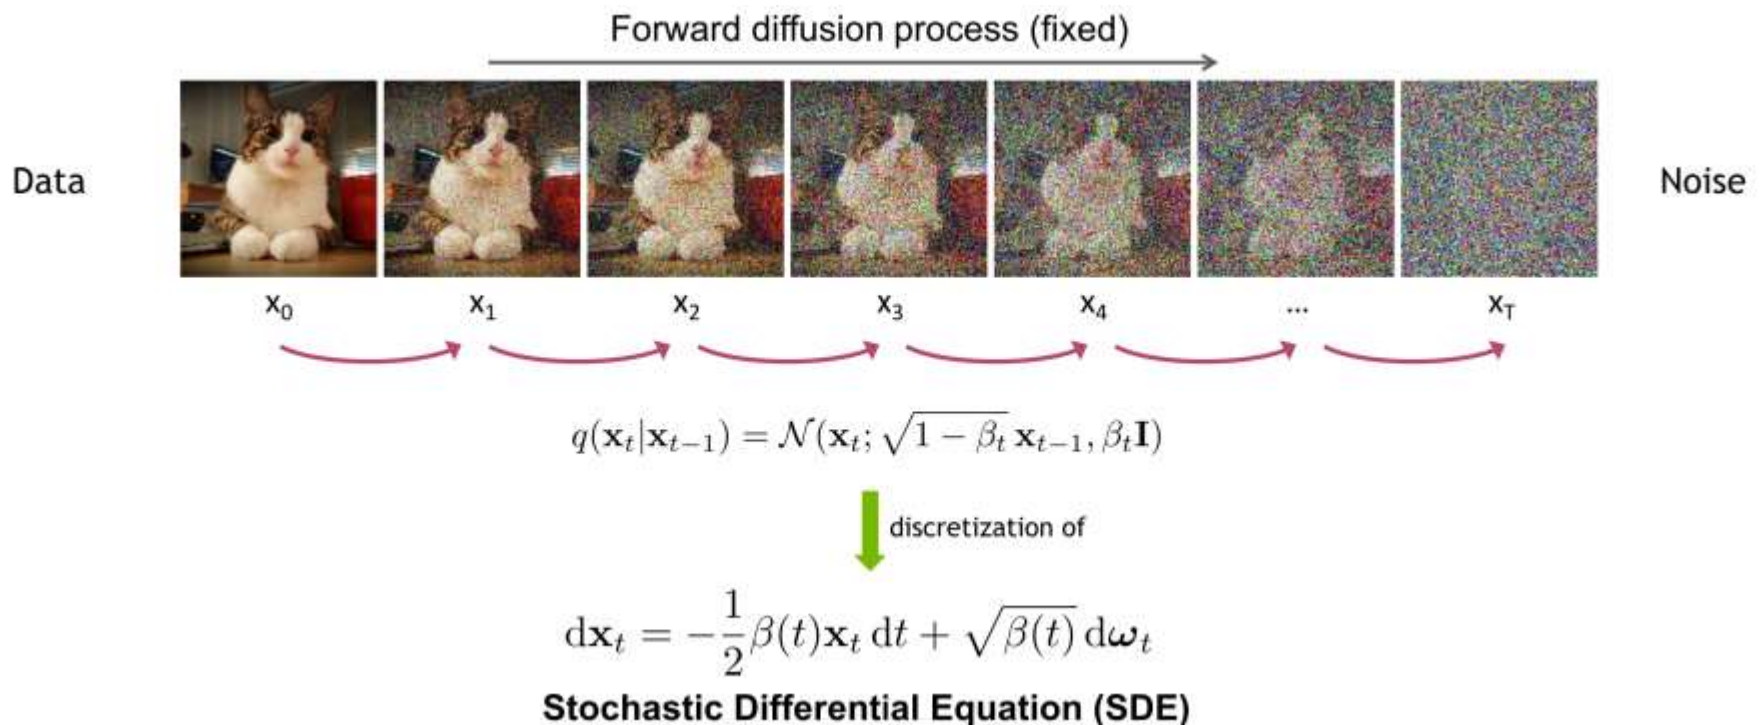
Wiener Process  
(Gaussian  
White Noise)





# Score-based Diffusion Models

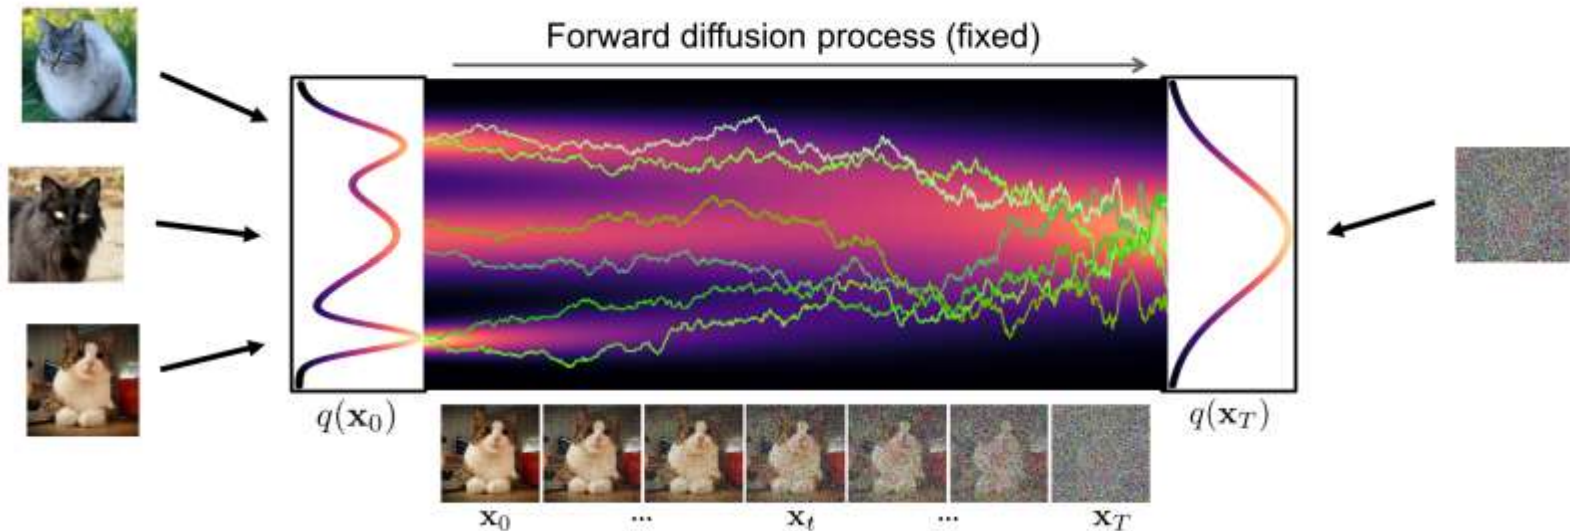
- **Forward Diffusion Process** as Stochastic Differential Equation (continuous form)
- Consider the limit of many small steps:



Song et al., “Score-Based Generative Modeling through Stochastic Differential Equations” , ICLR, 2021

# Score-based Diffusion Models

- **Forward Diffusion Process** as Stochastic Differential Equation (continuous form)

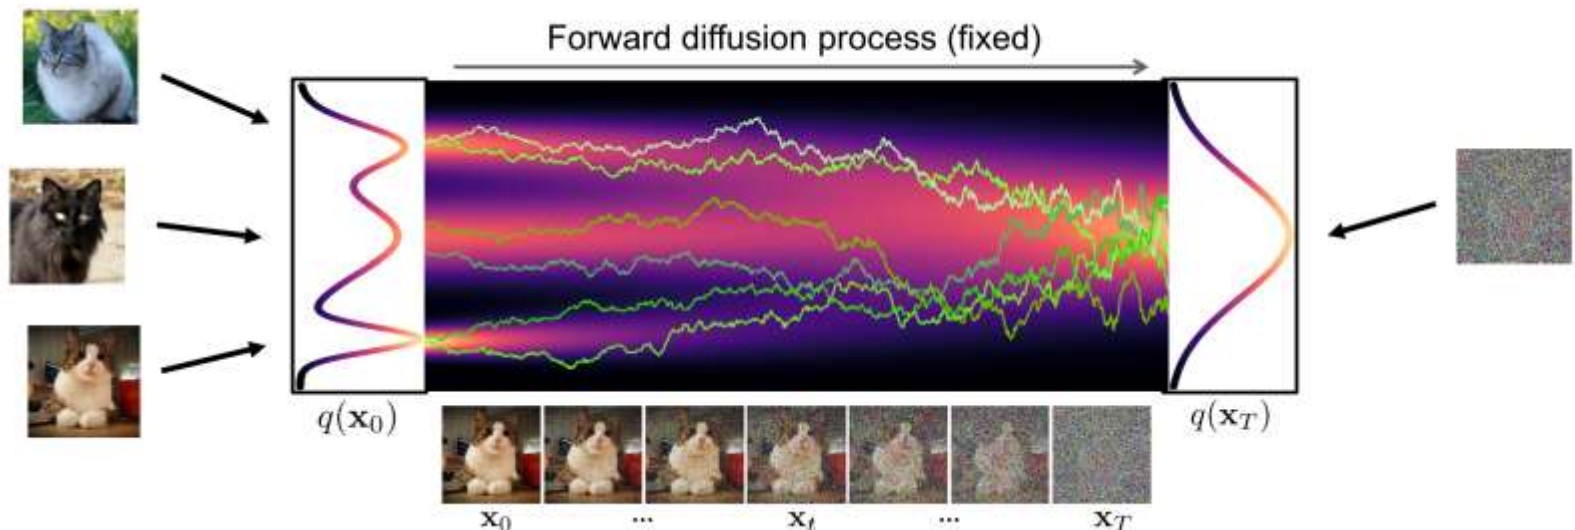


**Forward Diffusion SDE:**

$$d\mathbf{x}_t = \underbrace{-\frac{1}{2}\beta(t)\mathbf{x}_t dt}_{\text{drift term (pulls towards mode)}} + \underbrace{\sqrt{\beta(t)} d\omega_t}_{\text{diffusion term (injects noise)}}$$

# Score-based Diffusion Models

- The Generative Reverse Stochastic Differential Equation



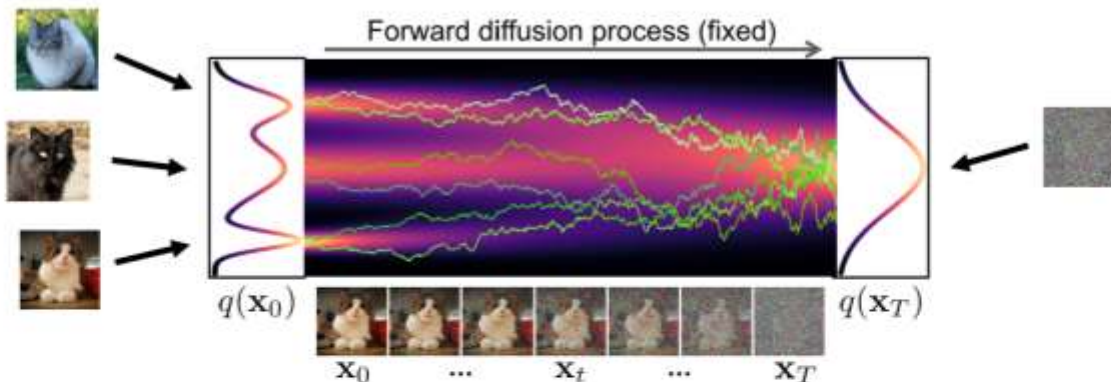
Forward Diffusion SDE:

$$d\mathbf{x}_t = -\frac{1}{2}\beta(t)\mathbf{x}_t dt + \sqrt{\beta(t)} d\omega_t$$

- But what about the **reverse direction**, necessary for generation?

# Score-based Diffusion Models

## ■ The Generative Reverse Stochastic Differential Equation



Forward Diffusion SDE:

$$dx_t = -\frac{1}{2}\beta(t)x_t dt + \sqrt{\beta(t)} d\omega_t$$

Reverse Generative  
Diffusion SDE:

$$dx_t = \underbrace{\left[ -\frac{1}{2}\beta(t)x_t - \beta(t) \underbrace{\nabla_{\mathbf{x}_t} \log q_t(\mathbf{x}_t)}_{\text{"Score Function"}} \right]}_{\text{drift term}} dt + \underbrace{\sqrt{\beta(t)} d\bar{\omega}_t}_{\text{diffusion term}}$$

➡ Simulate reverse diffusion process: Data generation from random noise!

# Score function example

- Probability density function (pdf)


$$p(\mathbf{X})$$

- Score function

$$\nabla_{\mathbf{x}} \log p(\mathbf{X})$$

- e.g. Gaussian distribution

$$p_{\theta}(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$
$$\longrightarrow s_{\theta}(x) = -\frac{x-\mu}{\sigma^2}$$



Density function: **Contours**  
Score function: **Vector field**



# Score-based Diffusion Models

## ■ The Generative Reverse Stochastic Differential Equation



**But how to get the score function  $\nabla_{\mathbf{x}_t} \log q_t(\mathbf{x}_t)$ ?**

Forward Diffusion SDE:

$$d\mathbf{x}_t = -\frac{1}{2}\beta(t)\mathbf{x}_t dt + \sqrt{\beta(t)} d\omega_t$$

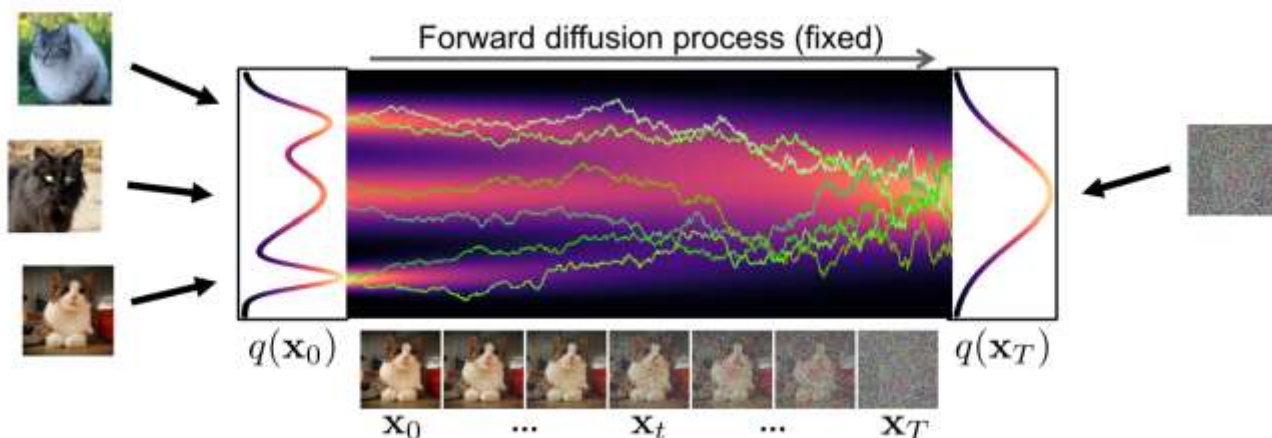
Reverse Generative  
Diffusion SDE:

$$d\mathbf{x}_t = \underbrace{\left[ -\frac{1}{2}\beta(t)\mathbf{x}_t - \beta(t) \underbrace{\nabla_{\mathbf{x}_t} \log q_t(\mathbf{x}_t)}_{\text{"Score Function"}} \right]}_{\text{drift term}} dt + \underbrace{\sqrt{\beta(t)} d\bar{\omega}_t}_{\text{diffusion term}}$$

➡ Simulate reverse diffusion process: Data generation from random noise!

# Score Matching

- Naïve idea, learn model for the score function by direct regression?



$$\min_{\theta} \underbrace{\mathbb{E}_{t \sim \mathcal{U}(0, T)}}_{\text{diffusion time } t} \underbrace{\mathbb{E}_{\mathbf{x}_t \sim q_t(\mathbf{x}_t)}}_{\text{diffused data } \mathbf{x}_t} \underbrace{\|s_{\theta}(\mathbf{x}_t, t) - \nabla_{\mathbf{x}_t} \log q_t(\mathbf{x}_t)\|_2^2}_{\substack{\text{neural} \\ \text{network} \quad \text{score of} \\ \text{diffused data} \\ \text{(marginal)}}}$$


➡ But  $\nabla_{\mathbf{x}_t} \log q_t(\mathbf{x}_t)$  (score of the *marginal diffused density*  $q_t(\mathbf{x}_t)$ ) is not tractable!



# Denoising Score Matching

- Instead, diffuse individual data points  $\mathbf{x}_0$

Diffused  $q_t(\mathbf{x}_t|\mathbf{x}_0)$  is Trackable!



"Variance Preserving" SDE:

$$d\mathbf{x}_t = -\frac{1}{2}\beta(t)\mathbf{x}_t dt + \sqrt{\beta(t)} d\omega_t$$

$$q_t(\mathbf{x}_t|\mathbf{x}_0) = \mathcal{N}(\mathbf{x}_t; \gamma_t\mathbf{x}_0, \sigma_t^2\mathbf{I})$$

- Denoising Score Matching:

$$\min_{\theta} \underbrace{\mathbb{E}_{t \sim \mathcal{U}(0,T)}}_{\text{diffusion time } t} \underbrace{\mathbb{E}_{\mathbf{x}_0 \sim q_0(\mathbf{x}_0)}}_{\text{data sample } \mathbf{x}_0} \underbrace{\mathbb{E}_{\mathbf{x}_t \sim q_t(\mathbf{x}_t|\mathbf{x}_0)}}_{\text{diffused data sample } \mathbf{x}_t} \underbrace{\|\mathbf{s}_{\theta}(\mathbf{x}_t, t) - \nabla_{\mathbf{x}_t} \log q_t(\mathbf{x}_t|\mathbf{x}_0)\|_2^2}_{\substack{\text{neural} \\ \text{network}} \quad \text{score of diffused} \\ \text{data sample}}$$

➔ After expectations,  $\mathbf{s}_{\theta}(\mathbf{x}_t, t) \approx \nabla_{\mathbf{x}_t} \log q_t(\mathbf{x}_t)$ !

# Denoising Score Matching

- Matching the score of a noise-perturbed distribution



$\mathbf{x}$

$p_{\text{data}}(\mathbf{x})$

$q_{\sigma}(\tilde{\mathbf{x}} \mid \mathbf{x})$

$q_{\sigma}(\tilde{\mathbf{x}})$

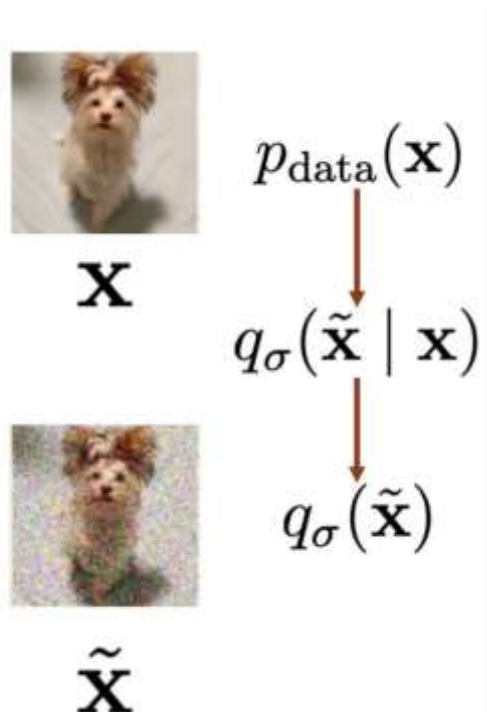


$\tilde{\mathbf{x}}$

$$\begin{aligned}
 & \frac{1}{2} E_{\tilde{\mathbf{x}} \sim q_{\sigma}} [\|\nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}}) - \mathbf{s}_{\theta}(\tilde{\mathbf{x}})\|_2^2] \\
 &= \frac{1}{2} \int q_{\sigma}(\tilde{\mathbf{x}}) \|\nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}}) - \mathbf{s}_{\theta}(\tilde{\mathbf{x}})\|_2^2 d\tilde{\mathbf{x}} \\
 &= \frac{1}{2} \int q_{\sigma}(\tilde{\mathbf{x}}) \|\nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}})\|_2^2 d\tilde{\mathbf{x}} + \frac{1}{2} \int q_{\sigma}(\tilde{\mathbf{x}}) \|\mathbf{s}_{\theta}(\tilde{\mathbf{x}})\|_2^2 d\tilde{\mathbf{x}} \\
 &\quad - \int q_{\sigma}(\tilde{\mathbf{x}}) \nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}})^{\top} \mathbf{s}_{\theta}(\tilde{\mathbf{x}}) d\tilde{\mathbf{x}} \\
 &= \text{const.} + \frac{1}{2} E_{\tilde{\mathbf{x}} \sim q_{\sigma}} [\|\mathbf{s}_{\theta}(\tilde{\mathbf{x}})\|_2^2] - \int q_{\sigma}(\tilde{\mathbf{x}}) \nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}})^{\top} \mathbf{s}_{\theta}(\tilde{\mathbf{x}}) d\tilde{\mathbf{x}}
 \end{aligned}$$

# Denoising Score Matching

- Matching the score of a noise-perturbed distribution



$$\begin{aligned}
 & - \int q_{\sigma}(\tilde{\mathbf{x}}) \nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}})^{\top} \mathbf{s}_{\theta}(\tilde{\mathbf{x}}) d\tilde{\mathbf{x}} \\
 &= - \int q_{\sigma}(\tilde{\mathbf{x}}) \frac{1}{q_{\sigma}(\tilde{\mathbf{x}})} \nabla_{\tilde{\mathbf{x}}} q_{\sigma}(\tilde{\mathbf{x}})^{\top} \mathbf{s}_{\theta}(\tilde{\mathbf{x}}) d\tilde{\mathbf{x}} \\
 &= - \int \nabla_{\tilde{\mathbf{x}}} q_{\sigma}(\tilde{\mathbf{x}})^{\top} \mathbf{s}_{\theta}(\tilde{\mathbf{x}}) d\tilde{\mathbf{x}} \\
 &= - \int \nabla_{\tilde{\mathbf{x}}} \left( \int p_{\text{data}}(\mathbf{x}) q_{\sigma}(\tilde{\mathbf{x}} | \mathbf{x}) d\mathbf{x} \right)^{\top} \mathbf{s}_{\theta}(\tilde{\mathbf{x}}) d\tilde{\mathbf{x}} \\
 &= - \int \left( \int p_{\text{data}}(\mathbf{x}) \nabla_{\tilde{\mathbf{x}}} q_{\sigma}(\tilde{\mathbf{x}} | \mathbf{x}) d\mathbf{x} \right)^{\top} \mathbf{s}_{\theta}(\tilde{\mathbf{x}}) d\tilde{\mathbf{x}} \\
 &= - \int \left( \int p_{\text{data}}(\mathbf{x}) q_{\sigma}(\tilde{\mathbf{x}} | \mathbf{x}) \nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}} | \mathbf{x}) d\mathbf{x} \right)^{\top} \mathbf{s}_{\theta}(\tilde{\mathbf{x}}) d\tilde{\mathbf{x}} \\
 &= - \iint p_{\text{data}}(\mathbf{x}) q_{\sigma}(\tilde{\mathbf{x}} | \mathbf{x}) \nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}} | \mathbf{x})^{\top} \mathbf{s}_{\theta}(\tilde{\mathbf{x}}) d\mathbf{x} d\tilde{\mathbf{x}} \\
 &= - E_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x}), \tilde{\mathbf{x}} \sim q_{\sigma}(\tilde{\mathbf{x}} | \mathbf{x})} [\nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}} | \mathbf{x})^{\top} \mathbf{s}_{\theta}(\tilde{\mathbf{x}})]
 \end{aligned}$$

# Denoising Score Matching

- Matching the score of a noise-perturbed distribution



$\mathbf{x}$

$p_{\text{data}}(\mathbf{x})$

$q_{\sigma}(\tilde{\mathbf{x}} | \mathbf{x})$

$q_{\sigma}(\tilde{\mathbf{x}})$



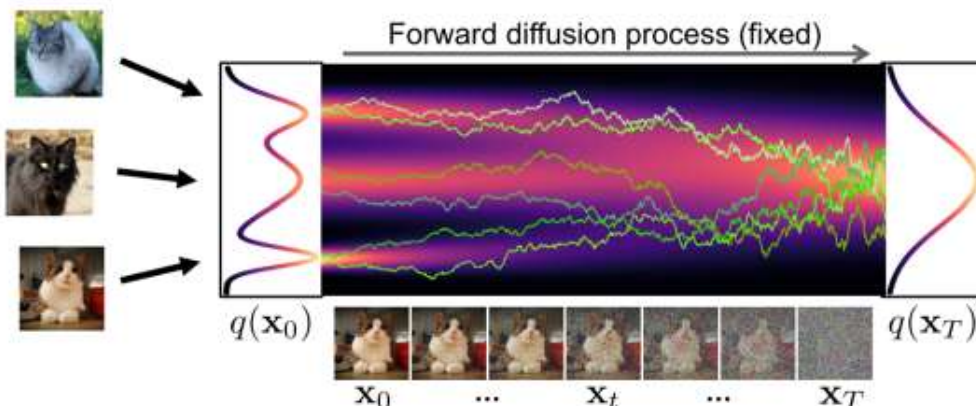
$\tilde{\mathbf{x}}$

$$\begin{aligned}
 & \frac{1}{2} E_{\tilde{\mathbf{x}} \sim q_{\sigma}} [\|\nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}}) - \mathbf{s}_{\theta}(\tilde{\mathbf{x}})\|_2^2] \\
 &= \text{const.} + \frac{1}{2} E_{\tilde{\mathbf{x}} \sim q_{\sigma}} [\|\mathbf{s}_{\theta}(\tilde{\mathbf{x}})\|_2^2] - \int q_{\sigma}(\tilde{\mathbf{x}}) \nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}})^{\top} \mathbf{s}_{\theta}(\tilde{\mathbf{x}}) d\tilde{\mathbf{x}} \\
 &= \text{const.} + \frac{1}{2} E_{\tilde{\mathbf{x}} \sim q_{\sigma}} [\|\mathbf{s}_{\theta}(\tilde{\mathbf{x}})\|_2^2] - E_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x}), \tilde{\mathbf{x}} \sim q_{\sigma}(\tilde{\mathbf{x}}|\mathbf{x})} [\nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}} | \mathbf{x})^{\top} \mathbf{s}_{\theta}(\tilde{\mathbf{x}})] \\
 &= \text{const.} + \frac{1}{2} E_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x}), \tilde{\mathbf{x}} \sim q_{\sigma}(\tilde{\mathbf{x}}|\mathbf{x})} [\|\mathbf{s}_{\theta}(\tilde{\mathbf{x}}) - \nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}} | \mathbf{x})\|_2^2] \\
 &\quad - \frac{1}{2} E_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x}), \tilde{\mathbf{x}} \sim q_{\sigma}(\tilde{\mathbf{x}}|\mathbf{x})} [\|\nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}} | \mathbf{x})\|_2^2] \\
 &= \text{const.} + \frac{1}{2} E_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x}), \tilde{\mathbf{x}} \sim q_{\sigma}(\tilde{\mathbf{x}}|\mathbf{x})} [\|\mathbf{s}_{\theta}(\tilde{\mathbf{x}}) - \nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}} | \mathbf{x})\|_2^2] + \text{const.} \\
 &= \frac{1}{2} E_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x}), \tilde{\mathbf{x}} \sim q_{\sigma}(\tilde{\mathbf{x}}|\mathbf{x})} [\|\mathbf{s}_{\theta}(\tilde{\mathbf{x}}) - \nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}} | \mathbf{x})\|_2^2] + \text{const.}
 \end{aligned}$$

# Denoising Score Matching

- Instead, diffuse individual data points  $\mathbf{x}_0$

Diffused  $q_t(\mathbf{x}_t|\mathbf{x}_0)$  is Trackable!



"Variance Preserving" SDE:

$$d\mathbf{x}_t = -\frac{1}{2}\beta(t)\mathbf{x}_t dt + \sqrt{\beta(t)} d\omega_t$$

$$q_t(\mathbf{x}_t|\mathbf{x}_0) = \mathcal{N}(\mathbf{x}_t; \gamma_t\mathbf{x}_0, \sigma_t^2\mathbf{I})$$

- Denoising Score Matching:

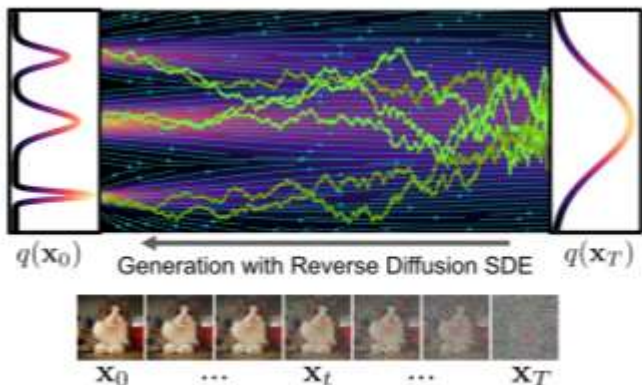
$$\min_{\theta} \underbrace{\mathbb{E}_{t \sim \mathcal{U}(0,T)}}_{\text{diffusion time } t} \underbrace{\mathbb{E}_{\mathbf{x}_0 \sim q_0(\mathbf{x}_0)}}_{\text{data sample } \mathbf{x}_0} \underbrace{\mathbb{E}_{\mathbf{x}_t \sim q_t(\mathbf{x}_t|\mathbf{x}_0)}}_{\text{diffused data sample } \mathbf{x}_t} \underbrace{\|\mathbf{s}_{\theta}(\mathbf{x}_t, t) - \nabla_{\mathbf{x}_t} \log q_t(\mathbf{x}_t|\mathbf{x}_0)\|_2^2}_{\text{neural network score of diffused data sample}}$$

→ 
$$\min_{\theta} \mathbb{E}_{t \sim \mathcal{U}(0,T)} \mathbb{E}_{\mathbf{x}_0 \sim q_0(\mathbf{x}_0)} \mathbb{E}_{\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})} \frac{1}{\sigma_t^2} \|\epsilon - \epsilon_{\theta}(\mathbf{x}_t, t)\|_2^2$$



# Synthesis with SDE vs. ODE

- How to solve the generative SDE or ODE in practice?




**Generative Diffusion SDE:**

$$dx_t = -\frac{1}{2}\beta(t) [x_t + 2s_\theta(x_t, t)] dt + \sqrt{\beta(t)} d\bar{\omega}_t$$

→ **Euler-Maruyama:**

$$x_{t-1} = x_t + \frac{1}{2}\beta(t) [x_t + 2s_\theta(x_t, t)] \Delta t + \sqrt{\beta(t)\Delta t} \mathcal{N}(0, \mathbf{I})$$

→ **Ancestral Sampling (Part 1)** is also a generative SDE sampler!



**Probability Flow ODE:**

$$dx_t = -\frac{1}{2}\beta(t) [x_t + s_\theta(x_t, t)] dt$$

→ **Euler's Method:**

$$x_{t-1} = x_t + \frac{1}{2}\beta(t) [x_t + s_\theta(x_t, t)] \Delta t$$

→ In practice: Higher-Order ODE solvers (Runge-Kutta, linear multistep methods, exponential integrators, ...)

# Synthesis with SDE vs. ODE

- How to solve the generative SDE or ODE in practice?

- Runge-Kutta adaptive step-size ODE solver
- Higher-Order adaptive step-size SDE solver
- Reparametrized, smoother ODE
- Higher-Order ODE solver with linear multisteping
- Exponential ODE Integrators
- Higher-order ODE solver with Heun's Method

⇒ Euler-Maruyama:

$$\mathbf{x}_{t-1} \quad \mathbf{x}_t + \frac{1}{2} \beta(t) [\mathbf{x}_t + 2s_\theta(\mathbf{x}_t, t)] \Delta t + \sqrt{\beta(t) \Delta t} \mathcal{N}(\mathbf{0}, \mathbf{I})$$

⇒ Ancestral Sampling (Part 1) is also a generative SDE sampler!

⇒ Euler's Method:

$$\mathbf{x}_{t-1} = \mathbf{x}_t + \frac{1}{2} \beta(t) [\mathbf{x}_t + s_\theta(\mathbf{x}_t, t)] \Delta t$$

⇒ In practice: Higher-Order ODE solvers (Runge-Kutta, linear multistep methods, exponential integrators, ...)






# Outline

- Diffusion: Denoising Diffusion Probabilistic Models
- Score-based Diffusion Models
- **Accelerated Sampling**
- Conditional Generation and Guidance

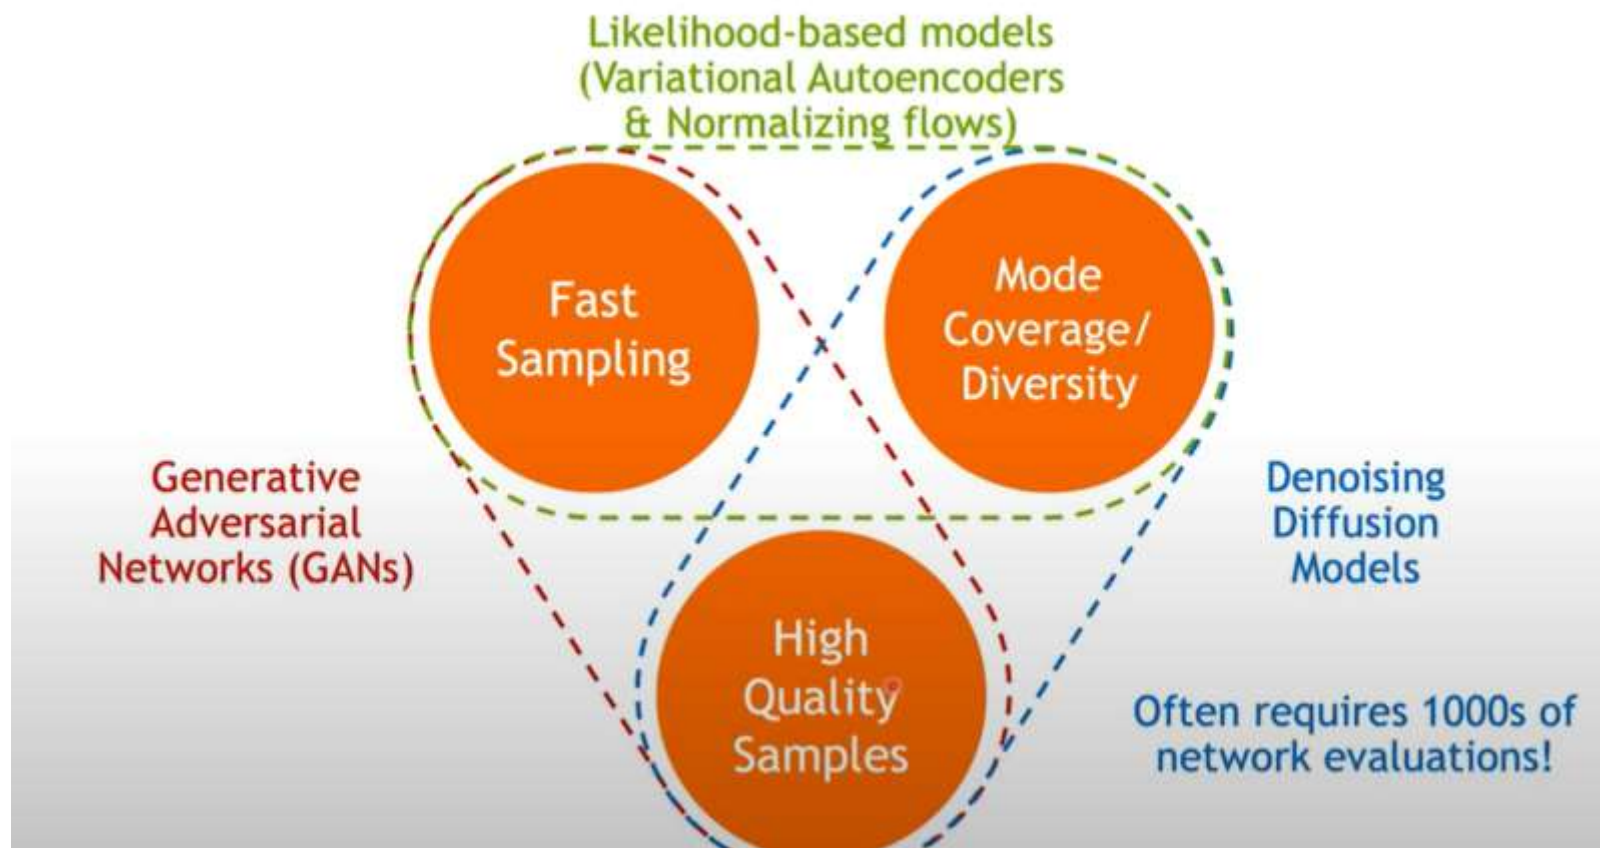
# What makes a good generative model?

- The generative learning trilemma
- Tackle the trilemma by accelerating diffusion models



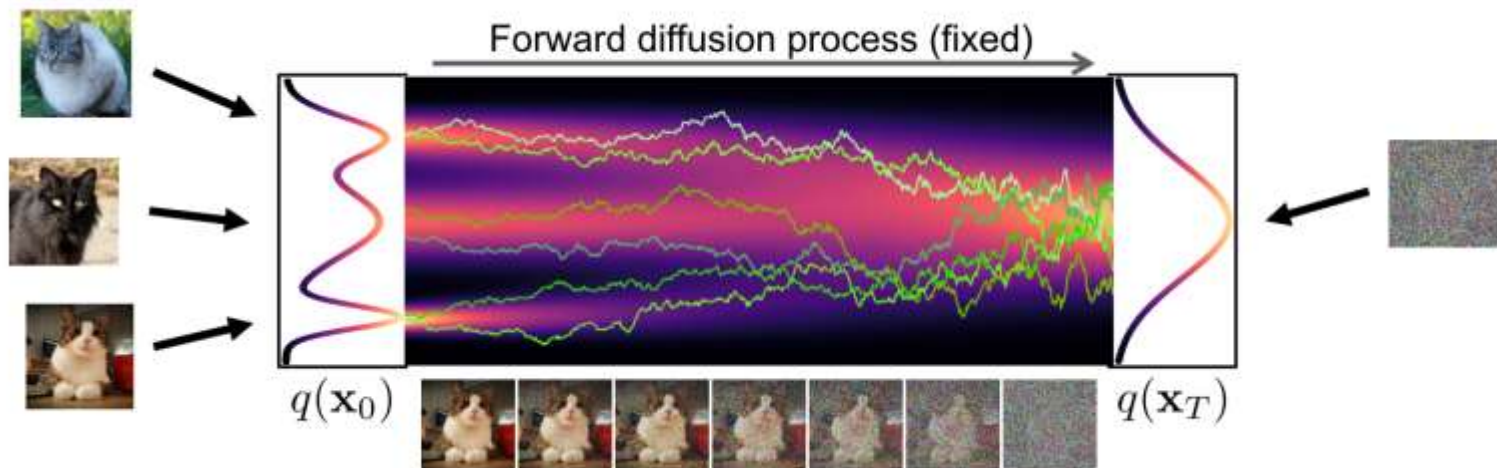
# What makes a good generative model?

- The generative learning trilemma
- Tackle the trilemma by accelerating diffusion models



# Accelerated Sampling


- Advanced ODE/SDE Solvers
- Distillation Techniques
- Low-dim Diffusion Processes
- . . . . .



# Generative ODEs

- Solve ODEs with as little function evaluations as possible!

$$dx = \epsilon_{\theta}(x, t)dt$$





Song et al., "Denoising Diffusion Implicit Models (DDIM)", ICLR 2021

<https://arxiv.org/abs/2010.02502>

# Generative ODEs

- Solve ODEs with as little function evaluations as possible!

$$dx = \epsilon_{\theta}(x, t)dt$$



First Truncated Taylor method  
(Euler Method)

Higher order methods  
(RK4, Multistepping, Heun)



Approximate the higher order terms numerically



# Generative ODEs


- Solve ODEs with as little function evaluations as possible!

$$dx = \epsilon_{\theta}(x, t)dt$$

- Runge-Kutta adaptive step-size ODE solver:  
<https://arxiv.org/abs/2011.13456>
- Higher-Order adaptive step-size SDE solver:  
<https://arxiv.org/abs/2105.14080>
- Reparametrized, smoother ODE: gDDIM  
<https://arxiv.org/abs/2206.05564>
- Higher-Order ODE solver with linear multisteping:  
<https://arxiv.org/abs/2202.09778>
- Exponential ODE Integrators: DPM, DPM-Solver++  
<https://arxiv.org/abs/2206.00927>
- Higher-Order ODE solver with Heun's Method:  
<https://arxiv.org/abs/2206.00364>
- Many more: <https://arxiv.org/abs/2305.19947>

# Distillation Techniques

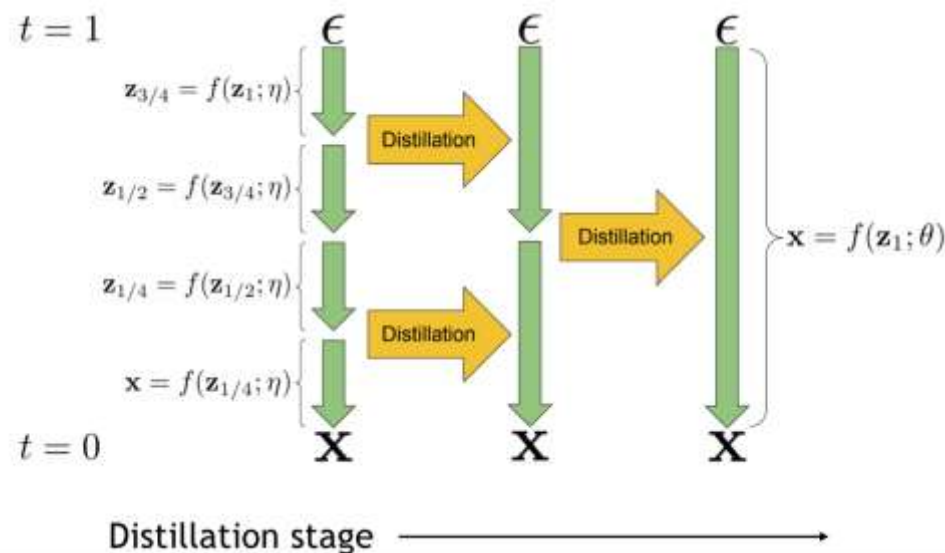
- ODE Distillation



Can we train a neural network to directly predict  $\mathbf{x}_{t'}$  given  $\mathbf{x}_t$  ?

# Progressive Distillation

- Distill a deterministic ODE sampler to the same model architecture
- At each stage, a “**student**” model is learned to distill two adjacent sampling steps of the “**teacher**” model to one sampling step.
- At next stage, the “**student**” model from previous stage will serve as the new “**teacher**” model.



Salimans & Ho, “Progressive distillation for fast sampling of diffusion models”, ICLR 2022. <https://arxiv.org/abs/2202.00512>

# Progressive Distillation in Latent Space




(a) 2 denoising steps

(b) 4 denoising steps

(c) 8 denoising steps

Meng et al., “On Distillation of Guided Diffusion Models” , CVPR 2023 (Award Candidate). <https://arxiv.org/abs/2210.03142>

# Consistency Distillation



Points on the same trajectory should generate the same  $\mathbf{x}_0$

Assume  $f_\theta(\mathbf{x}_t, t)$  is the current estimation of  $\mathbf{x}_0$

Basic idea:

- Find  $\mathbf{x}_t$  and  $\mathbf{x}_{t'}$  on a trajectory by solving generative ODE in  $[t, t']$
- Minimize:

$$\min_{\theta} \|f_{\text{EMA}}(\mathbf{x}_t, t) - f_\theta(\mathbf{x}_{t'}, t')\|_2^2$$


Song et al., Consistency Models, ICML 2023,

<https://arxiv.org/abs/2303.01469>



# Distillation Techniques

- SDE Distillation

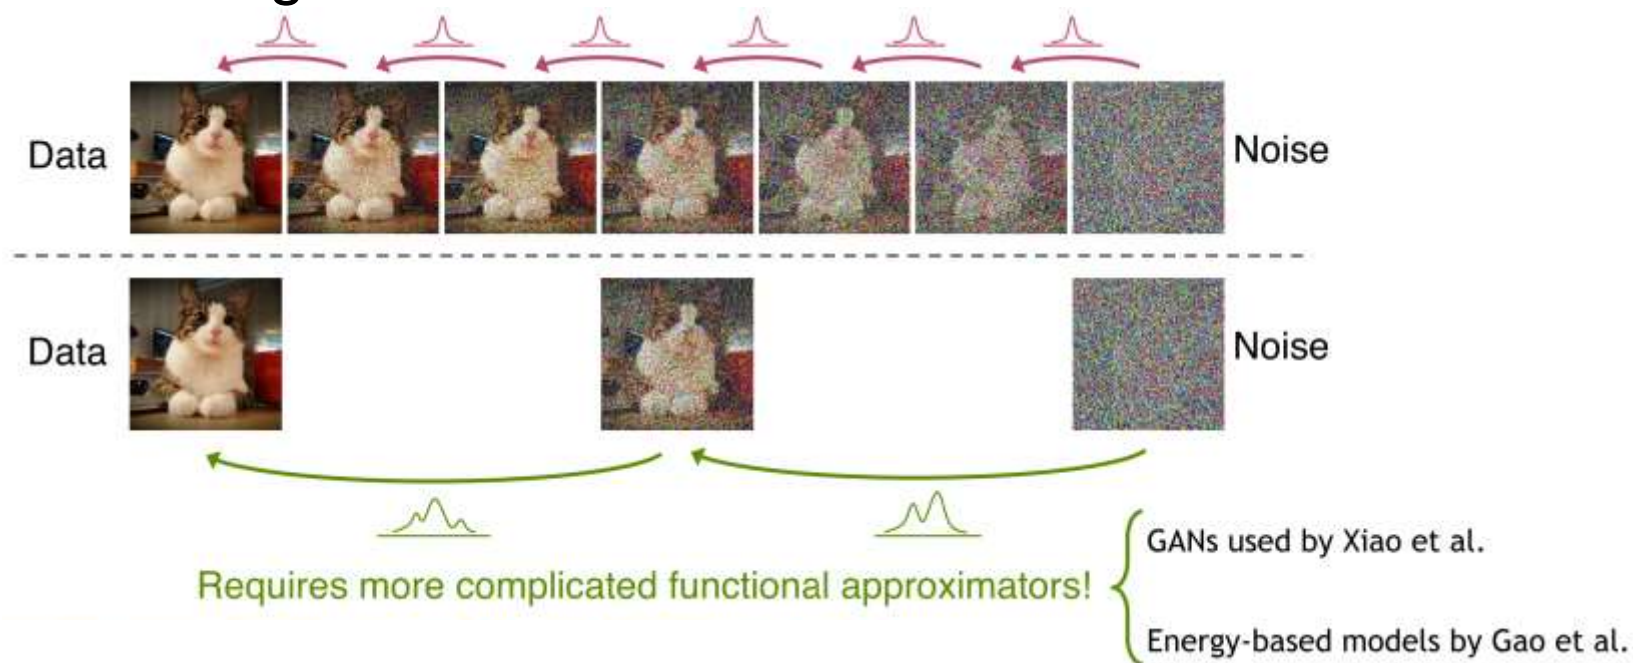


Can we train a neural network to directly predict  $x_{t'}$  given  $x_t$  ?



# Approximation of Reverse Process

- Normal assumption in denoising distribution holds only for small step
- Denoising Process with Uni-modal Normal Distribution



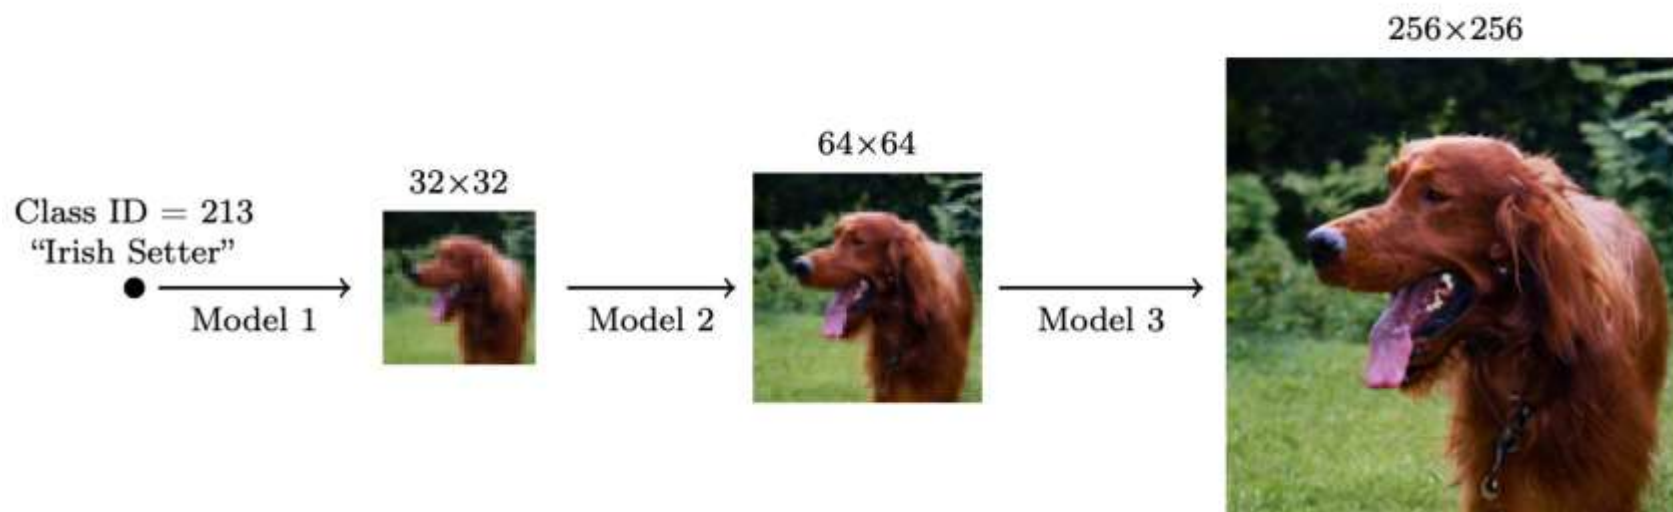
Xiao et al., "Tackling the Generative Learning Trilemma with Denoising Diffusion GANs", ICLR 2022.

Gao et al., "Learning energy-based models by diffusion recovery likelihood", ICLR 2021.

<https://arxiv.org/abs/2112.07804>

# Low-dim. Diffusion Processes

- Cascaded Generation
- Cascaded Diffusion Models outperform Big-GAN in FID and IS and VQ-VAE2 in Classification Accuracy Score.



Ho et al., "Cascaded Diffusion Models for High Fidelity Image Generation", 2021. <https://arxiv.org/abs/2106.15282>

Ramesh et al., "Hierarchical Text-Conditional Image Generation with CLIP Latents", 2022,

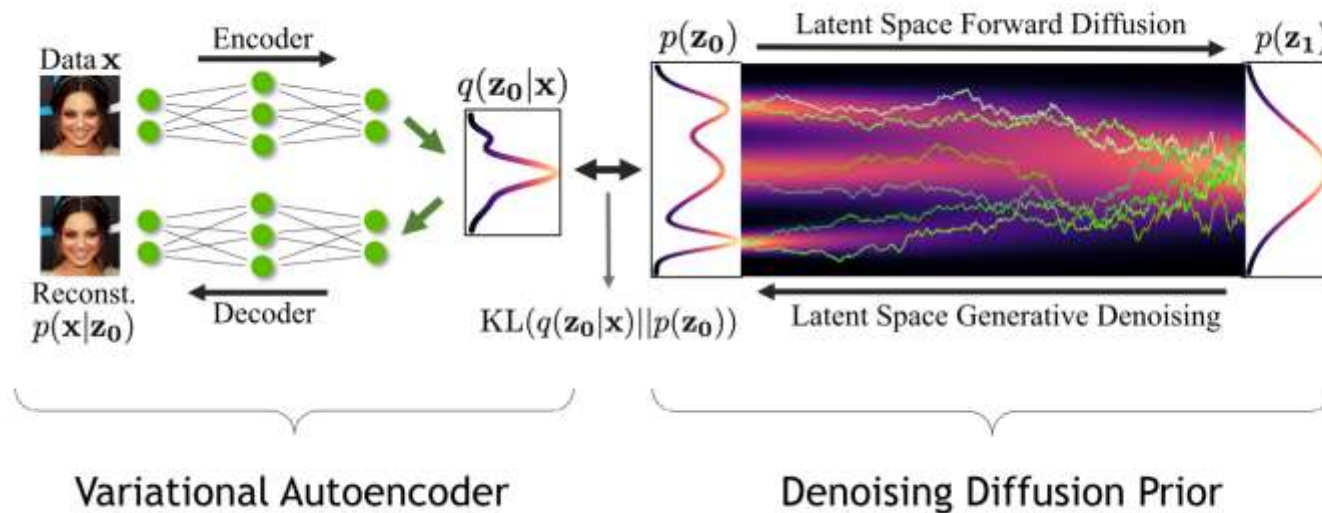
<https://arxiv.org/abs/2204.06125>

Saharia et al., "Photorealistic Text-to-Image Diffusion Models with Deep Language Understanding", 2022,

<https://arxiv.org/abs/2205.11487>

# Latent Diffusion Models

- VAE + score-based prior



## Main Idea

Encoder maps the input data to an embedding space

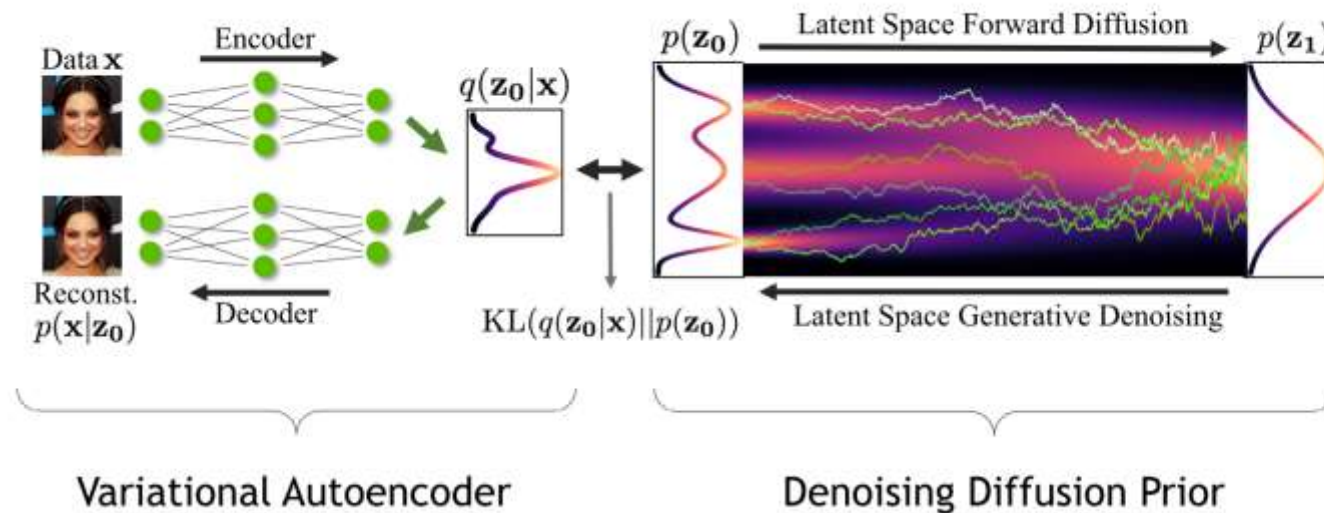
Denoising diffusion models are applied in the latent space



Vahdat et al., "Score-based generative modeling in latent space",  
NeurIPS 2021. <https://arxiv.org/abs/2106.05931>

# Latent Diffusion Models

## ■ VAE + score-based prior



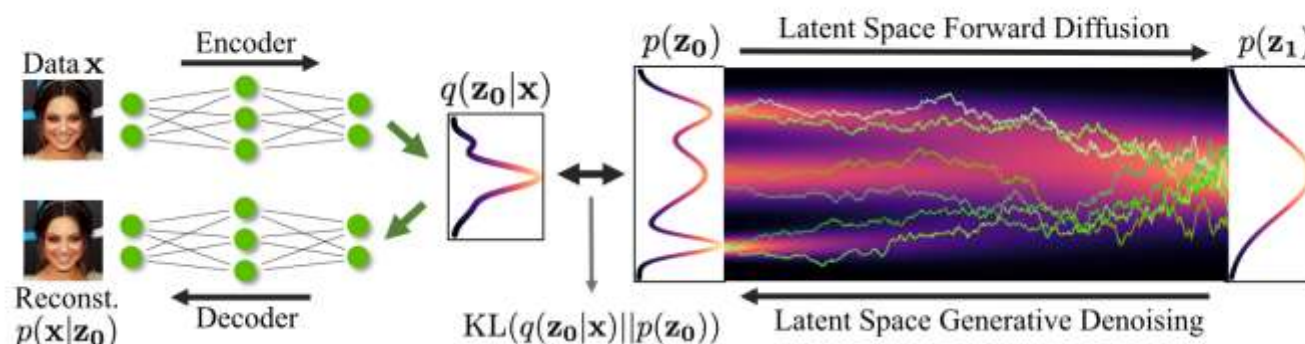
### Advantages:

- (1) The distribution of latent embeddings close to Normal distribution → *Simpler denoising, Faster synthesis!*
- (2) Latent space → *More expressivity and flexibility in design!*
- (3) Tailored Autoencoders → *More expressivity, Application to any data type (graphs, text, 3D data, etc.) !*

Vahdat et al., "Score-based generative modeling in latent space",  
NeurIPS 2021. <https://arxiv.org/abs/2106.05931>

# Latent Diffusion Models

## ■ End-to-End Training objective



$$\begin{aligned}\mathcal{L}(\mathbf{x}, \phi, \theta, \psi) &= \mathbb{E}_{q_\phi(\mathbf{z}_0|\mathbf{x})} [-\log p_\psi(\mathbf{x}|\mathbf{z}_0)] + \text{KL}(q_\phi(\mathbf{z}_0|\mathbf{x})||p_\theta(\mathbf{z}_0)) \\ &= \underbrace{\mathbb{E}_{q_\phi(\mathbf{z}_0|\mathbf{x})} [-\log p_\psi(\mathbf{x}|\mathbf{z}_0)]}_{\text{reconstruction term}} + \underbrace{\mathbb{E}_{q_\phi(\mathbf{z}_0|\mathbf{x})} [\log q_\phi(\mathbf{z}_0|\mathbf{x})]}_{\text{negative encoder entropy}} + \underbrace{\mathbb{E}_{q_\phi(\mathbf{z}_0|\mathbf{x})} [-\log p_\theta(\mathbf{z}_0)]}_{\text{cross entropy}}\end{aligned}$$

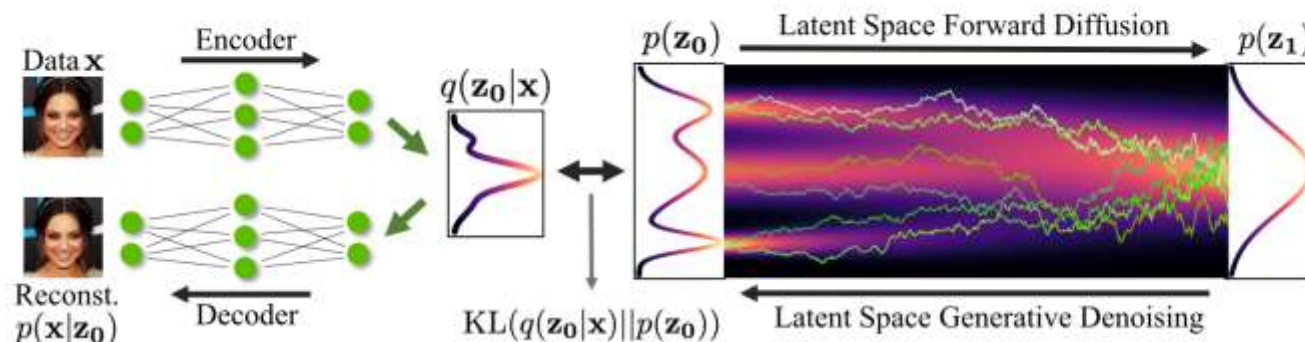
$$\begin{aligned}CE(q(\mathbf{z}_0|\mathbf{x})||p(\mathbf{z}_0)) &= \underbrace{\mathbb{E}_{t \sim \mathcal{U}[0,1]}}_{\text{time sampling}} \left[ \underbrace{\frac{g(t)^2}{2}}_{\text{Forward diffusion}} \underbrace{\mathbb{E}_{q(\mathbf{z}_t, \mathbf{z}_0|\mathbf{x})}}_{\text{Diffusion kernel}} \left[ \underbrace{\|\nabla_{\mathbf{z}_t} \log q(\mathbf{z}_t|\mathbf{z}_0) - \nabla_{\mathbf{z}_t} \log p(\mathbf{z}_t)\|_2^2}_{\text{Trainable score function}} \right] \right] + \underbrace{\frac{D}{2} \log(2\pi e \sigma_0^2)}_{\text{Constant}}\end{aligned}$$

Vahdat et al., “Score-based generative modeling in latent space”,  
NeurIPS 2021. <https://arxiv.org/abs/2106.05931>



# Latent Diffusion Models

- Stable Diffusion: the efficiency and expressivity of LDM+ open-source access fueled a large body of work in the community



- Two stage training: train autoencoder first, then train the diffusion prior
- Focus on compression without of any loss in reconstruction quality
- Demonstrated the expressivity of latent diffusion models on many conditional problems

Rombach et al., “High-Resolution Image Synthesis with Latent Diffusion Models”, CVPR 2022. <https://arxiv.org/abs/2112.10752>




# Latent Diffusion Models

- Training models in the pixel space is excessively computationally expensive (can easily multiple days on a V100 GPU)
  - Even image synthesis is very slow compared to GANs
  - Images are high dimensional → more things to model
- Researchers observed that most “bits” of an image contribute to its perceptual characteristics since aggressively compressing it usually maintains its semantic and conceptual composition
  - In layman’s terms, there are more bits for describing pixel-level details while less bits for describing “the meaning” within an image
  - Generative models should learn the latter
- Can we separate these two components?

# Huge success of text-2-img!

- Stable Diffusion model (CVPR2022)



## High-Resolution Image Synthesis with Latent Diffusion Models

Robin Rombach<sup>1</sup> \*    Andreas Blattmann<sup>1</sup> \*    Dominik Lorenz<sup>1</sup>    Patrick Esser<sup>℞</sup>    Björn Ommer<sup>1</sup>

<sup>1</sup>Ludwig Maximilian University of Munich & IWR, Heidelberg University, Germany    <sup>℞</sup>Runway ML

<https://github.com/CompVis/latent-diffusion>

# Latent Diffusion Models

Latent Diffusion Models can be divided **into two stages**:

- 1) Training perceptual compression models that strip away irrelevant high-level details and learn a latent space that is semantically equivalent to the high level image pixel-space
  - a. The loss is a combination of a reconstruction loss, an adversarial loss (remember GANs?) that promotes high quality decoder reconstruction, and regularization terms

$$L_{\text{Autoencoder}} = \min_{\mathcal{E}, \mathcal{D}} \max_{\psi} \left( L_{\text{rec}}(x, \mathcal{D}(\mathcal{E}(x))) - L_{\text{adv}}(\mathcal{D}(\mathcal{E}(x))) + \log D_{\psi}(x) + L_{\text{reg}}(x; \mathcal{E}, \mathcal{D}) \right)$$

- 2) Performing a diffusion process *in this latent space*. There are several benefits to this:
  - a. The diffusion process is only focusing on the relevant semantic bits of the data
  - b. Performing diffusion in a low dimensional space is significantly more efficient

# Huge success of text-2-img!

- Stable Diffusion model (CVPR2022)
- Long story between Stability AI, Runway ML and LAION-5B
- AI paradigm: data + algorithm + computing resource



Computer Vision & Learning Group  
Ludwig Maximilian University of Munich  
(LMU)



~4000 A100 from  
Stability AI



Huge text-image  
dataset from  
LAION

# Huge success of text-2-img!

- The multi-modality framework is important
- The trend continues: big data, big modal .....
- Enjoy better text-encoder and suitable generator



**Google: Parti Model**, “Scaling  
Autoregressive Models for Content-  
Rich Text-to-Image Generation”



**Stability AI: DeepFloyd IF Model**  
T5-XXL as Text-encoder; pixel-level  
Diffusion



# Huge success of text-2-img!

- Enjoy cross-modality abilities
- Enjoy downstream conditioning abilities



Conditioning using **ControlNet**



Subject-Driven Generation using **DreamBooth**



Editing Instructions using **InstructPix2Pix** (based on GPT-3)



# Huge success of text-2-img-3D!

- Use NeRF as inherent representation to bridge 2D-DM with 3D scene
- More explicit disentanglement towards geometry, color, lighting ...



3D Editing Instructions  
using **InstructNeRF2NeRF**



NVIDIA **Magic3D** 18 Nov 2022



OpenAI **Point-E** 21 Dec 2022

# Summary

- Preliminary theory of diffusion (don't worry if this is confusing!)
- Some tricks that modern diffusion models employ for image generation:
  - A U-Net architecture equipped with all kinds of modifications
  - Other architecture improvements
  - Several implementation tricks (different noise schedules, covariance parametrizations)
- Latent diffusion models for improving diffusion quality and efficiency